


Inside the Java Virtual Machine
by
Bill Venners

Specially Made for [OR]
by
Outlawtorn

Application development

Inside the
JAVA 2
Virtual Machine



Bill Venners

Inside the Java Virtual Machine

Acknowledgments

Introduction

Part One: Java's Architecture

1 Introduction to Java's Architecture

Why Java?

The Architecture

The Java Virtual Machine

The Class Loader Architecture

The Java Class File

The Java API

The Java Programming Language

Architectural Tradeoffs

Future Trends

On the CD-ROM

The Resources Page

2 Platform independence

Why Platform Independence?

Java's Architectural Support for Platform Independence

Scalability

Factors that Influence Platform Independence

The Java Platform

Native Methods

Other Factors

Seven Steps to Platform Independence

The Politics of Platform Independence

The Resources Page

3 Security

- Why Security?
- The Sandbox
- The Class Loader Architecture
- The Class File Verifier

- Phase One: Internal Checks
 - Phase Two: Verification of Symbolic References

- Safety Features Built Into the Java Virtual Machine
- The Security Manager and the Java API
- The Security API
- Security Beyond the Architecture
- The Resources Page

4 Network-mobility

- Why Network Mobility?

- A New Software Paradigm
- Java's Architectural Support for Network-Mobility
- The Applet: An Example of Network-Mobile Java
- The Resources Page

Part Two: Java Internals

5 The Java Virtual Machine

- What is a Java Virtual Machine?
- The Lifetime of a Java Virtual Machine
- The Architecture of the Java Virtual Machine

- Data Types
 - Word Size
 - The Class Loader Subsystem

- Loading, Linking and Initialization
 - The Primordial Class Loader
 - Class Loader Objects
 - Name Spaces

- The Method Area

- Type Information
 - The Constant Pool
 - Field Information
 - Method Information Class Variables
 - A Reference to Class ClassLoader
 - A Reference to Class Class

Method Tables
An Example of Method Area Use

The Heap

Garbage Collection
Object Representation
Array Representation

The Program Counter
The Java Stack
The Stack Frame

Local Variables
Operand Stack
Frame Data
Possible Implementations of the Java Stack

Native Method Stacks
Execution Engine

The Instruction Set
Execution Techniques
Threads

Native Method Interface

The Real Machine
Eternal Math: A Simulation
On the CD-ROM
The Resources Page

6 The Java Class File

What is a Java Class File?
What's in a Class File?
Special Strings

Fully Qualified Names
Simple Names
Descriptors

The Constant Pool

The CONSTANT_Utf8_info Table
The CONSTANT_Integer_info Table
The CONSTANT_Float_info Table
The CONSTANT_Long_info Table
The CONSTANT_Double_info Table
The CONSTANT_Class_info Table
The CONSTANT_String_info Table

- The CONSTANT_Fieldref_info Table
- The CONSTANT_Methodref_info Table
- The CONSTANT_InterfaceMethodref_info Table
- The CONSTANT_NameAndType_info Table

- Fields
- Methods
- Attributes

- The Code Attribute
- The ConstantValue Attribute
- The Exceptions Attribute
- The InnerClasses Attribute
- The LineNumberTable Attribute
- The LocalVariableTable Attribute
- The SourceFile Attribute
- The Synthetic Attribute

- Getting Loaded: A Simulation
- On the CD-ROM
- The Resources Page

7 The Lifetime of a Class

- Class Loading, Linking, and Initialization

- Loading
- Verification
- Preparation
- Resolution
- Initialization

- The Class Initialization Method
- Active versus Passive Use

- The Lifetime of an Object

- Class Instantiation
- Garbage Collection and Finalization of Objects

- Unloading and Finalization of Classes
- On the CD-ROM
- The Resources Page

8 The Linking Model

- Dynamic Linking and Resolution
- Resolution and Dynamic Extension
- Constant Pool Resolution

Resolution of CONSTANT_Class_info Entries

Array Classes

Non-Array Classes and Interfaces

Resolution of CONSTANT_Fieldref_info Entries

Resolution of CONSTANT_Methodref_info Entries

Resolution of CONSTANT_InterfaceMethodref_info Entries

Resolution of CONSTANT_String_info Entries

Resolution of Other Types of Entries

Compile-Time Resolution of Constants

Direct References

_quick Instructions

Example: The Linking of the Salutation Application

Example: The Dynamic Extension of the Greet Application

Example: Unloading Unreachable Greeters

On the CD-ROM

The Resources Page

9 Garbage Collection

Why Garbage Collection?

Garbage Collection Algorithms

Reference Counting Collectors

Tracing Collectors

Compacting Collectors

Copying Collectors

Generational Collectors

Adaptive Collectors

Finalization

Heap of Fish: A Simulation

Allocate Fish

Assign References

Garbage Collect

Compact Heap

On the CD-ROM

The Resources Page

10 Stack and Local Variable Operations

Pushing Constants Onto the Stack

Generic Stack Operations

Pushing Local Variables Onto the Stack

Popping to Local Variables

The wide Instruction

Fibonacci Forever: A Simulation

On the CD-ROM

The Resources Page

11 Type Conversion

The Conversion Opcodes
Conversion Diversion: A Simulation
On the CD-ROM
The Resources Page

12 Integer Arithmetic

Two's Complement Arithmetic
Inner Int: A Java int Reveals its Inner Nature
Arithmetic Opcodes
Prime Time: A Simulation
On the CD-ROM
The Resources Page

13 Logic

The Logic Opcodes
Logical Results: A Simulation
On the CD-ROM
The Resources Page

14 Floating Point Arithmetic

Floating Point Numbers
Inner Float: A Java float Reveals its Inner Nature
The Floating Point Opcodes
Circle of Squares: A Simulation
On the CD-ROM
The Resources Page

15 Objects and Arrays

A Refresher on Objects and Arrays
Opcodes for Objects
Opcodes for Arrays
Three-Dimensional Array: A Simulation
On the CD-ROM
The Resources Page

16 Control Flow

Conditional Branching
Unconditional Branching

Conditional Branching with Tables
Saying Tomato: A Simulation
On the CD-ROM
The Resources Page

17 Exceptions

Throwing and Catching Exceptions
The Exception Table
Play Ball!: A Simulation
On the CD-ROM
The Resources Page

18 Finally Clauses

Miniature Subroutines
Asymmetrical Invocation and Return
Hop Around: A Simulation
On the CD-ROM
The Resources Page

19 Method Invocation and Return

Method Invocation

Invoking a Java Method
Invoking a Native Method

Other Forms of Method Invocation
The invokespecial instruction

invokespecial and <init()
invokespecial and Private Methods
invokespecial and super

The invokeinterface Instruction
Invocation Instructions and Speed
Examples of Method Invocation
Returning from Methods
On the CD-ROM
The Resources Page

20 Thread Synchronization

Monitors
Object Locking
Synchronization Support in the Instruction Set

Synchronized Statements
Synchronized Methods

Coordination Support in Class Object
On the CD-ROM
The Resources Page

Appendix A. Instructions by Opcode Mnemonic

Appendix B. Opcode Mnemonic by Functional Group

Appendix C. Opcode Mnemonic by Opcode

Appendix D. Slices of Pi: A Simulation of the Java Virtual Machine

Preface

My primary goal in writing this book was to explain the Java Virtual Machine, and the software technologies upon which it is based, to Java programmers. Although the Java Virtual Machine incorporates technologies that had been tried and proven in other programming languages, prior to Java many of these technologies had not yet entered into common use. As a consequence, many programmers will be encountering these technologies for the first time as they begin to program in Java. Garbage collection, multi-threading, exception handling, dynamic extension--even the use of a virtual machine itself--may be new to many programmers. The aim of this book is to help programmers understand how all these things work, and in the process, to help them become more adept with the Java programming language.

Another goal I had in mind as I wrote this book was to experiment a bit with the changing nature of text. Web pages have three interesting characteristics that differentiate them from paper-based text: they're dynamic (they can evolve over time), they're interactive (especially if you embed Java applets in them), and they're interconnected (you can easily navigate from one to another). Besides the traditional text and figures, this book includes several Java applets (in a mini-website on the CD-ROM) that serve as "interactive illustrations" of the concepts presented in the text. In addition, I maintain a website on the internet that serves as a launching point for readers to find more (and more current) information on the topics covered in the book. This book is composed of all of these components: text, figures, interactive illustrations, and constantly evolving links to further reading.

Bill Venners

Sunnyvale, California

August, 1997

Introduction

This book describes the Java Virtual Machine, the abstract computer on which all Java programs run. Through a combination of tutorial explanations, working examples, reference material, and applets that interactively illustrate the concepts presented in the text, this book provides an in-depth, technical survey of Java as a technology.

The Java programming language seems poised to be the next popular language for mainstream commercial software development, the next step after C and C++. One of the fundamental reasons Java is a likely candidate for this role is that Java's architecture helps programmers deal with emerging hardware realities. Java has features that the shifting hardware environment is demanding, features that are made possible by the Java Virtual Machine.

The evolution of programming languages has to a great extent been driven by changes in the hardware being programmed. As hardware has grown faster, cheaper, and more powerful, software has become larger and more complex. The migration from assembly languages to procedural languages, such as C, and to object oriented languages, such as C++, was largely driven by a need to manage ever greater complexity--complexity made possible by increasingly powerful hardware.

Today the progression towards cheaper, faster, and more powerful hardware continues, as does the need for managing increasing software complexity. Building on C and C++, Java helps programmers deal with complexity by rendering impossible certain kinds of bugs that frequently plague C and C++ programmers. Java's inherent memory safety--garbage collection, lack of pointer arithmetic, run-time checks on the use of references--prevents most memory bugs from ever occurring. Java's memory safety makes programmers more productive and helps them manage complexity.

In addition, besides the ongoing increase in the capabilities of hardware, there is another fundamental shift taking place in the hardware environment--the network. As networks interconnect more and more computers and devices, new demands are being made on software. With the rise of the network, platform independence and security have become much more important than they were in the past.

The Java Virtual Machine is responsible for the memory safety, platform neutrality, and security features of the Java programming language. Although virtual machines have been around for a long time, prior to Java they hadn't quite entered the mainstream. But given today's emerging hardware realities, software developers needed a programming language with a virtual machine, and Sun hit the market window with Java.

Thus, the Java Virtual Machine embodies the right software "stuff" for the coming years of computing. This book will help you get to know this virtual machine, and armed with this knowledge, you'll know how best to put the Java Virtual Machine to use in your programs.

Who Should Read the Book

This book is aimed primarily at professional software developers and students who want to understand Java technology. I assume you are familiar, though not necessarily proficient, with the Java language. Reading this book should help you add a depth to your knowledge of Java programming. If you are one of the elite few who are actually writing Java compilers or creating implementations of the Java Virtual Machine, this book can serve as a companion to the Java Virtual Machine specification. Where the specification specifies, this book explains.

How to Use the Book

This book has five basic parts:

1. An introduction to Java's architecture (Chapters 1 through 4)
2. An in-depth, technical tutorial of Java internals (Chapters 5 through 20)
3. A class file and instruction set reference (Chapter 6, Appendices A through C)
4. Interactive illustrations, example source code, and the JDK (On the CD-ROM)
5. The Resources Web Site (<http://www.artima.com/insidejvm>)

An Introduction to Java's Architecture

Chapters 1 through 4 (Part I of this book) give an overview of Java's architecture, including the motivations behind--and the implications of--Java's architectural design. These chapters show how the Java Virtual Machine relates to the other components of Java's architecture: the class file, API, and language. If you want a basic understanding of Java as a technology, consult these chapters. Here are some specific points of interest from this portion of the book:

- For an overview of Java's architecture and a discussion of its inherent tradeoffs, see Chapter 1, "Introduction to Java's Architecture."
- For a discussion of what "platform independence" really means, how Java's architecture supports it, and seven steps to take to create a platform independent Java program, see Chapter 2, "Platform Independence."
- For a description of the security model built into Java's core architecture, including a tutorial explaining how to write a security-minded class loader, see Chapter 3, "Security."
- For a discussion of the new paradigm of network-mobile software, see Chapter 4, "Network Mobility."

A Tutorial of Java Internals

Chapters 5 through 20 (Part II of this book) give an in-depth technical description of the inner workings of the Java Virtual Machine. These chapters will help you understand how Java programs actually work. All the material in Part II is presented in a tutorial manner, with lots of examples. Here are some specific points of interest from this portion of the book:

- For a comprehensive overview of the inner workings of the Java Virtual Machine, see Chapter 5,

"The Java Virtual Machine."

- If you are parsing, generating, or simply peering into Java class files, see Chapter 6, "The Java Class File," for a complete tutorial and reference on the class file format.
- For a discussion of the lifetime of a class inside the Java Virtual Machine, including the circumstances in which classes can be unloaded, see Chapter 7, "The Lifetime of a Class."
- For a thorough explanation of Java's linking model, including a tutorial and examples on writing your own class loaders, see Chapter 8, "The Linking Model."
- For a discussion of garbage collection and finalization, and suggestions on how to use finalizers, see Chapter 9, "Garbage Collection."
- For a tutorial on the Java Virtual Machine's instruction set, read Chapters 10 through 20.
- For an explanation of monitors and how you can use them to write thread-safe Java code, see Chapter 20, "Thread Synchronization."

A Class File and Instruction Set Reference

In addition to being a tutorial on the Java class file, Chapter 6, "The Java Class File," serves as a complete reference of the class file format. Similarly, Chapters 10 through 20 form a tutorial of the Java Virtual Machine's instruction set, and Appendices A through C serve as a complete reference of the instruction set. If you need to look something up, check out these chapters and the appendices.

Interactive Illustrations, Example Source Code, and the JDK

For most of this book's chapters, material associated with the chapter--such as example code or simulation applets--appears on the CD-ROM.

The `applets` directory of the CD-ROM contains a mini-website that includes 14 Java applets that illustrate the concepts presented in the text. These "interactive illustrations" form an integral part of this book. Ten of the applets simulate the Java Virtual Machine executing bytecodes. The other applets illustrate garbage collection, two's-complement and IEEE 754 floating point numbers, and the loading of class files. The applets can be viewed on any platform by any Java-capable browser. The source code for the simulation applets is also included on the CD-ROM.

The copyright notice accompanying the HTML, `.java`, and `.class` files for the mini-website enables you to post the mini-website on any network, including the internet, providing you adhere to a few simple rules. For example, you must post the mini-website in its entirety, you can't make any changes to it, and you can't charge people to look at it. The full text of the copyright notice is given later in this introduction.

All the example source code shown in this book appears on the CD-ROM in both source and compiled (class files) form. If some example code in the text strikes you as interesting (or dubious), you can try it out for yourself.

Most of the example code is illustrative and not likely to be of much practical use besides helping you to understand Java. Nevertheless, you are free to cut and paste from the example code, use it in your own programs, and distribute it in binary (such as Java class file) form. The full text of the copyright notice for the example source code is shown later in this introduction.

Besides the interactive illustrations and example source code, the CD-ROM contains one last item: a full distribution of version 1.1.3 of Sun's JDK. This is contained in the CD-ROM's `jdk` directory.

The Resources Web Site

To help you find more information and keep abreast of changes, I maintain a "Resources Web Site" with links to further reading about the material presented in this book. There is at least one "resources page" for each chapter in the book. The main URL of the Resources Web Site is <http://www.artima.com/insidejvm>. The URL for each chapter's individual resources page is given at the end of each chapter in the "The Resources Page" section.

Chapter by Chapter Summary

Part One: Java's Architecture

Chapter 1. Introduction to Java's Architecture

This chapter gives an introduction to Java as a technology. It gives an overview of Java's architecture, discusses why Java is important, and looks at Java's pros and cons.

Chapter 2. Platform independence

This chapter shows how Java's architecture enables programs to run on any platform, discusses the factors that determine the true portability of Java programs, and looks at the relevant tradeoffs.

Chapter 3. Security

This chapter gives an overview of the security model built into Java's core architecture.

Chapter 4. Network-mobility

This chapter examines the new paradigm of network-mobile software heralded by the arrival of Java, and shows how Java's architecture makes it possible.

Part Two: Java Internals

Chapter 5. The Java Virtual Machine

This chapter gives a detailed overview of the Java Virtual Machine's internal architecture. Accompanying the chapter on the CD-ROM is an applet, named *Eternal Math*, that simulates the Java Virtual Machine executing a short sequence of bytecodes.

Chapter 6. The Java Class File

This chapter describes the contents of the class file, including the structure and format of the constant pool, and serves as both a tutorial and a complete reference of the Java class file format. Accompanying the chapter on the CD-ROM is an applet, named *Getting Loaded*, that simulates the Java Virtual Machine loading a Java class file.

Chapter 7. The Lifetime of a Class

This chapter follows the lifetime of a type (class or interface) from the type's initial entrance into the virtual machine to its ultimate exit. It discusses the processes of loading, linking, and initialization; object instantiation, garbage collection, and finalization; and type finalization and unloading.

Chapter 8. The Linking Model

This chapter takes an in-depth look at Java's linking model. It describes constant pool resolution and shows how to write class loaders to enable a Java application to dynamically extend itself at run-time.

Chapter 9. Garbage Collection

This chapter describes various garbage collection techniques and explains how garbage collection works in Java Virtual Machines. Accompanying this chapter on the CD-ROM is an applet, named Heap of Fish, that simulates a compacting, mark-and-sweep garbage-collected heap.

Chapter 10. Stack and Local Variable Operations

This chapter describes the Java Virtual Machine instructions that focus most exclusively on the operand stack--those that push constants onto the operand stack, perform generic stack operations, and transfer values back and forth between the operand stack and local variables. Accompanying this chapter on the CD-ROM is an applet, named Fibonacci Forever, that simulates the Java Virtual Machine executing a method that generates the Fibonacci sequence.

Chapter 11. Type Conversion

This chapter describes the instructions that convert values from one primitive type to another. Accompanying the chapter on the CD-ROM is an applet, named Conversion Diversion, that simulates the Java Virtual Machine executing a method that performs type conversion.

Chapter 12. Integer Arithmetic

This chapter describes integer arithmetic in the Java Virtual Machine. It explains two's complement arithmetic and describes the instructions that perform integer arithmetic. Accompanying this chapter on the CD-ROM are two applets that interactively illustrate the material presented in the chapter. One applet, named Inner Int, allows you to view and manipulate a two's complement number. The other applet, named Prime Time, simulates the Java Virtual Machine executing a method that generates prime numbers.

Chapter 13. Logic

This chapter describes the instructions that perform bitwise logical operations inside the Java Virtual Machine. These instructions include opcodes to perform shifting and boolean operations on integers. Accompanying this chapter on the CD-ROM is an applet, named Logical Results, that simulates the Java Virtual Machine executing a method that includes uses several of the logic opcodes.

Chapter 14. Floating Point Arithmetic

This chapter describes the floating point numbers and the instructions that perform floating point arithmetic inside the Java Virtual Machine. Accompanying this chapter on the CD-ROM are two applets

that interactively illustrate the material presented in the chapter. One applet, named Inner Float, allows you to view and manipulate the individual components that make up a floating point number. The other applet, named Circle of Squares, simulates the Java Virtual Machine executing a method that uses several of the floating point opcodes.

Chapter 15. Objects and Arrays

This chapter describes the Java Virtual Machine instructions that create and manipulate objects and arrays. Accompanying this chapter on the CD-ROM is an applet, named Three-Dimensional Array, that simulates the Java Virtual Machine executing a method that allocates and initializes a three-dimensional array.

Chapter 16. Control Flow

This chapter describes the instructions that cause the Java Virtual Machine to conditionally or unconditionally branch to a different location within the same method. Accompanying this chapter on the CD-ROM is an applet, named Saying Tomato, that simulates the Java Virtual Machine executing a method that includes bytecodes that perform table jumps (the compiled version of a Java `switch` statement).

Chapter 17. Exceptions

This chapter shows how exceptions are implemented in bytecodes. It describes the instruction for throwing an exception explicitly, explains exception tables, and shows how catch clauses work. Accompanying this chapter on the CD-ROM is an applet, named Play Ball!, that simulates the Java Virtual Machine executing a method that throws and catches exceptions.

Chapter 18. Finally Clauses

This chapter shows how finally clauses are implemented in bytecodes. It describes the relevant instructions and gives examples of their use. The chapter also describes some surprising behaviors exhibited by finally clauses in Java source code and explains this behavior at the bytecode level. Accompanying this chapter on the CD-ROM is an applet, named Hop Around, that simulates the Java Virtual Machine executing a method that includes finally clauses.

Chapter 19. Method Invocation and Return

This chapter describes the four instructions that the Java Virtual Machine uses to invoke methods and the situations in which each instruction is used.

Chapter 20. Thread Synchronization

This chapter describes monitors--the mechanism that Java uses to support synchronization--and shows how they are used by the Java Virtual Machine. It shows how one aspect of monitors, the locking and unlocking of data, is supported in the instruction set.

The Appendices

Appendix A. Instruction Set by Opcode Mnemonic

This appendix lists the opcodes alphabetically by mnemonic. For each opcode, it gives the mnemonic, opcode byte value, instruction format (the operands, if any), a snapshot image of the stack before and after the instruction is executed, and a description of the execution of the instruction. Appendix A serves as the primary instruction set reference of the book.

Appendix B. Opcode Mnemonic by Functional Group

This appendix organizes the instructions by functional group. The organization used in this appendix corresponds to the order the instructions are described in Chapters 10 through 20.

Appendix C. Opcode Mnemonic by Opcode

This appendix organizes the opcodes in numerical order. For each numerical value, this appendix gives the mnemonic.

Chapter One

Introduction to Java's Architecture

At the heart of Java technology lies the Java Virtual Machine--the abstract computer on which all Java programs run. Although the name "Java" is generally used to refer to the Java programming language, there is more to Java than the language. The Java Virtual Machine, Java API, and Java class file work together with the Java language to make the Java phenomenon possible.

The first four chapters of this book (Part I. "Java's Architecture") show how the Java Virtual Machine fits into the big picture. They show how the virtual machine relates to the other components of Java's architecture: the class file, API, and language. They describe the motivation behind--and the implications of--the overall design of Java technology.

This chapter gives an introduction to Java as a technology. It gives an overview of Java's architecture, discusses why Java is important, and looks at Java's pros and cons.

Why Java?

Over the ages people have used tools to help them accomplish tasks, but lately their tools have been getting smarter and interconnected. Microprocessors have appeared inside many commonly used items, and increasingly, they have been connected to networks. As the heart of personal computers and workstations, for example, microprocessors have been routinely connected to networks. They have also appeared inside devices with more specific functionality than the personal computer or the workstation. Televisions, VCRs, audio components, fax machines, scanners, printers, cell phones, personal digital assistants, pagers, and wrist-watches--all have been enhanced with microprocessors; most have been connected to networks.

Given the increasing capabilities and decreasing costs of information processing and data networking technologies, the network is rapidly extending its reach. The emerging infrastructure of smart devices and computers interconnected by networks represents a new environment for software--an environment that presents new challenges and offers new opportunities to software developers.

Java technology is a tool well suited to help you meet the challenges and seize the opportunities presented by the emerging computing environment. Java was designed for networks. Its suitability for networked environments is inherent in its architecture, which enables secure, robust, platform-independent programs to be delivered across networks and run on a great variety of computers and devices.

The Challenges and Opportunities of Networks

One challenge presented to developers by a networked computing environment is the wide range of devices that networks interconnect. A typical network usually has many different kinds of attached devices, with diverse hardware architectures, operating systems, and purposes. Java addresses this challenge by enabling the creation of platform-independent programs. A single Java program can run unchanged on a wide range of computers and devices. Compared with programs compiled for a specific hardware and operating system, platform-independent programs written in Java can be easier and cheaper to develop, administer, and maintain.

Another challenge the network presents to software developers is security. In addition to their potential for good, networks represent an avenue for malicious programmers to steal or destroy information, steal computing resources, or simply be a nuisance. Virus writers, for example, can place their wares on the network for unsuspecting users to download. Java addresses the security challenge by providing an environment in which programs downloaded across a network can be run with customizable degrees of security. A downloaded program can do anything it wants inside the boundaries of the secure environment, but can't read or write data outside those boundaries.

One aspect of security is simple program robustness. Java's architecture guarantees a certain level of program robustness by preventing certain types of pernicious bugs, such as memory corruption, from ever occurring in Java programs. This establishes trust that downloaded code will not inadvertently (or intentionally) crash, but it also has an important benefit unrelated to networks: it makes programmers more productive. Because Java prevents many types of bugs from ever occurring, Java programmers need not spend time trying to find and fix them.

One opportunity created by an omnipresent network is online software distribution. Java takes advantage of this opportunity by enabling the transmission of binary code in small pieces across networks. This capability can make Java programs easier and cheaper to deliver than programs that are not network-mobile. It can also simplify version control. Because the most recent version of a Java program can be delivered on-demand across a network, you needn't worry about what version your end-users are running. They will always get the most recent version each time they use your program.

Platform independence, security, and network-mobility--these three facets of Java's architecture work together to make Java suitable for the emerging networked computing environment. Because Java programs are platform independent, network-delivery of software is more practical. The same version of a program can be delivered to all the computers and devices the network interconnects. Java's built-in security framework also helps make network-delivery of software more practical. By reducing risk, the security framework helps to build trust in a new paradigm of network-mobile code.

The Architecture

Java's architecture arises out of four distinct but interrelated technologies, each of which is defined by a separate specification from Sun Microsystems:

- the Java programming language
- the Java class file format
- the Java Application Programming Interface
- the Java Virtual Machine

When you write and run a Java program, you are tapping the power of these four technologies. You express the program in source files written in the Java programming language, compile the source to Java class files, and run the class files on a Java Virtual Machine. When you write your program, you

access system resources (such as I/O, for example) by calling methods in the classes that implement the Java Application Programming Interface, or Java API. As your program runs, it fulfills your program's Java API calls by invoking methods in class files that implement the Java API. You can see the relationship between these four parts in Figure 1-1.




Figure 1-1. The Java programming environment.

Together, the Java Virtual Machine and Java API form a "platform" for which all Java programs are compiled. In addition to being called the *Java runtime system*, the combination of the Java Virtual Machine and Java API is called the **Java Platform**. Java programs can run on many different kinds of computers because the Java Platform can itself be implemented in software. As you can see in Figure 1-2, a Java program can run anywhere the Java Platform is present.




Figure 1-2. Java programs run on top of the Java Platform.

The Java Virtual Machine

At the heart of Java's network-orientation is the Java Virtual Machine, which supports all three prongs of Java's network-oriented architecture: platform independence, security, and network-mobility.

The Java Virtual Machine is an abstract computer. Its specification defines certain features every Java Virtual Machine must have, but leaves many choices to the designers of each implementation. For example, although all Java Virtual Machines must be able to execute Java bytecodes, they may use any technique to execute them. Also, the specification is flexible enough to allow a Java Virtual Machine to be implemented either completely in software or to varying degrees in hardware. The flexible nature of the Java Virtual Machine's specification enables it to be implemented on a wide variety of computers and devices.

A Java Virtual Machine's main job is to load class files and execute the bytecodes they contain. As you can see in Figure 1-3, the Java Virtual Machine contains a *class loader*, which loads class files from both the program and the Java API. Only those class files from the Java API that are actually needed by a running program are loaded into the virtual machine. The bytecodes are executed in an *execution engine*, which is one part of the virtual machine that can vary in different implementations. On a Java Virtual Machine implemented in software, the simplest kind of execution engine just interprets the bytecodes one at a time. Another kind of execution engine, one that is faster but requires more memory, is a *just-in-time compiler*. In this scheme, the bytecodes of a method are compiled to native machine code the first time the method is invoked. The native machine code for the method is then cached, so it can be re-used the next time that same method is invoked. On a Java Virtual Machine built on top of a chip that executes Java bytecodes natively, the execution engine is actually embedded in the chip.




Figure 1-3. A basic block diagram of the Java Virtual Machine.

Sometimes the Java Virtual Machine is called the *Java interpreter*; however, given the various ways in which bytecodes can be executed, this term can be misleading. While "Java interpreter" is a reasonable name for a Java Virtual Machine that interprets bytecodes, virtual machines also use other techniques (such as just-in-time compiling) to execute bytecodes. Therefore, although all Java interpreters are Java Virtual Machines, not all Java Virtual Machines are Java interpreters.

When running on a Java Virtual Machine that is implemented in software on top of a host operating system, a Java program interacts with the host by invoking *native methods*. In Java, there are two kinds of methods: Java and native. A Java method is written in the Java language, compiled to bytecodes, and

stored in class files. A native method is written in some other language, such as C, C++, or assembly, and compiled to the native machine code of a particular processor. Native methods are stored in a dynamically linked library whose exact form is platform specific. While Java methods are platform independent, native methods are not. When a running Java program calls a native method, the virtual machine loads the dynamic library that contains the native method and invokes it. As you can see in Figure 1-4, native methods are the connection between a Java program and an underlying host operating system.




Figure 1-4. A Java Virtual Machine implemented in software on top of a host operating system.

You can use native methods to give your Java programs direct access to the resources of the underlying operating system. Their use, however, will render your program platform specific. This is because the dynamic libraries containing the native methods are platform specific. In addition, the use of native methods may render your program specific to a particular implementation of the Java Platform. One native method interface--the *Java Native Interface*, or *JNI*--enables native methods to work with any Java Platform implementation on a particular host computer. Vendors of the Java Platform, however, are not required to support JNI. They may provide their own proprietary native method interfaces in addition to (or in place of) JNI.

Java gives you a choice. If you want to access resources of a particular host that are unavailable through the Java API, you can write a platform-specific Java program that calls native methods. If you want to keep your program platform independent, however, you must call only Java methods and access the system resources of the underlying operating system through the Java API.

The Class Loader Architecture

One aspect of the Java Virtual Machine that plays an important role in both security and network-mobility is the class loader architecture. In the block diagrams of Figures 1-3 and 1-4, a single mysterious cube identifies itself as "the class loader," but in reality there may be more than one class

loader inside a Java Virtual Machine. Thus the class loader cube of the block diagram actually represents a subsystem that may involve many class loaders. The Java Virtual Machine has a flexible class loader architecture that allows a Java application to load classes in custom ways.

A Java application can use two types of class loaders: a "primordial" class loader and class loader objects. The primordial class loader (there is only one of them) is a part of the Java Virtual Machine implementation. For example, if a Java Virtual Machine is implemented as a C program on top of an existing operating system, then the primordial class loader will be part of that C program. The primordial class loader loads trusted classes, including the classes of the Java API, usually from the local disk.

At run-time, a Java application can install class loader objects that load classes in custom ways, such as by downloading class files across a network. The Java Virtual Machine considers any class it loads through the primordial class loader to be trusted, regardless of whether or not the class is part of the Java API. Classes it loads through class loader objects, however, it views with suspicion--by default, it considers them to be untrusted. While the primordial class loader is an intrinsic part of the virtual machine implementation, class loader objects are not. Instead, class loader objects are written in Java, compiled to class files, loaded into the virtual machine, and instantiated just like any other object. They are really just another part of the executable code of a running Java application. You can see a graphical depiction of this architecture in Figure 1-5.




Figure 1-5. Java's class loader architecture.

Because of class loader objects, you don't have to know at compile-time all the classes that may ultimately take part in a running Java application. They enable you to dynamically extend a Java application at run-time. As it runs, your application can determine what extra classes it needs and load them through one or more class loader objects. Because you write the class loader in Java, you can load classes in any manner. You can download them across a network, get them out of some kind of database, or even calculate them on the fly.

For each class it loads, the Java Virtual Machine keeps track of which class loader--whether primordial or object--loaded the class. When a loaded class first refers to another class, the virtual machine requests the referenced class from the same class loader that originally loaded the referencing class. For example, if the virtual machine loads class `Volcano` through a particular class loader, it will attempt to load any classes `Volcano` refers to through the same class loader. If `Volcano` refers to a class named `Lava`, perhaps by invoking a method in class `Lava`, the virtual machine will request `Lava` from the class loader object that loaded `Volcano`. The `Lava` class returned by the class loader is dynamically linked with class `Volcano`.

Because the Java Virtual Machine takes this approach to loading classes, classes can by default only see other classes that were loaded by the same class loader. This is how Java's architecture enables you to create multiple *name-spaces* inside a single Java application. Each class loader in your running Java program maintains its own name-space, which is populated by the names of all the classes it has loaded.

A Java application can instantiate multiple class loader objects either from the same class or from multiple classes. It can, therefore, create as many (and as many different kinds of) class loader objects as it needs. Classes loaded by different class loaders are in different name-spaces and cannot gain access to each other unless the application explicitly allows it. When you write a Java application, you can segregate classes loaded from different sources into different name-spaces. In this way, you can use Java's class loader architecture to control any interaction between code loaded from different sources. You can prevent hostile code from gaining access to and subverting friendly code.

One example of dynamic extension is the web browser, which uses class loader objects to download the class files for an applet across a network. A web browser fires off a Java application that installs a class loader object--usually called an *applet class loader*--that knows how to request class files from an HTTP server. Applets are an example of dynamic extension, because the Java application doesn't know when it starts which class files the browser will ask it to download across the network. The class files to download are determined at run-time, as the browser encounters pages that contain Java applets.

The Java application started by the web browser usually creates a different applet class loader object for each location on the network from which it retrieves class files. As a result, class files from different sources are loaded by different class loader objects. This places them into different name-spaces inside the host Java application. Because the class files for applets from different sources are placed in separate name-spaces, the code of a malicious applet is restricted from interfering directly with class files downloaded from any other source.

By allowing you to instantiate class loader objects that know how to download class files across a network, Java's class loader architecture supports network-mobility. It supports security by allowing you to load class files from different sources through different class loader objects. This puts the class files from different sources into different name-spaces, which allows you to restrict or prevent access between code loaded from different sources.

The Java Class File

The Java class file helps make Java suitable for networks mainly in the areas of platform-independence and network-mobility. Its role in platform independence is serving as a binary form for Java programs that is expected by the Java Virtual Machine but independent of underlying host platforms. This approach breaks with the tradition followed by languages such as C or C++. Programs written in these languages are most often compiled and linked into a single binary executable file specific to a particular hardware platform and operating system. In general, a binary executable file for one platform won't

work on another. The Java class file, by contrast, is a binary file that can be run on any hardware platform and operating system that hosts the Java Virtual Machine.

When you compile and link a C++ program, the executable binary file you get is specific to a particular target hardware platform and operating system because it contains machine language specific to the target processor. A Java compiler, by contrast, translates the instructions of the Java source files into bytecodes, the "machine language" of the Java Virtual Machine.

In addition to processor-specific machine language, another platform-dependent attribute of a traditional binary executable file is the byte order of integers. In executable binary files for the Intel X86 family of processors, for example, the byte order is *little-endian*, or lower order byte first. In executable files for the PowerPC chip, however, the byte order is *big-endian*, or higher order byte first. In a Java class file, byte order is big-endian irrespective of what platform generated the file and independent of whatever platforms may eventually use it.

In addition to its support for platform independence, the Java class file plays a critical role in Java's architectural support for network-mobility. First, class files were designed to be compact, so they can more quickly move across a network. Also, because Java programs are dynamically linked and dynamically extensible, class files can be downloaded as needed. This feature helps a Java application manage the time it takes to download class files across a network, so the end-user's wait time can be kept to a minimum.

The Java API

The Java API helps make Java suitable for networks through its support for platform independence and security. The Java API is set of runtime libraries that give you a standard way to access the system resources of a host computer. When you write a Java program, you assume the class files of the Java API will be available at any Java Virtual Machine that may ever have the privilege of running your program. This is a safe assumption because the Java Virtual Machine and the class files for the Java API are the required components of any implementation of the Java Platform. When you run a Java program, the virtual machine loads the Java API class files that are referred to by your program's class files. The combination of all loaded class files (from your program and from the Java API) and any loaded dynamic libraries (containing native methods) constitute the full program executed by the Java Virtual Machine.

The class files of the Java API are inherently specific to the host platform. The API's functionality must be implemented expressly for a particular platform before that platform can host Java programs. In a system where bytecodes are executed directly in silicon (on a "Java chip") the API will likely be implemented as part of a Java-based operating system. On a system where the virtual machine is implemented in software on top of a non-Java operating system, the Java API will access the host resources through native methods. As you can see in Figure 1-6, the class files of the Java API invoke native methods so your Java program doesn't have to. In this manner, the Java API's class files provide a Java program with a standard, platform-independent interface to the underlying host. To the Java program, the Java API looks the same and behaves predictably no matter what platform happens to be underneath. Precisely because the Java Virtual Machine and Java API are implemented specifically for each particular host platform, Java programs themselves can be platform independent.




Figure 1-6. A platform-independent Java program.

The internal design of the Java API is also geared towards platform independence. For example, the graphical user interface library of the Java API, called the Abstract Windows Toolkit (or AWT), is designed to facilitate the creation of user interfaces that work on all platforms. Creating platform-independent user interfaces is inherently difficult, given that the native look and feel of user interfaces vary greatly from one platform to another. The AWT library's architecture does not coerce implementations of the Java API to give Java programs a user interface that looks exactly the same everywhere. Instead, it encourages implementations to adopt the look and feel of the underlying platform. Also, because the size of fonts, buttons, and other user interface components will vary from platform to platform, the AWT includes *layout managers* to position the elements of a window or dialog box at run-time. Rather than forcing you to indicate exact X and Y coordinates for the various elements that constitute, say, a dialog box, the layout manager positions them when your dialog box is displayed. With the aim of making the dialog look its best on each platform, the layout manager will very likely position the dialog box elements slightly differently on different platforms. In these ways and others, the internal architecture of the Java API is aimed at facilitating the platform independence of the Java programs that use it.

In addition to facilitating platform independence, the Java API contributes to Java's security model. The methods of the Java API, before they perform any action that could potentially be harmful (such as writing to the local disk), check for permission from the *security manager*. The security manager is a special object that a Java application can instantiate that defines a custom security policy for the application. A security manager could, for example, forbid access to the local disk. If the application requested a local disk write by invoking a method from the Java API, that method would first check with the security manager. Upon learning from the security manager that disk access is forbidden, the Java API would refuse to perform the write. By enforcing the security policy established by the security manager, the Java API helps to establish a safe environment in which you can run potentially unsafe code.

The Java Programming Language

Although Java was designed for the network, its utility is not restricted to networks.

Platform independence, network-mobility, and security are of prime importance in a networked computing environment, but you may not always find yourself facing network-oriented problems. As a result, you may not always want to write programs that are platform independent. You may not always want to deliver your programs across networks or limit their capabilities with security restrictions. There may be times when you use Java technology primarily because you want to get the advantages of the Java programming language.

As a whole, Java technology leans heavily in the direction of networks, but the Java programming language is quite general-purpose. The Java language allows you to write programs that take advantage of many software technologies:

- object-orientation
- multi-threading
- structured error-handling
- garbage collection
- dynamic linking
- dynamic extension

Instead of serving as a test bed for new and experimental software technologies, the Java language combines in a new way concepts and techniques that have already been tried and proven in other languages. These features make the Java programming language a powerful general-purpose tool that you can apply to a variety of situations, independent of whether or not they involve a network.

At the beginning of a new project, you may be faced with the question, "Should I use C++ (or some other language) for my next project, or should I use Java?" As an implementation language, Java has some advantages and some disadvantages over other languages. One of the most compelling reasons for using Java as a language is that it can enhance developer productivity. The main disadvantage is slower execution speed.

Java is, first and foremost, an object-oriented language. One promise of object-orientation is that it promotes the re-use of code, resulting in better productivity for developers. This may make Java more attractive than a procedural language such as C, but doesn't add much value to Java over C++, the object-oriented language that Java most closely resembles. Yet compared to C++, Java has some significant differences that can improve a developer's productivity. This productivity boost comes mostly from Java's restrictions on direct memory manipulation.

In Java, there is no way to directly access memory by arbitrarily casting pointers to a different type or by using pointer arithmetic, as there is in C++. Java requires that you strictly obey rules of type when working with objects. If you have a *reference* (similar to a pointer in C++) to an object of type **Mountain**, you can only manipulate it as a **Mountain**. You can't cast the reference to type **Lava** and manipulate the memory as if it were a **Lava**. Neither can you simply add an arbitrary offset to the reference, as pointer arithmetic allows you to do in C++. You can, in Java, cast a reference to a different type, but only if the object really is of the new type. For example, if the **Mountain** reference actually referred to an instance of class **Volcano** (a specialized type of **Mountain**), you could cast the **Mountain** reference to a **Volcano** reference. Because Java enforces strict type rules at run-time, you are not able to directly manipulate memory in ways that can accidentally corrupt it. As a result, you can't ever create certain kinds of bugs in Java programs that regularly harass C++ programmers and reduce their productivity.

Another way Java prevents you from inadvertently corrupting memory is through automatic garbage collection. Java has a `new` operator, just like C++, that you use to allocate memory on the heap for a new object. But unlike C++, Java has no corresponding `delete` operator, which C++ programmers use to free the memory for an object that is no longer needed by the program. In Java, you merely stop referencing an object, and at some later time, the garbage collector will reclaim the memory occupied by the object.

The garbage collector prevents Java programmers from needing to explicitly indicate which objects should be freed. As a C++ project grows in size and complexity, it often becomes increasingly difficult for programmers to determine when an object should be freed, or even whether an object has already been freed. This results in memory leaks, in which unused objects are never freed, and memory corruption, in which the same object is accidentally freed multiple times. Both kinds of memory troubles cause C++ programs to crash, but in ways that make it difficult to track down the exact source of the problem. You can be more productive in Java in part because you don't have to chase down memory corruption bugs. But perhaps more significantly, you can be more productive because when you no longer have to worry about explicitly freeing memory, program design becomes easier.

A third way Java protects the integrity of memory at run-time is array bounds checking. In C++, arrays are really shorthand for pointer arithmetic, which brings with it the potential for memory corruption. C++ allows you to declare an array of ten items, then write to the eleventh item, even though that tramples on memory. In Java, arrays are full-fledged objects, and array bounds are checked each time an array is used. If you create an array of ten items in Java and try to write to the eleventh, Java will throw an exception. Java won't let you corrupt memory by writing beyond the end of an array.

One final example of how Java ensures program robustness is by checking object references, each time they are used, to make sure they are not `null`. In C++, using a null pointer usually results in a program crash. In Java, using a null reference results in an exception being thrown.

The productivity boost you can get just by using the Java language results in quicker development cycles and lower development costs. You can realize further cost savings if you take advantage of the potential platform independence of Java programs. Even if you are not concerned about a network, you may still want to deliver a program on multiple platforms. Java can make support for multiple platforms easier, and therefore, cheaper.

As you might expect, however, all this good news about productivity, quick development cycles, and lower development costs does not come without a catch. The designers of Java made tradeoffs. They designed an architecture that favors network-oriented features--such as platform-independence, program robustness, security, and network-mobility--over other concerns. The primary tradeoff, and thus the primary hit you will take if you use Java, is execution speed.

Java's extra run-time housekeeping--array bounds checking, type-safe reference casting, checking for null references, and garbage-collection--will cause your Java program to be slower than an equivalent C++ program. Yet often the tradeoff in speed is made up for in productivity increases enjoyed by the developer and robustness enjoyed by the end-user. And often, the Java program simply runs quickly *enough* to satisfy end-users.

Another speed hit, and one that can be far more substantial, arises from the interpreted nature of Java programs. Whereas C++ programs are usually compiled to native machine code, which is stored in a monolithic executable file, Java programs are usually compiled to Java bytecodes, which are stored in class files. When the Java program runs, a virtual machine loads the class files and executes the

bytecodes they contain. When running on a virtual machine that interprets bytecodes, a Java program may be 10 to 30 times slower than an equivalent C++ program compiled to native machine code.

This performance degradation is primarily a tradeoff in exchange for platform independence. Instead of compiling a Java program to platform-specific native machine code, you compile it to platform independent Java bytecodes. Native machine code can run fast, but only on the native platform. Java bytecodes (when interpreted) run slowly, but can be executed on any platform that hosts the Java Virtual Machine.

Fortunately, other techniques can improve the performance of bytecode execution. For example, just-in-time compiling can speed up program execution 7 to 10 times over interpreting. Rather than merely interpreting a method's bytecodes, a virtual machine can compile the bytecodes to native machine code the first time the method is invoked. (The method is compiled "just-in-time" for its first use by the virtual machine.) The native machine code version of the method is then cached by the virtual machine, and re-used the next time the method is invoked by the program. Execution techniques such as just-in-time compiling allows Java programs to be delivered as platform-independent class files, and still, in many cases, run quickly enough to satisfy end-users.

Raw execution speed is not always the most important factor determining an end-user's perception of a program's performance. In some situations, programs spend much of their time waiting for data to come across a network or waiting for the user to hit another key on the keyboard. In such cases, even executing the program via an interpreter may be adequate. For more demanding applications, a just-in-time compiler may be sufficient to satisfy the end-user's need for speed.

The simulation applets incorporated into Part II of this book are an example of a type of program for which execution speed is not that critical. Most of time in these programs is spent waiting for the user to click a button. For many programs, however, execution speed is extremely important. For such programs, if you want to use the Java language, you may have to execute part or all of your program natively. One way to do that is to run the class files on a virtual machine built on top of a chip that executes bytecodes directly in silicon. If you (or your end-users) don't have such a chip handy, another possibility is to identify time-critical portions of your program and implement them as native methods. Using native methods yields a program that is delivered as a combination of platform independent class files and platform-specific dynamic libraries. The bytecodes from the class files are executed by interpreting or just-in-time compiling, but the time-critical code stored in the dynamic libraries is executed natively.

One final alternative is to compile the Java program to a platform-specific, monolithic native executable, as is usually done with C++ programs. Such a strategy bypasses class files entirely, and generates a platform-specific binary. A monolithic native executable can be faster than the same program just-in-time compiled for several reasons. First, just-in-time compilers do not usually do as much optimization as native compilers because of the time trade-off. When compiling a Java program to a monolithic native executable, you have plenty of time to spend performing optimization. When just-in-time compiling, however, time is more scarce. The whole point of just-in-time compiling is to speed up program execution on the fly, but at some stage the speedup gained by certain optimizations will not be worth the time spent doing the optimization. Another reason using just-in-time compiler is slower than a native executable is the just-in-time compiled program will likely occupy a larger memory footprint. The larger footprint could require more paging (or swapping) on a virtual memory system.

So when you compile your Java program to a monolithic native executable, you give up binary platform independence in return for speed. In cases where platform independence is not important to you, or speed is more important, compiling to a native executable can give you both fast execution and the

productivity benefits of the Java language.

One way to get the best of both the platform independence and speed execution worlds is by install-time compiling. In this scheme, you deliver platform-independent class files, which are compiled at install time to a platform-specific, monolithic native executable. The binary form that you deliver (Java class files) is platform independent, but the binary form that the end-user executes (a monolithic native executable) is platform specific. Because the translation from class files to native executable is done during installation on the end-user's system, optimizations can be made for the user's particular system setup.

Java, therefore, gives you many options of program delivery and execution. Moreover, if you write your program in the Java language, you need not choose just one option. You can use several or all methods of program delivery and execution made possible by Java. You can deliver the same program to some users over a network, where they are executed via interpreting or just-in-time compiling. To other users you can deliver class files that are install-time compiled. To still other users you can deliver a monolithic native executable.

Although program speed is a concern when you use Java, there are ways you can address it. By appropriate use of the various techniques for developing, delivering, and executing Java programs, you can often satisfy end-user's expectations for speed. As long as you are able to address the speed issue successfully, you can use the Java language and realize its benefits: productivity for the developer and program robustness for the end-user.

Architectural Tradeoffs

Although Java's network-oriented features are desirable, especially in a networked environment, they did not come for free. They required tradeoffs against other desirable features. Whenever a potential tradeoff between desirable characteristics arose, the designers of Java made the architectural choice that made better sense in a networked world. Hence, Java is not the right tool for every job. It is suitable for solving problems that involve networks and has utility in many problem that don't involve networks, but its architectural tradeoffs will disqualify it for certain types of jobs.

As mentioned before, one of the prime costs of Java's network-oriented features is the potential reduction in program execution speed compared to other technologies such as C++. Java programs can run slower than an equivalent C++ program for many reasons:

- Interpreting bytecodes is 10 to 30 times slower than native execution.
- Just-in-time compiling bytecodes can be 7 to 10 times faster than interpreting, but still not quite as fast as native execution.
- Java programs are dynamically linked.
- The Java Virtual Machine may have to wait for class files to download across a network.
- Array bounds are checked on each array access.
- All objects are created on the heap (no objects are created on the stack).
- All uses of object references are checked at run-time for `null`.
- All reference casts are checked at run-time for type safety.
- The garbage collector is likely less efficient (though often more effective) at managing the heap than you could be if you managed it directly as in C++.
- Primitive types in Java are the same on every platform, rather than adjusting to the most efficient size on each platform as in C++.
- Strings in Java are always UNICODE. When you really need to manipulate just an ASCII string, a

Java program will be slightly less efficient than an equivalent C++ program.

Although many of Java's speed hits are manageable through techniques such as just-in-time compiling, some--such as those that result from run-time checking--can't be eliminated even by compilation to native executable. Still, you get something, such as platform independence or program robustness, for all of the speed hits associated with Java programs. In many cases the end-user will not be able to perceive any speed deficit. In many other cases, the benefits of platform independence and improved program robustness will be worth the speed degradation. Sometimes, however, Java may be disqualified as a tool to help you solve a problem because that problem requires the utmost in speed and Java can't deliver it.

Another tradeoff is loss of control of memory management. Garbage collection can help make programs more robust and easier to design, but adds a level of uncertainty to the runtime performance of the program. You can't always be sure when a garbage collector will decide it is time to collect garbage, nor how long it will take. This loss of control of memory management makes Java a questionable candidate for software problems that require a real-time response to events. While it is possible to create a garbage collector that attempts to meet real-time requirements, for many real-time problems, robustness and platform independence are simply not important enough to justify using Java.

Still another tradeoff arises from Java's goal of platform independence. One difficulty inherent in any API that attempts to provide cross-platform functionality is the lowest-common-denominator problem. Although there is much overlap between operating systems, each operating system usually has a handful of traits all its own. An API that aims to give programs access to the system services of any operating system has to decide which capabilities to support. If a feature exists on only one operating system, the designers of the API may decide not to include support for that feature. If a feature exists on most operating systems, but not all, the designers may decide to support it anyway. This will require an implementation of something similar in the API on operating systems that lack the feature. Both of these lowest-common-denominator kinds of choices may to some degree offend developers and end-users on the affected operating systems.

What's worse, not only does the lowest-common-denominator problem afflict the designers of a platform independent API, it also affects the designer of a program that uses that API. Take user interface as an example. The AWT attempts to give your program a user interface that adopts the native look on each platform. You might find it difficult, however, to design a user interface in which the components interact in a way that *feels* native on every platform, even though the individual components may have the native look. So on top of the lowest-common-denominator choices that were made when the AWT was designed, you may find yourself faced with your own lowest-common-denominator choices when you use the AWT.

One last tradeoff stems from the dynamically linked nature of Java programs combined with the close relationship between Java class files and the Java programming language. Because Java programs are dynamically linked, the references from one class file to another are symbolic. In a statically-linked executable, references between classes are direct pointers or offsets. Inside a Java class file, by contrast, a reference to another class spells out the name of the other class in a text string. If the reference is to a field, the field's name and *descriptor* (the field's type) are also specified. If the reference is to a method, the method's name and descriptor (the method's return type, number and types of its arguments) are specified. Moreover, not only do Java class files contain symbolic references to the fields and methods of other classes, they also contain symbolic references to their own fields and methods. Java class files also may contain optional debugging information that includes the names and types of local variables. A class file's symbolic information, and the close relationship between the bytecode instruction set and the Java language, make it quite easy to decompile Java class files back into Java source. This in turn makes

it quite easy for your competitors to borrow heavily from your hard work.

While it has always been possible for competitors to decompile a statically-linked binary executable and glean insights into your program, by comparison decompilation is far easier with an intermediate (not yet linked) binary form such as Java class files. Decompilation of statically-linked binary executables is more difficult not only because the symbolic information (the original class, field, method, and local variable names) is missing, but also because statically-linked binaries are usually heavily optimized. The more optimized a statically-linked binary is, the less it corresponds to the original source code. Still, if you have an algorithm buried in your binary executable, and it is worth the trouble to your competitors, they can peer into your binary executable and retrieve that algorithm.

Fortunately, there is a way to combat the easy borrowing of your intellectual property: you can obfuscate your class files. Obfuscation alters your class files by changing the names of classes, fields, methods, and local variables, but without altering the operation of the program. Your program can still be decompiled, but will no longer have the (hopefully) meaningful names you originally gave to all of your classes, fields, methods, and local variables. For large programs, obfuscation can make the code that comes out of the decompiler so cryptic as to require nearly the same effort to steal your work as would be required by a statically-linked executable.

Future Trends

As Java matures, some of the tradeoffs described in this chapter may change. One area in which you can expect improvement over time is in the execution speed of Java programs. Sun, for example, is currently working on a technology they call "hot-spot compiling," which is a hybrid of interpreting and just-in-time compiling. They claim this technique will yield Java programs that run as fast as natively compiled C++. Although this seems like a rash claim, when you look at the approach, it makes sense that speeds very close to natively compiled C++ could be achievable.

As a programmer, you may sometimes be faced with the task of speeding up a program by looking through your code for ways to optimize. Often, programmers waste time optimizing code that is rarely executed when the program runs. The proper approach is usually to profile the program to discover exactly where the program spends most of its time. Programs often spend 80 or 90 percent of their time in 10 to 20 percent of the code. To be most effective, you should focus your optimization efforts on just the 10 to 20 percent of the code that really matters to execution speed.

In a sense, a Java Virtual Machine that does just-in-time compiling is like a programmer who spends time to optimize all the code in a program. 80 to 90 percent of the time such a virtual machine spends just-in-time compiling is probably spent on code that only runs 10 to 20 percent of the time. Because all the code is just-in-time compiled, the memory footprint of the program grows much larger than that of an interpreted program, where all the code remains in bytecode form. Also, because so much time is spent just-in-time compiling everything, the virtual machine doesn't have enough time left over to do a thorough job of optimization.

A Java Virtual Machine that does hot-spot compiling, by contrast, is like a programmer who profiles the code and only optimizes the code's time-critical portions. In this approach, the virtual machine begins by interpreting the program. As it interprets bytecodes, it analyzes the execution of the program to determine the program's "hot spot" --that part of the code where the program is spending most of its time. When it identifies the hot spot, it just-in-time compiles only that part of the code that makes up the hot spot. As the program continues to run, the virtual machine continues to analyze it. If the hot spot moves, the virtual machine can just-in-time compile and optimize new areas as they move into the hot

spot. Also, it can revert back to using bytecodes for areas that move out of the hot spot back, to keep the memory footprint at a minimum.

Because only a small part of the program is just-in-time compiled, the memory footprint of the program remains small and the virtual machine has more time to do optimizations. On systems with virtual memory, a smaller memory footprint means less paging. On systems that lack virtual memory--such as many embedded devices--a smaller memory footprint may mean the difference between a program fitting or not fitting in memory at all. More time for optimizations yields hot-spot code that could potentially be optimized as much as natively compiled C++.

In the hot-spot compiling approach, the Java Virtual Machine loads platform-independent class files, just-in-time compiles and heavily optimizes only the most time-critical code, and interprets the rest of the code. Such a program could spend 80 to 90 percent of its time executing native code that is optimized as heavily as natively compiled C++. At the same time, it could keep a memory footprint that is not much larger than a Java program that is 100 percent interpreted. It makes sense that a Java program running on such a virtual machine could achieve speeds very close to the speed of natively compiled C++.

If emerging technologies, such as hot-spot compiling, fulfill their promise, the speed tradeoff of Java programs could eventually become much less significant. It remains to be seen, however, what execution speeds such technologies will actually be able to achieve. For links to the latest information about emerging virtual machine technologies, visit the resources page for this chapter.

Another area in which much work is being done is user interface. One of the tradeoffs listed in this chapter for writing platform-independent programs is the lowest-common-denominator problem. A major area in which this problem reveals itself is user interfaces. In an effort to provide a user-interface library that could map to native components on most platforms, Sun filled the AWT library in Java 1.0 and 1.1 with a lowest-common-denominator subset of components. The 1.0 AWT library included a button class, for example, because every platform had a native button. The library did not include more advanced components such as tab controls or spin controls, however, in part because of schedule constraints, but also because these kinds of controls weren't native to enough platforms.

The Java programmer was faced with an AWT library that directly supported the creation of rather simple user interfaces. With work, however, the programmer could build a fancier user interface on top of the AWT primitives. Many third party vendors built more advanced user interface libraries on top of AWT to help ease the programmer's burden. Microsoft's AFC (Application Foundation Classes) and Netscape's IFC (Internet Foundation Classes) are two good examples. These libraries add support for more advanced user-interface components and functionality to those directly supported by the AWT. They are, however, built on top of AWT, so programs that use them are still platform independent.

Sun has announced JFC (Java Foundation Classes), which is their approach to solving the lowest-common-denominator problem with the 1.0 and 1.1 AWT libraries. Rather than attempting to map more components to native counterparts, Sun's strategy is to provide what they call "lightweight components." A lightweight component doesn't directly map to a native component. Instead, it is built out of the AWT primitives. So for example, instead of providing a tab control that maps to a native tab control on each platform that supports one, JFC would provide a "Java Platform tab control". When such a control is used on Windows95, which supports tab controls natively, the control would not necessarily have the native Windows look and feel. It would have the Java Platform look and feel.

As Java user interface libraries evolve, they will reduce the pain of writing platform-independent user

interfaces. Sun's lightweight component approach could enable the Java Platform to become more of a driving force in the evolution of user-interface. Rather than just trying to catch up with the user interfaces available on native platforms, Sun can develop the user interface of the Java Platform. With lightweight components, they need not be restrained by lowest-common-denominator choices between native user interfaces.

It's not clear to what extent users will accept a Java Platform look and feel over a native one, but user interface does seem to be evolving towards more heterogeneity. Back in the eighties, the Apple Macintosh established a principle that stated all Macintosh applications should adhere to certain user-interface guidelines. The theory was that software would be easier to use if all a user's applications were homogeneous: if they all used familiar metaphors and exhibited the same look and feel. Today, however, when a Macintosh user browses the World Wide Web, they don't expect every web page to look like a Macintosh page. When they go to the IBM site, they expect it to look like IBM. When they go to the Disney site, they expect it to look like Disney. This is similar to the real world in that when you go to New York, you expect it to look and feel like New York. When you go to Paris, you expect it to look and feel like Paris. You don't expect all cities to have the same look and feel. As users are exposed to the web, they are becoming accustomed to working with more heterogeneous user interfaces than they might have encountered on an isolated personal computer.

As Java user-interface libraries evolve, the lowest-common-denominator problem inherent in platform-independent user interfaces may gradually become less painful. It remains to be seen, however, the extent to which users will accept interfaces that do not look and feel 100% native. For links to the latest information about the evolution of user interface technologies for Java, visit the resources page for this chapter.

The Resources Page

For links to more information about the material presented in this chapter, visit the resources page at <http://www.artima.com/insidejvm/intro.html>

Chapter Two

Platform Independence

The last chapter showed how Java's architecture makes it a useful tool for developing software solutions in a networked environment. The next three chapters take a closer look at how Java's architecture accomplishes its suitability for networks. This chapter examines platform independence in detail. It shows how Java's architecture enables programs to run on any platform, discusses the factors that determine the true portability of Java programs, and looks at the relevant tradeoffs.

Why Platform Independence?

One of the key reasons Java technology is useful in a networked environment is that Java makes it possible to create binary executables that will run unchanged on multiple platforms. This is important in a networked environment because networks usually interconnect many different kinds of computers and devices. An internal network at a medium-sized company might connect Macintoshes in the art department, UNIX workstations in engineering, and PCs running Windows everywhere else. Also, various kinds of embedded devices, such as printers, scanners, and fax machines, would typically be connected to the same network. Although this arrangement enables various kinds of computers and devices within the company to share data, it requires a great deal of administration. Such a network presents a system administrator with the task of keeping different platform-specific editions of programs up to date on many different kinds of computers. Programs that can run without change on any networked computer, regardless of the computer's type, make the system administrator's job simpler, especially if those programs can actually be delivered across the network.

On the developer's side, Java can reduce the cost and time required to develop and deploy applications on multiple platforms. Even though historically, many (or most) applications have been supported on only one platform, often the reason was that the cost involved in supporting multiple platforms wasn't worth the added return. Java can help make multi-platform support affordable for more types of programs.

For software developers, Java's platform independence can be both an advantage and a disadvantage. If you are developing and selling a software product, Java's support for platform independence can help you to compete in more markets. Instead of developing a product that runs only on Windows, for example, you can write one that runs on Windows, Macintosh, UNIX, and OS/2. With Java, you can have more potential customers. The trouble is, so can everyone else. Imagine, for example, that you have focused your efforts on writing great software for OS/2. Java makes it easier for others to write software that competes in your chosen market niche. With Java, therefore, you may not only end up with more potential customers, but also with more potential competitors.

Java's Architectural Support for Platform Independence

Java's architecture facilitates the creation of platform-independent software, but also allows you to create software that is platform-specific. When you write a Java program, platform independence is an *option*.

Support for platform independence, like support for security and network-mobility, is spread throughout Java's architecture. All the components of the architecture--the language, the class file, the API, and the virtual machine--play a role in enabling platform independence.

The Java Platform

Java's architecture supports the platform independence of Java programs in several ways, but primarily through the Java Platform itself. The Java Platform acts as a buffer between a running Java program and the underlying hardware and operating system. Java programs are compiled to run on a Java Virtual Machine, with the assumption that the class files of the Java API will be available at run-time. The virtual machine runs the program; the API gives the program access to the underlying computer's resources. No matter where a Java program goes, it need only interact with the Java Platform. It needn't worry about the underlying hardware and operating system. As a result, it can run on any computer that hosts a Java Platform.

The Java Language

The Java programming language reflects Java's platform independence in one principal way: the ranges and behavior of its primitive types are defined by the language. In languages such as C or C++, the range of the primitive type `int` is determined by its size, and its size is determined by the target platform. The size of an `int` in C or C++ is generally chosen by the compiler to match the word size of the platform for which the program is compiled. This means that a C++ program might have different behavior when compiled for different platforms merely because the ranges of the primitive types are not consistent across the platforms. For example, no matter what underlying platform might be hosting the program, an `int` in Java behaves as a signed 32-bit two's complement number. A `float` adheres to the 32-bit IEEE 754 floating point standard. This consistency is also reflected in the internals of the Java Virtual Machine, which has primitive data types that match those of the language, and in the class file, where the same primitive data types appear. By guaranteeing that primitive types behave the same on all platforms, the Java language itself promotes the platform independence of Java programs.

The Java Class File

As mentioned in the previous chapter, the class file defines a binary format that is specific to the Java Virtual Machine. Java class files can be generated on any platform. They can be loaded and run by a Java Virtual Machine that sits on top of any platform. Their format, including the big-endian order of multi-byte values, is strictly defined and independent of any platform that hosts a Java Virtual Machine.

Scaleability

One aspect of Java's support for platform independence is its scaleability. The Java Platform can be implemented on a wide range of hosts with varying levels of resources, from embedded devices to mainframe computers.

Even though Java first came to prominence by riding on top of a wave that was crashing through the desktop computer industry, the World Wide Web, Java was initially envisioned as a technology for embedded devices, not desktop computers. Part of the early reasoning behind Java was that although Microsoft and Intel had a dominant clutch on the desktop market, no such dominance existed in the embedded systems market. Microprocessors had been appearing in device after device for years--audio-video equipment, cell phones, printers, fax machines, copiers--and the coming trend was that, increasingly, embedded microprocessors would be connected to networks. An original design goal of Java, therefore, was to provide a way for software to be delivered across networks to any kind of embedded device--independent of its microprocessor and operating system.

To accomplish this goal, the Java runtime system (the Java Platform) had to be compact enough to be implemented in software using the resources available to a typical embedded system. Embedded microprocessors often have special constraints, such as small memory footprint, no hard disk, a non-graphical display, or no display.

Given the special requirements of embedded systems, several incarnations of the Java Platform exist just for embedded systems:

- the Java Embedded Platform
- the Java Personal Platform
- the Java Card Platform

These Java Platforms are composed of a Java Virtual Machine and a smaller shell of runtime libraries than are available in the Java Core Platform. The difference between the Core and the Embedded Platform, therefore, is that the Embedded Platform guarantees the availability of fewer Java API runtime libraries. The Personal Platform guarantees fewer APIs than the Embedded Platform, and the Card Platform fewer than the Personal.

In addition to guaranteeing the smallest set of APIs, the Card Platform, which is targeted at SmartCards, uses only a subset of the full Java Virtual Machine instruction set. Only a subset of the features of the Java language are supported by this smaller instruction set. As a result, only Java programs that restrict themselves to features available on the Card Platform can run on a SmartCard.

Because the Java Platform is compact, it can be implemented on a wide variety of embedded systems. The compactness of the Java Platform, however, does not restrict implementation at the opposite end of the spectrum. The Java Platform also scales up to personal computers, workstations, and mainframes.

Factors that Influence Platform Independence

When you write a Java program, its degree of platform independence depends on several factors. As a developer, some of these factors are beyond your control, but most are within your control. Primarily, the degree of platform independence of any Java program you write depends on how you write it.

Java Platform Deployment

The most basic factor determining the a Java program's platform independence is the extent to which the Java Platform has been deployed on multiple platforms. Java programs will only run on computers and devices that host a Java Platform. Thus, before one of your Java programs will run on a particular computer owned by, say, your friend Alicia, two things must happen. First, the Java Platform must be ported to Alicia's particular type of hardware and operating system. Once the port has been done by

some Java Platform vendor, that port must in some way get installed on Alicia's computer. So a critical factor determining the true extent of platform independence of Java programs--and one that is beyond the control of the average developer--is the availability of Java Platform implementations and their distribution.

Fortunately for the Java developer, the deployment of the Java Platform has proceeded with great momentum, starting with Web-browsers, then moving on to desktop, workstation, and network operating systems. With the advent of chips optimized to execute Java bytecodes efficiently, the Java Platform will to some extent work its way into many different kinds of embedded devices. It is increasingly likely, therefore, that your friend Alicia will have a Java Platform implementation on her computer.

The Java Platform Version and Edition

The deployment of the Java Platform is a bit more complicated, however, because not all standard runtime libraries are guaranteed to be available at every Java Platform. The basic set of libraries guaranteed to be available at a Java Platform is called the *Java Core API*. A Java Virtual Machine accompanied by the class files that constitute the Core API is called the *Java Core Platform*. This edition of the Java Platform has the minimum set of Java API libraries that you can assume will be available at network computers, desktop computers, and workstations. As mentioned earlier, three other editions of the Java Platform--the Embedded, Personal, and Card Platforms--provide subsets of the Core API for embedded systems. The standard runtime libraries not guaranteed to be available at the Core Platform are collectively called the *Java Standard Extension API*. These libraries include such services as telephony, commerce, and media such as audio, video, or 3D. If your program uses libraries from the Standard Extension API, it will run anywhere those standard extension API libraries are available, but not on a computer that implements only the basic Java Core Platform.

Another complicating factor is that in a sense the Java Platform is a moving target--it evolves over time. Although the Java Virtual Machine is likely to evolve very gradually, the Java API will probably change more frequently. Over time, features will be added to and removed from both the Core and Standard Extension APIs, and parts of the Standard Extension API may migrate into the Core API. The changes made to the Java Platform should for the most part be upwards compatible, meaning they won't break existing Java programs, but some changes may not be. As obsolete features are removed in a new version of the Java Platform, existing Java programs that depend upon those features won't run on the new version. Also, changes may not be downwards compatible, meaning programs that are compiled for a new version of the Java Platform won't necessarily work on an old version. The dynamic nature of the Java Platform complicates things somewhat for the developer wishing to write a Java program that will run on any computer.

In theory, your program should run on all computers that host a Java Core Platform so long as you depend only upon the runtime libraries in the Core API. In practice, however, new versions of the Core API will take time to percolate everywhere. When your program depends on newly added features of the latest version of the Java Core API, there may be some hosts that can't run it because they have an older version. This is not a new problem to software developers--programs written for Windows 95, for example, don't work on the previous version of the operating system, Windows 3.1--but because Java enables the network delivery of software, it becomes a more acute problem. The promise of Java is not only that it is easy to port programs from one platform to another, but that one version of a binary executable Java program placed on a server can be delivered across a network and run on all computers.

As a developer, you can't control the release cycles or deployment schedules of the Java Platform, but

you can choose the Java Platform edition and version that your programs depend upon. In practice, therefore, you will have to decide when a new version of the Java Platform has been distributed to a great enough extent to justify writing programs for that version.

Native Methods

Besides the Java Platform version and edition your program depends on, the other major factor determining the extent of platform independence of your Java program is whether or not you call native methods. The most important rule to follow when you are writing a platform independent Java program is: don't directly or indirectly invoke any native methods that aren't part of the Java API. As you can see in Figure 2-1, calling native methods outside the Java API renders your program platform-specific.




Figure 2-1. A platform-specific Java program.

Calling native methods directly is appropriate in situations where you don't desire platform independence. In general, native methods are useful in three cases:

- for accessing features of an underlying host platform that are not accessible through the Java API
- for accessing a legacy system or using an already existing library that isn't written in Java
- for speeding up the performance of a program by implementing time-critical code as native methods

If you need to use native methods and also need your program to run on several platforms, you'll have to port the native methods to all the required platforms. This porting must be done the old fashioned way, and once you've done it, you'll have to figure out how to deliver the platform-specific native method libraries to the appropriate hosts. Because Java's architecture was designed to simplify multi-platform support, your initial goal in writing a platform-independent Java program should be to avoid native methods altogether and interact with the host only through the Java API.

Non-Standard Runtime Libraries

Native methods aren't inherently incompatible with platform independence. What is important is whether or not the methods you invoke are implemented "everywhere." Implementations of the Java API on operating systems such as Windows or Solaris use native methods to access the host. When you call a method in the Java API, you are certain it will be available everywhere. It doesn't matter if in some places the method is implemented as a native method.

Java Platform implementations can come from a variety of vendors, and although every vendor must supply the standard runtime libraries of the Java API, individual vendors may also supply extra libraries. If you are interested in platform independence, you must remain aware of whether any non-standard runtime libraries you use call native methods. Non-standard libraries that don't call native methods don't degrade your program's platform independence. Using non-standard libraries that do call native methods, however, yields the same result as calling native methods directly—it renders your program platform-specific.

For example, Microsoft offers several sets of non-standard runtime libraries with their Java Platform implementation. One set is the Application Foundation Classes (or AFC). The AFC library is delivered along with the Java Platform implementation in Microsoft's Internet Explorer (Version 4.0 and beyond) web browser. The AFC library extends the capabilities provided by the standard runtime libraries of the Java API, but doesn't call native methods outside the Java API. The AFC library accesses the host only through the Java API. As a result, programs that use AFC should run on all implementations of the Java Platform.

One concern surrounding the use of a non-standard class library (such as AFC) that doesn't call native methods, is you must deliver it to Java Platforms that don't support it directly. For instance, imagine you write a Java program that uses AFC and make it available for download across a network. If you want the program to run on all platforms, you'll have to make the AFC classes available for download too. Compared to other browsers, the program might get started sooner at Microsoft's browser because it already has the AFC classes close at hand. Other browsers that don't store the AFC library locally would have to download the needed AFC classes across the network. Other than this potential difference in download time, using a non-standard library that interacts with the host only through the Java API won't reduce the platform independence of a Java program. To keep your program platform independent, however, you must deliver the library along with your program.

One other Microsoft extension to the standard libraries of the Java Platform gives you a way to generate Java class files that grant access to COM (Component Object Model) objects. If you want to interact with a particular COM object, you use Microsoft's tool to generate a Java class file that gives your program a Java interface to the COM object. In effect, the generated class files form an extra runtime library that is available only at Microsoft's Java Platform implementation. Although this extra library offers more capabilities to Java programs, it reduces the platform independence of any Java program that takes advantage of it. Why? Because the extra libraries call native methods that are, initially at least, only available on Windows 95 and Windows NT. (In the future, Microsoft plans to port the Java/COM interface to non-Microsoft platforms.) Because the extra libraries provided by Microsoft use native methods that are specific to Microsoft operating systems, Java programs that use the extra libraries will only work on Microsoft operating systems.

Another potential ramification of using a vendor's non-standard runtime library that calls native methods directly is that your program will only work on that vendor's Java Platform implementation. For example, the class files described above that give your Java program access to COM objects use a

special native method interface (the Java/COM interface) of Microsoft's Java Virtual Machine. Currently, this special native method interface is available only on Microsoft's virtual machine implementation on Windows 95 and Windows NT. Microsoft may port the Java/COM interface to other platforms, to increase the platform independence of programs that take advantage of it. But even if Microsoft is able to port it to every major platform, Java programs that use it will most likely only work properly when running on a Microsoft implementation of the Java Virtual Machine.

Because Java Platforms can come from different vendors, there can be different Java Platform implementations from different vendors for the same hardware and operating system. To run a Java program on Windows 95, for example, you could use a Java Platform from Sun, Microsoft, Borland, Symantec, or Asymetrix. The level of platform independence that a Java program has depends not only on how many different host computers it can run on, but also on how many different Java Platform implementations it can run on each host.

Virtual Machine Dependencies

Two other rules to follow when writing a platform independent Java program involve portions of the Java Virtual Machine that can be implemented differently by different vendors. The rules are:

1. don't depend upon timely finalization for program correctness, and
2. don't depend upon thread prioritization for program correctness.

These two rules address the variations allowed in the Java Virtual Machine specification for garbage collection and threads.

All Java Virtual Machine must have a garbage-collected heap, but different implementations can use different garbage collection techniques. This flexibility in the Java Virtual Machine specification means that the objects of a particular Java program can be garbage collected at completely different times on different virtual machines. This in turn means that finalizers, which are run by the garbage collector before an object is freed, can run at different times on different virtual machines. If you use a finalizer to free finite memory resources, such as file handles, your program may run on some virtual machine implementations but not others. On some implementations, your program could run out of the finite resource before the garbage collector gets around to invoking the finalizers that free the resource.

Another variation allowed in different implementations of the Java Virtual Machine involves thread prioritization. The Java Virtual Machine specification guarantees that all runnable threads at the highest priority in your program will get some CPU time. The specification also guarantees that lower priority threads will run when higher priority threads are blocked. The specification does not, however, prohibit lower priority threads from running when higher priority threads aren't blocked. On some virtual machine implementations, therefore, lower priority threads may get some CPU time even when the higher priority threads aren't blocked. If your program depends for correctness on this behavior, however, it may work on some virtual machine implementations but not others. To keep your multi-threaded Java program platform independent, you must rely on synchronization--not prioritization--to coordinate inter-activity between threads.

User Interface Dependencies

Another major variation between different Java Platform implementations is user interface. User interface is one of the more difficult issues in writing platform independent Java programs. The AWT

user interface library gives you a set of basic user-interface components that map to native components on each platform. Libraries such as Microsoft's AFC, Netscape's IFC, and Sun's JFC, give you advanced components that don't map directly to native components. From this raw material, you must build an interface that users on many different platforms will feel comfortable with. This is not always an easy task.

Users on different platforms are accustomed to different ways of interacting with their computer. They metaphors are different. The components are different. The interaction between the components is different. Although the AWT library makes it fairly easy to create a user interface that runs on multiple platforms, it doesn't necessarily make it easy to devise an interface that keeps users happy on multiple platforms.

Bugs in Java Platform Implementations

One final source of variation among different implementations of the Java Platform is bugs. Although Sun has developed a comprehensive suite of tests that Java Platform implementations must pass, it is still possible that some implementations will be distributed with bugs in them. The only way you can defend yourself against this possibility is through testing. If there is a bug, you can determine through testing whether the bug affects your program, and if so, attempt to find a work-around.

Testing

Given the allowable differences between Java Platform implementations, the platform dependent ways you can potentially write a Java program, and the simple possibility of bugs in any particular Java Platform implementation, you should test your Java programs on all platforms you plan to claim it runs on. Java programs are not platform independent to a great enough extent that you only need test them on one platform. You still need to test a Java program on multiple platforms, and you should probably test it on the various Java Platform implementations that are likely to be found on each host computer you claim your program runs on. In practice, therefore, testing your Java program on the various host computers and Java Platform implementations that you plan to claim your program works on is a key factor in making your program platform independent.

Seven Steps to Platform Independence

Java's architecture allows you to choose between platform independence and other concerns. You make your choice by the way in which you write your program. If your goal is to take advantage of platform-specific features not available through the Java API, to interact with a legacy system, to use an existing library written not written in Java, or to maximize the execution speed of your program, you can use native methods to help you achieve that goal. In such cases, your programs will have reduced platform independence, and that will usually be acceptable. If, on the other hand, your goal is platform independence, then you should follow certain rules when writing your program. The following seven steps outline one path you can take to maximize your program's portability:

1. Choose a set of host computers that you will claim your program runs on (your "target hosts").
2. Choose a version of the Java Platform that you feel is well enough distributed among your target hosts. Write your program to run on this version of the Java Platform.
3. For each target host, choose a set of Java Platform implementations that you will claim your program runs on (your "target runtimes").

4. Write your program so that it accesses the host computer only through the standard runtime libraries of the Java API. (Don't invoke native methods, or use vendor-specific libraries that invoke native methods.)
5. Write your program so that it doesn't depend for correctness on timely finalization by the garbage collector or on thread prioritization.
6. Strive to design a user interface that works well on all of your target hosts.
7. Test your program on all of your target runtimes and all of your target hosts.

If you follow the seven steps outlined above, your Java program will definitely run on all your target hosts. If your target hosts cover most major Java Platform vendors on most major host computers, there is a very good chance your program will run everywhere else as well.

If you wish, you can have your program certified as "100% Pure Java." There are several reasons that you may wish to do this if you are writing a program that you intend to be platform independent. For example, if your program is certified 100% Pure, you can brand it with the 100% Pure Java coffee cup icon. You can also potentially participate in co-marketing programs with Sun. You may, however, wish to go through the certification process simply as an added check of the platform independence of your program. In this case, you have the option of just running 100% Pure verification tools you can download for free. These tools will report problems with your program's "purity" without requiring you to go through the full certification process.

The 100% Pure certification is not quite a full measure of platform independence. Part of platform independence is that user's expectations are fulfilled on multiple platforms. The 100% Pure testing process does not attempt to measure user fulfillment. It only checks to make certain your program depends only on the Java Core Platform. You could write a Java program that passes the 100% Pure tests, but still doesn't work well on all platforms from the perspective of users. Nonetheless, running your code through the 100% Pure testing process can be a worthwhile step on the road to creating a platform independent Java program.

The Politics of Platform Independence

As illustrated in Figure 2-2, Java Platform vendors are allowed to extend the standard components of the Java Platform in non-standard and platform-specific ways, but they must always support the standard components. In the future, Sun Microsystems intends to prevent the standard components of the Java Platform from splitting into several competing, slightly incompatible systems, as happened, for instance, with UNIX. The license that all Java Platform vendors must sign requires compatibility at the level of the Java Virtual Machine and the Java API, but permits differentiation in the areas of performance and extensions. There is some flexibility, as mentioned above, in the way vendors are allowed to implement threads, garbage-collection, and user interface look and feel. If things go as Sun plans, the core components of the Java Platform will remain a standard to which all vendors faithfully adhere, and the ubiquitous nature of the standard Java Platform will enable you to write programs that really are platform independent.




Figure 2-2. Java Platform implementations from different vendors.

You can rely on the standard components of the Java Platform because every Java Platform vendor must support them. If you write a program that only depends on these components, it should "run anywhere," but may suffer to some extent from the lowest-common-denominator problem. Yet because vendors are allowed to extend the Java Platform, they can give you a way to write platform-specific programs that take full advantage of the features of the underlying host operating system. The presence of both required standard components and permitted vendor extensions at any Java Platform implementation gives developers a choice. This arrangement allows developers to balance platform independence with other concerns.

There is currently a marketing battle raging for the hearts and minds of software developers over how they will write Java programs--in particular, whether or not they will choose to write platform-independent or platform-specific programs. The choice that Java graciously gives to developers also potentially threatens some vested interests in the software industry.

Java's support for platform independence threatens to weaken the "lock" enjoyed by operating system vendors. If all of your software runs on only one operating system, then your next computer will also very likely run that same operating system. You are "locked in" to one operating system vendor because your investment in software depends on an API proprietary to that vendor. You are also likely locked into one hardware architecture because the binary form of your programs requires a particular kind of microprocessor. If instead, much of your software is written to the Java API and stored as bytecodes in class files, it becomes easier for you to migrate to a different operating system vendor the next time you buy a computer. Because the Java Platform can be implemented in software on top of existing operating systems, you can switch operating systems and take all of your old platform-independent Java-based software with you.

Microsoft dominates the desktop operating system market largely because most available software runs only on Microsoft operating systems. It is in Microsoft's strategic interest to continue this status quo, so

they are encouraging developers to use Java as a language to write programs that run only on Microsoft platforms. It is in just about every other operating system vendor's strategic interest to weaken Microsoft's lock on the operating system market, so the other players are encouraging developers to write Java programs that are platform independent. Sun, Netscape, IBM, and many others have banded together to promote Sun's "100% Pure Java initiative," through which they hope to educate and persuade developers to go the platform independence route.

Microsoft's approach to Java is to make Windows the best platform on which to develop and run Java programs. They want developers to use Microsoft's tools and libraries whether the developer chooses platform independence or not. Their AWT library, for example, enables Java developers to build advanced platform independent user interfaces. Microsoft also provides the Java/COM native method interface, which allows developers to use Java to write full-fledged, platform-specific Windows programs. Still, in the "spin" Microsoft gives to Java in promotional material to developers, they strongly favor the platform-specific Windows path. They extol the virtues of using Java to write programs that take full advantage of the Windows platform.

Sun and the other operating system vendors behind the 100% Pure Java initiative are attempting to counter Microsoft's spin with some of their own. The promotional material from these companies focuses on the benefits of writing platform-independent Java programs.

On one level, it is a battle between two icons. If you write your Java program Microsoft's way, you get to brand your product with a Windows 95/NT icon that displays the famous four-paneled Windows logo. If you go the 100% Pure Java route, you get to brand your product with a 100% Pure Java icon that displays the famous steaming coffee cup logo.

As a developer, the politics and propaganda swirling around the software industry need not be a major factor when you decide how to write a particular Java program. For some programs you write, platform independence may be the right approach; for others, a platform-specific program may make more sense. In each individual case, you can make a decision based on what you feel your customers want and how you want to position yourself in the marketplace with respect to your competitors.

The Resources Page

For links to more information about Java and platform independence, visit the resources page for this chapter: <http://www.artima.com/insidejvm/platindep.html>



Chapter Three

Security

Aside from platform independence, discussed in the previous chapter, the other major technical challenge a network-oriented software technology must deal with is security. Networks, because they allow computers to share data and distribute processing, can potentially serve as a way to break into a computer system, enabling someone to steal information, destroy information, or steal computing resources. As a consequence, connecting a computer to a network raises many security issues.

To address the security concerns raised by networks, Java's architecture comes with an extensive built-in security model. This chapter gives an overview of the security model built into Java's core architecture.

Why Security?

Java's security model is one of the key architectural features that makes it an appropriate technology for networked environments. Security is important because networks represent a potential avenue of attack to any computer hooked to them. This concern becomes especially strong in an environment in which software is downloaded across the network and executed locally, as is done, for example, with Java applets. Because the class files for an applet are automatically downloaded when a user goes to the containing web page in a browser, it is likely that a user will encounter applets from untrusted sources. Without any security, this would be a convenient way to spread viruses. Thus, Java's security mechanisms help make Java suitable for networks because they establish a needed trust in the safety of network-mobile code.

Java's security model is focused on protecting end-users from hostile programs downloaded across a network from untrusted sources. To accomplish this goal, Java provides a customizable "sandbox" in which Java programs run. A Java program must play only inside its sandbox. It can do anything within the boundaries of its sandbox, but can't take any action outside those boundaries. The sandbox for untrusted Java applets, for example, prohibits many activities, including:

- reading or writing to the local disk,
- making a network connection to any but the host from which the applet came,
- creating a new process, and
- loading a new dynamic library and directly calling a native method.

By making it impossible for downloaded code to perform certain actions, Java's security model protects the end-user from the threat of hostile code.

The Sandbox

Traditionally, you had to trust software before you ran it. You achieved security by being careful only to use software from trusted sources, and by regularly scanning for viruses just to make sure. Once some software got access to your system, it had full reign. If it was malicious, it could do a great deal of damage because there were no restrictions placed on it by the runtime environment of your computer. So in the traditional security scheme, you tried to prevent malicious code from ever gaining access to your computer in the first place.

The sandbox security model makes it easier to work with software that comes from sources you don't fully trust. Instead of approaching security by requiring you to prevent any code you don't trust from ever making its way onto your computer, the sandbox model allows you to welcome code from any source. But as code from an untrusted source runs, the sandbox restricts code from untrusted sources from taking any actions that could possibly harm your system. You don't need to figure out what code you can and can't trust. You don't need to scan for viruses. The sandbox itself prevents any viruses or other malicious code you may invite into your computer from doing any damage.

If you have a properly skeptical mind, you'll need to be convinced a sandbox has no leaks before you trust it to protect you. To make sure the sandbox must have no leaks, Java's security model involves every aspect of its architecture. If there were areas in Java's architecture where security was weak, a malicious programmer (a "cracker") could potentially exploit those areas to "go around" the sandbox. To understand the sandbox, therefore, you must look at several different parts of Java's architecture, and understand how they work together.

The fundamental components responsible for Java's sandbox are:

- the class loader architecture
- the class file verifier
- safety features built into the Java Virtual Machine (and the language)
- the security manager and the Java API

One of the greatest strengths of Java's security model is that two of these components, the class loader and the security manager, are customizable. By customizing these components, you can create a customized security policy for a Java application. As a developer, you may never need to create your own customized sandbox. You can often make use of sandboxes created by others. When you write and run a Java applet, for instance, you make use of a sandbox created by the developers of the web browser that hosts your applet.

The Class Loader Architecture

In Java's sandbox, the class loader architecture is the first line of defense. It is the class loader, after all, that brings code into the Java Virtual Machine--code that could be hostile. The class loader architecture contributes to Java's sandbox in two ways:

1. it guards the borders of the trusted class libraries, and
2. it prevents malicious code from interfering with benevolent code.

The class loader architecture guards the borders of the trusted class libraries by preventing untrusted classes from pretending to be trusted. If a malicious class could successfully trick the Java Virtual

Machine into believing it was a trusted class from the Java API, that malicious class could potentially break through the sandbox barrier. By preventing untrusted classes from impersonating trusted classes, the class loader architecture blocks one potential approach to compromising the security of the Java runtime.

The class loader architecture prevents malicious code from interfering with benevolent code by providing protected name-spaces for classes loaded by different class loaders. A *name-space* is a set of unique names for loaded classes that is maintained by the Java Virtual Machine. Once a Java Virtual Machine has loaded a class named `Volcano` into a particular name-space, for example, it is impossible to load a different class named `Volcano` into that same name-space. You can load multiple `Volcano` classes into a Java Virtual Machine, however, because you can create multiple name-spaces inside a Java application by creating multiple class loaders. If you create three separate name-spaces (one for each of three class loaders) in a running Java application, then, by loading one `Volcano` class into each name-space, your program could load three different `Volcano` classes into your application.

Name-spaces contribute to security because you can place a shield between classes loaded into different name-spaces. Inside the Java Virtual Machine, classes in the same name-space can interact with one another directly. Classes in different name-spaces, however, can't even detect each other's presence unless you explicitly provide a mechanism that allows them to interact. If a malicious class, once loaded, had guaranteed access to every other class currently loaded by the virtual machine, that class could potentially learn things it shouldn't know or interfere with the proper execution of your program.

Often, a class loader object relies on other class loaders--at the very least, upon the primordial class loader--to help it fulfill some of the class load requests that come its way. For example, imagine you write a Java application that installs a class loader whose particular manner of loading class files is by downloading them across a network. Assume that during the course of running the Java application, a request is made of your class loader to load a class named `Volcano`. One way you could write the class loader is to have it first asks the primordial class loader to find and load the class from its trusted repository. In this case, since `Volcano` is not a part of the Java API, assume the primordial class loader can't find a class named `Volcano`. When the primordial class loader responds that it can't load the class, your class loader could then attempt to load the `Volcano` class in its custom manner, by downloading it across the network. Assuming your class loader was able to download class `Volcano`, that `Volcano` class could then play a role in the application's future course of execution.

To continue with the same example, assume that at some time later a method of class `Volcano` is invoked for the first time, and that method references class `String` from the Java API. Because it is the first time the reference was used by the running program, the virtual machine asks your class loader (the one that loaded `Volcano`) to load `String`. As before, your class loader first passes the request to the primordial class loader, but in this case, the primordial class loader is able to return a `String` class back to your class loader. (The primordial class loader most likely didn't have to actually load `String` at this point because, given that `String` is such a fundamental class in Java programs, it was almost certainly used before and therefore already loaded. Most likely, the primordial class loader just returned the `String` class that it had previously loaded from the trusted repository.) Since the primordial class loader was able to find the class, your class loader doesn't attempt to download it across the network; it merely passes to the virtual machine the `String` class returned by the primordial class loader. From that point forward, the virtual machine uses that `String` class whenever class `Volcano` references a class named `String`.

When you write a class loader, you create a new environment in which the loaded code runs. If you want

the environment to be free of security holes, you must follow certain rules when you write your class loader. In general, you will want to write class loaders such that they protect the borders of trusted class libraries, such as those of the Java API.

Java allows classes in the same package to grant each other special access privileges that aren't granted to classes outside the package. So, if your class loader receives a request to load a class that by its name brazenly declares itself to be part of the Java API (for example, a class named `java.lang.Virus`), it could gain special access to the trusted classes of `java.lang` and could possibly use that special access for devious purposes. Consequently, you would normally write a class loader so that it simply refuses to load any class that claims to be part of the Java API (or any other trusted runtime library), but that doesn't exist in the local trusted repository. In other words, after your class loader passes a request to the primordial class loader, and the primordial class loader indicates it can't load the class, your class loader should check to make sure the class doesn't declare itself to be a member of a trusted package. If it does, your class loader, instead of trying to download the class across the network, should throw a security exception.

In addition, you may have installed some packages in the trusted repository that contain classes you want your application to be able to load through the primordial class loader, but that you don't want to be accessible to classes loaded through your class loader. For example, assume you have created a package named `absolutePower` and installed it on the local repository accessible by the primordial class loader. Assume also that you don't want classes loaded by your class loader to be able to load any class from the `absolutePower` package. In this case, you would write your class loader such that the very first thing it does is make sure the requested class doesn't declare itself as a member of the `absolutePower` package. If such a class is requested, your class loader, rather than passing the class name to the primordial class loader, should throw a security exception.

The only way a class loader can know whether or not a class is from a restricted package, such as `java.lang`, or a forbidden package, such as `absolutePower`, is by the class's name. Thus a class loader must be given a list of the names of restricted and forbidden packages. Because the name of class `java.lang.Virus` indicates it is from the `java.lang` package, and `java.lang` is on the list of restricted packages, your class loader should throw a security exception if the primordial class loader can't load it. Likewise, because the name of class `absolutePower.FancyClassLoader` indicates it is part of the `absolutePower` package, and the `absolutePower` package is on the list of forbidden packages, your class loader should throw a security exception absolutely.

A common way, therefore, to write a security-minded class loader is using the following four steps:

1. If packages exist that this class loader is not allowed to load from, the class loader checks whether the requested class is in one of those forbidden packages. If so, it throws a security exception. Else, it continues on to step two.
2. The class loader passes the request to the primordial class loader. If the primordial class loader successfully returns the class, the class loader returns that same class. Else, it continues on to step three.
3. If trusted packages exist that this class loader is not allowed to add classes to, the class loader checks whether the requested class is in one of those restricted packages. If so, it throws a security exception. Else, it continues on to step four.
4. Finally, the class loader attempts to load the class in the custom way, such as by downloading it

across a network. If successful, it returns the class. Else, it throws a "no class definition found" error.

By performing steps one and three as outlined above, the class loader guards the borders of the trusted packages. With step one, it prevents a class from a forbidden package to be loaded at all. With step three, it doesn't allow an untrusted class to insert itself into a trusted package.

The Class File Verifier

Working in conjunction with the class loader, the class file verifier ensures that loaded class files have a proper internal structure. If the class file verifier discovers a problem with a class file, it throws an exception. Although compliant Java compilers should not generate malformed class files, a Java Virtual Machine can't tell how a particular class file was created. Because a class file is just a sequence of binary data, a virtual machine can't know whether a particular class file was generated by a well-meaning Java compiler or by shady crackers bent on compromising the integrity of the virtual machine. As a consequence, all Java Virtual Machine implementations have a class file verifier that can be invoked on untrusted classes, to make sure the classes are safe to use.

One of the security goals that the class file verifier helps achieve is program robustness. If a buggy compiler or savvy cracker generated a class file that contained a method whose bytecodes included an instruction to jump beyond the end of the method, that method could, if it were invoked, cause the virtual machine to crash. Thus, for the sake of robustness, it is important that the virtual machine verify the integrity of the bytecodes it imports. Although Java Virtual Machine designers are allowed to decide when their virtual machines will perform these checks, many implementations will do most checking just after a class is loaded. Such a virtual machine, rather than checking every time it encounters a jump instruction as it executes bytecodes, analyzes bytecodes (and verifies their integrity) once, before they are ever executed. As part of its verification of bytecodes, the Java Virtual Machine makes sure all jump instructions cause a jump to another valid instruction in the bytecode stream of the method. In most cases, checking all bytecodes once, before they are executed, is a more efficient way to guarantee robustness than checking every bytecode instruction every time it is executed.

A class file verifier that performs its checking as early as possible most likely operates in two distinct phases. During phase one, which takes place just after a class is loaded, the class file verifier checks the internal structure of the class file, including verifying the integrity of the bytecodes it contains. During phase two, which takes place as bytecodes are executed, the class file verifier confirms the existence of symbolically referenced classes, fields, and methods.

Phase One: Internal Checks

During phase one, the class file verifier checks everything that's possible to check in a class file by looking at only the class file itself. In addition to verifying the integrity of the bytecodes during phase one, the verifier performs many checks for proper class file format and internal consistency. For example, every class file must start with the same four bytes, the magic number: 0xCAFEBAE. The purpose of magic numbers is to make it easy for file parsers to recognize a certain type of file. Thus, the first thing a class file verifier likely checks is that the imported file does indeed begin with 0xCAFEBAE.

The class file verifier also checks to make sure the class file is neither truncated nor enhanced with extra trailing bytes. Although different class files can be different lengths, each individual component contained inside a class file indicates its length as well as its type. The verifier can use the component

types and lengths to determine the correct total length for each individual class file. In this way, it can verify that the imported file has a length consistent with its internal contents.

The verifier also looks at individual components, to make sure they are well-formed instances of their type of component. For example, a method descriptor (its return type and the number and types of its parameters) is stored in the class file as a string that must adhere to a certain context-free grammar. One check the verifier performs on individual components is to make sure each method descriptor is a well-formed string of the appropriate grammar.

In addition, the class file verifier checks that the class itself adheres to certain constraints placed upon it by the specification of the Java programming language. For example, the verifier enforces the rule that all classes, except class `Object`, must have a superclass. Thus, the class file verifier checks at run-time some of the Java language rules that should have been enforced at compile-time. Because the verifier has no way of knowing if the class file was generated by a benevolent, bug-free compiler, it checks each class file to make sure the rules are followed.

Once the class file verifier has successfully completed the checks for proper format and internal consistency, it turns its attention to the bytecodes. During this part of phase one, which is commonly called the "bytecode verifier," the Java Virtual Machine performs a data-flow analysis on the streams of bytecodes that represent the methods of the class. To understand the bytecode verifier, you need to understand a bit about bytecodes and frames.

The bytecode streams that represent Java methods are a series of one-byte instructions, called *opcodes*, each of which may be followed by one or more *operands*. The operands supply extra data needed by the Java Virtual Machine to execute the opcode instruction. The activity of executing bytecodes, one opcode after another, constitutes a thread of execution inside the Java Virtual Machine. Each thread is awarded its own *Java Stack*, which is made up of discrete *frames*. Each method invocation gets its own frame, a section of memory where it stores, among other things, local variables and intermediate results of computation. The part of the frame in which a method stores intermediate results is called the method's *operand stack*. An opcode and its (optional) operands may refer to the data stored on the operand stack or in the local variables of the method's frame. Thus, the virtual machine may use data on the operand stack, in the local variables, or both, in addition to any data stored as operands following an opcode when it executes the opcode.

The bytecode verifier does a great deal of checking. It checks to make sure that no matter what path of execution is taken to get to a certain opcode in the bytecode stream, the operand stack always contains the same number and types of items. It checks to make sure no local variable is accessed before it is known to contain a proper value. It checks that fields of the class are always assigned values of the proper type, and that methods of the class are always invoked with the correct number and types of arguments. The bytecode verifier also checks to make sure that each opcode is valid, that each opcode has valid operands, and that for each opcode, values of the proper type are in the local variables and on the operand stack. These are just a few of the many checks performed by the bytecode verifier, which is able, through all its checking, to verify that a stream of bytecodes is safe for the Java Virtual Machine to execute.

Phase one of the class file verifier makes sure the imported class file is properly formed, internally consistent, adheres to the constraints of the Java programming language, and contains bytecodes that will be safe for the Java Virtual Machine to execute. If the class file verifier finds that any of these are not true, it throws an error, and the class file is never used by the program.

Phase Two: Verification of Symbolic References

Although phase one happens immediately after the Java Virtual Machine loads a class file, phase two is delayed until the bytecodes contained in the class file are actually executed. During phase two, the Java Virtual Machine follows the references from the class file being verified to the referenced class files, to make sure the references are correct. Because phase two has to look at other classes external to the class file being checked, phase two may require that new classes be loaded. Most Java Virtual Machine implementations will likely delay loading classes until they are actually used by the program. If an implementation does load classes earlier, perhaps in an attempt to speed up the loading process, then it must still give the impression that it is loading classes as late as possible. If, for example, a Java Virtual Machine discovers during early loading that it can't find a certain referenced class, it doesn't throw a "class definition not found" error until (and unless) the referenced class is used for the first time by the running program. Therefore, phase two, the checking of symbolic references, is usually delayed until each symbolic reference is actually used for the first time during bytecode execution.

Phase two of class file verification is really just part of the process of dynamic linking. When a class file is loaded, it contains symbolic references to other classes and their fields and methods. A symbolic reference is a character string that gives the name and possibly other information about the referenced item--enough information to uniquely identify a class, field, or method. Thus, symbolic references to other classes give the full name of the class; symbolic references to the fields of other classes give the class name, field name, and field descriptor; symbolic references to the methods of other classes give the class name, method name, and method descriptor.

Dynamic linking is the process of *resolving* symbolic references into direct references. As the Java Virtual Machine executes bytecodes and encounters an opcode that, for the first time, uses a symbolic reference to another class, the virtual machine must resolve the symbolic reference. The virtual machine performs two basic tasks during resolution:

1. find the class being referenced (loading it if necessary)
2. replace the symbolic reference with a direct reference, such as a pointer or offset, to the class, field, or method

The virtual machine remembers the direct reference so that if it encounters the same reference again later, it can immediately use the direct reference without needing to spend time resolving the symbolic reference again.

When the Java Virtual Machine resolves a symbolic reference, phase two of the class file verifier makes sure the reference is valid. If the reference is not valid--for instance, if the class cannot be loaded or if the class exists but doesn't contain the referenced field or method--the class file verifier throws an error.

As an example, consider again the `Volcano` class. If a method of class `Volcano` invokes a method in a class named `Lava`, the name and descriptor of the method in `Lava` are included as part of the binary data in the class file for `Volcano`. So, during the course of execution when the `Volcano`'s method first invokes the `Lava`'s method, the Java Virtual Machine makes sure a method exists in class `Lava` that has a name and descriptor that matches those expected by class `Volcano`. If the symbolic reference (class name, method name and descriptor) is correct, the virtual machine replaces it with a direct reference, such as a pointer, which it will use from then on. But if the symbolic reference from class `Volcano` doesn't match any method in class `Lava`, phase two verification fails, and the Java Virtual Machine throws a "no such method" error.

Binary Compatibility

The reason phase two of the class file verifier must look at classes that refer to one

another to make sure they are compatible is because Java programs are dynamically linked. Java compilers will often recompile classes that depend on a class you have changed, and in so doing, detect any incompatibility at compile-time. But there may be times when your compiler doesn't recompile a dependent class. For example, if you are developing a large system, you will likely partition the various parts of the system into packages. If you compile each package separately, then a change to one class in a package would cause a recompilation of affected classes within that same package, but not necessarily in any other package. Moreover, if you are using someone else's packages, especially if your program downloads class files from someone else's package across a network as it runs, it may be impossible for you to check for compatibility at compile-time. That's why phase two of the class file verifier must check for compatibility at run-time.

As an example of incompatible changes, imagine you compiled class `Volcano` (from the above example) with a Java compiler. Because a method in `Volcano` invokes a method in another class named `Lava`, the Java compiler would look for a class file or a source file for class `Lava` to make sure there was a method in `Lava` with the appropriate name, return type, and number and types of arguments. If the compiler couldn't find any `Lava` class, or if it encountered a `Lava` class that didn't contain the desired method, the compiler would generate an error and would not create a class file for `Volcano`. Otherwise, the Java compiler would produce a class file for `Volcano` that is compatible with the class file for `Lava`. In this case, the Java compiler refused to generate a class file for `Volcano` that wasn't already compatible with class `Lava`.

The converse, however, is not necessarily true. The Java compiler could conceivably generate a class file for `Lava` that isn't compatible with `Volcano`. If the `Lava` class doesn't refer to `Volcano`, you could potentially change the name of the method `Volcano` invokes from the `Lava` class, and then recompile only the `Lava` class. If you tried to run your program using the new version of `Lava`, but still using the old version of `Volcano` that wasn't recompiled since you made your change to `Lava`, the Java Virtual Machine would, as a result of phase two class file verification, throw a "no such method" error when `Volcano` attempted to invoke the now non-existent method in `Lava`.

In this case, the change to class `Lava` broke *binary compatibility* with the pre-existing class file for `Volcano`. In practice, this situation may arise when you update a library you have been using, and your existing code isn't compatible with the new version of the library. To make it easier to alter the code for libraries, the Java programming language was designed to allow you to make many kinds of changes to a class that don't require recompilation of classes that depend upon it. The changes you are allowed to make, which are listed in the Java Language Specification, are called the rules of binary compatibility. These rules clearly define what can be changed, added, or deleted in a class without breaking binary compatibility with pre-existing class files that depend on the changed class. For example, it is always a binary compatible change to add a new method to a class, but never to delete a method that other classes may be using. So in the case of `Lava`, you violated the rules of binary compatibility when you changed the name of the method used by `Volcano`, because you in effect deleted the old method and added a new. If you had, instead, added the new method and then rewritten the old method so it calls the new, that change would have been binary compatible with any pre-existing class file that already used `Lava`, including `Volcano`.

Safety Features Built Into the Java Virtual Machine

Once the Java Virtual Machine has loaded a class and performed phase one of class file verification, the bytecodes are ready to be executed. Besides the verification of symbolic references (phase two of class file verification), the Java Virtual Machine has several other built-in security mechanisms operating as bytecodes are executed. These are the same mechanisms listed in Chapter 1 as features of the Java programming language that make Java programs robust. They are, not surprisingly, also features of the Java Virtual Machine:

- type-safe reference casting
- structured memory access (no pointer arithmetic)
- automatic garbage collection (can't explicitly free allocated memory)
- array bounds checking
- checking references for `null`

By granting a Java program only safe, structured ways to access memory, the Java Virtual Machine makes Java programs more robust, but it also makes their execution more secure. Why? There are two reasons. First, a program that corrupts memory, crashes, and possibly causes other programs to crash represents one kind of security breach. If you are running a mission critical server process, it is critical that the process doesn't crash. This level of robustness is also important in embedded systems, such as a cell phone, which people don't usually expect to have to reboot. The second reason unrestrained memory access would be a security risk is because a wily cracker could potentially use it to subvert the security system. If, for example, a cracker could learn where in memory a class loader is stored, it could assign a pointer to that memory and manipulate the class loader's data. By enforcing structured access to memory, the Java Virtual Machine yields programs that are robust, but also frustrates crackers who dream of harnessing the internal memory of the Java Virtual Machine for their own devious plots.

Another safety feature built into the Java Virtual Machine—one that serves as a backup to structured memory access—is the unspecified manner in which the runtime data areas are laid out inside the Java Virtual Machine. The *runtime data areas* are the memory areas in which the Java Virtual Machine stores the data it needs to execute a Java application: Java stacks (one for each thread), a *method area*, where bytecodes are stored, and a *garbage-collected heap*, where the objects created by the running program are stored. If you peer into a class file, you won't find any memory addresses. When the Java Virtual Machine loads a class file, it decides where in its internal memory to put the bytecodes and other data it parses from the class file. When the Java Virtual Machine starts a thread, it decides where to put the Java stack it creates for the thread. When it creates a new object, it decides where in memory to put the object. Thus, a cracker cannot predict by looking at a class file where in memory the data representing that class, or objects instantiated from that class, will be kept. What's worse (for the cracker) is the cracker can't tell anything about memory layout by reading the Java Virtual Machine specification either. The manner in which a Java Virtual Machine lays out its internal data is not part of the specification. The designers of each Java Virtual Machine implementation decide which data structures their implementation will use to represent the runtime data areas, and where in memory their implementation will place them. As a result, even if a cracker were somehow able to break through the Java Virtual Machine's memory access restrictions, they would next be faced with the difficult task of finding something to subvert by looking around.

The prohibition on unstructured memory access is not something the Java Virtual Machine must actively enforce on a running program; rather, it is intrinsic to the bytecode instruction set itself. Just as there is no way to express an unstructured memory access in the Java programming language, there is also no way to express it in bytecodes—even if you write the bytecodes by hand. Thus, the prohibition on unstructured memory access is a solid barrier against the malicious manipulation of memory.

There is, however, a way to penetrate the security barriers erected by the Java Virtual Machine.

Although the bytecode instruction set doesn't give you an unsafe, unstructured way to access memory, there is a way you can go around bytecodes: native methods. Basically, when you call a native method, Java's security sandbox becomes dust in the wind. First of all, the robustness guarantees don't hold for native methods. Although you can't corrupt memory from a Java method, you can from a native method. But most importantly, native methods don't go through the Java API (they are how you go around the Java API) so the security manager isn't checked before a native method attempts to do something that could be potentially damaging. (This is, of course, often how the Java API itself gets anything done. But the native methods used by the Java API are "trusted.") Thus, once a thread gets into a native method, no matter what security policy was established inside the Java Virtual Machine, it doesn't apply anymore to that thread, so long as that thread continues to execute the native method. This is why the security manager includes a method that establishes whether or not a program can load dynamic libraries, which are necessary for invoking native methods. Applets, for example, aren't allowed to load a new dynamic library, therefore they can't install their own new native methods. They can, however, call methods in the Java API, methods which may be native, but which are always trusted. When a thread invokes a native method, that thread leaps outside the sandbox. The security model for native methods is, therefore, the same security model described earlier as the traditional approach to computer security: you have to trust a native method before you call it.

One final mechanism that is built into the Java Virtual Machine that contributes to security is structured error handling with exceptions. Because of its support for exceptions, the Java Virtual Machine has something structured to do when a security violation occurs. Instead of crashing, the Java Virtual Machine can throw an exception or an error, which may result in the death of the offending thread, but shouldn't crash the system. Throwing an error (as opposed to throwing an exception) almost always results in the death of the thread in which the error was thrown. This is usually a major inconvenience to a running Java program, but won't necessarily result in termination of the entire program. If the program has other threads doing useful things, those threads may be able to carry on without their recently departed colleague. Throwing an exception, on the other hand, may result in the death of the thread, but is often just used as a way to transfer control from the point in the program where the exception condition arose to the point in the program where the exception condition is handled.

The Security Manager and the Java API

By using class loaders, you can prevent code loaded by different class loaders from interfering with one another inside the Java Virtual Machine, but to protect assets external to the Java Virtual Machine, you must use a security manager. The security manager defines the outer boundaries of the sandbox. Because it is customizable, the security manager allows you to establish a custom security policy for an application. The Java API enforces the custom security policy by asking the security manager for permission before it takes any action that is potentially unsafe. For each potentially unsafe action, there is a method in the security manager that defines whether that action is allowed by the sandbox. Each method's name starts with "check," so for example, `checkRead()` defines whether or not a thread is allowed to read to a specified file, and `checkWrite()` defines whether or not a thread is allowed to write to a specified file. The implementation of these methods is what defines the custom security policy of the application.

Most of the activities that are regulated by a "check" method are listed below. The classes of the Java API check with the security manager before they:

- accept a socket connection from a specified host and port number
- modify a thread (change its priority, stop it, etc.)
- open a socket connection to a specified host and port number

- create a new class loader
- delete a specified file
- create a new process
- cause the application to exit
- load a dynamic library that contains native methods
- wait for a connection on a specified local port number
- load a class from a specified package (used by class loaders)
- add a new class to a specified package (used by class loaders)
- access or modify system properties
- access a specified system property
- read from a specified file
- write to a specified file

Because the Java API always checks with the security manager before it performs any of the activities listed above, the Java API will not perform any action forbidden under the security policy established by the security manager.

Two actions not present in the above list that could potentially be unsafe are allocation of memory and invocation of threads. Currently, a hostile applet that can possibly crash the browser by:

- allocating memory until it runs out
- firing off threads until everything slows to a crawl

These kinds of attacks are called *denial of service*, because they deny the end-users from using their own computers. The security manager does not allow you to enforce any kind of limit on allocated memory or thread creation. (There are no `checkAllocateMemory()` or `checkCreateThread()` methods in the security manager class.) The difficulty in attempting to thwart this kind of hostile code is that it is hard to tell the difference, for example, between a hostile applet allocating a lot of memory and an image processing applet attempting to do useful work. Other kinds of hostile applets that are currently possible are:

- applets that send unauthorized e-mail from the end-user's computer
- applets that make annoying noises even after you leave the web page
- applets that display offensive images or animations

So a security manager isn't enough to prevent every possible action that could possibly offend or inconvenience an end-user. Other than the attacks listed here, however, the security manager attempts to provide a check method that allows you to control access to any potentially unsafe action.

When a Java application starts, it has no security manager, but the application can install one at its option. If it does not install a security manager, there are no restrictions placed on any activities requested of the Java API--the Java API will do whatever it is asked. (This is why Java applications, by default, do not have any security restrictions such as those that limit the activities of untrusted applets.) If the application does install a security manager, then that security manager will be in charge for the entire remainder of the lifetime of that application. It can't be replaced, extended, or changed. From that point on, the Java API will only fulfill those requests that are sanctioned by the security manager.

In general, a "check" method of the security manager throws a security exception if the checked upon activity is forbidden, and simply returns if the activity is permitted. Therefore, the procedure a Java API method generally follows when it is about to perform a potentially unsafe activity involves two steps.

First, the Java API code checks whether a security manager has been installed. If not, it skips step two and goes ahead with the potentially unsafe action. Otherwise, as step two, it calls the appropriate "check" method in the security manager. If the action is forbidden, the "check" method will throw a security exception, which will cause the Java API method to immediately abort. The potentially unsafe action will never be taken. If, on the other hand, the action is permitted, the "check" method will simply return. In this case, the Java API method carries on and performs the potentially unsafe action.

Although you can only install one security manager, you can write the security manager so that it establishes multiple security policies. In addition to the "check" methods, the security manager also has methods that allow you to determine if a request is being made either directly or indirectly from a class loaded by a class loader object, and if so, which class loader object. This enables you to implement a security policy that varies depending upon which class loader loaded the classes making the request. You can also vary the security policy based on information about the class files loaded by the class loader, such whether or not the class files were down loaded across a network or imported from the local disk. So even though an application can only have one security manager, that security manager can establish a flexible security policy that varies depending upon the trustworthiness of the code requesting the potentially unsafe action.

Authentication

The support for authentication introduced in Java 1.1 in the `java.security` package expands your ability to establish multiple security policies by enabling you to implement a sandbox that varies depending upon who actually created the code. Authentication allows you to verify that a set of class files was blessed as trustworthy by some vendor, and that the class files were not altered en route to your virtual machine. Thus, to the extent you trust the vendor, you can ease the restrictions placed on the code by the sandbox. You can establish different security policies for code that comes from different vendors.

For links to more information about authentication and `java.security`, visit the resources page for this chapter.

Security Beyond the Architecture

Security is a tradeoff between cost and risk: the lower the security risk, the higher the cost of security. The costs associated with any computer or network security strategy must be weighed against the costs that would be associated with the theft or destruction of the information or computing resources being protected. The nature of a computer or network security strategy should be shaped by the value of the assets being protected.

To be effective, a computer or network security strategy must be comprehensive. It cannot consist exclusively of a sandbox for running downloaded Java code. For instance, it may not matter much that the Java applets you download from the internet and run on your computer can't read the word processing file of your top-secret business plan if you:

- routinely download untrusted native executables from the internet and run them
- throw away extra printed copies of your business plan without shredding them
- leave your doors unlocked when you're gone
- hire someone to help you who is actually a spy for your arch-rival

In the context of a comprehensive security strategy, however, Java's security model can play a useful role.

The nice thing about Java's security model is that once you set it up, it does most of the work for you. You don't have to worry about whether a particular program is trusted or not--the Java runtime will determine that for you; and if it is untrusted, the Java runtime will protect your assets by encasing the untrusted code in a sandbox.

End-users of Java software cannot rely only on the security mechanisms built into Java's architecture. They must have a comprehensive security policy appropriate to their actual security requirements.

Similarly, the security strategy of Java technology itself does not rely exclusively on the architectural security mechanisms described in this chapter. For example, one aspect of Java's security strategy is that anyone can sign a license agreement and get a copy of the source code of Sun's Java Platform implementation. Instead of keeping the internal implementation of Java's security architecture a secret "black box," it is open to anyone who wishes to look at it. This encourages security experts seeking a good technical challenge to try and find security holes in the implementation. When security holes are discovered, they can be patched. Thus, the openness of Java's internal implementation is part of Java's overall security strategy.

Besides openness, there are several other aspects to Java's overall security strategy that don't directly involve its architecture. You can find out more information about these on the resources page for this chapter.

The Resources Page

For more information about Java and security, see the resource page for this chapter:
<http://www.artima.com/insidejvm/security.html>

McGraw-Hill

A Division of The McGraw-Hill Companies



Chapter Four

Network mobility

The previous two chapters discussed how Java's architecture deals with the two major challenges presented to software developers by a networked computing environment. Platform independence deals with the challenge that many different kinds of computers and devices are usually connected to the same network. The sandbox security model deals with the challenge that networks represent a convenient way to transmit viruses and other forms of malicious code. This chapter describes not how Java's architecture deals with a challenge, but how it seizes an opportunity made possible by the network.

One of the fundamental reasons Java is a useful tool for networked software environments is that Java's architecture enables the network mobility of software. In fact, it was primarily this aspect of Java technology that was considered by many in the software industry to represent a paradigm shift. This chapter examines the nature of this new paradigm of network-mobile software, and how Java's architecture makes it possible.

Why Network Mobility?

Prior to the advent of the personal computer, the dominant computing model was the large mainframe computer serving multiple users. By time-sharing, a mainframe computer divided its attention among several users, who logged onto the mainframe at dumb terminals. Software applications were stored on disks attached to the mainframe computer, allowing multiple users to share the same applications while they shared the same CPU. A drawback of this model was that if one user ran a CPU-intensive job, all other users would experience degraded performance.

The appearance of the microprocessor led to the proliferation of the personal computer. This change in the hardware status quo changed the software paradigm as well. Rather than sharing software applications stored at a mainframe computer, individual users had individual copies of software applications stored at each personal computer. Because each user ran software on a dedicated CPU, this new model of computing addressed the difficulties of dividing CPU-time among many users attempting to share one mainframe CPU.

Initially, personal computers operated as unconnected islands of computing. The dominant software model was of isolated executables running on isolated personal computers. But soon, personal computers began to be connected to networks. Because a personal computer gave its user undivided attention, it addressed the CPU-time sharing difficulties of mainframes. But unless personal computers were connected to a network, they couldn't replicate the mainframe's ability to let multiple users view and manipulate a central repository of data.

As personal computers connected to networks became the norm, another software model began to increase in importance: client/server. The client/server model divided work between two processes running on two different computers: a client process ran on the end-user's personal computer, and a server process ran on some other computer hooked to the same network. The client and server processes communicated with one another by sending data back and forth across the network. The server process often simply accepted data query commands from clients across the network, retrieved the requested data from a central database, and sent the retrieved data back across the network to the client. Upon receiving the data, the client processed, displayed, and allowed the user to manipulate the data. This model allowed users of personal computers to view and manipulate data stored at a central repository, while not forcing them to share a central CPU for all of the processing of that data. Users did share the CPU running the server process, but to the extent that data processing was performed by the clients, the burden on the central CPU hosting the server process was lessened.

The client/server architecture was soon extended to include more than two processes. The original client/server model began to be called 2-tier client/server, to indicate two processes: one client and one server. More elaborate architectures were called 3-tier, to indicate three processes, 4-tier, to indicate four processes, or N-tier, to indicate people were getting tired of counting processes. Eventually, as more processes became involved, the distinction between client and server blurred, and people just started using the term *distributed processing* to encompass all of these schemes.

The distributed processing model leveraged the network and the proliferation of processors by dividing processing work loads among many processors while allowing those processors to share data. Although this model had many advantages over the mainframe model, there was one notable disadvantage: distributed processing systems were more difficult to administer than mainframe systems. On mainframe systems, software applications were stored on a disk attached to the mainframe. Even though an application could serve many users, it only needed to be installed and maintained in one place. When an application was upgraded, all users got the new version the next time they logged on and started the application. By contrast, the software executables for different components of a distributed processing system were usually stored on many different disks. In a client/server architecture, for example, each computer that hosted a client process usually had its own copy of the client software stored on its local disk. As a result, a system administrator had to install and maintain the various components of a distributed software system in many different places. When a software component was upgraded, the system administrator had to physically upgrade each copy of the component on each computer that hosted it. As a result, system administration was more difficult for the distributed processing model than for the mainframe model.

The arrival of Java, with an architecture that enabled the network-mobility of software, heralded yet another model for computing. Building on the prevailing distributed processing model, the new model added the automatic delivery of software across networks to computers that ran the software. This addressed the difficulties involved in system administration of distributed processing systems. For example, in a client/server system, client software could be stored at one central computer attached to the network. Whenever an end-user needed to use the client software, the binary executable would be sent from the central computer across the network to the end-user's computer, where the software would run.

So network-mobility of software represented another step in the evolution of the computing model. In particular, it addressed the difficulty of administering a distributed processing system. It simplified the job of distributing any software that was to be used on more than one CPU. It allowed data to be delivered together with the software that knows how to manipulate or display the data. Because code was sent along with data, end-users would always have the most up-to-date version of the code. Thus, because of network-mobility, software can be administered from a central computer, reminiscent of the

mainframe model, but processing can still be distributed among many CPUs.

A New Software Paradigm

The shift away from the mainframe model towards the distributed processing model was a consequence of the personal computer revolution, which was made possible by the rapidly increasing capabilities and decreasing costs of processors. Similarly, lurking underneath the latest software paradigm shift towards distributed processing with network-mobile code is another hardware trend--the increasing capabilities and decreasing costs of network bandwidth. As bandwidth, the amount of information that can be carried by a network, increases, it becomes practical to send new kinds of information across a network; and with each new kind of information a network carries, the network takes on a new character. Thus, as bandwidth grows, simple text sent across a network can become enhanced with graphics, and the network begins to take on an appearance reminiscent of newspapers or magazines. Once bandwidth expands enough to support live streams of audio data, the network begins to act like a radio, a CD-player, or a telephone. With still more bandwidth, video becomes possible, resulting in a network that competes with TV and VCRs for the attention of couch potatoes. But there is still one other kind of bandwidth-hungry content that becomes increasingly practical as bandwidth improves: computer software. Because networks by definition interconnect processors, one processor can, given enough bandwidth, send code across a network for another processor to execute. Once networks begin to move software as well as data, the network begins to look like a computer in its own right.

As software begins to travel across networks, not only does the network begin to take on a new character, but so does the software itself. Network-mobile code makes it easier to ensure that an end-user has the necessary software to view or manipulate some data sent across the network, because the software can be sent along with the data. In the old model, software executables from a local disk were invoked to view data that came across the network, thus the software application was usually a distinct entity, easily discernible from the data. In the new model, because software and data are both sent across the network, the distinction between software and data is not as stark--software and data blur together to become "content."

As the nature of software evolves, the end-user's relationship to software evolves as well. Prior to network-mobility, an end-user had to think in terms of software applications and version numbers. Software was generally distributed on media such as tapes, floppy disks, or CD-ROMs. To use an application, an end-user had to get the installation media, physically insert them into a drive or reader attached to the computer, and run an installation program that copied files from the installation media to the computer's hard disk. Moreover, the end-user often did this process multiple times for each application, because software applications were routinely replaced by new versions that fixed old bugs and added new features (and usually added new bugs too). When a new version was released, end-users had to decide whether or not to upgrade. If an end-user decided to upgrade, the installation process had to be repeated. Thus, end-users had to think in terms of software applications and version numbers, and take deliberate action to keep their software applications up-to-date.

In the new model, end-users think less in terms of software applications with discrete versions, and more in terms of self-evolving "content services." Whereas installing a traditional software application or an upgrade was a deliberate act on the part of the end-user, network-mobility of software enables installation and upgrading that is more automatic. Network-delivered software need not have discrete version numbers that are known to the end-user. The end-user need not decide whether to upgrade, and need not take any special action to upgrade. Network-delivered software can just evolve of its own accord. Instead of buying discrete versions of a software application, end-users can subscribe to a content service--software that is delivered across a network along with relevant data--and watch as both

the software and data evolve automatically over time.

Once you move away from delivering software in discrete versions towards delivering software as self-evolving streams of interactive content, your end-user loses some control. In the old model, if a new version appeared that had serious bugs, an end-user could simply opt not to upgrade. But in the new model, an end-user can't necessarily wait until the bugs are worked out of a new version before upgrading to the new version, because the end-user may have no control over the upgrading process.

For certain kinds of products, especially those that are large and full-featured, end-users may prefer to retain control over whether and when to upgrade. Consequently, in some situations software vendors may publish discrete versions of a content service over the network. At the very least, a vendor can publish two branches of a service: a beta branch and a released branch. End-users that want to stay on the bleeding edge can subscribe to the beta service, and the rest can subscribe to the released service that, although it may not have all the newest features, is likely more robust.

Yet for many content services, especially simple ones, most end-users won't want to have to worry about versions, because worrying about versions makes software harder to use. The end-user has to have knowledge about the differences between versions, make decisions about when and if to upgrade, and take deliberate action to cause an upgrade. Content services that are not chopped up into discrete versions are easier to use, because they evolve automatically. Such a content service, because the end-user doesn't have to maintain it but can just simply use it, takes on the feel of a "software appliance."

Many self-evolving content services will share two fundamental characteristics with common household appliances: a focused functionality and a simple user-interface. Consider the toaster. A toaster's functionality is focused exclusively on the job of preparing toast, and it has a simple user-interface. When you walk up to a toaster, you don't expect to have to read a manual. You expect to put the bread in at the top and push down a handle until it clicks. You expect to be able to peer in and see orange wires glowing, and after a moment, to hear that satisfying pop and see your bread transformed into toast. If the result is too light or too dark, you expect to be able to slide a knob to indicate to the toaster that the next time, you want your toast a bit darker or lighter. That's the extent of the functionality and user-interface of a toaster. Likewise, the functionality of many content services will be as focused and the user-interface will be as simple. If you want to order a movie through the network, for example, you don't want to worry whether you have the correct version of movie-ordering software. You don't want to have to install it. You just want to switch on the movie-ordering content service, and through a simple user-interface, order your movie. Then you can sit back and enjoy your network-delivered movie as you eat your toast.

A good example of a content service is a World Wide Web page. If you look at an HTML file, it looks like a source file for some kind of program. But if you see the browser as the program, the HTML file looks more like data. Thus, the distinction between code and data is blurred. Also, people who browse the World Wide Web expect web pages to evolve over time, without any deliberate action on their part. They don't expect to see discrete version numbers for web pages. They don't expect to have to do anything to upgrade to the latest version of a page besides simply revisiting the page in their browser.

In the coming years, many of today's media may to some extent be assimilated by the network and transformed into content services. (As with the Borg from Star Trek, resistance is futile.) Broadcast radio, broadcast and cable TV, telephones, answering machines, faxes, video rental stores, newspapers, magazines, books, computer software—all of these will be affected by the proliferation of networks. But just as TV didn't supplant radio entirely, content services will not entirely supplant existing media. Instead, content services will likely take over some aspects of existing media, leaving the existing media to adjust accordingly, and create some new forms that didn't previously exist.

In the computer software domain, the content service model will not completely replace the old models either. Instead, it will likely take over certain aspects of the old models that fit better in the new model, add new forms that didn't exist before, and leave the old models to adjust their focus slightly in light of the newcomer.

This book is an example of how the network can affect existing media. The book was not entirely replaced by a content service counterpart, but instead of including resource pointers (sources where you can find further information on topics presented in the book) as part of the book, they were placed on a web page. Because resource pointers change so often, it made sense to let the network assimilate that part of the book. Thus, the resource pointers portion of the book has become a content service.

The crux of the new software paradigm, therefore, is that software begins to act more like appliances. End-users no longer have to worry about installation, version numbers, or upgrading. As code is sent along with data across the network, software delivery and updating become automatic. In this way, simply by making code mobile, Java unleashes a whole new way to think about software development, delivery, and use.

Java's Architectural Support for Network-Mobility

Java's architectural support for network-mobility begins with its support for platform independence and security. Although they are not strictly required for network-mobility, platform independence and security help make network-mobility practical. Platform independence makes it easier to deliver a program across the network because you don't have to maintain a separate version of the program for different platforms, and you don't have to figure out how to get the right version to each computer. One version of a program can serve all computers. Java's security features help promote network-mobility because they give end-users confidence to download class files from untrusted sources. In practice, therefore, Java's architectural support for platform independence and security facilitate the network-mobility of its class files.

Beyond platform independence and security, Java's architectural support for network-mobility is focused on managing the time it takes to move software across a network. If you store a program on a server and download it across a network when you need it, it will likely take longer for your program to start than if you had started the same program from a local disk. Thus, one of the primary issues of network-mobile software is the time it takes to send a program across a network. Java's architecture addresses this issue by rejecting the traditional monolithic binary executable in favor of small binary pieces: Java class files. Class files can travel across networks independently, and because Java programs are dynamically linked and dynamically extensible, an end-user needn't wait until all of a program's class files are downloaded before the program starts. The program starts when the first class file arrives. Class files themselves are designed to be compact, so that they fly more quickly across networks. Therefore, the main way Java's architecture facilitates network-mobility directly is by breaking up the monolithic binary executable into compact class files, which can be loaded as needed.

The execution of a Java application begins at a `main()` method of some class, and other classes are loaded and dynamically linked as they are needed by the application. If a class is never actually used during one session, that class won't ever be loaded during that session. For example, if you are using a word processor that has a spelling checker, but during one session you never invoke the spelling checker, the class files for the spelling checker will not be loaded during that session.

In addition to dynamic linking, Java's architecture also enables dynamic extension. Dynamic extension is another way the loading of class files (and the downloading of them across a network) can be delayed

in a Java application. Using class loader objects, a Java program can load extra classes at run-time, which then become a part of the running program. Therefore, dynamic linking and dynamic extension give a Java programmer some flexibility in designing when class files for a program are loaded, and as a result, how much time an end-user must spend waiting for class files to come across the network.

Besides dynamic linking and dynamic extension, another way Java's architecture directly supports network mobility is through the class file format itself. To reduce the time it takes to send them across networks, class files are designed to be compact. In particular, the bytecode streams they contain are designed to be compact. They are called "bytecodes" because each instruction occupies only one byte. With only two exceptions, all opcodes and their ensuing operands are byte aligned to make the bytecode streams smaller. The two exceptions are opcodes that may have one to three bytes of padding after the opcode and before the start of the operands, so that the operands are aligned on word boundaries. Other than the two opcodes that may have a small amount of padding, all data in a class file is byte aligned.

One of the implications of the compactness goal for class files is that Java compilers are not likely to do any local optimization. Because of binary compatibility rules, Java compilers can't perform global optimizations such as inlining the invocation of another class's method. (Inlining means replacing the method invocation with the code performed by the method, which saves the time it takes to invoke and return from the method as the code executes.) Binary compatibility requires that a method's implementation can be changed without breaking compatibility with pre-existing class files that depend on the method. Inlining could be performed in some circumstances on methods within a single class, but in general that kind of optimization is not done by Java compilers, partly because it goes against the grain of class file compactness. Optimizations are often a tradeoff between execution speed and code size. Therefore, Java compilers generally leave optimization up to the Java Virtual Machine, which can optimize code as it loads classes for interpreting or just-in-time compiling.

Beyond the architectural features of dynamic linking, dynamic extension and class file compactness, there are some strategies that, although they are really not necessarily part of the architecture, help manage the time it takes to move class files across a network. Because HTTP protocols require that each class file of Java applet be requested individually, it turns out that often a large percentage of applet download time is due not to the actual transmission of class files across the network, but to the network handshaking of each class file request. The overhead for a file request is multiplied by the number of class files being requested. To address this problem, Java 1.1 included support for JAR (Java ARchive) files. JAR files enable many class files to be sent in one network transaction, which greatly reduces the overhead time required to move class files across a network compared with sending one class file at a time. Moreover, the data inside a JAR file can be compressed, which results in an even shorter download time. So sometimes it pays to send software across a network in one big chunk. If a set of class files is definitely needed by a program before that program can start, those class files can be more speedily transmitted if they are sent together in a JAR file.

One other strategy to minimize an end-user's wait time is to not download class files on-demand. Through various techniques, such as the subscription model used by Marimba Castanet, class files can be downloaded before they are needed, resulting in a program that starts up faster. You can obtain more information about Castanet's model from the resource page for this chapter.

Therefore, other than platform independence and security, which help make network-mobility practical, the main focus of Java's architectural support for network-mobility is managing the time it takes to send class files across a network. Dynamic linking and dynamic extension allow Java programs to be designed in small functional units that are downloaded as needed by the end-user. Class file compactness helps reduce the time it takes to move a Java program across the network. The JAR file enables compression and the sending of multiple class files across the network in a single network file-

transfer transaction.

The Applet: An Example of Network-Mobile Java

Java is a network-oriented technology that first appeared at a time when the network was looking increasingly like the next revolution in computing. The reason Java was adopted so rapidly and so widely, however, was not simply because it was a timely technology, but because it had timely marketing. Java was not the only network-oriented technology being developed in the early to mid 1990s. And although it was a good technology, it wasn't the necessarily the best technology--but it probably had the best marketing. Java was the one technology to hit a slim market window in early 1995, resulting in such a strong response that many companies developing similar technologies canceled their projects, including Microsoft, which canceled a project code-named Blackbird. Companies that carried on with their technologies, such as AT&T did with a network-oriented technology named Inferno, saw Java steal much of their potential thunder.

There were several important factors in how Java was initially unleashed on the world that contributed to its successful marketing. First, it had a cool name--one that could be appreciated by programmers and non-programmers alike. Second, it was, for all practical purposes, free--always a strong selling point among prospective buyers. But the most critical factor contributing to the successful marketing of Java, however, was that Sun's engineers hooked Java technology to the World Wide Web at the precise moment Netscape was looking to transform their web browser from a graphical hypertext viewer to a full-fledged computing platform. As the World Wide Web swept through the software industry (and the global consciousness) like an ever-increasing tidal wave, Java rode with it. Therefore, in a sense Java became a success because Java "surfing the web." It caught the wave at just the right time and kept riding it as one by one, its potential competitors dropped uneventfully into the cold, dark sea. The way the engineers at Sun hooked Java technology to the World Wide Web--and therefore, the key way Java was successfully marketed--was by creating a special flavor of Java program that ran inside a web browser: the Java applet.

The Java applet showed off all of Java's network-oriented features: platform independence, network-mobility, and security. Platform independence was one of the main tenets of the World Wide Web, and Java applets fit right in. Java applets can run on any platform so long as there is a Java-capable browser for that platform. Java applets also demonstrated Java's security capabilities, because they run inside a strict sandbox. But most significantly, Java applets demonstrated the promise of network-mobility. As shown in Figure 4-1, Java applets can be maintained on one server, from which they can travel across a network to many different kinds of computers. To update an applet, you only need to update the server. Users will automatically get the updated version the next time they use the applet. Thus, maintenance is localized, but processing is distributed.




Figure 4-1. The Java Applet and the network.

Java-capable browsers fire off a Java application that hosts the applets the browser displays. To display a web page, a web browser requests an HTML file from an HTTP server. If the HTML file includes an applet, the browser will see an HTML tag such as this:

begin

```
<applet CODE="HeapOfFish.class"
```

```
CODEBASE="gcsupport/classes"
```

```
WIDTH=525
```

```
HEIGHT=360</applet>
```

end

This "applet" tag provides enough information to enable the browser to display the applet. The `CODE` attribute indicates the name of the applet's starting class file, in this case: `HeapOfFish.class`. The `CODEBASE` attribute gives the location of the applet's class files relative to the base URL of the web page. The `WIDTH` and `HEIGHT` attributes indicate the size in pixels of the applet's panel, the visible portion of the applet that is displayed as part of the web page.

When a browser encounters a web page that includes an applet tag, it passes information from the tag to the running Java application. The Java application either creates a new class loader object, or re-uses an existing one, to download the starting class file for the applet. It then initializes the applet, by invoking the `init()` method of the applet's starting class. The other class files for the applet are

downloaded on an as needed basis, by the normal process of dynamic linking. For example, when a new class is first used by the applet's starting class, the symbolic reference to the new class must be resolved. During resolution, if the class has not already been loaded, the Java Virtual Machine will ask the same class loader object that loaded the applet's starting class to load the new class. If the class loader object is unable to load the class from the local trusted repository through the primordial class loader, the class loader object will attempt to download the class file across the network from the same location it retrieved the applet's starting class. Once initialization of the applet is complete, the applet appears as part of the web page inside the browser.

Network-mobile Java can take other forms besides Java applets, but the framework that supports other network-mobile forms will likely look very similar to the framework for applets. Like applets, other forms of network-mobile Java code will in general be implemented as special flavors of Java program running in the context of a host Java application. Network-mobile class files will generally be loaded by class loader objects for two reasons. The first reason is simply that a class loader object can download class files across a network in custom ways that a primordial class loader can't. But the second reason involves security. Because network-mobile class files are not always known to be trustworthy, the separate name-spaces provided by class loader objects are needed to protect a malicious applet from interfering with applets loaded from other sources. Also because network-mobile class files can't always be trusted, there will generally be a security manager establishing a sandbox for the network-mobile code.

The Resources Page

For links to information about other examples of network-mobile Java, such as Marimba Castanet Channels, Jeeves Servlets, and Aglets (Java-based autonomous software agents), see the resource page for this chapter: <http://www.artima.com/insidejvm/mobility.html>

McGraw-Hill

A Division of The McGraw-Hill Companies



Chapter Five

The Java Virtual Machine

The previous four chapters of this book gave a broad overview of Java's architecture. They showed how the Java Virtual Machine fits into the overall architecture relative to other components such as the language and API. The remainder of this book will focus more narrowly on the Java Virtual Machine. This chapter gives an overview of the Java Virtual Machine's internal architecture.

The Java Virtual Machine is called "virtual" because it is an abstract computer defined by a specification. To run a Java program, you need a concrete implementation of the abstract specification. This chapter describes primarily the abstract specification of the Java Virtual Machine. To illustrate the abstract definition of certain features, however, this chapter also discusses various ways in which those features could be implemented.

What is a Java Virtual Machine?

To understand the Java Virtual Machine you must first be aware that you may be talking about any of three different things when you say "Java Virtual Machine." You may be speaking of:

- the abstract specification,
- a concrete implementation, or
- a runtime instance.

The abstract specification is a concept, described in detail in the book: *The Java Virtual Machine Specification*, by Tim Lindholm and Frank Yellin. Concrete implementations, which exist on many platforms and come from many vendors, are either all software or a combination of hardware and software. A runtime instance hosts a single running Java application.

Each Java application runs inside a runtime instance of some concrete implementation of the abstract specification of the Java Virtual Machine. In this book, the term "Java Virtual Machine" is used in all three of these senses. Where the intended sense is not clear from the context, one of the terms "specification," "implementation," or "instance" is added to the term "Java Virtual Machine".

The Lifetime of a Java Virtual Machine

A runtime instance of the Java Virtual Machine has a clear mission in life: to run one Java application. When a Java application starts, a runtime instance is born. When the application completes, the instance dies. If you start three Java applications at the same time, on the same computer, using the same

concrete implementation, you'll get three Java Virtual Machine instances. Each Java application runs inside its own Java Virtual Machine.

A Java Virtual Machine instance starts running its solitary application by invoking the `main()` method of some initial class. The `main()` method must be public, static, return `void`, and accept one parameter: a `String` array. Any class with such a `main()` method can be used as the starting point for a Java application.

For example, consider an application that prints out its command line arguments:

begin

```
// On CD-ROM in file jvm/ex1/Echo.java
class Echo {

    public static void main(String[] args) {
        int len = args.length;
        for (int i = 0; i < len; ++i) {
            System.out.print(args[i] + " ");
        }
        System.out.println();
    }
}
```

end

You must in some implementation-dependent way give a Java Virtual Machine the name of the initial class that has the `main()` method that will start the entire application. One real world example of a Java Virtual Machine implementation is the `java` program from Sun's JDK. If you wanted to run the `Echo` application using Sun's `java` on Window95, for example, you would type in a command such as:

insert

```
java Echo Greetings, Planet.
```

end

The first word in the command, "`java`," indicates that the Java Virtual Machine from Sun's JDK should be run by the operating system. The second word, "`Echo`," is the name of the initial class. `Echo` must have a public static method named `main()` that returns `void` and takes a `String` array as its only parameter. The subsequent words, "`Greetings, Planet.`," are the command line arguments for the application. These are passed to the `main()` method in the `String` array in the order in which they appear on the command line. So, for the above example, the contents of the `String` array passed

to main in Echo are:

arg[0] is "Greetings, "

arg[1] is "Planet."

The `main()` method of an application's initial class serves as the starting point for that application's initial thread. The initial thread can in turn fire off other threads.

Inside the Java Virtual Machine, threads come in two flavors: *daemon* and *non-daemon*. A daemon thread is ordinarily a thread used by the virtual machine itself, such as a thread that performs garbage collection. The application, however, can mark any threads it creates as daemon threads. The initial thread of an application--the one that begins at `main()`--is a non-daemon thread.

A Java application continues to execute (the virtual machine instance continues to live) as long as any non-daemon threads are still running. When all non-daemon threads of a Java application terminate, the virtual machine instance will exit. If permitted by the security manager, the application can also cause its own demise by invoking the `exit()` method of class `Runtime` or `System`.

In the Echo application above, the `main()` method doesn't invoke any other threads. After it prints out the command line arguments, `main()` returns. This terminates the application's only non-daemon thread, which causes the virtual machine instance to exit.

The Architecture of the Java Virtual Machine

In the Java Virtual Machine specification, the behavior of a virtual machine instance is described in terms of subsystems, memory areas, data types, and instructions. These components describe an abstract inner architecture for the abstract Java Virtual Machine. The purpose of these components is not so much to dictate an inner architecture for implementations. It is more to provide a way to strictly define the external behavior of implementations. The specification defines the required behavior of any Java Virtual Machine implementation in terms of these abstract components and their interactions.

Figure 5-1 shows a block diagram of the Java Virtual Machine that includes the major subsystems and memory areas described in the specification. As mentioned in previous chapters, each Java Virtual Machine has a *class loader subsystem*: a mechanism for loading types (classes and interfaces) given fully qualified names. Each Java Virtual Machine also has an *execution engine*: a mechanism responsible for executing the instructions contained in the methods of loaded classes.




Figure 5-1. The internal architecture of the Java Virtual Machine.

When a Java Virtual Machine runs a program, it needs memory to store many things, including bytecodes and other information it extracts from loaded class files, objects the program instantiates, parameters to methods, return values, local variables, and intermediate results of computations. The Java Virtual Machine organizes the memory it needs to execute a program into several *runtime data areas*.

Although the same runtime data areas exist in some form in every Java Virtual Machine implementation, their specification is quite abstract. Many decisions about the structural details of the runtime data areas are left to the designers of individual implementations.

Different implementations of the virtual machine can have very different memory constraints. Some implementations may have a lot of memory in which to work, others may have very little. Some implementations may be able to take advantage of virtual memory, others may not. The abstract nature of the specification of the runtime data areas helps make it easier to implement the Java Virtual Machine on a wide variety of computers and devices.

Some runtime data areas are shared among all of an application's threads and others are unique to individual threads. Each instance of the Java Virtual Machine has one *method area* and one *heap*. These areas are shared by all threads running inside the virtual machine. When the virtual machine loads a class file, it parses information about a type from the binary data contained in the class file. It places this type information into the method area. As the program runs, the virtual machine places all objects the program instantiates onto the heap. See Figure 5-2 for a graphical depiction of these memory areas.




Figure 5-2. Runtime data areas shared by all threads.

As each new thread comes into existence, it gets its own *pc register* (program counter) and *Java stack*. If the thread is executing a Java method (not a native method), the value of the pc register indicates the next instruction to execute. A thread's Java stack stores the state of Java (not native) method invocations for the thread. The state of a Java method invocation includes its local variables, the parameters with which it was invoked, its return value (if any), and intermediate calculations. The state of native method invocations is stored in an implementation-dependent way in *native method stacks*, as well as possibly in registers or other implementation-dependent memory areas.

The Java stack is composed of *stack frames* (or *frames*). A stack frame contains the state of one Java method invocation. When a thread invokes a method, the Java Virtual Machine pushes a new frame onto that thread's Java stack. When the method completes, the virtual machine pops and discards the frame for that method.

The Java Virtual Machine has no registers to hold intermediate data values. The instruction set uses the Java stack for storage of intermediate data values. This approach was taken by Java's designers to keep the Java Virtual Machine's instruction set compact and to facilitate implementation on architectures with few or irregular general purpose registers.

See Figure 5-3 for a graphical depiction of the memory areas the Java Virtual Machine creates for each thread. These areas are private to the owning thread. No thread can access the pc register or Java stack of another thread.




Figure 5-3. Runtime data areas exclusive to each thread.

Figure 5-3 shows a snapshot of a virtual machine instance in which three threads are executing. At the instant of the snapshot, threads one and two are executing Java methods. Thread three is executing a native method.

In Figure 5-3, as in all graphical depictions of the Java stack in this book, the stacks are shown growing downwards. The "top" of each stack is shown at the bottom of the figure. Stack frames for currently executing methods are shown in a lighter shade. For threads that are currently executing a Java method, the pc register indicates the next instruction to execute. In Figure 5-3, such pc registers (the ones for threads one and two) are shown in a lighter shade. Because thread three is currently executing a native method, the contents of its pc register--the one shown in dark gray--is undefined.

Data Types

The Java Virtual Machine computes by performing operations on certain types of data. Both the data types and operations are strictly defined by the Java Virtual Machine specification. The data types can be divided into a set of *primitive types* and a *reference type*. Variables of the primitive types hold *primitive values*, and variables of the reference type hold *reference values*. Reference values refer to objects, but are not objects themselves. Primitive values, by contrast, do not refer to anything. They are the actual data themselves. You can see a graphical depiction of the Java Virtual Machine's families of data types in Figure 5-4.




Figure 5-4. Data types of the Java Virtual Machine.

All the primitive types of the Java programming language, except `boolean`, are primitive types of the Java Virtual Machine. When a compiler translates Java source code into bytecodes, it uses `ints` or `bytes` to represent `booleans`. In the Java Virtual Machine, `false` is represented by integer zero and `true` by any non-zero integer. Operations involving `boolean` values use `ints`. Arrays of `boolean` are accessed as arrays of `byte`, though they may be represented on the heap as arrays of `byte` or as bit fields.

The primitive types of the Java programming language other than `boolean` form the *numeric types* of the Java Virtual Machine. The numeric types are divided between the *integral types*: `byte`, `short`, `int`, `long`, and `char`, and the *floating-point types*: `float` and `double`. As with the Java programming language, the primitive types of the Java Virtual Machine have the same range everywhere. A `long` in the Java Virtual Machine always acts like a 64-bit signed twos complement number, independent of the underlying host platform.

The Java Virtual Machine works with one other primitive type that is unavailable to the Java programmer: the `returnValue` type. This primitive type is used to implement `finally` clauses of Java programs. The use of the `returnValue` type is described in detail in Chapter 18, "Finally Clauses."

The reference type of the Java Virtual Machine is cleverly named `reference`. Values of type `reference` come in three flavors: the *class type*, the *interface type*, and the *array type*. All three types have values that are references to dynamically created objects. The class type's values are references to class instances. The array type's values are references to arrays, which are full-fledged objects in the Java Virtual Machine. The interface type's values are references to class instances that implement an interface. One other reference value is the `null` value, which indicates the `reference` variable

doesn't refer to any object.

The Java Virtual Machine specification defines the range of values for each of the data types, but does not define their sizes. The number of bits used to store each data type value is a decision of the designers of individual implementations. The ranges of the Java Virtual Machine's data types are shown in Table 5-1. More information on the floating point ranges is given in Chapter 14, "Floating Point Arithmetic."

Table 5-1. Ranges of the Java Virtual Machine's data types

Type	Range
byte	8-bit signed two's complement integer (-2^7 to $2^7 - 1$, inclusive)
short	16-bit signed two's complement integer (-2^{15} to $2^{15} - 1$, inclusive)
int	32-bit signed two's complement integer (-2^{31} to $2^{31} - 1$, inclusive)
long	64-bit signed two's complement integer (-2^{63} to $2^{63} - 1$, inclusive)
char	16-bit unsigned Unicode character (0 to $2^{16} - 1$, inclusive)
float	32-bit IEEE 754 single-precision float
double	64-bit IEEE 754 double-precision float
returnValue	address of an opcode within the same method
reference	reference to an object on the heap, or null

Word Size

The basic unit of size for data values in the Java Virtual Machine is the *word*--a fixed size chosen by the designer of each Java Virtual Machine implementation. The word size must be large enough to hold a value of type `byte`, `short`, `int`, `char`, `float`, `returnValue`, or `reference`. Two words must be large enough to hold a value of type `long` or `double`. An implementation designer must therefore choose a word size that is at least 32 bits, but otherwise can pick whatever word size will yield the most efficient implementation. The word size is often chosen to be the size of a native pointer on the host platform.

The specification of many of the Java Virtual Machine's runtime data areas are based upon this abstract concept of a word. For example, two sections of a Java stack frame--the local variables and operand stack--are defined in terms of words. These areas can contain values of any of the virtual machine's data types. When placed into the local variables or operand stack, a value occupies either one or two words.

As they run, Java programs cannot determine the word size of their host virtual machine implementation. The word size does not affect the behavior of a program. It is only an internal attribute of a virtual machine implementation.

The Class Loader Subsystem

The part of a Java Virtual Machine implementation that takes care of finding and loading types is the

class loader subsystem. Chapter 1, "Introduction to Java's Architecture," gives an overview of this subsystem. Chapter 3, "Security," shows how the subsystem fits into Java's security model. This chapter describes the class loader subsystem in more detail and show how it relates to the other components of the virtual machine's internal architecture.

As mentioned in Chapter 1, the Java Virtual Machine contains two kinds of class loaders: a *primordial class loader* and *class loader objects*. The primordial class loader is a part of the virtual machine implementation, and class loader objects are part of the running Java application. Classes loaded by different class loaders are placed into separate *name spaces* inside the Java Virtual Machine.

The class loader subsystem involves many other parts of the Java Virtual Machine and several classes from the `java.lang` library. For example, class loader objects are regular Java objects whose class descends from `java.lang.ClassLoader`. The methods of class `ClassLoader` allow Java applications to access the virtual machine's class loading machinery. Also, for every type a Java Virtual Machine loads, it creates an instance of class `java.lang.Class` to represent that type. Like all objects, class loader objects and instances of class `Class` reside on the heap. Data for loaded types resides in the method area.

Loading, Linking and Initialization

The class loader subsystem is responsible for more than just locating and importing the binary data for classes. It must also verify the correctness of imported classes, allocate and initialize memory for class variables, and assist in the resolution of symbolic references. These activities are performed in a strict order:

1. Loading: finding and importing the binary data for a type
2. Linking: performing verification, preparation, and (optionally) resolution
 - a. Verification: ensuring the correctness of the imported type
 - b. Preparation: allocating memory for class variables and initializing the memory to default values
 - c. Resolution: transforming symbolic references from the type into direct references.
- Initialization: invoking Java code that initializes class variables to their proper starting values.

The details of these processes are given Chapter 7, "The Lifetime of a Class."

The Primordial Class Loader

Java Virtual Machine implementations must be able to recognize and load classes and interfaces stored in binary files that conform to the Java class file format. An implementation is free to recognize other binary forms besides class files, but it must recognize class files. One example of an alternative binary format recognized by a particular Java Virtual Machine implementation is the CAB file. This file format, which is an archive of class files and other data files, is defined by Microsoft and recognized by their implementation of the Java Virtual Machine.

Every Java Virtual Machine implementation has a primordial class loader, which knows how to load

trusted classes, including the classes of the Java API. The Java Virtual Machine specification doesn't define how the primordial loader should locate classes. That is another decision the specification leaves to implementation designers.

Given a fully qualified type name, the primordial class loader must *in some way* attempt to locate a file with the type's simple name plus ".class". One common approach is demonstrated by the Java Virtual Machine implementation in Sun's 1.1 JDK on Windows95. This implementation searches a user-defined directory path stored in an environment variable named CLASSPATH. The primordial loader looks in each directory, in the order the directories appear in the CLASSPATH, until it finds a file with the appropriate name: the type's simple name plus ".class". Unless the type is part of the unnamed package, the primordial loader expects the file to be in a subdirectory of one of the directories in the CLASSPATH. The path name of the subdirectory is built from the package name of the type. For example, if the primordial class loader is searching for class `java.lang.Object`, it will look for `Object.class` in the `java\lang` subdirectory of each CLASSPATH directory.

Class Loader Objects

Although class loader objects themselves are part of the Java application, three of the methods in class `ClassLoader` are gateways into the Java Virtual Machine:

```
begin
// Three of the methods declared in class java.lang.ClassLoader:
protected final Class defineClass(byte data[], int offset,
    int length);
protected final Class findSystemClass(String name);
protected final void resolveClass(Class c);
end
```

Any Java Virtual Machine implementation must take care to connect these methods of class `ClassLoader` to the internal class loader subsystem.

The `defineClass()` method accepts a byte array, `data[]`, as input. Starting at position `offset` in the array and continuing for `length` bytes, class `ClassLoader` expects binary data conforming to the Java class file format--binary data that represents a new type for the running application. Every Java Virtual Machine implementation must make sure the `defineClass()` method of class `ClassLoader` can cause a new type to be imported into the method area.

The `findSystemClass()` method accepts a `String` representing a fully qualified name of a type. When a class loader object invokes this method, it is requesting that the virtual machine attempt to load the named type via its primordial class loader. If the primordial class loader has already loaded or successfully loads the type, it returns a reference to the `Class` object representing the type. If it can't locate the binary data for the type, it throws `ClassNotFoundException`. Every Java Virtual Machine implementation must make sure the `findSystemClass()` method can invoke the primordial class loader in this way.

The `resolveClass()` method accepts a reference to a `Class` instance. This method causes the type represented by the `Class` instance to be linked and initialized (if it hasn't already been linked and initialized). The `defineClass()` method, described above, only takes care of loading. (See the above section, "Loading, Linking, and Initialization" for definitions of these terms.) When `defineClass()` returns a `Class` instance, the binary file for the type has definitely been located and imported into the method area, but not necessarily linked and initialized. Java Virtual Machine implementations make sure the `resolveClass()` method of class `ClassLoader` can cause the class loader subsystem to perform linking and initialization.

The details of how a Java Virtual Machine performs class loading, linking, and initialization, with class loader objects is given in Chapter 8, "The Linking Model."

Name Spaces

As mentioned in Chapter 3, "Security," each class loader maintains its own name space populated by the types it has loaded. Because each class loader has its own name space, a single Java application can load multiple types with the same fully qualified name. A type's fully qualified name, therefore, is not always enough to uniquely identify it inside a Java Virtual Machine instance. If multiple types of that same name have been loaded into different name spaces, the identity of the class loader that loaded the type (the identity of the name space it is in) will also be needed to uniquely identify that type.

Name spaces arise inside a Java Virtual Machine instance as a result of the process of resolution. As part of the data for each loaded type, the Java Virtual Machine keeps track of the class loader that imported the type. When the virtual machine needs to resolve a symbolic reference from one class to another, it requests the referenced class from the same class loader that loaded the referencing class. This process is described in detail in Chapter 8, "The Linking Model."

The Method Area

Inside a Java Virtual Machine instance, information about loaded types is stored in a logical area of memory called the method area. When the Java Virtual Machine loads a type, it uses a class loader to locate the appropriate class file. The class loader reads in the class file--a linear stream of binary data--and passes it to the virtual machine. The virtual machine extracts information about the type from the binary data and stores the information in the method area. Memory for class (static) variables declared in the class is also taken from the method area.

The manner in which a Java Virtual Machine implementation represents type information internally is a decision of the implementation designer. For example, multi-byte quantities in class files are stored in big-endian (most significant byte first) order. When the data is imported into the method area, however, a virtual machine can store the data in any manner. If an implementation sits on top of a little-endian processor, the designers may decide to store multi-byte values in the method area in little-endian order.

The virtual machine will search through and use the type information stored in the method area as it executes the application it is hosting. Designers must attempt to devise data structures that will facilitate speedy execution of the Java application, but must also think of compactness. If designing an implementation that will operate under low memory constraints, designers may decide to trade off some execution speed in favor of compactness. If designing an implementation that will run on a virtual memory system, on the other hand, designers may decide to store redundant information in the method area to facilitate execution speed. (If the underlying host doesn't offer virtual memory, but does offer a hard disk, designers could create their own virtual memory system as part of their implementation.)

Designers can choose whatever data structures and organization they feel optimize their implementations performance, in the context of its requirements.

All threads share the same method area, so access to the method area's data structures must be designed to be thread-safe. If two threads are attempting to find a class named `Lava`, for example, and `Lava` has not yet been loaded, only one thread should be allowed to load it while the other one waits.

The size of the method area need not be fixed. As the Java application runs, the virtual machine can expand and contract the method area to fit the application's needs. Also, the memory of the method area need not be contiguous. It could be allocated on a heap—even on the virtual machine's own heap. Implementations may allow users or programmers to specify an initial size for the method area, as well as a maximum or minimum size.

The method area can also be garbage collected. Because Java programs can be dynamically extended via class loader objects, classes can become "unreferenced" by the application. If a class becomes unreferenced, a Java Virtual Machine can unload the class (garbage collect it) to keep the memory occupied by the method area at a minimum. The unloading of classes—including the conditions under which a class can become "unreferenced"—is described in Chapter 7, "The Lifetime of a Class."

Type Information

For each type it loads, a Java Virtual Machine must store the following kinds of information in the method area:

- The fully qualified name of the type
- The fully qualified name of the type's direct superclass (unless the type is an interface or class `java.lang.Object`, neither of which have a superclass)
- Whether or not the type is a class or an interface
- The type's modifiers (some subset of `public`, `abstract`, `final`)
- An ordered list of the fully qualified names of any direct superinterfaces

Inside the Java class file and Java Virtual Machine, type names are always stored as *fully qualified names*. In Java source code, a fully qualified name is the name of a type's package, plus a dot, plus the type's *simple name*. For example, the fully qualified name of class `Object` in package `java.lang` is `java.lang.Object`. In class files, the dots are replaced by slashes, as in `java/lang/Object`. In the method area, fully qualified names can be represented in whatever form and data structures a designer chooses.

In addition to the basic type information listed above, the virtual machine must also store for each loaded type:

- The constant pool for the type
- Field information
- Method information
- All class (static) variables declared in the type, except constants
- A reference to class `ClassLoader`
- A reference to class `Class`

This data is described in the following sections.

The Constant Pool

For each type it loads, a Java Virtual Machine must store a *constant pool*. A constant pool is an ordered set of constants used by the type, including literals (string, integer, and floating point constants) and symbolic references to types, fields, and methods. Entries in the constant pool are referenced by index, much like the elements of an array. Because it holds symbolic references to all types, fields, and methods used by a type, the constant pool plays a central role in the dynamic linking of Java programs. The constant pool is described in more detail later in this chapter and in Chapter 6, "The Java Class File."

Field Information

For each field declared in the type, the following information must be stored in the method area. In addition to the information for each field, the order in which the fields are declared by the class or interface must also be recorded. Here is the list for fields:

- The field's name
- The field's type
- The field's modifiers (some subset of `public`, `private`, `protected`, `static`, `final`, `volatile`, `transient`)

Method Information

For each method declared in the type, the following information must be stored in the method area. As with fields, the order in which the methods are declared by the class or interface must be recorded as well as the data. Here is the list:

- The method's name
- The method's return type (or `void`)
- The number and types (in order) of the method's parameters
- The method's modifiers (some subset of `public`, `private`, `protected`, `static`, `final`, `synchronized`, `native`, `abstract`)

In addition to the items listed above, the following information must also be stored with each method that is not abstract or native:

- The method's bytecodes
- The sizes of the operand stack and local variables sections of the method's stack frame (these are described in a later section of this chapter)
- An exception table (this is described in Chapter 17, "Exceptions")

Class Variables

Class variables are shared among all instances of a class and can be accessed even in the absence of any instance. These variables are associated with the class--not with instances of the class--so they are logically part of the class data in the method area. Before a Java Virtual Machine uses a class, it must allocate memory from the method area for each non-final class variable declared in the class.

Constants (class variables declared `final`) are not treated in the same way as non-final class variables. Every type that uses a final class variable gets a copy of the constant value in its own constant pool. As

part of the constant pool, final class variables are stored in the method area--just like non-final class variables. But whereas non-final class variables are stored as part of the data for the type that *declares* them, final class variables are stored as part of the data for any type that *uses* them. This special treatment of constants is explained in more detail in Chapter 6, "The Java Class File."

*A Reference to Class **ClassLoader***

For each type it loads, a Java Virtual Machine must keep track of whether or not the type was loaded via the primordial class loader or a class loader object. For those types loaded via a class loader object, the virtual machine must store a reference to the class loader object that loaded the type. This information is stored as part of the type's data in the method area.

The virtual machine uses this information during dynamic linking. When one type refers to another type, the virtual machine requests the referenced type from the same class loader that loaded the referencing type. This process of dynamic linking is also central to the way the virtual machine forms separate name spaces. To be able to properly perform dynamic linking and maintain multiple name spaces, the virtual machine needs to know what class loader loaded each type in its method area. The details of dynamic linking and name spaces are given in Chapter 8, "The Linking Model."

*A Reference to Class **Class***

An instance of class `java.lang.Class` is created by the Java Virtual Machine for every type it loads. The virtual machine must in some way associate a reference to the `Class` instance for a type with the type's data in the method area.

Your Java programs can obtain and use references to `Class` objects. One static method in class `Class`, `forName()`, allows you to get a reference to the `Class` instance for any loaded class:

```
begin
```

```
// A method declared in class java.lang.Class:
```

```
public static Class forName(String className);
```

```
end
```

If you invoke `forName("java.lang.Object")`, for example, you will get a reference to the `Class` object that represents `java.lang.Object`. If you invoke `forName("java.util Enumeration")`, you will get a reference to the `Class` object that represents the `Enumeration` interface from the `java.util` package. You can use `forName()` to get a `Class` reference for any loaded type from any package, so long as the type can be (or already has been) loaded into the current name space. If the virtual machine is unable to load the requested type into the current name space, `forName()` will throw `ClassNotFoundException`.

An alternative way to get a `Class` reference is to invoke `getClass()` on any object reference. This method is inherited by every object from class `Object` itself:

```
begin
```

```
// A method declared in class java.lang.Object:
```

```
public final Class getClass();
```

```
end
```

If you have a reference to an object of class `java.lang.Integer`, for example, you could get the `Class` object for `java.lang.Integer` simply by invoking `getClass()` on your reference to the `Integer` object.

Given a reference to a `Class` object, you can find out information about the type by invoking methods declared in class `Class`. If you look at these methods, you will quickly realize that class `Class` gives the running application access to the information stored in the method area. Here are some of the methods declared in class `Class`:

```
begin
```

```
// Some of the methods declared in class java.lang.Class:
```

```
public String getName();
```

```
public Class getSuperClass();
```

```
public boolean isInterface();
```

```
public Class[] getInterfaces();
```

```
public ClassLoader getClassLoader()
```

```
end
```

These methods just return information about a loaded type. `getName()` returns the fully qualified name of the type. `getSuperClass()` returns the `Class` instance for the type's direct superclass. If the type is class `java.lang.Object` or an interface, none of which have a superclass, `getSuperClass()` returns `null`. `isInterface()` returns `true` if the `Class` object describes an interface, `false` if it describes a class. `getInterfaces()` returns an array of `Class` objects, one for each direct superinterface. The superinterfaces appear in the array in the order they are declared as superinterfaces by the type. If the type has no direct superinterfaces, `getInterfaces()` returns an array of length zero. `getClassLoader()` returns a reference to the `ClassLoader` object that loaded this type, or `null` if the type was loaded by the primordial class loader. All this information comes straight out of the method area.

Method Tables

The type information stored in the method area must be organized to be quickly accessible. In addition to the raw type information listed above, implementations may include other data structures that speed up access to the raw data. One example of such a data structure is a *method table*. For each non-abstract class a Java Virtual Machine loads, it could generate a method table and include it as part of the class information it stores in the method area. A method table is an array of direct references to all the

instance methods that may be invoked on a class instance, including instance methods inherited from superclasses. (A method table isn't helpful in the case of abstract classes or interfaces, because the program will never instantiate these.) A method table allows a virtual machine to quickly locate an instance method invoked on an object. Method tables are described in detail in Chapter 8, "The Linking Model."

An Example of Method Area Use

As an example of how the Java Virtual Machine uses the information it stores in the method area, consider these classes:

begin

```
// On CD-ROM in file jvm/ex2/Lava.java
class Lava {
```

```
    private int speed = 5; // 5 kilometers per hour
```

```
    void flow() {
    }
}
```

```
// On CD-ROM in file jvm/ex2/Volcano.java
class Volcano {
```

```
    public static void main(String[] args) {
        Lava lava = new Lava();
        lava.flow();
    }
}
```

end

The following paragraphs describe how an implementation might execute the first instruction in the bytecodes for the `main()` method of the `Volcano` application. Different implementations of the Java Virtual Machine can operate in very different ways. The following description illustrates one way--but not the only way--a Java Virtual Machine could execute the first instruction of `Volcano`'s `main()` method.

To run the `Volcano` application, you give the name "`Volcano`" to a Java Virtual Machine in an implementation-dependent manner. Given the name `Volcano`, the virtual machine finds and reads in

file `Volcano.class`. It extracts the definition of class `Volcano` from the binary data in the imported class file and places the information into the method area. The virtual machine then invokes the `main()` method, by interpreting the bytecodes stored in the method area. As the virtual machine executes `main()`, it maintains a pointer to the constant pool (a data structure in the method area) for the current class (class `Volcano`).

Note that this Java Virtual Machine has already begun to execute the bytecodes for `main()` in class `Volcano` even though it hasn't yet loaded class `Lava`. Like many (probably most) implementations of the Java Virtual Machine, this implementation doesn't wait until all classes used by the application are loaded before it begins executing `main()`. It loads classes only as it needs them.

`main()`'s first instruction tells the Java Virtual Machine to allocate enough memory for the class listed in constant pool entry one. The virtual machine uses its pointer into `Volcano`'s constant pool to look up entry one and finds a symbolic reference to class `Lava`. It checks the method area to see if `Lava` has already been loaded.

The symbolic reference is just a string giving the class's fully qualified name: `"Lava"`. Here you can see that the method area must be organized so a class can be located--as quickly as possible--given only the class's fully qualified name. Implementation designers can choose whatever algorithm and data structures best fit their needs--a hash table, a search tree, anything. This same mechanism can be used by the static `forName()` method of class `Class`, which returns a `Class` reference given a fully qualified name.

When the virtual machine discovers that it hasn't yet loaded a class named `"Lava"`, it proceeds to find and read in file `Lava.class`. It extracts the definition of class `Lava` from the imported binary data and places the information into the method area.

The Java Virtual Machine then replaces the symbolic reference in `Volcano`'s constant pool entry one, which is just the string `"Lava"`, with a pointer to the class data for `Lava`. If the virtual machine ever has to use `Volcano`'s constant pool entry one again, it won't have to go through the relatively slow process of searching through the method area for class `Lava` given only a symbolic reference, the string `"Lava"`. It can just use the pointer to more quickly access the class data for `Lava`. This process of replacing symbolic references with direct references (in this case, a native pointer) is called *constant pool resolution*. The symbolic reference is *resolved* into a direct reference by searching through the method area until the referenced entity is found, loading new classes if necessary.

Finally, the virtual machine is ready to actually allocate memory for a new `Lava` object. Once again, the virtual machine consults the information stored in the method area. It uses the pointer (which was just put into `Volcano`'s constant pool entry one) to the `Lava` data (which was just imported into the method area) to find out how much heap space is required by a `Lava` object.

A Java Virtual Machine can always determine the amount of memory required to represent an object by looking into the class data stored in the method area. The actual amount of heap space required by a particular object, however, is implementation-dependent. The internal representation of objects inside a Java Virtual Machine is another decision of implementation designers. Object representation is discussed in more detail later in this chapter.

Once the Java Virtual Machine has determined the amount of heap space required by a `Lava` object, it

allocates that space on the heap and initializes the instance variable `speed` to zero, its default initial value. If class `Lava` is superclass, `Object`, has any instance variables, those are also initialized to default initial values. (The details of initialization of both classes and objects are given in Chapter 7, "The Lifetime of a Class.")

The first instruction of `main()` completes by pushing a reference to the new `Lava` object onto the stack. A later instruction will use the reference to invoke Java code that initializes the `speed` variable to its proper initial value, five. Another instruction will use the reference to invoke the `flow()` method on the referenced `Lava` object.

The Heap

Whenever a class instance or array is created in a running Java application, the memory for the new object is allocated from a single heap. As there is only one heap inside a Java Virtual Machine instance, all threads share it. Because a Java application runs inside its "own" exclusive Java Virtual Machine instance, there is a separate heap for every individual running application. There is no way two different Java applications could trample on each other's heap data. Two different threads of the same application, however, could trample on each other's heap data. This is why you must be concerned about proper synchronization of multi-threaded access to objects (heap data) in your Java programs.

The Java Virtual Machine has an instruction that allocates memory on the heap for a new object, but has no instruction for freeing that memory. Just as you can't explicitly free an object in Java source code, you can't explicitly free an object in Java bytecodes. The virtual machine itself is responsible for deciding whether and when to free memory occupied by objects that are no longer referenced by the running application. Usually, a Java Virtual Machine implementation uses a *garbage collector* to manage the heap.

Garbage Collection

A garbage collector's primary function is to automatically reclaim the memory used by objects that are no longer referenced by the running application. It may also move objects as the application runs to reduce heap fragmentation.

A garbage collector is not strictly required by the Java Virtual Machine specification. The specification only requires that an implementation manage its own heap *in some manner*. For example, an implementation could simply have a fixed amount of heap space available and throw an `OutOfMemory` exception when that space fills up. While this implementation may not win many prizes, it does qualify as a Java Virtual Machine. The Java Virtual Machine specification does not say how much memory an implementation must make available to running programs. It does not say how an implementation must manage its heap. It says to implementation designers only that the program will be allocating memory from the heap, but not freeing it. It is up to designers to figure out how they want to deal with that fact.

No garbage collection technique is dictated by the Java Virtual Machine specification. Designers can use whatever techniques seem most appropriate given their goals, constraints, and talents. Because references to objects can exist in many places--Java Stacks, the heap, the method area, native method stacks--the choice of garbage collection technique heavily influences the design of an implementation's runtime data areas. Various garbage collection techniques are described in Chapter 9, "Garbage Collection."

As with the method area, the memory that makes up the heap need not be contiguous, and may be

expanded and contracted as the running program progresses. An implementation's method area could, in fact, be implemented on top of its heap. In other words, when a virtual machine needs memory for a freshly loaded class, it could take that memory from the same heap on which objects reside. The same garbage collector that frees memory occupied by unreferenced objects could take care of finding and freeing (unloading) unreferenced classes. Implementations may allow users or programmers to specify an initial size for the heap, as well as a maximum and minimum size.

Object Representation

The Java Virtual Machine specification is silent on how objects should be represented on the heap. Object representation--an integral aspect of the overall design of the heap and garbage collector--is a decision of implementation designers

The primary data that must in some way be represented for each object is the instance variables declared in the object's class and all its superclasses. Given an object reference, the virtual machine must be able to quickly locate the instance data for the object. In addition, there must be some way to access an object's class data (stored in the method area) given a reference to the object. For this reason, the memory allocated for an object usually includes some kind of pointer into the method area.

One possible heap design divides the heap into two parts: a handle pool and an object pool. An object reference is a native pointer to a handle pool entry. A handle pool entry has two components: a pointer to instance data in the object pool and a pointer to class data in the method area. The advantage of this scheme is that it makes it easy for the virtual machine to combat heap fragmentation. When the virtual machine moves an object in the object pool, it need only update one pointer with the object's new address: the relevant pointer in the handle pool. The disadvantage of this approach is that every access to an object's instance data requires dereferencing two pointers. This approach to object representation is shown graphically in Figure 5-5. This kind of heap is demonstrated interactively by the HeapOfFish applet, described in Chapter 9, "Garbage Collection."




Figure 5-5. Splitting an object across handle and object pools.

Another design makes an object reference a native pointer to a bundle of data that contains the object's instance data and a pointer to the object's class data. This approach requires dereferencing only one pointer to access an object's instance data, but makes moving objects more complicated. When the virtual machine moves an object to combat fragmentation of this kind of heap, it must update every reference to that object anywhere in the runtime data areas. This approach to object representation is shown graphically in Figure 5-6.




Figure 5-6. Keeping object data all in one place.

The virtual machine needs to get from an object reference to that object's class data for several reasons. When a running program attempts to cast an object reference to another type, the virtual machine must check to see if the type being cast to is the actual class of the referenced object or one of its supertypes. . It must perform the same kind of check when a program performs an `instanceof` operation. In either case, the virtual machine must look into the class data of the referenced object. When a program invokes an instance method, the virtual machine must perform dynamic binding: it must choose the method to invoke based not on the type of the reference but on the class of the object. To do this, it must once again have access to the class data given only a reference to the object.

No matter what object representation an implementation uses, it is likely that a method table is close at hand for each object. Method tables, because they speed up the invocation of instance methods, can play an important role in achieving good overall performance for a virtual machine implementation. Method tables are not required by the Java Virtual Machine specification and may not exist in all implementations. Implementations that have extremely low memory requirements, for instance, may not be able to afford the extra memory space method tables occupy. If an implementation does use method tables, however, an object's method table will likely be quickly accessible given just a reference to the object.

One way an implementation could connect a method table to an object reference is shown graphically in Figure 5-7. This figure shows that the pointer kept with the instance data for each object points to a special structure. The special structure has two components:

- A pointer to the full the class data for the object
- The method table for the object

The method table is an array of pointers to the data for each instance method that can be invoked on objects of that class. The method data pointed to by method table includes:

- The sizes of the operand stack and local variables sections of the method's stack
- The method's bytecodes
- An exception table

This gives the virtual machine enough information to invoke the method. The method table include pointers to data for methods declared explicitly in the object's class or inherited from superclasses. In other words, the pointers in the method table may point to methods defined in the object's class or any of its superclasses. More information on method tables is given in Chapter 8, "The Linking Model."




Figure 5-7. Keeping the method table close at hand.

If you are familiar with the inner workings of C++, you may recognize the method table as similar to the VTBL or virtual table of C++ objects. In C++, objects are represented by their instance data plus an array of pointers to any virtual functions that can be invoked on the object. This approach could also be taken by a Java Virtual Machine implementation. An implementation could include a copy of the method table for a class as part of the heap image for every instance of that class. This approach would consume more heap space than the approach shown in Figure 5-7, but might yield slightly better performance on a systems that enjoy large quantities of available memory.

One other kind of data that is not shown in Figures 5-5 and 5-6, but which is logically part of an object's data on the heap, is the object's *lock*. Each object in a Java Virtual Machine is associated with a lock (or *mutex*) that a program can use to coordinate multi-threaded access to the object. Only one thread at a time can "own" an object's lock. While a particular thread owns a particular object's lock, only that thread can access that object's instance variables. All other threads that attempt to access the object's

variables have to wait until the owning thread releases the object's lock. If a thread requests a lock that is already owned by another thread, the requesting thread has to wait until the owning thread releases the lock. Once a thread owns a lock, it can request the same lock again multiple times, but then has to release the lock the same number of times before it is made available to other threads. If a thread requests a lock three times, for example, that thread will continue to own the lock until it has released it three times.

Many objects will go through their entire lifetimes without ever being locked by a thread. The data required to implement an object's lock is not needed unless the lock is actually requested by a thread. As a result, many implementations, such as the ones shown in Figure 5-5 and 5-6, may not include a pointer to "lock data" within the object itself. Such implementations must create the necessary data to represent a lock when the lock is requested for the first time. In this scheme, the virtual machine must associate the lock with the object in some indirect way, such as by placing the lock data into a search tree based on the object's address.

Along with data that implements a lock, every Java object is logically associated with data that implements a *wait set*. Whereas locks help threads to work independently on shared data without interfering with one another, wait sets help threads to cooperate with one another--to work together towards a common goal.

Wait sets are used in conjunction with wait and notify methods. Every class inherits from `Object` three "wait methods" (overloaded forms of a method named `wait()`) and two "notify methods" (`notify()` and `notifyAll()`). When a thread invokes a wait method on an object, the Java Virtual Machine suspends that thread and adds it to that object's wait set. When a thread invokes a notify method on an object, the virtual machine will at some future time wake up one or more threads from that object's wait set. As with the data that implements an object's lock, the data that implements an object's wait set is not needed unless a wait or notify method is actually invoked on the object. As a result, many implementations of the Java Virtual Machine may keep the wait set data separate from the actual object data. Such implementations could allocate the data needed to represent an object's wait set when a wait or notify method is first invoked on that object by the running application. For more information about locks and wait sets, see Chapter 20, "Thread Synchronization."

One last example of a type of data that may be included as part of the image of an object on the heap is any data needed by the garbage collector. The garbage collector must in some way keep track of which objects are referenced by the program. This task invariably requires data to be kept for each object on the heap. The kind of data required depends upon the garbage collection technique being used. For example, if an implementation uses a *mark and sweep* algorithm, it must be able to mark an object as referenced or unreferenced. For each unreferenced object, it may also need to indicate whether or not the object's finalizer has been run. As with thread locks, this data may be kept separate from the object image. Some garbage collection techniques only require this extra data while the garbage collector is actually running. A mark and sweep algorithm, for instance, could potentially use a separate bitmap for marking referenced and unreferenced objects. More detail on various garbage collection techniques, and the data that is required by each of them, is given in Chapter 9, "Garbage Collection."

In addition to data that a garbage collector uses to distinguish between reference and unreferenced objects, a garbage collector needs data to keep track of which objects on which it has already executed a finalizer. Garbage collectors must run the finalizer of any object whose class declares one before it reclaims the memory occupied by that object. The Java language specification states that a garbage collector will only execute an object's finalizer once, but allows that finalizer to "resurrect" the object: to make the object referenced again. When the object becomes unreferenced for a second time, the garbage

collector must not finalize it again. Because most objects will likely not have a finalizer, and very few of those will resurrect their objects, this scenario of garbage collecting the same object twice will probably be extremely rare. As a result, the data used to keep track of objects that have already been finalized, though logically part of the data associated with an object, will likely not be part of the object representation on the heap. In most cases, garbage collectors will keep this information in a separate place. Chapter 9, "Garbage Collection," gives more information about finalization.

Array Representation

In Java, arrays are full-fledged objects. Like objects, arrays are always stored on the heap. Also like objects, implementation designers can decide how they want to represent arrays on the heap.

Arrays have a `CLASS` instance associated with their class, just like any other object. All arrays of the same dimension and type have the same class. The length of an array (or the lengths of each dimension of a multidimensional array) does not play any role in establishing the array's class. For example, an array of three `ints` has the same class as an array of three hundred `ints`. The length of an array is considered part of its instance data.

The name of an array's class has one open square bracket for each dimension plus a letter or string representing the array's type. For example, the class name for an array of `ints` is "`[I`". The class name for a three-dimensional array of `bytes` is "`[[[B`". The class name for a two-dimensional array of `Objects` is "`[[Ljava.lang.Object`". The full details of this naming convention for array classes is given in Chapter 6, "The Java Class File."

Multi-dimensional arrays are represented as arrays of arrays. A two dimensional array of `ints`, for example, would be represented by a one dimensional array of references to several one dimensional arrays of `ints`. This is shown graphically in Figure 5-8.




Figure 5-8. One possible heap representation for arrays.

The data that must be kept on the heap for each array is the array's length, the array data, and some kind of reference to the array's class data. Given a reference to an array, the virtual machine must be able to determine the array's length, to get and set its elements by index (checking to make sure the array bounds are not exceeded), and to invoke any methods declared by `Object`, the direct superclass of all arrays.

The Program Counter

Each thread of a running program has its own pc register, or program counter, which is created when the thread is started. The pc register is one word in size, so it can hold both a native pointer and a `returnValue`. As a thread executes a Java method, the pc register contains the address of the current instruction being executed by the thread. An "address" can be a native pointer or an offset from the beginning of a method's bytecodes. If a thread is executing a native method, the value of the pc register is undefined.

The Java Stack

When a new thread is launched, the Java Virtual Machine creates a new Java stack for the thread. As mentioned earlier, a Java stack stores a thread's state in discrete frames. The Java Virtual Machine only performs two operations directly on Java Stacks: it pushes and pops frames.

The method that is currently being executed by a thread is the thread's *current method*. The stack frame for the current method is the *current frame*. The class in which the current method is defined is called the *current class*, and the current class's constant pool is the *current constant pool*. As it executes a method, the Java Virtual Machine keeps track of the current class and current constant pool. When the virtual machine encounters instructions that operate on data stored in the stack frame, it performs those operations on the current frame.

When a thread invokes a Java method, the virtual machine creates and pushes a new frame onto the thread's Java stack. This new frame then becomes the current frame. As the method executes, it uses the frame to store parameters, local variables, intermediate computations, and other data.

A method can complete in either of two ways. If a method completes by returning, it is said to have *normal completion*. If it completes by throwing an exception, it is said to have *abrupt completion*. When a method completes, whether normally or abruptly, the Java Virtual Machine pops and discards the method's stack frame. The frame for the previous method then becomes the current frame.

All the data on a thread's Java stack is private to that thread. There is no way for a thread to access or alter the Java stack of another thread. Because of this, you need never worry about synchronizing multi-threaded access to local variables in your Java programs. When a thread invokes a method, the method's local variables are stored in a frame on the invoking thread's Java stack. Only one thread can ever access those local variables: the thread that invoked the method.

Like the method area and heap, the Java stack and stack frames need not be contiguous in memory. Frames could be allocated on a contiguous stack, or they could be allocated on a heap, or some combination of both. The actual data structures used to represent the Java stack and stack frames is a decision of implementation designers. Implementations may allow users or programmers to specify an initial size for Java stacks, as well as a maximum or minimum size.

The Stack Frame

The stack frame has three parts: local variables, operand stack, and frame data. The sizes of the local variables and operand stack, which are measured in words, depend upon the needs of each individual method. These sizes are determined at compile time and included in the class file data for each method. The size of the frame data is implementation dependent.

When the Java Virtual Machine invokes a Java method, it checks the class data to determine the number of words required by the method in the local variables and operand stack. It creates a stack frame of the proper size for the method and pushes it onto the Java stack.

Local Variables

The local variables section of the Java stack frame is organized as a zero-based array of words. Instructions that use a value from the local variables section provide an index into the zero-based array. Values of type `int`, `float`, `reference`, and `returnValue` occupy one entry in the local variables array. Values of type `byte`, `short`, and `char` are converted to `int` before being stored into the local variables. Values of type `long` and `double` occupy two consecutive entries in the array.

To refer to a long or double in the local variables, instructions provide the index of the first of the two consecutive entries occupied by the value. For example, if a `long` occupies array entries three and four, instructions would refer to that `long` by index three. All values in the local variables are word-aligned. Dual-entry `longs` and `doubles` can start at any index.

The local variables section contains a method's parameters and local variables. Compilers place the parameters into the local variable array first, in the order in which they are declared. Figure 5-9 shows the local variables section for the following two methods:

begin

```
// On CD-ROM in file jvm/ex3/Example3a.java
class Example3a {
```

```
    public static int runClassMethod(int i, long l, float f,
        double d, Object o, byte b) {

        return 0;
    }
```

```
    public int runInstanceMethod(char c, double d, short s,
        boolean b) {

        return 0;
    }
```

}

end

runClassMethod()			runInstanceMethod()		
index	type	parameter	index	type	parameter
0	int	int i	0	reference	hidden this
1	long	long l	1	int	char c
3	float	float f	2	double	double d
4	double	double d	4	int	short s
6	reference	Object o	5	int	boolean b
7	int	byte b			

Figure 5-9. Method parameters on the local variables section of a Java stack.

Note that Figure 5-9 shows that the first parameter in the local variables for `runInstanceMethod()` is of type `reference`, even though no such parameter appears in the source code. This is the hidden `this` reference passed to every instance method. Instance methods use this reference to access the instance data of the object upon which they were invoked. As you can see by looking at the local variables for `runClassMethod()` in Figure 5-9, class methods do not receive a hidden `this`. Class methods are not invoked on objects. You can't directly access a class's instance variables from a class method, because there is no instance associated with the method invocation.

Note also that types `byte`, `short`, `char`, and `boolean` in the source code become `ints` in the local variables. This is also true of the operand stack. As mentioned earlier, the `boolean` type is not supported directly by the Java Virtual Machine. The Java compiler always uses `ints` to represent `boolean` values in the local variables or operand stack. Data types `byte`, `short`, and `char`, however, are supported directly by the Java Virtual Machine. These can be stored on the heap as instance variables or array elements, or in the method area as class variables. When placed into local variables or the operand stack, however, values of type `byte`, `short`, and `char` are converted into `ints`. They are manipulated as `ints` while on the stack frame, then converted back into `byte`, `short`, or `char` when stored back into heap or method area.

Also note that `Object o` is passed as a reference to `runClassMethod()`. In Java, all objects are passed by reference. As all objects are stored on the heap, you will never find an image of an object in the local variables or operand stack, only object references.

Aside from a method's parameters, which compilers must place into the local variables array first and in

order of declaration, Java compilers can arrange the local variables array as they wish. Compilers can place the method's local variables into the array in any order, and they can use the same array entry for more than one local variable. For example, if two local variables have limited scopes that don't overlap, such as the `i` and `j` local variables in `Example3b` below, compilers are free to use the same array entry for both variables. During the first half of the method, before `j` comes into scope, entry zero could be used for `i`. During the second half of the method, after `i` has gone out of scope, entry zero could be used for `j`.

begin

```
// On CD-ROM in file jvm/ex3/Example3b.java
class Example3b {
```

```
    public static void runtwoLoops() {
```

```
        for (int i = 0; i < 10; ++i) {
            System.out.println(i);
        }
```

```
        for (int j = 9; j >= 0; --j) {
            System.out.println(j);
        }
```

```
    }
```

```
}
```

end

As with all the other runtime memory areas, implementation designers can use whatever data structures they deem most appropriate to represent the local variables. The Java Virtual Machine specification does not indicate how `long`s and `double`s should be split across the two array entries they occupy. Implementations that use a word size of 64 bits could, for example, store the entire long or double in the lower of the two consecutive entries, leaving the higher entry unused.

Operand Stack

Like the local variables, the operand stack is organized as an array of words. But unlike the local variables, which are accessed via array indices, the operand stack is accessed by pushing and popping values. If an instruction pushes a value onto the operand stack, a later instruction can pop and use that value.

The virtual machine stores the same data types in the operand stack that it stores in the local variables: `int`, `long`, `float`, `double`, `reference`, and `returnType`. It converts values of type `byte`,

`short`, and `char` to `int` before pushing them onto the operand stack.

Other than the program counter, which can't be directly accessed by instructions, the Java Virtual Machine has no registers. The Java Virtual Machine is stack-based rather than register-based because its instructions take their operands from the operand stack rather than from registers. Instructions can also take operands from other places, such as immediately following the opcode (the byte representing the instruction) in the bytecode stream, or from the constant pool. The Java Virtual Machine instruction set's main focus of attention, however, is the operand stack.

The Java Virtual Machine uses the operand stack as a work space. Many instructions pop values from the operand stack, operate on them, and push the result. For example, the `iadd` instruction adds two integers by popping two `ints` off the top of the operand stack, adding them, and pushing the `int` result. Here is how a Java Virtual Machine would add two local variables that contain `ints` and store the `int` result in a third local variable:

```
begin
  iload_0    // push the int in local variable 0 onto the stack
  iload_1    // push the int in local variable 1 onto the stack
  iadd       // pop two ints, add them, push result
  istore_2   // pop int, store into local variable 2
end
```

In this sequence of bytecodes, the first two instructions, `iload_0` and `iload_1`, push the `ints` stored in local variable positions zero and one onto the operand stack. The `iadd` instruction pops those two `int` values, adds them, and pushes the `int` result back onto the operand stack. The fourth instruction, `istore_2`, pops the result of the add off the top of the operand stack and stores it into local variable position two. In Figure 5-10, you can see a graphical depiction of the state of the local variables and operand stack while executing the above instructions. In this figure, unused slots of the local variables and operand stack are left blank.




Figure 5-10. Adding two local variables.

Frame Data

In addition to the local variables and operand stack, the Java stack frame includes data to support constant pool resolution, normal method return, and exception dispatch. This data is stored in the *frame data* portion of the Java stack frame.

Many instructions in the Java Virtual Machine's instruction set refer to entries in the constant pool. Some instructions merely push constant values of type `int`, `long`, `float`, `double`, or `String` from the constant pool onto the operand stack. Some instructions use constant pool entries to refer to classes or arrays to instantiate, fields to access, or methods to invoke. Other instructions determine whether a particular object is a descendant of a particular class or interface specified by a constant pool entry.

Whenever the Java Virtual Machine encounters any of the instructions that refer to an entry in the constant pool, it uses the frame data's pointer to the constant pool to access that information. As mentioned earlier, references to types, fields, and methods in the constant pool are initially symbolic. When the virtual machine looks up a constant pool entry that refers to a class, interface, field, or method, that reference may still be symbolic. If so, the virtual machine must resolve the reference at that time.

Aside from constant pool resolution, the frame data must assist the virtual machine in processing a normal or abrupt method completion. If a method completes normally (by returning), the virtual machine must restore the stack frame of the invoking method. It must set the pc register to point to the instruction in the invoking method that follows the instruction that invoked the completing method. If the completing method returns a value, the virtual machine must push that value onto the operand stack of the invoking method.

The frame data must also contain some kind of reference to the method's exception table, which the virtual machine uses to process any exceptions thrown during the course of execution of the method. An exception table, which is described in detail in Chapter 17, "Exceptions," defines ranges within the bytecodes of a method that are protected by catch clauses. Each entry in an exception table gives a starting and ending position of the range protected by a catch clause, an index into the constant pool that gives the exception class being caught, and a starting position of the catch clause's code.

When a method throws an exception, the Java Virtual Machine uses the exception table referred to by the frame data to determine how to handle the exception. If the virtual machine finds a matching catch clause in the method's exception table, it transfers control to the beginning of that catch clause. If the virtual machine doesn't find a matching catch clause, the method completes abnormally. The virtual machine uses the information in the frame data to restore the invoking method's frame. It then rethrows the same exception in the context of the invoking method.

In addition to data to support constant pool resolution, normal method return, and exception dispatch, the stack frame may also include other information that is implementation dependent, such as data to support debugging.

Possible Implementations of the Java Stack

Implementation designers can represent the Java stack in whatever way they wish. As mentioned earlier, one potential way to implement the stack is by allocating each frame separately from a heap. As an example of this approach, consider the following class:

begin

```
// On CD-ROM in file jvm/ex3/Example3c.java
class Example3c {

    public static void addAndPrint() {
        double result = addTwoTypes(1, 88.88);
        System.out.println(result);
    }

    public static double addTwoTypes(int i, double d) {
        return i + d;
    }
}
```

end

Figure 5-11 shows three snapshots of the Java stack for a thread that invokes the `addAndPrint()` method. In the implementation of the Java Virtual Machine represented in this figure, each frame is allocated separately from a heap. To invoke the `addTwoTypes()` method, the `addAndPrint()` method first pushes an `int` one and `double` 88.88 onto its operand stack. It then invokes the `addTwoTypes()` method.




Figure 5-11. Allocating frames from a heap.

The instruction to invoke `addTwoTypes()` refers to a constant pool entry. The Java Virtual Machine looks up the entry and resolves it if necessary.

Note that the `addAndPrint()` method uses the constant pool to identify the `addTwoTypes()` method, even though it is part of the same class. Like references to fields and methods of other classes, references to the fields and methods of the *same* class are initially symbolic and must be resolved before they are used.

The resolved constant pool entry points to information in the method area about the `addTwoTypes()` method. The virtual machine uses this information to determine the sizes required by `addTwoTypes()` for the local variables and operand stack. In the class file generated by Sun's `javac` compiler from the JDK 1.1, `addTwoTypes()` requires three words in the local variables and four words in the operand stack. (As mentioned earlier, the size of the frame data portion is implementation dependent.) The virtual machine allocates enough memory for the `addTwoTypes()` frame from a heap. It then pops the `double` and `int` parameters (`88.88` and `one`) from `addAndPrint()`'s operand stack and places them into `addTwoType()`'s local variable slots one and zero.

When `addTwoTypes()` returns, it first pushes the `double` return value (in this case, `89.88`) onto its operand stack. The virtual machine uses the information in the frame data to locate the stack frame of the invoking method, `addAndPrint()`. It pushes the `double` return value onto `addAndPrint()`'s operand stack and frees the memory occupied by `addTwoType()`'s frame. It makes `addAndPrint()`'s frame current and continues executing the `addAndPrint()` method at the first instruction past the `addTwoType()` method invocation.

Figure 5-12 shows snapshots of the Java stack of a different virtual machine implementation executing the same methods. Instead of allocating each frame separately from a heap, this implementation

allocates frames from a contiguous stack. This approach allows the implementation to overlap the frames of adjacent methods. The portion of the invoking method's operand stack that contains the parameters to the invoked method become the base of the invoked method's local variables. In this example, `addAndPrint()`'s entire operand stack becomes `addTwoTypes()`'s entire local variables section.




Figure 5-12. Allocating frames from a contiguous stack.

This approach saves memory space because the same memory is used by the calling method to store the parameters as is used by the invoked method to access the parameters. It saves time because the Java Virtual Machine doesn't have to spend time copying the parameter values from one frame to another.

Note that the operand stack of the current frame is always at the "top" of the Java stack. Although this may be easier to visualize in the contiguous memory implementation of Figure 5-12, it is true no matter how the Java stack is implemented. (As mentioned earlier, in all the graphical images of the stack shown in this book, the stack grows downwards. The "top" of the stack is always shown at the bottom of the picture.) Instructions that push values onto (or pop values off of) the operand stack always operate on the current frame. Thus, pushing a value onto the operand stack can be seen as pushing a value onto the top of the entire Java stack. In the remainder of this book, "pushing a value onto the stack" refers to pushing a value onto the operand stack of the current frame.

One other possible approach to implementing the Java stack is a hybrid of the two approaches shown in Figure 5-11 and Figure 5-12. A Java Virtual Machine implementation can allocate a chunk of contiguous memory from a heap when a thread starts. In this memory, the virtual machine can use the overlapping frames approach shown in Figure 5-12. If the stack outgrows the contiguous memory, the virtual machine can allocate another chunk of contiguous memory from the heap. It can use the separate frames approach shown in Figure 5-11 to connect the invoking method's frame sitting in the old chunk with the invoked method's frame sitting in the new chunk. Within the new chunk, it can once again use the contiguous memory approach.

Native Method Stacks

In addition to all the runtime data areas defined by the Java Virtual Machine specification and described above, a running Java application may use other data areas created by or for native methods. When a thread invokes a native method, it enters a new world in which the structures and security restrictions of the Java Virtual Machine no longer hamper its freedom. A native method can likely access the runtime data areas of the virtual machine (it depends upon the native method interface), but can also do anything else it wants. It may use registers inside the native processor, allocate memory on any number of native heaps, or use any kind of stack.

Native methods are inherently implementation dependent. Implementation designers are free to decide what mechanisms they will use to enable a Java application running on their implementation to invoke native methods.

Any native method interface will use some kind of native method stack. When a thread invokes a Java method, the virtual machine creates a new frame and pushes it onto the Java stack. When a thread invokes a native method, however, that thread leaves the Java stack behind. Instead of pushing a new frame onto the thread's Java stack, the Java Virtual Machine will simply dynamically link to and directly invoke the native method. One way to think of it is that the Java Virtual Machine is dynamically extending itself with native code. It is as if the Java Virtual Machine implementation is just calling another (dynamically linked) method within itself, at the behest of the running Java program.

If an implementation's native method interface uses a C-linkage model, then the native method stacks are C stacks. When a C program invokes a C function, the stack operates in a certain way. The arguments to the function are pushed onto the stack in a certain order. The return value is passed back to the invoking function in a certain way. This would be the behavior of the native method stacks in that implementation.

A native method interface will likely (once again, it is up to the designers to decide) be able to call back into the Java Virtual Machine and invoke a Java method. In this case, the thread leaves the native method stack and enters another Java stack.

Figure 5-13 shows a graphical depiction of a thread that invokes a native method that calls back into the virtual machine to invoke another Java method. This figure shows the full picture of what a thread can expect inside the Java Virtual Machine. A thread may spend its entire lifetime executing Java methods, working with frames on its Java stack. Or, it may jump back and forth between the Java stack and native method stacks.




Figure 5-13. The stack for a thread that invokes Java and native methods.

As depicted in Figure 5-13, a thread first invoked two Java methods, the second of which invoked a native method. This act caused the virtual machine to use a native method stack. In this figure, the native method stack is shown as a finite amount of contiguous memory space. Assume it is a C stack. The stack area used by each C-linkage function is shown in gray and bounded by a dashed line. The first C-linkage function, which was invoked as a native method, invoked another C-linkage function. The second C-linkage function invoked a Java method through the native method interface. This Java method invoked another Java method, which is the current method shown in the figure.

As with the other runtime memory areas, the memory they occupied by native method stacks need not be of a fixed size. It can expand and contract as needed by the running application. Implementations may allow users or programmers to specify an initial size for the method area, as well as a maximum or minimum size.

Execution Engine

At the core of any Java Virtual Machine implementation is its execution engine. In the Java Virtual Machine specification, the behavior of the execution engine is defined in terms of an instruction set. For each instruction, the specification describes in detail *what* an implementation should do when it encounters the instruction as it executes bytecodes, but says very little about *how*. As mentioned in previous chapters, implementation designers are free to decide how their implementations will execute bytecodes. Their implementations can interpret, just-in-time compile, execute natively in silicon, use a combination of these, or dream up some brand new technique.

Similar to the three senses of the term "Java Virtual Machine" described at the beginning of this chapter, the term "execution engine" can also be used in any of three senses: an abstract specification, a concrete implementation, or a runtime instance. The abstract specification defines the behavior of an execution engine in terms of the instruction set. Concrete implementations, which may use a variety of techniques, are either software, hardware, or a combination of both. A runtime instance of an execution engine is a thread.

Each thread of a running Java application is a distinct instance of the virtual machine's execution engine. From the beginning of its lifetime to the end, a thread is either executing bytecodes or native methods. A thread may execute bytecodes directly, by interpreting or executing natively in silicon, or indirectly, by just-in-time compiling and executing the resulting native code. A Java Virtual Machine implementation may use other threads invisible to the running application, such as a thread that performs garbage collection. Such threads need not be "instances" of the implementation's execution engine. All threads that belong to the running application, however, are execution engines in action.

The Instruction Set

A method's bytecode stream is a sequence of instructions for the Java Virtual Machine. Each instruction consists of a one-byte *opcode* followed by zero or more *operands*. The opcode indicates the operation to be performed. Operands supply extra information needed by the Java Virtual Machine to perform the operation specified by the opcode. The opcode itself indicates whether or not it is followed by operands, and the form the operands (if any) take. Many Java Virtual Machine instructions take no operands, and therefore consist only of an opcode. Depending upon the opcode, the virtual machine may refer to data stored in other areas in addition to (or instead of) operands that trail the opcode. When it executes an instruction, the virtual machine may use entries in the current constant pool, entries in the current frame's local variables, or values sitting on the top of the current frame's operand stack.

The abstract execution engine runs by executing bytecodes one instruction at a time. This process takes place for each thread (execution engine instance) of the application running in the Java Virtual Machine. An execution engine fetches an opcode and, if that opcode has operands, fetches the operands. It executes the action requested by the opcode and its operands, then fetches another opcode. Execution of bytecodes continues until a thread completes either by returning from its starting method or by not catching a thrown exception.

From time to time, the execution engine may encounter an instruction that requests a native method invocation. On such occasions, the execution engine will dutifully attempt to invoke that native method. When the native method returns (if it completes normally, not by throwing an exception), the execution engine will continue executing the next instruction in the bytecode stream.

One way to think of native methods, therefore, is as programmer-customized extensions to the Java Virtual Machine's instruction set. If an instruction requests an invocation of a native method, the execution engine invokes the native method. Running the native method is how the Java Virtual Machine executes the instruction. When the native method returns, the virtual machine moves on to the next instruction. If the native method completes abnormally (by throwing an exception), the virtual machine follows the same steps to handle the exception as it does when any instruction throws an exception.

Part of the job of executing an instruction is determining the next instruction to execute. An execution engine determines the next opcode to fetch in one of three ways. For many instructions, the next opcode to execute directly follows the current opcode and its operands, if any, in the bytecode stream. For some instructions, such as `goto` or `return`, the execution engine determines the next opcode as part of its execution of the current instruction. If an instruction throws an exception, the execution engine determines the next opcode to fetch by searching for an appropriate catch clause.

Several instructions can throw exceptions. The `athrow` instruction, for example, throws an exception explicitly. This instruction is the compiled form of the `throw` statement in Java source code. Every time the `athrow` instruction is executed, it will throw an exception. Other instructions throw exceptions

only when certain conditions are encountered. For example, if the Java Virtual Machine discovers, to its chagrin, that the program is attempting to perform an integer divide by zero, it will throw an `ArithmeticException`. This can occur while executing any of four instructions--`idiv`, `ldiv`, `irem`, and `lrem`--which perform divisions or calculate remainders on `ints` or `longs`.

Each type of opcode in the Java Virtual Machine's instruction set has a mnemonic. In the typical assembly language style, streams of Java bytecodes can be represented by their mnemonics followed by (optional) operand values.

For an example of method's bytecode stream and mnemonics, consider the `doMathForever()` method of this class:

begin

```
// On CD-ROM in file jvm/ex4/Act.java
class Act {
```

```
    public static void doMathForever() {
        int i = 0;
        for (;;) {
            i += 1;
            i *= 2;
        }
    }
}
```

}

end

The stream of bytecodes for `doMathForever()` can be disassembled into mnemonics as shown

begin

```
// Bytecode stream: 03 3b 84 00 01 1a 05 68 3b a7 ff f9
```

```
// Disassembly:
```

```
// Method void doMathForever()
```

```
// Left column: offset of instruction from beginning of method
```

```
// |   Center column: instruction mnemonic and any operands
```

```
// |   |               Right column: comment
```

```
0   iconst_0           // 03
```

```
1   istore_0           // 3b
```



```

2   iinc 0, 1          // 84 00 01
5   iload_0            // 1a
6   iconst_2          // 05
7   imul               // 68
8   istore_0           // 3b
9   goto 2             // a7 ff f9

```

end

This way of representing mnemonics is very similar to the output of the `javap` program of Sun's JDK. `javap` allows you to look at the bytecode mnemonics of the methods of any class file. Note that jump addresses are given as offsets from the beginning of the method. The `goto` instruction causes the virtual machine to jump to the instruction at offset two (an `iinc`). The actual operand in the stream is minus seven. To execute this instruction, the virtual machine adds the operand to the current contents of the pc register. The result is the address of the `iinc` instruction at offset two. To make the mnemonics easier to read, the operands for jump instructions are shown as if the addition has already taken place. Instead of saying "`goto -7`," the mnemonics say, "`goto 2`."

The central focus of the Java Virtual Machine's instruction set is the operand stack. Values are generally pushed onto the operand stack before they are used. Although the Java Virtual Machine has no registers for storing arbitrary values, each method has a set of local variables. The instruction set treats the local variables, in effect, as a set of registers that are referred to by indexes. Nevertheless, other than the `iinc` instruction, which increments a local variable directly, values stored in the local variables must be moved to the operand stack before being used.

For example, to divide one local variable by another, the virtual machine must push both onto the stack, perform the division, and then store the result back into the local variables. To move the value of an array element or object field into a local variable, the virtual machine must first push the value onto the stack, then store it into the local variable. To set an array element or object field to a value stored in a local variable, the virtual machine must follow the reverse procedure. First, it must push the value of the local variable onto the stack, then pop it off the stack and into the array element or object field on the heap.

Several goals--some conflicting--guided the design of the Java Virtual Machine's instruction set. These goals are basically the same as those described in Part I of this book as the motivation behind Java's entire architecture: platform independence, network mobility, and security.

The platform independence goal was a major influence in the design of the instruction set. The instruction set's stack-centered approach, described above, was chosen over a register-centered approach to facilitate efficient implementation on architectures with few or irregular registers, such as the Intel 80X86. This feature of the instruction set--the stack-centered design--make it easier to implement the Java Virtual Machine on a wide variety of host architectures.

As mentioned in Chapter 4, "Network Mobility," one major design consideration was class file compactness. Compactness is important because it facilitates speedy transmission of class files across networks. In the bytecodes stored in class files, all instructions--except two that deal with table jumping--

-are aligned on byte boundaries. The total number of opcodes is small enough so that opcodes occupy only one byte. This design strategy favors class file compactness possibly at the cost of some performance when the program runs. In some Java Virtual Machine implementations, especially those executing bytecodes in silicon, the single-byte opcode may preclude certain optimizations that could improve performance. Also, better performance may have been possible on some implementations if the bytecode streams were word-aligned instead of byte-aligned. (An implementation could always realign bytecode streams, or translate opcodes into a more efficient form as classes are loaded. Bytecodes are byte-aligned in the class file and in the specification of the abstract method area and execution engine. Concrete implementations can store the loaded bytecode streams any way they wish.)

Another goal that guided the design of the instruction set was the ability to do bytecode verification, especially all at once by a data flow analyzer. The verification capability is needed as part of Java's security framework. The ability to use a data flow analyzer on the bytecodes when they are loaded, rather than verifying each instruction as it is executed, facilitates execution speed. One way this design goal manifests itself in the instruction set is that most opcodes indicate the type they operate on.

For example, instead of simply having one instruction that pops a word from the operand stack and stores it in a local variable, the Java Virtual Machine's instruction set has two. One instruction, `istore`, pops and stores an `int`. The other instruction, `fstore`, pops and stores a `float`. Both of these instructions perform the exact same function when executed: they pop a word and store it. Distinguishing between popping and storing an `int` versus a `float` is important only to the verification process.

For many instructions, the virtual machine needs to know the types being operated on to know how to perform the operation. For example, the Java Virtual Machine supports two ways of adding two words together, yielding a one-word result. One addition treats the words as `ints`, the other as `floats`. The difference between these two instructions facilitates verification, but also tells the virtual machine whether it should perform integer or floating point arithmetic.

A few instructions operate on any type. The `dup` instruction, for example, duplicates the top word of a stack irrespective of its type. Some instructions, such as `goto`, don't operate on typed values. The majority of the instructions, however, operate on a specific type. The mnemonics for most of these "typed" instructions indicate their type by a single character prefix that starts their mnemonic. Table 5-2 shows the prefixes for the various types. A few instructions, such as `arraylength` or `instanceof`, don't include a prefix because their type is obvious. The `arraylength` opcode requires an array reference. The `instanceof` opcode requires an object reference.

Table 5-2. Type prefixes of bytecode mnemonics

Type	Code	Example	Description
byte	b	ba load	load byte from array
short	s	sa load	load short from array
int	i	ia load	load int from array
long	l	la load	load long from array
char	c	ca load	load char from array
float	f	fa load	load float from array
double	d	da load	load double from array

reference a aaload load reference from array

Values on the operand stack must be used in a manner appropriate to their type. It is illegal, for example, to push four `ints`, then add them as if they were two `longs`. It is illegal to push a `float` value onto the operand stack from the local variables, then store it as an `int` in an array on the heap. It is illegal to push a `double` value from an object field on the heap, then store the topmost of its two words into the local variables as an value of type `reference`. The strict type rules that are enforced by Java compilers must also be enforced by Java Virtual Machine implementations.

Implementations must also observe rules when executing instructions that perform generic stack operations independent of type. As mentioned above, the `dup` instruction pushes a copy of the top word of the stack, irrespective of type. This instruction can be used on any value that occupies one word: an `int`, `float`, `reference`, or `returnValue`. It is illegal, however, to use `dup` when the top of the stack contains either a `long` or `double`, the data types that occupy two consecutive operand stack locations. A `long` or `double` sitting on the top of the operand stack can be duplicated in their entirety by the `dup2` instruction, which pushes a copy of the top two words onto the operand stack. The generic instructions cannot be used to split up dual-word values.

To keep the instruction set small enough to enable each opcode to be represented by a single byte, not all operations are supported on all types. Most operations are not supported for types `byte`, `short`, and `char`. These types are converted to `int` when moved from the heap or method area to the stack frame. They are operated on as `ints`, then converted back to `byte`, `short`, or `char` before being stored back into the heap or method area.

Table 5-3 shows the computation types that correspond to each storage type in the Java Virtual Machine. As used here, a *storage type* is the manner in which values of the type are represented on the heap. The storage type corresponds to the type of the variable in Java source code. A *computation type* is the manner in which the type is represented on the Java stack frame.

Table 5-3. Storage and computation types inside the Java Virtual Machine

Storage Type	Minimum Bits in Heap or Method Area	Computation Type Words in the Java Stack Frame
byte	8int	load byte from array1
short	16int	load short from array1
int	32int	load int from array1
long	64long	load long from array2
char	16int	load char from array1
float	32float	load float from array1
double	64double	load double from array2
reference	32reference	load reference from array1

Implementations of the Java Virtual Machine must in some way ensure that values are operated on by instructions appropriate to their type. They can verify bytecodes up front as part of the class verification process, on the fly as the program executes, or some combination of both. Bytecode verification is described in more detail in Chapter 7, "The Lifetime of a Class." The entire instruction set is covered in

detail in Chapters 10 through 20.

Execution Techniques

Various execution techniques that may be used by an implementation--interpreting, just-in-time compiling, hot-spot compiling, native execution in silicon--were described in Chapter 1, "Introduction to Java's Architecture." The main point to remember about execution techniques is that an implementation can use any technique to execute bytecodes so long as it adheres to the semantics of the Java Virtual Machine instruction set.

Threads

The Java Virtual Machine specification defines a threading model that aims to facilitate implementation on a wide variety of architectures. One goal of the Java threading model is to enable implementation designers, where possible and appropriate, to use native threads. Alternatively, designers can implement a thread mechanism as part of their virtual machine implementation. One advantage to using native threads on a multi-processor host is that different threads of a Java application could run simultaneously on different processors.

One tradeoff of Java's threading model is that the specification of priorities is lowest-common-denominator. A Java thread can run at any one of ten priorities. Priority one is the lowest, and priority ten is the highest. If designers use native threads, they can map the ten Java priorities onto the native priorities however seems most appropriate. The Java Virtual Machine specification defines the behavior of threads at different priorities only by saying that all threads at the highest priority will get some CPU time. Threads at lower priorities are guaranteed to get CPU time only when all higher priority threads are blocked. Lower priority threads *may* get some CPU time when higher priority threads aren't blocked, but there are no guarantees.

The specification doesn't assume time-slicing between threads of different priorities, because not all architecture's time-slice. (As used here, *time-slicing* means that all threads at all priorities will be guaranteed some CPU time, even when no threads are blocked.) Even among those architectures that do time-slice, the algorithms used to allot time slots to threads at various priorities can differ greatly.

As mentioned in Chapter 2, "Platform Independence," you must not rely on time-slicing for program correctness. You should use thread priorities only to give the Java Virtual Machine hints at what it should spend more time on. To coordinate the activities of multiple threads, you should use synchronization.

The thread implementation of any Java Virtual Machine must support two aspects of *synchronization*: object locking and thread wait and notify. Object locking helps keep threads from interfering with one another while working independently on shared data. Thread wait and notify helps threads to cooperate with one another while working together toward some common goal. Running applications access the Java Virtual Machine's locking capabilities via the instruction set, and its wait and notify capabilities via the `wait()`, `notify()`, and `notifyAll()` methods of class `Object`. For more details, see Chapter 20, "Thread Synchronization."

In the Java Virtual Machine Specification, the behavior of Java threads is defined in terms of *variables*, a *main memory*, and *working memories*. Each Java Virtual Machine instance has a main memory, which contains all the program's variables: instance variables of objects, components of arrays, and class variables. Each thread has a working memory, in which the thread stores "working copies" of variables

it uses or assigns. Local variables and parameters, because they are private to individual threads, can be logically seen as part of either the working memory or main memory.

The Java Virtual Machine Specification defines many rules that govern the low-level interactions of threads with main memory. For example, one rule states that all operations on primitive types, except in some cases `longs` and `doubles`, are *atomic*. For example, if two threads compete to write two different values to an `int` variable, even in the absence of synchronization, the variable will end up with one value or the other. The variable will not contain a corrupted value. In other words, one thread will win the competition and write its value to the variable first. The losing thread need not sulk, however, because it will write its value the variable second, overwriting the "winning" thread's value.

The exception to this rule is any `long` or `double` variable that is not declared `volatile`. Rather than being treated as a single atomic 64-bit value, such variables may be treated by some implementations as two atomic 32-bit values. Storing a non-volatile `long` to memory, for example, could involve two 32-bit write operations. This non-atomic treatment of `longs` and `doubles` means that two threads competing to write two different values to a `long` or `double` variable can legally yield a corrupted result.

Although implementation designers are not required to treat operations involving non-volatile `longs` and `doubles` atomically, the Java Virtual Machine specification encourages them to do so anyway. This non-atomic treatment of `longs` and `doubles` is an exception to the general rule that operations on primitive types are atomic. This exception is intended to facilitate efficient implementation of the threading model on processors that don't provide efficient ways to transfer 64-bit values to and from memory. In the future, this exception may be eliminated. For the time being, however, Java programmers must be sure to synchronize access to shared `longs` and `doubles`.

Fundamentally, the rules governing low-level thread behavior specify when a thread may and when it must:

1. copy values of variables from the main memory to its working memory, and
2. write values from its working memory back into the main memory.

For certain conditions, the rules specify a precise and predictable order of memory reads and writes. For other conditions, however, the rules do not specify any order. The rules are designed to enable Java programmers to build multi-threaded programs that exhibit predictable behavior, while giving implementation designers some flexibility. This flexibility enables designers of Java Virtual Machine implementations to take advantage of standard hardware and software techniques that can improve the performance of multi-threaded applications.

The fundamental high-level implication of all the low-level rules that govern the behavior of threads is this: If access to certain variables isn't synchronized, threads are allowed update those variables in main memory in any order. Without synchronization, your multi-threaded applications may exhibit surprising behavior on some Java Virtual Machine implementations. With proper use of synchronization, however, you can create multi-threaded Java applications that behave in a predictable way on any implementation of the Java Virtual Machine.

Native Method Interface

Java Virtual Machine implementations aren't required to support any particular native method interface. Some implementations may support no native method interfaces at all. Others may support several, each geared towards a different purpose.

Sun's Java Native Interface, or JNI, is geared towards portability. JNI is designed so it can be supported by any implementation of the Java Virtual Machine, no matter what garbage collection technique or object representation the implementation uses. This in turn enables developers to link the same (JNI compatible) native method binaries to any JNI-supporting virtual machine implementation on a particular host platform.

Implementation designers can choose to create proprietary native method interfaces in addition to, or instead of, JNI. To achieve its portability, the JNI uses a lot of indirection through pointers to pointers and pointers to functions. To obtain the ultimate in performance, designers of an implementation may decide to offer their own low-level native method interface that is tied closely to the structure of their particular implementation. Designers could also decide to offer a higher-level native method interface than JNI, such as one that brings Java objects into a component software model.

To do useful work, a native method must be able to interact to some degree with the internal state of the Java Virtual Machine instance. For example, a native method interface may allow native methods to do some or all of the following:

- Pass and return data
- Access instance variables or invoke methods in objects on the garbage-collected heap
- Access class variables or invoke class methods
- Accessing arrays
- Lock an object on the heap for exclusive use by the current thread
- Create new objects on the garbage-collected heap
- Load new classes
- Throw new exceptions
- Catch exceptions thrown by Java methods that the native method invoked
- Catch asynchronous exceptions thrown by the virtual machine
- Indicate to the garbage collector that it no longer needs to use a particular object

Designing a native method interface that offers these services can be complicated. The design needs to ensure that the garbage collector doesn't free any objects that are being used by native methods. If an implementation's garbage collector moves objects to keep heap fragmentation at a minimum, the native method interface design must make sure that either:

1. an object can be moved after its reference has been passed to a native method, or
2. any objects whose references have been passed to a native method are pinned until the native method returns or otherwise indicates it is done with the objects

As you can see, native method interfaces are very intertwined with the inner workings of a Java Virtual Machine.

The Real Machine

As mentioned at the beginning of this chapter, all the subsystems, runtime data areas, and internal behaviors defined by the Java Virtual Machine specification are abstract. Designers aren't required to

organize their implementations around "real" components that map closely to the abstract components of the specification. The abstract internal components and behaviors are merely a vocabulary with which the specification defines the required external behavior of any Java Virtual Machine implementation.

In other words, an implementation can be anything on the inside, so long as it behaves like a Java Virtual Machine on the outside. Implementations must be able to recognize Java class files and must adhere to the semantics of the Java code the class files contain. But otherwise, anything goes. How bytecodes are executed, how the runtime data areas are organized, how garbage collection is accomplished, how threads are implemented, how the primordial class loader finds classes, what native method interfaces are supported--these are some of the many decisions left to implementation designers.

The flexibility of the specification gives designers the freedom to tailor their implementations to fit their circumstances. In some implementations, minimizing usage of resources may be critical. In other implementations, where resources are plentiful, maximizing performance may be the one and only goal.

By clearly marking the line between the external behavior and the internal implementation of a Java Virtual Machine, the specification preserves compatibility among all implementations while promoting innovation. Designers are encouraged to apply their talents and creativity towards building ever-better Java Virtual Machines.

Eternal Math: A Simulation

The CD-ROM contains several simulation applets that serve as interactive illustrations for the material presented in this book. The applet shown in Figure 5-14 simulates a Java Virtual Machine executing a few bytecodes. You can run this applet by loading `applets/EternalMath.html` from the CD-ROM into any Java enabled web browser or applet viewer that supports JDK 1.1.

The instructions in the simulation represent the body of the `doMathForever()` method of class `Act`, shown above in the "Instruction Set" section of this chapter. This simulation shows the local variables and operand stack of the current frame, the pc register, and the bytecodes in the method area. It also shows an optop register, which you can think of as part of the frame data of this particular implementation of the Java Virtual Machine. The optop register always points to one word beyond the top of the operand stack.

The applet has four buttons: Step, Reset, Run, and Stop. Each time you press the Step button, the Java Virtual Machine simulator will execute the instruction pointed to by the pc register. Initially, the pc register points to an `iconst_0` instruction. The first time you press the Step button, therefore, the virtual machine will execute `iconst_0`. It will push a zero onto the stack and set the pc register to point to the next instruction to execute. Subsequent presses of the Step button will execute subsequent instructions and the pc register will lead the way. If you press the Run button, the simulation will continue with no further coaxing on your part until you press the Stop button. To start the simulation over, press the Reset button.

The value of each register (pc and optop) is shown two ways. The contents of each register, an integer offset from the beginning of either the method's bytecodes or the operand stack, is shown in an edit box. Also, a small arrow (either "pc" or "optop") indicates the location contained in the register.

In the simulation the operand stack is shown growing down the panel (up in memory offsets) as words are pushed onto it. The top of the stack recedes back up the panel as words are popped from it.

The `doMathForever()` method has only one local variable, `i`, which sits at array position zero. The first two instructions, `iconst_0` and `istore_0` initialize the local variable to zero. The next instruction, `iinc`, increments `i` by one. This instruction implements the `i += 1` statement from `doMathForever()`. The next instruction, `iload_0`, pushes the value of the local variable onto the operand stack. `iconst_2` pushes an `int 2` onto the operand stack. `imul` pops the top two `ints` from the operand stack, multiplies them, and pushes the result. The `istore_0` instruction pops the result of the multiply and puts it into the local variable. The previous four instructions implement the `i *= 2` statement from `doMathForever()`. The last instruction, `goto`, sends the program counter back to the `iinc` instruction. The `goto` implements the `for (;;) loop` of `doMathForever()`.

With enough patience and clicks of the Step button (or a long enough run of the Run button), you can get an arithmetic overflow. When the Java Virtual Machine encounters such a condition, it just truncates, as is shown by this simulation. It does not throw any exceptions.

For each step of the simulation, a panel at the bottom of the applet contains an explanation of what the next instruction will do. Happy clicking.

Applet Viewer: EternalMath.class

Applet

ETERNAL MATH

Local Variables			pc: 7		The Method		
index	hex value	value	offset	bytecode	mnemonic		
0	00000003	3	0	03	iconst_0		
			1	3b	istore_0		
			2	84	iinc 0 1		
			3	00			
			4	01			
			5	1a	iload_0		
			6	05	iconst_2		
			7	68	imul		
			8	3b	istore_0		
			9	a7	goto 2		
			10	ff			
			11	f9			

Operand Stack		
offset	hex value	value
0	00000003	3
1	00000002	2

optop: 2

optop >

Step Run

Reset Stop

imul will pop two integers, multiply them, and push the result.

On the CD-ROM

The CD-ROM contains the source code examples from this chapter in the `jvm` directory. The *Eternal Math* applet is contained in a web page on the CD-ROM in file `applets/EternalMath.html`. The source code for this applet is found alongside its class files, in the `applets/JVMSimulators` and `applets/JVMSimulators/COM/artima/jvmsim` directories.

The Resources Page

For links to more information about the Java Virtual Machine, visit the resources page for this chapter:
<http://www.artima.com/insidejvm/jvm.html>

Chapter Six

The Java Class File

The previous chapter, the first of Part II, "Java Internals," gave an overview of the Java Virtual Machine. The next four chapters will focus on different aspects of the Java Virtual Machine. This chapter takes a look at the Java class file. It describes the contents of the class file, including the structure and format of the constant pool. This chapter serves as a complete reference of the Java class file format.

Accompanying this chapter on the CD-ROM is an applet that interactively illustrates the material presented in the chapter. The applet, named *Getting Loaded*, simulates the Java Virtual Machine loading a Java class file. At the end of this chapter, you will find a description of this applet and an instructions on how to use it.

What is a Java Class File?

The Java class file is a precisely defined binary file format for Java programs. Each Java class file represents a complete description of one Java class or interface. There is no way to put more than one class or interface into a single class file. The precise definition of the class file format ensures that any Java class file can be loaded and correctly interpreted by any Java Virtual Machine, no matter what system produced the class file or what system hosts the virtual machine.

Although the class file is related to the Java language architecturally, it is not inextricably linked to the Java language. As shown in Figure 6-1, you could write programs in other languages and compile them to class files, or you could compile your Java programs to a different binary file format. Nevertheless, most Java programmers will likely use the class file as the primary vehicle for delivering their programs to Java Virtual Machines.




Figure 6-1. The non-exclusive relationship of the Java language and class file.

As mentioned in earlier chapters, the Java class file is a binary stream of 8-bit bytes. Data items are stored sequentially in the class file with no padding between adjacent items. The lack of padding helps keep class files compact. Items that occupy more than one byte are split up into several consecutive bytes that appear in *big-endian* (higher bytes first) order.

Just as your Java classes can contain varying numbers of fields, methods, method parameters, local variables, and so on, the Java class file can contain many items that vary in size or number from one class file to another. In the class file, the size or length of a variable-length item precedes the actual data for the item. This allows class file streams to be parsed from beginning to end, reading in the size of an item first followed by the item data.

What's in a Class File?

The Java class file contains everything a Java Virtual Machine needs to know about one Java class or interface. The remainder of this chapter describes the class file format using tables. Each table has a name and shows an ordered list of items that can appear in a class file. Items appear in the table in the order in which they appear in the class file. Each item has a type, a name, and a count. The type is either a table name or one of the "primitive types" shown in Table 6-1. All values stored in items of type u2, u4, and u8 appear in the class file in big-endian order.

Table 6-1. Class file "primitive types"

u1	a single unsigned byte
u2	two unsigned bytes
u4	four unsigned bytes
u8	eight unsigned bytes

The major components of the class file, in their order of appearance in the class file, are shown in Table 6-2. Each of these components is described in more detail below.

Table 6-2. Format of a ClassFile Table

Type	Name	Count
u4	magic	1
u2	minor_version	1
u2	major_version	1
u2	constant_pool_count	1
cp_info	constant_pool	constant_pool_count - 1
u2	access_flags	1
u2	this_class	1
u2	super_class	1
u2	interfaces_count	1
u2	interfaces	interfaces_count
u2	fields_count	1
field_info	fields	fields_count
u2	methods_count	1
method_info	methods	methods_count
u2	attributes_count	1
attribute_info	attributes	attributes_count

magic

The first four bytes of every Java class file are its *magic number*, 0xCAFEBAFE. The magic number makes non-Java class files easier to identify. If a file doesn't start with 0xCAFEBAFE, it definitely isn't a Java class file. A magic number can be chosen by a file format's designers to be any arbitrary number that isn't already in widespread use. The magic number for the Java class file was chosen back in the days when "Java" was called "Oak." According to Patrick Naughton, a key member of the original Java team, the magic number was chosen "long before the name Java was ever uttered in reference to this language. We were looking for something fun, unique, and easy to remember. It is only a coincidence that 0xCAFEBAFE, an oblique reference to the cute baristas at Peet's Coffee, was foreshadowing for the name Java."

minor_version and major_version

The second four bytes of the class file contain the minor and major version numbers. As Java technology evolves, new features may occasionally be added to the Java class file format. Each time the

class file format changes, the version numbers will change as well. To the Java Virtual Machine, the version numbers identify the format to which a particular class file adheres. Java Virtual Machines will generally be able to load class files with a given major version number and a range of minor version numbers. Java Virtual Machines must reject class files with version numbers outside their valid range.

For class files generated by 1.0 or 1.1 compilers, the major version number is 45. The minor version number is 3.

constant_pool_count and constant_pool

Following the magic and version numbers in the class file is the *constant pool*. As mentioned in Chapter 5, "The Java Virtual Machine," the constant pool contains the constants associated with the class or interface defined by the file. Constants such as literal strings, final variable values, class names, and method names are stored in the constant pool. The constant pool is organized as a list of entries. A count of the number of entries in the list, `constant_pool_count`, precedes the actual list, `constant_pool`.

Many entries in the constant pool refer to other entries in the constant pool, and many items that follow the constant pool in the class file refer back to entries in the constant pool. Throughout the class file, constant pool entries are referred to by the integer index that indicates their position in the `constant_pool` list. The first entry in the list has an index of one, the second has an index of two, and so on. Although there is no entry in the `constant_pool` list that has an index of zero, the missing zeroeth entry is included in the `constant_pool_count`. For example, if a `constant_pool` list includes fourteen entries (with indexes one through fourteen), the `constant_pool_count` would be fifteen.

Each constant pool entry starts with a one-byte tag that indicates the type of constant making its home at that position in the list. Once a Java Virtual Machine grabs and interprets this tag, it knows what to expect after the tag. Table 6-3 shows the names and values of the constant pool tags.

Table 6-3. Constant pool tags

Entry Type	Tag Value	Description
<code>CONSTANT_Utf8</code>	1	A UTF-8 encoded Unicode string
<code>CONSTANT_Integer</code>	3	An <code>int</code> literal value
<code>CONSTANT_Float</code>	4	A <code>float</code> literal value
<code>CONSTANT_Long</code>	5	A <code>long</code> literal value
<code>CONSTANT_Double</code>	6	A <code>double</code> literal value
<code>CONSTANT_Class</code>	7	A symbolic reference to a class or interface
<code>CONSTANT_String</code>	8	A <code>String</code> literal value
<code>CONSTANT_Fieldref</code>	9	A symbolic reference to a field
<code>CONSTANT_Methodref</code>	10	A symbolic reference to a method declared in a class

CONSTANT_InterfaceMethodref	11	A symbolic reference to a method declared in an interface
CONSTANT_NameAndType	12	Part of a symbolic reference to a field or method

For each tag shown in Table 6-3, there is a corresponding table. The name of the table is formed by appending "_info" to the tag name. For example, the table that corresponds to the `CONSTANT_Class` tag is called `CONSTANT_Class_info`. The `CONSTANT_Utf8_info` table stores a compressed form of Unicode strings. The tables for the various kinds of constant pool entries are described in detail later in this chapter.

The constant pool plays an important role in the dynamic linking of Java programs. In addition to literal constant values, the constant pool contains the following kinds of symbolic references:

- fully qualified names of classes and interfaces
- field names and descriptors
- method names and descriptors

A *field* is an instance or class variable of the class or interface. A *field descriptor* is a string that indicates the field's type. A *method descriptor* is a string that indicates the method's return type and the number, order, and types of its parameters. The constant pool's fully qualified names and method and field descriptors are used at run time to link code in this class or interface with code and data in other classes and interfaces. The class file contains no information about the eventual memory layout of its components, so classes, fields, and methods cannot be referenced directly by the bytecodes in the class file. The Java Virtual Machine resolves the actual address of any referenced item at run time given a symbolic reference from the constant pool. For example, bytecode instructions that invoke a method give constant pool index of a symbolic reference to the method to invoke. This process of using the symbolic references in the constant pool is described in more detail in Chapter 8, "The Linking Model."

access_flags

The first two bytes after the constant pool, the *access flags*, reveal several pieces of information about the class or interface defined in the file. To start with, the access flags indicate whether the file defines a class or an interface. The access flags also indicate what modifiers were used in the declaration of the class or interface. Classes and interfaces can be public or abstract. Classes can be final, though final classes cannot be abstract. Interfaces can't be final. The bits used for the various flags are shown in Table 6-4.

Table 6-4. Flag bits in the `access_flags` item of `ClassFile` tables

Flag Name	Value	Meaning if Set	Set By
ACC_PUBLIC	0x0001	Type is public	Classes and interfaces
ACC_FINAL	0x0010	Class is final	Classes only
ACC_SUPER	0x0020	Use new	<code>invokespecial</code> semantics Classes and interfaces

ACC_INTERFACE 0x0200 Type is an interface, not a class All interfaces, no classes

ACC_ABSTRACT 0x0400 Type is abstract All interfaces, some classes

The ACC_SUPER flag exists for backwards compatibility with Sun's older Java compilers. In Sun's older Java Virtual Machines, the `invokespecial` instruction had more relaxed semantics. All new compilers should set the ACC_SUPER flag. All new implementations of the Java Virtual Machine should implement the newer, stricter `invokespecial` semantics. (See the `invokespecial` instruction in Appendix A for a description of these semantics.) Sun's older compilers generate class files with the ACC_SUPER flag set to zero. Sun's older Java Virtual Machines ignore the flag if it is set.

All unused bits in `access_flags` must be set to zero and ignored by Java Virtual Machine implementations.

this_class

The next two bytes are the `this_class` item, an index into the constant pool. The constant pool entry at position `this_class` must be a `CONSTANT_Class_info` table, which has two parts: a `tag` and a `name_index`. The tag will have the value `CONSTANT_Class`. The constant pool entry at position `name_index` will be a `CONSTANT_Utf8_info` table containing the fully qualified name of the class or interface.

The `this_class` item provides a glimpse of how the constant pool is used. By itself, the `this_class` item is just an index into the constant pool. When a Java Virtual Machine looks up the constant pool entry at position `this_class`, it will find an entry that identifies itself via its `tag` as a `CONSTANT_Class_info`. The Java Virtual Machine knows `CONSTANT_Class_info` entries always have an index into the constant pool, called `name_index`, following their `tag`. So the virtual machine looks up the constant pool entry at position `name_index`, where it should find a `CONSTANT_Utf8_info` entry that contains the fully qualified name of the class or interface. See Figure 6-2 for a graphical depiction of this process.




Figure 6-2. Example of constant pool usage.

super_class

Following `this_class` in the class file is the `super_class` item, another two-byte index into the constant pool. The constant pool entry at position `super_class` will be a `CONSTANT_Class_info` entry that refers to the fully qualified name of this class's superclass. Because the base class of every object in Java programs is the `java.lang.Object` class, the `super_class` constant pool index will be valid for every class except `Object`. For `Object`, `super_class` is a zero. For interfaces, the constant pool entry at position `super_class` is `java.lang.Object`.

interfaces_count and interfaces

The component that follows `super_class` starts with `interfaces_count`, a count of the number of superinterfaces directly implemented by the class or interface defined in this file. Immediately following the count is `interfaces`, an array that contains one index into the constant pool for each superinterface directly implemented by this class or interface. Each superinterface is represented by a `CONSTANT_Class_info` entry in the constant pool that refers to the fully qualified name of the interface. Only direct superinterfaces, those that appear in the `implements` clause of the class or the `extends` clause of the interface declaration, appear in this array. The superinterfaces appear in the array in the order in which they appear (left to right) in the `implements` or `extends` clause.

fields_count and fields

Following the `interfaces` component in the class file is a description of the fields declared by this class or interface. This component starts with `fields_count`, a count of the number of fields, including both

class and instance variables. Following the count is a list of variable-length `field_info` tables, one for each field. (The `fields_count` indicates the number of `field_info` tables in the list.) The only fields that appear in the `fields` list are those that were declared by the class or interface defined in the file. No fields inherited from superclasses or superinterfaces appear in the `fields` list.

Each `field_info` table reveals information about one field. The table contains the field's name, descriptor, and modifiers. If the field is declared as final, the `field_info` table also reveals the field's constant value. Some of this information is contained in the `field_info` table itself, and some is contained in constant pool locations referred to by the table. The `field_info` table is described in more detail later in this chapter.

methods_count and methods

Following the fields in the class file is a description of the methods declared by the class or interface. This component starts with `methods_count`, a two-byte count of the number of methods in the class or interface. The count includes only those methods that are explicitly defined by this class or interface. (It does not include any methods inherited from superclasses or superinterfaces.) Following the method count are the methods themselves, described in a list of `method_info` tables. (The `methods_count` indicates the number of `method_info` tables in the list.)

The `method_info` table contains several pieces of information about the method, including the method's name and descriptor (its return type and argument types). If the method is not abstract, the `method_info` table includes the number of stack words required for the method's local variables, the maximum number of stack words required for the method's operand stack, a table of exceptions caught by the method, the bytecode sequence, and optional line number and local variable tables. If the method can throw any checked exceptions, the `method_info` table includes a list of those checked exceptions. The `method_info` table is described in detail later in this chapter.

attributes_count and attributes

The last component in the class file are the attributes, which give general information about the particular class or interface defined by the file. The attributes component starts with `attributes_count`, a count of the number of `attribute_info` tables appearing in the subsequent `attributes` list. The first item in each `attribute_info` table is an index into the constant pool of a `CONSTANT_Utf8_info` table that gives the attribute's name.

Attributes come in many varieties. Several varieties are defined by the Java Virtual Machine specification, but anyone can create their own varieties of attributes (following certain rules) and place them into class files. Java Virtual Machine implementations must silently ignore any attributes they don't recognize. The rules surrounding the creation of new varieties of attributes are described later in this chapter.

Attributes appear in several places in the class file, not just in the `attributes` item of the top-level `ClassFile` table. The attributes that appear in the `ClassFile` table give more information about the class or interface defined by the file. Attributes that give more information about a field may be included as part of `field_info` table. Attributes that give more information about a method may be included as part of a `method_info` table.

The Java Virtual Machine specification defines two kinds of attributes that may appear in the attributes list of the `ClassFile` table: `SourceCode` and `InnerClasses`. These two attributes are described in detail later in this chapter.

Special Strings

The symbolic references contained in the constant pool involve three special kinds of strings: fully qualified names, simple names, and descriptors. All symbolic references include the fully qualified name of a class or interface. Symbolic references to fields include a simple field name and field descriptor in addition to a fully qualified type name. Symbolic references to methods include a simple method name and method descriptor in addition to a fully qualified name.

The same special strings are used to describe the class or interface that is defined by the class file. The class or interface name, the superclass name (if any), and the names of any superinterfaces are all given as fully qualified names. For each field declared by the class or interface, the constant pool contains a simple name and field descriptor. For each method declared by the class or interface, the constant pool contains a simple name and method descriptor.

Fully Qualified Names

Whenever constant pool entries refer to classes and interfaces, they give the fully qualified name of the class or interface. In the class file, fully qualified names have their dots replaced with slashes. For example, the representation of the fully qualified name of `java.lang.Object` in the class file is `java/lang/Object`. The fully qualified name of `java.util.Hashtable` in the class file is `java/util/Hashtable`.

Simple Names

The names of fields and methods appear in constant pool entries as simple (not fully qualified) names. For example, a constant pool entry that refers to the `toString()` method of class `java.lang.Object` would give its method name as `"toString"`. A constant pool entry that refers to the `java.io.PrintStream out` field of class `java.lang.System` would specify the field name simply as `"out"`.

Descriptors

Symbolic references to fields and methods include a descriptor string in addition to a fully qualified class or interface name and a simple field or method name. A field descriptor gives the field's type. A method descriptor gives the method's return type and the number and types of the method's parameters.

Field and method descriptors are defined by the context free grammar shown below. Nonterminals of this grammar, such as *FieldType*, are shown in italic font. Terminals, such as `B` or `V`, are shown in fixed width font. The asterisk character (*) stands for zero or more occurrences of the item that precedes it placed side by side (with no intervening white space).

FieldDescriptor:

FieldType

ComponentType:

FieldType

FieldType:

BaseType

ObjectType

ArrayType

BaseType:

B

C

D

F

I

J

S

Z

ObjectType:

L<classname>;

ArrayType:

[ComponentType

ParameterDescriptor:

FieldType

MethodDescriptor:

(ParameterDescriptor) ReturnDescriptor*

ReturnDescriptor:

FieldType

V

The meaning of each of the *BaseType* terminals is shown in Table 6-5. The *V* terminal represents methods that return *void*.

Table 6-5. *BaseType* terminals

Terminal	Type
B	byte
C	char
D	double
F	float
I	int
J	long
S	short
Z	boolean

Some examples of field descriptors are shown in Table 6-6.

Table 6-6. Examples of field descriptors

Descriptor	Field Declaration
I	int i;
[[J	long[][] windingRoad;
[Ljava/lang/Object;	java.lang.Object[] stuff;
Ljava/util/Hashtable;	java.util.Hashtable ht;
[[[Z	boolean[][][] isReady;

Some examples of method descriptors are shown in Table 6-7. Note that method descriptors don't include the hidden *this* parameter passed as the first argument to all instance methods.

Table 6-7. Examples of method descriptors

Descriptor	Method Declaration
()I	int getSize();
()Ljava/lang/String;	String toString();
([Ljava/lang/String;)V	void main(String[] args);
()V	void wait()
(JI)V	void wait(long timeout, int nanos)
(ZILjava/lang/String;II)	boolean regionMatches
Z	(boolean ignoreCase, int toOffset, String other, int ooffset, int len);

```
([BII])I                                int read(byte[] b, int
                                         off, int len);
```

The Constant Pool

The constant pool is an ordered list of *cp_info* tables, each of which follows the general form shown in Table 6-8. The *tag* item of a *cp_info* table, an unsigned byte, indicates the table's variety and format. *cp_info* tables come in eleven varieties, each of which is described in detail in the following sections.

Table 6-8. General form of a *cp_info* table

Type	Name	Count
u1	tag	1
u1	info	depends on tag value

The *CONSTANT_Utf8_info* Table

A *CONSTANT_Utf8_info* table stores one constant string value in a modified UTF-8 format. This table is used to store many different kinds of strings, including:

- string literals that get instantiated as *String* objects
- the fully qualified name of the class or interface being defined
- the fully qualified name of the superclass (if any) of the class being defined
- the fully qualified names of any superinterfaces of the class or interface being defined
- the simple names and descriptors of any fields declared by the class or interface
- the simple names and descriptors of any methods declared by the class or interface
- fully qualified names of any referenced classes and interfaces
- simple names and descriptors of any referenced fields
- simple names and descriptors of any referenced methods
- strings associated with attributes

As you can see from the above list, there are four basic kinds of information stored in *CONSTANT_Utf8_info* tables: string literals, descriptions of the class or interface being defined, symbolic references to other classes and interfaces, and strings associated with attributes. Some examples of strings associated with attributes are: the name of the attribute, the name of the source file from which the class file was generated, and the names and descriptors of local variables.

The UTF-8 encoding scheme allows all two-byte Unicode characters to be represented in a string, but enables ASCII characters to be represented by just one byte. Table 6-9 shows the format of a *CONSTANT_Utf8_info* table.

Table 6-9. Format of a *CONSTANT_Utf8_info* table

Type	Name	Count
u1	tag	1
u2	length	1

u1 bytes length

tag

The *tag* item has the value *CONSTANT_Utf8* (1).

length

The *length* item gives the length in bytes of the subsequent *bytes* item.

bytes

The *bytes* item contains the characters of the string stored in a modified UTF-8 format. Characters in the range '0001' through '007f' (all the ASCII characters except the null character) are represented by one byte:




Figure 6-Single

The *null* character, '0000', and the characters in the range '0080' through '07ff' are represented by two bytes:




Figure 6-Double

Characters in the range '0800' through 'ffff' are represented by three bytes:




Figure 6-Triple

The encoding of UTF-8 strings in the *bytes* item of *CONSTANT_Utf8_info* tables differs from the standard UTF-8 format in two ways. First, in the standard UTF-8 encoding scheme, the *null* character is represented by one byte. In a *CONSTANT_Utf8_info* table, *null* characters are represented by two bytes. This two-byte encoding of *null*s means that the *bytes* item never contains any byte equal to zero. The second way the *bytes* item of a *CONSTANT_Utf8_info* departs from the standard UTF-8 encoding is that only one-, two-, and three-byte encodings are used in the *bytes* item. The standard UTF-8 includes longer formats that aren't used in *CONSTANT_Utf8_info* tables.

The *CONSTANT_Integer_info* Table

The *CONSTANT_Integer_info* table stores a constant *int* value. This table is used only to store *int* literals. It is not used in symbolic references. Table 6-10 shows the format of a *CONSTANT_Integer_info* table.

Table 6-10. Format of a *CONSTANT_Integer_info* table

Type	Name	Count
u1	tag	1
u4	bytes	1

tag

The *tag* item has the value *CONSTANT_Integer* (3).

bytes

The *bytes* item contains the *int* value stored in big-endian order.

The *CONSTANT_Float_info* Table

The *CONSTANT_Float_info* table stores a constant *float* value. This table is used only to store *float* literals. It is not used in symbolic references. Table 6-11 shows the format of a *CONSTANT_Float_info* table.

Table 6-11. Format of a *CONSTANT_Float_info* table

Type	Name	Count
u1	tag	1
u4	bytes	1

tag

The *tag* item has the value *CONSTANT_Float* (4).

bytes

The *bytes* item contains the *float* value stored in big-endian order. For the details of the representation of *float* in the Java class file, see Chapter 14, "Floating Point Arithmetic."

The *CONSTANT_Long_info* Table

The *CONSTANT_Long_info* table stores a constant *long* value. This table is used only to store *long* literals. It is not used in symbolic references. Table 6-12 shows the format of a *CONSTANT_Long_info* table.

As noted above, a *long* occupies two slots in the constant pool table. In the class file, a *long* entry is just followed by the next entry, but the index of the next entry is two more than that of the *long* entry.

Table 6-12. Format of a *CONSTANT_Long_info* table

Type	Name	Count
u1	tag	1
u8	bytes	1

tag

The *tag* item has the value *CONSTANT_Long* (5).

bytes

The *bytes* item contains the *long* value stored in big-endian order.

The *CONSTANT_Double_info* Table

The *CONSTANT_Double_info* table stores a constant *double* value. This table is used only to store *double* literals. It is not used in symbolic references. Table 6-13 shows the format of a *CONSTANT_Double_info* table.

As noted above, a *double* occupies two slots in the constant pool table. In the class file, a *double* entry is just followed by the next entry, but the index of the next entry is two more than that of the *double* entry.

Table 6-13. Format of a *CONSTANT_Double_info* table

Type	Name	Count
u1	tag	1

u8 bytes 1

tag

The *tag* item has the value *CONSTANT_Double* (6).

bytes

The *bytes* item contains the *double* value stored in big-endian order. For the details of the representation of *double* in the Java class file, see Chapter 14, "Floating Point Arithmetic."

The *CONSTANT_Class_info* Table

The *CONSTANT_Class_info* table represents a class or interface in symbolic references. All symbolic references, whether they refer to a class, interface, field, or method, include a *CONSTANT_Class_info* table. Table 6-14 shows the format of a *CONSTANT_Class_info* table.

Table 6-14. Format of a *CONSTANT_Class_info* table

Type	Name	Count
u1	tag	1
u2	name_index	1

tag

The *tag* item has the value *CONSTANT_Class* (7).

name_index

The *name_index* item gives the index of a *CONSTANT_Utf8_info* table that contains a fully qualified name of a class or interface.

The *CONSTANT_String_info* Table

A *CONSTANT_String_info* represents a literal string value, which will be represented as an instance of class *java.lang.String*. This table is only used to represent literal strings. It is not used in symbolic references. Table 6-15 shows the format of a *CONSTANT_String_info* table.

Table 6-15. Format of a *CONSTANT_String_info* table

Type	Name	Count
u1	tag	1
u2	string_index	1

tag

The **tag** item has the value **CONSTANT_String** (8).

string_index

The **string_index** item gives the index of a **CONSTANT_Utf8_info** entry that contains the value of the literal string.

The CONSTANT_Fieldref_info Table

The **CONSTANT_Fieldref_info** table represents a symbolic reference to a field. Table 6-16 shows the format of a **CONSTANT_Fieldref_info** table.

Table 6-16. Format of a **CONSTANT_Fieldref_info** table

Type	Name	Count
u1	tag	1
u2	class_index	1
u2	name_and_type_index	1

tag

The **tag** item has the value **CONSTANT_Fieldref** (9).

class_index

The **class_index** gives the index of the **CONSTANT_Class_info** entry for the class that declares the referenced field.

Note that the **CONSTANT_Class_info** specified by **class_index** must represent a class and not an interface. Although interfaces can declare fields, those fields are by definition public, static, and final. As mentioned in earlier chapters, class files do not contain symbolic references to static final fields of other classes. Instead, class files contain a copy of the constant value of any static final fields it uses. For example, if a class uses a static final field of type **float** that is declared in an interface, the class would have a **CONSTANT_Float_info** table in its own constant pool that stores the **float** value. For more information about this special treatment of static final fields, see Chapter 8, "The Linking Model."

name_and_type_index

The **name_and_type_index** gives the index of a **CONSTANT_NameAndType_info** entry that gives the field's simple name and descriptor.

The CONSTANT_Methodref_info Table

The **CONSTANT_Methodref_info** table represents a symbolic reference to a method declared in a class (not in an interface). Table 6-17 shows the format of a **CONSTANT_Methodref_info** table.

Table 6-17. Format of a *CONSTANT_Methodref_info* table

Type	Name	Count
u1	tag	1
u2	class_index	1
u2	name_and_type_index	1

tag

The *tag* item has the value *CONSTANT_Methodref* (10).

class_index

The *class_index* gives the index of a *CONSTANT_Class_info* entry for the class that declares the referenced method. The *CONSTANT_Class_info* table specified by *class_index* must be a class and not an interface. Symbolic references to methods declared in interfaces use *CONSTANT_InterfaceMethodref*.

name_and_type_index

The *name_and_type_index* gives the index of a *CONSTANT_NameAndType_info* entry that gives the method's simple name and descriptor.

The *CONSTANT_InterfaceMethodref_info* Table

The *CONSTANT_InterfaceMethodref_info* table represents a symbolic reference to a method declared in an interface (not in a class). Table 6-18 shows the format of a *CONSTANT_InterfaceMethodref_info* table.

Table 6-18. Format of a *CONSTANT_InterfaceMethodref_info* table

Type	Name	Count
u1	tag	1
u2	class_index	1
u2	name_and_type_index	1

tag

The *tag* item has the value *CONSTANT_InterfaceMethodref* (11).

class_index

The *class_index* gives the index of a *CONSTANT_Class_info* entry for the interface that declares the referenced method. The *CONSTANT_Class_info* table specified by *class_index* must be an interface and not a class. Symbolic references to methods declared in classes use *CONSTANT_Methodref*.

name_and_type_index

The *name_and_type_index* gives the index of a *CONSTANT_NameAndType_info* entry that gives the method's simple name and descriptor.

The *CONSTANT_NameAndType_info* Table

The *CONSTANT_NameAndType_info* table forms part of a symbolic reference to a field or method. This table gives constant pool entries of the simple name and the descriptor of the referenced field or method. Table 6-19 shows the format of a *CONSTANT_NameAndType_info* table.

Table 6-19. Format of a *CONSTANT_NameAndType_info* table

Type	Name	Count
u1	tag	1
u2	name_index	1
u2	descriptor_index	1

tag

The *tag* item has the value *CONSTANT_NameAndType* (12).

name_index

The *name_index* gives the index of a *CONSTANT_Utf8_info* entry that gives the name of the field or method.

descriptor_index

The *descriptor_index* gives the index of a *CONSTANT_Utf8_info* entry that gives the descriptor of the field or method.

Fields

Each field (class variable and instance variable) declared in a class or interface is described by a *field_info* table in the class file. The format of the *field_info* table is shown in Table 6-20.

Table 6-20. Format of a *field_info* table

Type	Name	Count
u2	access_flags	1
u2	name_index	1
u2	descriptor_index	1
u2	attributes_count	1
attribute_info	attributes	attributes_count

access_flags

The modifiers used in declaring the field are placed into the field's ***access_flags*** item. Table 6-21 shows the bits used by each flag.

Table 6-21. Flags in the ***access_flags*** item of ***field_info*** tables

Flag Name	Value	Meaning if Set	Set By
ACC_PUBLIC	0x0001	Field is public	Classes and interfaces
ACC_PRIVATE	0x0002	Field is private	Classes only
ACC_PROTECTED	0x0004	Field is protected	Classes only
ACC_STATIC	0x0008	Field is static	Classes and interfaces
ACC_FINAL	0x0010	Field is final	Classes and interfaces
ACC_VOLATILE	0x0040	Field is volatile	Classes only
ACC_TRANSIENT	0x0080	Field is transient	Classes only

For fields declared in a class (not an interface), at most one of ***ACC_PUBLIC***, ***ACC_PRIVATE***, and ***ACC_PROTECTED*** may be set. ***ACC_FINAL*** and ***ACC_VOLATILE*** must not both be set. All fields declared in interfaces must have the ***ACC_PUBLIC***, ***ACC_STATIC***, and ***ACC_FINAL*** flags set.

All unused bits in ***access_flags*** must be set to zero and ignored by Java Virtual Machine implementations.

name_index

The ***name_index*** gives the index of a ***CONSTANT_Utf8_info*** entry that gives the simple (not fully qualified) name of the field.

descriptor_index

The ***descriptor_index*** gives the index of a ***CONSTANT_Utf8_info*** entry that gives the descriptor of the field.

attributes_count* and *attributes

The ***attributes*** item is a list of ***attribute_info*** tables. The ***attributes_count*** indicates the number of ***attribute_info*** tables in the list. Two kinds of attributes defined by the Java Virtual Machine specification that may appear in this item are ***ConstantValue*** and ***Synthetic***. These two attributes are described in detail later in this chapter.

Methods

Each method declared in a class or interface or generated by the compiler is described in the class file by a ***method_info*** table. The two types of compiler-generated methods that may appear in class files are instance initialization methods (named ***<init>***) and class initialization methods (named ***<clinit>***). For more information on the compiler-generated methods, see Chapter 7, "The Lifetime of a Class." The

format of the *method_info* table is shown in Table 6-22.

Table 6-22. Format of a *method_info* table

Type	Name	Count
u2	access_flags	1
u2	name_index	1
u2	descriptor_index	1
u2	attributes_count	1
attribute_info	attributes	attributes_count

access_flags

The modifiers used in declaring the method are placed into the method's *access_flags* item. Table 6-23 shows the bits used by each flag.

Table 6-23. Flags in the *access_flags* item of *method_info* tables

Flag Name	Value	Meaning if Set	Set By
ACC_PUBLIC	0x0001	Method is public	Classes and all methods of interfaces
ACC_PRIVATE	0x0002	Method is private	Classes only
ACC_PROTECTED	0x0004	Method is protected	Classes only
ACC_STATIC	0x0008	Method is static	Classes only
ACC_FINAL	0x0010	Method is final	Classes only
ACC_SYNCHRONIZED	0x0020	Method is synchronized	Classes only
ACC_NATIVE	0x0100	Method is native	Classes only
ACC_ABSTRACT	0x0400	Method is abstract	Classes and all methods of interfaces

For methods declared in a class (not an interface), at most one of *ACC_PUBLIC*, *ACC_PRIVATE*, and *ACC_PROTECTED* may be set. If a method's *ACC_FINAL* flag is set, then its *ACC_SYNCHRONIZED*, *ACC_NATIVE*, and *ACC_ABSTRACT* flags must not be set. If a method's *ACC_PRIVATE* or *ACC_STATIC* flag is set, then its *ACC_ABSTRACT* flag must not be set. All methods declared in interfaces must have their *ACC_PUBLIC* and *ACC_ABSTRACT* flags set.

Instance initialization (*<init*) methods may only use flags *ACC_PUBLIC*, *ACC_PRIVATE*, and *ACC_PROTECTED*. Because class initialization (*<clinit*) methods are invoked by the Java Virtual Machine, never directly by Java bytecodes, the the *access_flags* for *<clinit* methods is ignored.

All unused bits in *access_flags* must be set to zero and ignored by Java Virtual Machine implementations.

name_index

The ***name_index*** gives the index of a ***CONSTANT_Utf8_info*** entry that gives the simple (not fully qualified) name of the method.

descriptor_index

The ***descriptor_index*** gives the index of a ***CONSTANT_Utf8_info*** entry that gives the descriptor of the method.

attributes_count and attributes

The ***attributes*** item is a list of ***attribute_info*** tables. The ***attributes_count*** indicates the number of ***attribute_info*** tables in the list. Three kinds of attributes that are defined by the Java Virtual Machine specification that may appear in this item are ***Code***, ***Exceptions***, and ***Synthetic***. These three attributes are described in detail later in this chapter.

Attributes

As mentioned above, attributes appear in several places inside a Java class file. They can appear in the ***ClassFile***, ***field_info***, ***method_info***, and ***Code_attribute*** tables. The ***Code_attribute*** table, an attribute itself, is described later in this section.

Every attribute follows the same general format of the ***attribute_info*** table, shown in Table 6-24. The first two bytes of an attribute, the ***attribute_name_index***, form an index into the constant pool of a ***CONSTANT_Utf8_info*** table that contains the string name of the attribute. Each ***attribute_info***, therefore, identifies its "type" by the first item in its table much like ***cp_info*** tables identify their type by the initial ***tag*** byte. The difference is that whereas the type of a ***cp_info*** table is indicated by an unsigned byte value, such as 3 (***CONSTANT_Integer_info***), the type of an ***attribute_info*** table is indicated by a string.

Following the ***attribute_name_index*** is a four-byte ***attribute_length*** item, which gives the length of the entire ***attribute_info*** table minus the initial six bytes. This length is necessary because anyone, following certain rules (outlined below), is allowed to add attributes to a Java class file. Java Virtual Machine implementations are allowed to recognize new attributes. Implementations must ignore any attributes they don't recognize. The ***attribute_length*** allows virtual machines to skip unrecognized attributes as they parse the class file.

Anyone who wishes to add a new attribute to a Java class file must follow these two rules:

1. Any attribute that is not predefined by the specification must not affect the semantics of class or interface types. New attributes can only add more information to the class file, such as information used during debugging.
2. The attribute must be named using the reverse Internet domain name scheme that is defined for package naming in the Java Language Specification. For example, if your Internet domain name were ***artima.com*** and you wished to create a new attribute named ***CompilerVersion***, you would name the attribute: ***COM.artima.CompilerVersion***.

Table 6-24. Format of an *attribute_info* table

Type	Name	Count
u2	attribute_name_index	1
u4	attribute_length	1
u1	info	attribute_length

attribute_name_index

The *attribute_name_index* gives the index in the constant pool of a *CONSTANT_Utf8_info* entry that contains the name of the attribute.

attribute_length

The *attribute_length* item indicates the length in bytes of the attribute data excluding the initial six bytes that contain the *attribute_name_index* and *attribute_length*.

info

The *info* item contains the attribute data.

Attributes

The Java Virtual Machine specification defines eight types of attributes, shown in Table 6-25. All Java Virtual Machine implementations must recognize three of these attributes: **Code**, **ConstantValue**, and **Exceptions**. Implementations can choose whether to recognize or ignore the other predefined attributes. (The **InnerClasses** and **Synthetic** attributes were added in Java 1.1). All of these predefined attributes are described in detail later in this chapter.

Table 6-25. Types of *attribute_info* tables defined by the specification

Name	Used By	Description
Code	method_info	The bytecodes and other data for one method
ConstantValue	field_info	The value of a final variable
Exceptions	method_info	The checked exceptions that a method may throw
InnerClasses	ClassFile	A description of the inner classes defined inside a class
LineNumberTable	Code_attribute	A mapping of line numbers to bytecodes for one method
LocalVariableTable	Code_attribute	A description of the local variables for one

SourceFile	ClassFile	method The name of the source file
Synthetic	field_info, method_info	An indicator that a field or method was generated by the compiler

The Code Attribute

One *Code_attribute* table appears in the *method_info* table of every method that is not abstract or native. The format of a *Code_attribute* table is shown in Table 6-26.

Table 6-26. Format of a *Code_attribute* table

Type	Name	Count
u2	attribute_name_index	1
u4	attribute_length	1
u2	max_stack	1
u2	max_locals	1
u4	code_length	1
u1	code	code_length
u2	exception_table_length	1
exception_info	exception_table	exception_table_length
u2	attributes_count	1
attribute_info	attributes	attributes_count

attribute_name_index

The *attribute_name_index* gives the index in the constant pool of a *CONSTANT_Utf8_info* entry that contains the string "Code".

attribute_length

The *attribute_length* item gives the length in bytes of the *Code* attribute excluding the initial six bytes that contain the *attribute_name_index* and *attribute_length* items.

max_stack

The *max_stack* item gives the maximum number of words that will be on the operand stack of this method at any point during its execution.

max_locals

The *max_locals* item gives the number of words in the local variables that are required by this method.

code_length and ***code***

The ***code_length*** item gives the length in bytes of the bytecode stream for this method. The bytecodes themselves appear in the ***code*** item.

exception_table_length and ***exception_table***

The ***exception_table*** item is a list of ***exception_info*** tables. Each ***exception_info*** table describes one exception table entry. The ***exception_table_length*** item gives the number of ***exception_info*** tables that appear in the ***exception_table*** list. The order in which the ***exception_info*** tables appear in the list is the order in which the Java Virtual Machine will check for a matching exception handler (catch clause) if an exception is thrown while this method executes. The format of a ***exception_info*** table is shown in Table 6-27. For more information about exception tables, see Chapter 17, "Exceptions."

Table 6-27. Format of an ***exception_info*** table

Type	Name	Count
u2	start_pc	1
u2	end_pc	1
u2	handler_pc	1
u2	catch_type	1

start_pc

The ***start_pc*** item gives the offset from the beginning of the code array for the beginning of the range covered by this exception handler.

end_pc

The ***end_pc*** item gives the offset from the beginning of the code array for one byte past the end of the range covered by this exception handler.

handler_pc

The ***handler_pc*** item gives the offset from the beginning of the code array for the instruction to jump to (the first instruction of the exception handler) if a thrown exception is caught by this entry.

catch_type

The ***catch_type*** item gives the constant pool index of a ***CONSTANT_Class_info*** entry for the type of exception caught by this exception handler. The ***CONSTANT_Class_info*** entry must represent class ***java.lang.Throwable*** or one of its subclasses.

If the value of ***catch_type*** is zero (which isn't a valid index into the constant pool, because the constant pool starts at index one), the exception handler handles all exceptions. A ***catch_type*** of zero is used to implement finally clauses. See Chapter 18, "Finally Clauses," for more information about

how finally clauses are implemented.

attributes_count and attributes

The ***attributes*** item is a list of ***attribute_info*** tables. The ***attributes_count*** indicates the number of ***attribute_info*** tables in the list. The two kinds of attributes defined by the Java Virtual Machine specification that may appear in this item are ***LineNumberTable*** and ***LocalVariableTable***. These two attributes are described in detail later in this chapter.

The ConstantValue Attribute

The ***ConstantValue*** attribute appears in ***field_info*** tables for fields that have a constant value. In the ***access_flags*** of a ***field_info*** table that includes a ***ConstantValue*** attribute, the ***ACC_STATIC*** flag must be set. The ***ACC_FINAL*** flag may also be set, though it is not required. When the virtual machine initializes a field that has a ***ConstantValue*** attribute, it assigns the constant value to the field. The format of a ***ConstantValue_attribute*** table is shown in Table 6-28.

Table 6-28. Format of a ***ConstantValue_attribute*** table

Type	Name	Count
u2	attribute_name_index	1
u4	attribute_length	1
u2	constantvalue_index	1

attribute_name_index

The ***attribute_name_index*** gives the index in the constant pool of a ***CONSTANT_Utf8_info*** entry that contains the string "***ConstantValue***".

attribute_length

The ***attribute_length*** item of a ***ConstantValue_attribute*** is always 2.

constantvalue_index

The ***constantvalue_index*** item gives the index in the constant pool of an entry that gives a constant value. Table 6-29 shows the type of entry for each type of field.

Table 6-29. Constant pool entry types for constant value attributes

byte, short, char, int, boolean	CONSTANT_Integer_info
long	CONSTANT_Long_info
float	CONSTANT_Float_info
double	CONSTANT_Double_info
java.lang.String	CONSTANT_String_info

The Exceptions Attribute

The **Exceptions** attribute lists the checked exceptions that a method may throw. One **Exceptions_attribute** table appears in the **method_info** table of every method that may throw checked exceptions. The format of a **Exceptions_attribute** table is shown in Table 6-30.

Table 6-30. Format of an **Exceptions_attribute** table

Type	Name	Count
u2	attribute_name_index	1
u4	attribute_length	1
u2	number_of_exceptions	1
u2	exception_index_table	number_of_exceptions

attribute_name_index

The **attribute_name_index** gives the index in the constant pool of a **CONSTANT_Utf8_info** entry that contains the string "Exceptions".

attribute_length

The **attribute_length** item gives the length in bytes of the **Exceptions_attribute** excluding the initial six bytes that contain the **attribute_name_index** and **attribute_length** items.

number_of_exceptions and **exception_index_table**

The **exception_index_table** is an array of indexes into the constant pool of **CONSTANT_Class_info** entries for the exceptions declared in this method's **throws** clause. In other words, the **exception_index_table** lists all the checked exceptions that this method may throw. The **number_of_exceptions** item indicates the number of indexes in the array.

The InnerClasses Attribute

The **InnerClasses** attribute lists any inner classes declared in a class. The **InnerClasses_attribute** table appears in the attributes of the **ClassFile** table for those classes that have inner classes. The format of a **InnerClasses_attribute** table is shown in Table 6-31.

Table 6-31. Format of a **InnerClasses_attribute** table

Type	Name	Count
u2	attribute_name_index	1
u4	attribute_length	1
u2	number_of_classes	1
classes_info	classes	number_of_classes

attribute_name_index

The *attribute_name_index* gives the index in the constant pool of a *CONSTANT_Utf8_info* entry that contains the string "InnerClasses".

attribute_length

The *attribute_length* item gives the length in bytes of the *Exceptions_attribute* excluding the initial six bytes that contain the *attribute_name_index* and *attribute_length* items.

number_of_classes and *classes*

The *classes* item is an array of *inner_class_info* tables. The *number_of_classes* gives the number of *inner_class_info* tables that appear in the *classes* array. The format of the *inner_class_info* table is shown in Table 6-32.

Table 6-32. Format of an *inner_class_info* table

Type	Name	Count
u2	<i>inner_class_info_index</i>	1
u2	<i>outer_class_info_index</i>	1
u2	<i>inner_name_index</i>	1
u2	<i>inner_class_access_flags</i>	1

inner_class_info_index

The *inner_class_info_index* gives an index into the constant pool for a *CONSTANT_Class_info* entry for an inner class.

outer_class_info_index

The *outer_class_info_index* gives an index into the constant pool for a *CONSTANT_Class_info* entry for the class in which the inner class described by this *inner_class_info* table is defined. If the inner class is not a member, the *outer_class_info_index* is zero.

inner_name_index

The *inner_name_index* gives an index in the constant pool for a *CONSTANT_Utf8_info* entry that gives the simple name of the inner class. If the inner class is anonymous, the *inner_name_index* is zero.

inner_class_access_flags

The *inner_class_access_flags* item gives the access flags for the inner class. The flags used in this item are the same shown in Table 6-4 above for top-level classes.

The *LineNumberTable* Attribute

The *LineNumberTable* attribute maps offsets in a method's bytecode stream to line numbers in the source file. One *LineNumberTable_attribute* table may appear (it is optional) in the attributes component of *Code_attribute* tables. The format of a *LineNumberTable_attribute* table is shown in Table 6-33.

Table 6-33. Format of a *LineNumberTable_attribute* table

Type	Name	Count
u2	attribute_name_index	1
u4	attribute_length	1
u2	line_number_table_length	1
line_number_info	line_number_table	line_number_table_length

attribute_name_index

The *attribute_name_index* gives the index in the constant pool of a *CONSTANT_Utf8_info* entry that contains the string "*LineNumberTable*".

attribute_length

The *attribute_length* item gives the length in bytes of the *LineNumberTable_attribute* excluding the initial six bytes that contain the *attribute_name_index* and *attribute_length* items.

line_number_table_length and *line_number_table*

The *line_number_table* item is an array of *line_number_info* tables. The *line_number_table_length* gives the number of *line_number_info* tables that appear in the *line_number_table* array. The tables in this array may appear in any order, and there may be more than one table for the same line number. The format of a *line_number_info* is shown in Table 6-34.

Table 6-34. Format of a *line_number_info*

Type	Name	Count
u2	start_pc	1
u2	line_number	1

start_pc

The *start_pc* item gives an offset from the beginning of the code array where a new line begins.

line_number

The *line_number* item gives the line number of the line that begins at *start_pc*.

The LocalVariableTable Attribute

The *LocalVariableTable* attribute maps words in the local variables portion of the method's stack frame to names and descriptors of local variables in the source code. One *LocalVariableTable_attribute* table may appear (it is optional) in the attributes component of *Code_attribute* tables. The format of a *LocalVariableTable_attribute* table is shown in Table 6-35.

Table 6-35. Format of a *LocalVariableTable_attribute* table

Type	Name	Count
u2	attribute_name_index	1
u4	attribute_length	1
u2	local_variable_table_length	1
local_variable_info	local_variable_table	local_variable_table_length

attribute_name_index

The *attribute_name_index* gives the index in the constant pool of a *CONSTANT_Utf8_info* entry that contains the string "*LocalVariableTable*".

attribute_length

The *attribute_length* item gives the length in bytes of the *LocalVariableTable_attribute* excluding the initial six bytes that contain the *attribute_name_index* and *attribute_length* items.

local_variable_table_length and *local_variable_table*

The *local_variable_table* item is an array of *local_variable_info* tables. The *local_variable_table_length* gives the number of *local_variable_info* tables that appear in the *local_variable_table* array. The format of a *local_variable_info* table is shown in Table 6-36.

Table 6-36. Format of a *local_variable_info* table

Type	Name	Count
u2	start_pc	1
u2	length	1
u2	name_index	1
u2	descriptor_index	1
u2	index	1

start_pc and *length*

The **start_pc** item gives an offset in the code array of the start of an instruction. The **length** item gives the length of the range of code that starts with **start_pc** for which a local variable is valid. The byte at offset **start_pc + length** from the beginning of the code array must either be the first byte of an instruction or the first byte past the end of the code array.

name_index

The **name_index** item gives an index in the constant pool of a **CONSTANT_Utf8_info** entry for the name of the local variable.

descriptor_index

The **descriptor_index** item gives an index in the constant pool of a **CONSTANT_Utf8_info** entry that contains the descriptor for this local variable. (A local variable descriptor adheres to the same grammar as a field descriptor.)

index

The **index** item gives the index in the local variable portion of this method's stack frame where the data for this local variable is kept as the method executes. If the local variable is of type **long** or **double**, the data occupies two words at positions **index** and **index + 1**. Otherwise the data occupies one word at position **index**.

The SourceFile Attribute

The optional **SourceFile** attribute, which may appear in the attributes component of a **ClassFile** table, gives the name of the source file from which the class file was generated. The format of a **SourceFile_attribute** table is shown in Table 6-37.

Table 6-37. Format of a **SourceFile_attribute** table

Type	Name	Count
u2	attribute_name_index	1
u4	attribute_length	1
u2	sourcefile_index	1

attribute_name_index

The **attribute_name_index** gives the index in the constant pool of a **CONSTANT_Utf8_info** entry that contains the string "**SourceFile**".

attribute_length

The **attribute_length** item of a **SourceFile_attribute** is always 2.

sourcefile_index

The *sourcefile_index* item gives the index in the constant pool of a *CONSTANT_Utf8_info* entry that contains the name of the source file. The source file name never includes a directory path.

The Synthetic Attribute

The Synthetic attribute, which may optionally appear in the attributes components of *field_info* and *method_info* tables, indicates that a field or method was generated by the compiler. The Synthetic attribute was added in Java 1.1 to support inner classes. The format of a *Synthetic_attribute* is shown in Table 6-38.

Table 6-38. Format of a *Synthetic_attribute* table

Type	Name	Count
u2	attribute_name_index	1
u4	attribute_length	1

attribute_name_index

The *attribute_name_index* gives the index in the constant pool of a *CONSTANT_Utf8_info* entry that contains the string "Synthetic".

attribute_length

The *attribute_length* must be zero.

Getting Loaded: A Simulation

The Getting Loaded applet, shown in Figure 6-3, simulates a Java Virtual Machine loading a class file. The class file being loaded in the simulation was generated by the 1.1 **javac** compiler from the Java source code shown below. Although the snippet of code used in the simulation may not be very useful in the real world, it does compile to a real class file and provides a reasonably simple example of the class file format. This is the same class used in the Eternal Math simulation applet described in Chapter 5, "The Java Virtual Machine."

```
// On CD-ROM in file classfile/ex1/Act.java
class Act {
```

```
    public static void doMathForever() {
        int i = 0;
        for (;;) {
            i += 1;
            i *= 2;
        }
    }
```

```

    }
}

end

```

The *Getting Loaded* applet allows you to drive the class load simulation one step at a time. For each step along the way you can read about the next chunk of bytes that is about to be consumed and interpreted by the Java Virtual Machine. Just press the "Step" button to cause the Java Virtual Machine to consume the next chunk. Pressing "Back" will undo the previous step, and pressing "Reset" will return the simulation to its original state, allowing you to start over from the beginning.

The Java Virtual Machine is shown at the bottom left consuming the stream of bytes that makes up the class file **Act.class**. The bytes are shown in hex streaming out of a server on the bottom right. The bytes travel right to left, between the server and the Java Virtual Machine, one chunk at a time. The chunk of bytes to be consumed by the Java Virtual Machine on the next "Step" button press are shown in red. These highlighted bytes are described in the large text area above the Java Virtual Machine. Any remaining bytes beyond the next chunk are shown in black.

As mentioned in previous sections, many items in the class file refer to constant pool entries. To make it easier for you to look up constant pool entries as you step through the simulation, a list of the contents of **Act**'s constant pool is shown in Table 6-39.

Table 6-39. Class **Act**'s constant pool

Index	Type	Value
1	CONSTANT_Class_info	7
2	CONSTANT_Class_info	16
3	CONSTANT_Methodref_info	2, 4
4	CONSTANT_NameAndType_info	6, 5
5	CONSTANT_Utf8_info	"()V"
6	CONSTANT_Utf8_info	"<init"
7	CONSTANT_Utf8_info	"Act"
8	CONSTANT_Utf8_info	"Act.java"
9	CONSTANT_Utf8_info	"Code"
10	CONSTANT_Utf8_info	"ConstantValue"
11	CONSTANT_Utf8_info	"Exceptions"
12	CONSTANT_Utf8_info	"LineNumberTable"
13	CONSTANT_Utf8_info	"LocalVariables"
14	CONSTANT_Utf8_info	"SourceFile"
15	CONSTANT_Utf8_info	"doMathForever"
16	CONSTANT_Utf8_info	"java/lang/Object"

Each chunk of bytes is fully explained in the text area. Because there is a lot of detail in the text area, you may wish to skim through all the steps first to get the general idea, then look back for more details. Happy clicking.

On the CD-ROM

*The CD-ROM contains the source code examples from this chapter in the **classfile** directory. The Getting Loaded applet is contained in a web page on the CD-ROM in file **applets/GettingLoaded.html**. The source code for this applet is found alongside its class files, in the **applets/GettingLoaded** directory.*

The Resources Page

For more information about class files visit the resources page for this chapter:
<http://www.artima.com/insidejvm/classfile.html>

Chapter Seven

The Lifetime of a Class

The previous chapter described in detail the format of the Java class file, the standard binary form for representing Java types. This chapter looks at what happens when binary type data is imported into a Java Virtual Machine. The chapter follows the lifetime of a type (class or interface) from the type's initial entrance into the virtual machine to its ultimate exit. It discusses the processes of loading, linking, and initialization that occur at the beginning of a class's lifetime; the processes of object instantiation, garbage collection, and finalization that can occur in the prime of a class's lifetime; and the finalization and unloading of types that can occur at the end of a class's lifetime.

Class Loading, Linking, and Initialization

The Java Virtual Machine makes types available to the running program through a process of *loading*, *linking*, and *initialization*. Loading is the process of bringing a binary form for a type into the Java Virtual Machine. Linking is the process of incorporating the binary type data into the runtime state of the virtual machine. Linking is divided into three sub-steps: *verification*, *preparation*, and *resolution*. Verification ensures the type is properly formed and fit for use by the Java Virtual Machine. Preparation involves allocating memory needed by the type, such as memory for any class variables. Resolution is the process of transforming symbolic references in the constant pool into direct references. Implementations may delay the resolution step until each symbolic reference is actually used by the running program. After verification, preparation, and (optionally) resolution are completed, the type is ready for initialization. During initialization, the class variables are given their proper initial values. See Figure 7-1 for a graphical depiction of this process.




Figure 7-1. The beginning of a class's lifetime.

As you can see from Figure 7-1, the processes of (1) loading, (2) linking, and (3) initialization must take place in that order. The only exception to this required ordering is the third phase of linking, resolution, which may optionally take place after initialization.

The Java Virtual Machine specification gives implementations flexibility in the timing of class and interface loading and linking, but strictly defines the timing of initialization. All implementations must initialize each class and interface on its first *active use*. An active use of a class is:

- The invocation of a constructor on a new instance of the class
- The creation of an array that has the class as its an element type
- The invocation of a method declared by the class (not inherited from a superclass)
- The use or assignment of a field declared by the class (not inherited from a superclass or superinterface), except for fields that are both `static` and `final`, and are initialized by a compile-time constant expression

An active use of an interface is:

- The use or assignment of a field declared by the interface (not inherited from a superinterface), except for fields that are initialized by a compile-time constant expression

All other uses of a type besides the five listed above are *passive uses* of the type. Several examples illustrating the difference between active and passive uses are given later in this chapter.

Aside from its own initial active use, there is one other situation that will cause the initialization of a class: the initial active use of one of its subclasses. Initialization of a class requires prior initialization of all its superclasses.

The same is not true of interfaces, however. An interface is initialized only because a non-constant field declared by the interface is used, never because a subinterface or class that implements the interface needs to be initialized. Thus, initialization of a class requires prior initialization of all its superclasses, but not its superinterfaces. Initialization of an interface does not require initialization of its superinterfaces.

The "initialize on first active use" rule drives the mechanism that loads, links, and initializes classes. On its first active use, a type must be initialized. Before it can be initialized, however, it must be linked. And before it can be linked, it must be loaded. At their option, implementations may load and link types early. They need not wait until the type's first active use to load and link the type. If a type hasn't been loaded and linked before its first active use, however, it must be loaded and linked at that time, so that it can be initialized.

Loading

The loading process consists of three basic activities. To load a type, the Java Virtual Machine must:

- produce a stream of binary data that represents the type
- parse the stream of binary data into internal data structures in the method area
- create an instance of class `java.lang.Class` that represents the type

The stream of binary data may adhere to the Java class file format, but could alternatively follow some other format. As mentioned in previous chapters, all Java Virtual Machine implementations must recognize the Java class file format, but individual implementations may also recognize other binary formats.

The Java Virtual Machine specification does not say how the binary data for a type must be produced. Some potential ways to produce binary data for a type are:

- load a Java class file from the local file system
- download a Java class file across a network
- extract a Java class file from a ZIP, JAR, CAB, or other archive file
- extract a Java class file from a proprietary database
- compile a Java source file on the fly into the class file format
- compute the class file data for a type on the fly
- any of the above, but using a binary file format other than the Java class file

Given the binary data for a type, the Java Virtual Machine must process that data to a great enough extent that it can create an instance of class `java.lang.Class`. The virtual machine must parse the binary data into implementation-dependent internal data structures. (See Chapter 5, "The Java Virtual Machine," for a discussion of potential internal data structures for storing class data.) The `Class` instance, the end product of the loading step, serves as an interface between the program and the internal data structures. To access information about a type that is stored in the internal data structures, the program invokes methods on the `Class` instance for that type.

As described in previous chapters, types are loaded either through the primordial class loader or through class loader objects. The primordial class loader, a part of the virtual machine implementation, loads types (including the classes and interfaces of the Java API) in an implementation-dependent way. Class loader objects, instances of subclasses of `java.lang.ClassLoader`, load classes in custom ways. The inner workings of class loader objects are described in more detail later in Chapter 8, "The Linking

Model."

Class loaders (primordial or object) need not wait until a type's first active use before they load the type. Class loaders are allowed to cache binary representations of types, load types early in anticipation of eventual use, or load types together in related groups. If a class loader encounters a problem during early loading, however, it must report that problem (by throwing a subclass of `LinkageError`) only upon the type's first active use. In other words, if a class loader encounters a missing or malformed class file during early loading, it must wait to report that error until the class's first active use by the program. If the class is never actively used by the program, the class loader will never report the error.

Verification

After a type is loaded, it is ready to be linked. The first step of the linking process is verification--ensuring that the type obeys the semantics of the Java language and that it won't violate the integrity of the virtual machine.

Verification is another area in which implementations of the Java Virtual Machine have some flexibility. Implementation designers can decide how and when to verify types. The Java Virtual Machine specification lists all the exceptions that a virtual machine can throw and under what circumstances it must throw them. No matter what kind of trouble a Java Virtual Machine might encounter, there is an exception or error it is supposed to throw. The specification says what exception or error should be thrown in each situation. In some cases, the specification says exactly when the exception or error should be thrown, but usually doesn't dictate precisely how or when the error condition should be detected.

Nevertheless, certain kinds of checks are very likely to take place at certain times in most Java Virtual Machine implementations. For example, during the loading process, the virtual machine must parse the stream of binary data that represents the type and build internal data structures. At this point, certain checks will have to be done just to ensure the initial act of parsing the binary data won't crash the virtual machine. During this parsing, implementations will likely check the binary data to make sure it has the expected overall format. Parsers of the Java class file format might check the magic number, make sure each component is in the right place and of the proper length, verify that the file isn't too short or too long, and so on. Although these checks take place during loading, before the official verification phase of linking, they are still logically part of the verification phase. The entire process of detecting any kind of problem with loaded types is placed under the category of verification.

Another check that likely occurs during loading is making sure that every class except `Object` has a superclass. This may be done during loading because when the virtual machine loads a class, it must also make sure all of the class's superclasses are loaded also. The only way a virtual machine can know the name of a given class's superclass is by peering into the binary data for the class. Since the virtual machine is looking at every class's superclass data during loading anyway, it may as well make this check during the loading phase.

Another check--one that likely occurs after the official verification phase in most implementations--is the verification of symbolic references. As described in earlier chapters, the process of dynamic linking involves locating classes, interfaces, fields, and methods referred to by symbolic references stored in the constant pool, and replacing the symbolic references with direct references. When the virtual machine searches for a symbolically referenced entity (type, field, or method), it must first make sure the entity exists. If the virtual machine finds that the entity exists, it must further check that the referencing type has permission to access the entity, given the entity's access permissions. These checks for existence and

access permission are logically a part of verification, the first phase of linking, but most likely happen during resolution, the third phase of linking. Resolution itself can be delayed until each symbolic reference is first used by the program, so these checks may even take place after initialization.

So what gets checked during the official verification phase? Anything that hasn't already been checked before the official verification phase and that won't get checked after it. Here two lists of some of the things that are good candidates for checking during the official verification phase. This first list is composed of checks that ensure classes are binary compatible with each other:

- checking that final classes are not subclassed
- checking that final methods are not overridden
- if the type being checked is a non-abstract class, checking that all the methods declared in any interfaces implemented by the class are indeed implemented by the class
- making sure no incompatible method declarations (such as two methods that have the same name, the same number, order, and types of parameters, but different return types) appear between the type and its supertypes

Note that while these checks require looking at other types, they only require looking at supertypes. Superclasses need to be initialized before subclasses, so these classes are likely already loaded. Superinterfaces do not need to be initialized when a class that implements them is initialized. However, this verification step will require their loading. (They won't be initialized, just loaded and possibly linked at the option of the virtual machine implementation.) All a class's supertypes will have to be loaded to make sure they are all still binary compatible.

- checking that all constant pool entries are consistent with each other. (For example, the `string_index` item of a `CONSTANT_String_info` entry must be the index of a `CONSTANT_Utf8_info` entry.)
- checking that all special strings contained in the constant pool (class names, field and method names, field and method descriptors) are well-formed
- verifying the integrity of the bytecodes

The most complicated task in the above list is the last one: bytecode verification. All Java Virtual Machines must in some way verify the integrity of the bytecodes for every method they execute. For example, implementations are not allowed to crash because a jump instruction sends the virtual machine beyond the end of a method. They must detect that the jump instruction is invalid through some process of bytecode verification, and throw an error.

Java Virtual Machine implementations are not required to verify bytecodes during the official verification phase of linking. Implementations are free, for example, to verify individual instructions as each instruction is executed. One of the design goals of the Java Virtual Machine instruction set, however, was that it yield bytecodes streams that can be verified all at once by a data flow analyzer. The ability to verify bytecode streams all at once during linking, rather than on the fly as the program runs, gives a big boost to the potential execution speed of Java programs.

When verifying bytecodes via a data flow analyzer, the virtual machine may have to load other classes to ensure that the semantics of the Java language are being followed. For example, imagine a class contained a method that assigned a reference to an instance of `java.lang.Float` to a field of type `java.lang.Number`. In this case, the virtual machine would have to load class `Float` during bytecode verification to make sure it was a subclass of class `Number`. It would have to load `Number` to make sure it wasn't declared final. The virtual machine must not initialize class `Float` at this time, just

load it. `Float` will be initialized only upon its first active use.

For more information on the class verification process, see Chapter 3, "Security."

Preparation

After a Java Virtual Machine has loaded a class and performed whatever verification it chooses to do up front, the class is ready for preparation. During the preparation phase, the Java Virtual Machine allocates memory for the class variables and sets them to default initial values. The class variables are not initialized to their proper initial values until the initialization phase. (No Java code is executed during the preparation step.) During preparation, the Java Virtual Machine sets the newly allocated memory for the class variables to a default value determined by the type of the variable. The default values for the various types are shown in Table 7-1.

Table 7-1. Default initial values for the primitive types

<code>int</code>	<code>0</code>
<code>long</code>	<code>0L</code>
<code>short</code>	<code>(short) 0</code>
<code>char</code>	<code>'\u0000'</code>
<code>byte</code>	<code>(byte) 0</code>
<code>reference</code>	<code>null</code>
<code>float</code>	<code>0.0f</code>
<code>double</code>	<code>0.0d</code>

There is no `boolean` in Table 7-1, because `boolean` is not a primitive type in the Java Virtual Machine. Internally, `boolean` is implemented as an `int`, which gets set to zero (`boolean false`) by default. Therefore, `boolean` class variables are in effect initialized to `boolean false`.

During the preparation phase, Java Virtual Machine implementations may also allocate memory for data structures that are intended to improve the performance of the running program. An example of such a data structure is a method table, which contains a pointer to the data for every method in a class, including those inherited from its superclasses. A method table enables an inherited method to be invoked on an object without a search of superclasses at the point of invocation. Method tables are described in more detail in Chapter 8, "The Linking Model."

Resolution

After a type has been through the first two phases of linking: verification and preparation, it is ready for the third and final phase of linking: resolution. Resolution is the process of locating classes, interfaces, fields, and methods referenced symbolically from a type's constant pool, and replacing those symbolic references with direct references. As mentioned above, this phase of linking is optional until (and unless) each symbolic reference is first used by the program. Constant pool resolution is described in detail in Chapter 8, "The Linking Model."

Initialization

The final step required to ready a class or interface for its first active use is initialization, the process of setting class variables to their proper initial values. As used here, a "proper" initial value is the programmer's desired starting value for a class variable. A proper initial value contrasts with the default initial value given to class variables during preparation. As described above, the virtual machine assigns default values based only on each variable's type. Proper initial values, by contrast, are based on some master plan known only to the programmer.

In Java code, a proper initial value is specified via a class variable initializer or static initializer. A class variable initializer is an equals sign and expression next to a class variable declaration, as in:

begin

```
// On CD-ROM in file classlife/ex1/Example1a.java
class Example1a {

    // "= 3 * (int) (Math.random() * 5.0)" is the class variable
    // initializer
    static int size = 3 * (int) (Math.random() * 5.0);
}
```

end

A static initializer is a block of code introduced by the `static` keyword, as in:

begin

```
// On CD-ROM in file classlife/ex1/Example1b.java
class Example1b {

    static int size;

    // This is the static initializer
    static {

        size = 3 * (int) (Math.random() * 5.0);
    }
}
```

end

All the class variable initializers and static initializers of a class are collected by the Java compiler and

placed into one special method, the *class initialization method*. In the Java class file, the class initialization method is named "`<clinit`". Regular methods of a Java application cannot invoke a class initialization method. This kind of method can only be invoked by the Java Virtual Machine, which invokes it to set a class's static variables to their proper initial values.

Initialization of a class consists of two steps:

1. Initializing the class's direct superclass (if any), if the direct superclass hasn't already been initialized
2. Executing the class's class initialization method, if it has one

When initializing a class's direct superclass, the same two steps listed above must be followed. As a result, the first class that will be initialized will always be `Object`, then all the classes on down the inheritance hierarchy to the class being actively used. Superclasses will be initialized before subclasses.

Initialization of an interface does not require initialization of its superinterfaces. Initialization of an interface consists of only one step:

- Executing the interface's class initialization method, if it has one

The code of a `<clinit()` method does not explicitly invoke a superclass's `<clinit()` method. Before a Java Virtual Machine invokes the `<clinit()` method of a class, therefore, it must make certain the `<clinit()` methods of superclasses have been executed.

Java Virtual Machines must also make sure the initialization process is properly synchronized. If multiple threads need to initialize a class, only one thread should be allowed to perform the initialization while the other threads wait. After the active thread completes the initialization process, it must notify any waiting threads. See Chapter 20, "Thread Synchronization," for information about synchronization, wait and notify.

The Class Initialization Method

As mentioned above, Java compilers place the code for class variable initializers and static initializers into the `<clinit()` method of the class file in the order in which they appear in the class declaration. For example, given this class:

```
begin
```

```
// On CD-ROM in file classlife/ex1/Example1c.java
class Example1c {
```

```
    static int width;
```

```
    static int height = (int) (Math.random() * 2.0);
```

```
    // This is the static initializer
```

```

static {

    width = 3 * (int) (Math.random() * 5.0);

}

}

end

```

The Java compiler generates the following `<clinit()` method:

```

begin

// The code for height's class variable initializer begins here
// Invoke Math.random(), which will push
// a double return value

0 invokestatic #6 <Method double random()

3 ldc2_w #8 <Double 2.0 // Push double constant 2.0

6 dmul // Pop two doubles, multiply, push result

7 d2i // Pop double, convert to int, push int
// Pop int, store into class variable
// height

8 putstatic #5 <Field int height

// The code for the static initializer begins here

11 iconst_3 // Push int constant 3
// Invoke Math.random(), which will push
// a double return value

12 invokestatic #6 <Method double random()

15 ldc2_w #10 <Double 5.0 // Push double constant 5.0

18 dmul // Pop two doubles, multiply, push result

19 d2i // Pop double, convert to int, push int

20 imul // Pop two ints, multiply, push int result
// Pop int, store into class variable
// width

```

```

21 putstatic #7 <Field int width
24 return                // Return void from <clinit method

end

```

This `<clinit()` method first executes the code for `Example1c`'s only class variable initializer, which initializes `height`, then executes the code for the static initializer, which initializes `width`. The initialization is done in this order because the class variable initializer appears textually before the static initializer in the source code of the `Example1c` class.

Not all classes will necessarily have a `<clinit()` method in their class file. If a class declares no class variables, it won't have a `<clinit()` method. If a class declares class variables, but doesn't explicitly initialize them with class variable initializers or static initializers, it won't have a `<clinit()` method. If a class contains only class variable initializers for static final variables, and those class variable initializers use compile-time constant expressions, that class won't have a `<clinit()` method. Only those classes that actually require Java code to be executed to initialize class variables to proper initial values will have a class initialization method.

Here's an example of a class that won't be awarded a `<clinit()` method by the Java compiler:

```

begin

// On CD-ROM in file classlife/ex1/Example1d.java
class Example1d {

    static final int angle = 35;

    static final int length = angle * 2;

}

end

```

Class `Example1d` declares two constants, `angle` and `length`, and initializes them with expressions that are compile-time constants. The compiler knows that `angle` represents the value 35 and `length` represents the value 70. When the `Example1d` class is loaded by a Java Virtual Machine, `angle` and `length` are not stored as class variables in the method area. As a result, no `<clinit()` method is needed to initialize them. The `angle` and `length` fields are not class variables, they are constants, which are treated specially by the Java compiler.

Instead of treating `Example1d`'s `angle` and `length` fields as class variables, the Java compiler places the constant `int` values they represent into the constant pool or bytecode streams of any class that uses them. For example, if a class uses `Example1d`'s `angle` field, that class will not have in its constant pool a symbolic reference to the `angle` field of class `Example1d`. Instead, the class will have operands embedded in its bytecode streams that have the value 35. If the constant value of `angle` were outside the range of a `short` (-32,768 to 32,767), say 35,000, the class would have a

CONSTANT_Integer_info entry in its constant pool with the value of 35,000.

Here is a class that uses both a constant and a class variable from other classes:

begin

```
// On CD-ROM in file classlife/ex1/Example1e.java
class Example1e {
```

```
    // The class variable initializer for symbolicRef uses a symbolic
    // reference to the size class variable of class Example1a
    static int symbolicRef = Example1a.size;
```

```
    // The class variable initializer for localConst doesn't use a
    // symbolic reference to the length field of class Example1d.
    // Instead, it just uses a copy of the constant value 70.
    static int localConst = Example1d.length * (int) (Math.random()
        * 3.0);
```

```
}
```

end

The Java compiler generates the following <clinit() method for class Example1e:

begin

```
// The code for symbolicRef's class variable initializer begins here:
    // Push int value from Example1a.size.
```

```
    // This getstatic instruction refers to a
    // symbolic reference to Example1a.size.
```

```
0 getstatic #9 <Field int size
```

```
    // Pop int, store into class variable
    // symbolicRef
```

```
3 putstatic #10 <Field int symbolicRef
```

```
// The code for localConst's class variable initializer begins here:
```

```
    // Expand byte operand to int, push int
```

```

        // result. This is the local copy of
6  bipush 70          // Example1d's length constant, 70.
        // Invoke Math.random(), which will push
        // a double return value

8  invokestatic #8 <Method double random()
11 ldc2_w #11 <Double 3.0 // Push double constant 3.0
14 dmul              // Pop two doubles, multiply, push result
15 d2i              // Pop double, convert to int, push int
16 imul              // Pop two ints, multiply, push int result
        // Pop int, store into class variable
        // localConst

17 putstatic #7 <Field int localConst
20 return            // Return void from <clinit method

end

```

The `getstatic` instruction at offset zero uses a symbolic reference (in constant pool entry nine) to the `size` field of class `Example1a`. The `bipush` instruction at offset six is followed by a byte that contains the constant value represented by `Example1d.length`. `Example1e`'s constant pool contains no symbolic reference to anything in class `Example1d`.

Interfaces may also be awarded a `<clinit()` method in the class file. All fields declared in an interface are implicitly public, static, and final and must be initialized with a field initializer. If an interface has any field initializers that don't resolve at compile-time to a constant, that interface will have a `<clinit()` method. Here's an example:

```

begin

// On CD-ROM in file classlife/ex1/Example1f.java
interface Example1f {

    int ketchup = 5;

    int mustard = (int) (Math.random() * 5.0);

}

end

```

The Java compiler generates the following `<clinit()` method for interface `Example1f`:

```

begin
// The code for mustard's class variable initializer begins here
// Invoke Math.random(), which will push
// a double return value

0 invokestatic #6 <Method double random()
3 ldc2_w #7 <Double 5.0 // Push double constant 2.0

6 dmul // Pop two doubles, multiply, push result
7 d2i // Pop double, convert to int, push int
// Pop int, store into class variable
// mustard

8 putstatic #5 <Field int mustard

11 return // Return void from <clinit method

end

```

Note that only the `mustard` field is initialized by this `<clinit()` method. Because the `ketchup` field is initialized to a compile-time constant, it is treated specially by the compiler. Although types that use `Example1f.mustard` will contain a symbolic reference to the field, types that use `Example1f.ketchup` will contain a local copy of `ketchup`'s constant value, 5.

Active versus Passive Use

As mentioned above, the Java Virtual Machine initializes types on their first active use or, in the case of classes, upon the first active use of a subclass. Only four activities constitute an active use: invoking a class initialization method on a new instance of a class, creating an array whose element type is the class, invoking a method declared in a class, and accessing a non-constant field declared in a class or interface.

A use of a non-constant field is an active use of only the class or interface that actually declares the field. For example, a field declared in a class may be referred to via a subclass. A field declared in an interface may be referred to via a subinterface or class that implements the interface. These are passive uses of the subclass, subinterface, or class that implements the interface--uses that won't trigger their initialization. They are an active use only of the class or interface in which the field is actually declared. Here's an example that illustrates this principle:

```

begin

// On CD-ROM in file classlife/ex2/NewParent.java
class NewParent {

    static int hoursOfSleep = (int) (Math.random() * 3.0);

```



```

static {
    System.out.println("NewParent was initialized.");
}
}

// On CD-ROM in file classlife/ex2/NewbornBaby.java
class NewbornBaby extends NewParent {

    static int hoursOfCrying = 6 + (int) (Math.random() * 2.0);

    static {
        System.out.println("NewbornBaby was initialized.");
    }
}

// On CD-ROM in file classlife/ex2/Example2.java
class Example2 {

    // Invoking main() is an active use of Example2
    public static void main(String[] args) {

        // Using hoursOfSleep is an active use of Parent, but a
        // passive use of NewbornBaby
        int hours = NewbornBaby.hoursOfSleep;
        System.out.println(hours);
    }

    static {
        System.out.println("Example2 was initialized.");
    }
}

```

end

In the above example, executing `main()` of `Example2` causes only `Example2` and `Parent` to be initialized. `NewbornBaby` is not initialized and need not be loaded. The following text is printed to the standard output:

begin

Example2 was initialized.

NewParent was initialized.

2

end

A use of a field that is both `static` and `final`, and initialized by a compile-time constant expression, is not an active use of the type that declares the field. As mentioned above, the Java compiler resolves references to such fields to a local copy of the constant value that resides either in the referring classes constant pool, in its bytecode streams, or both. Here is an example that illustrates this special treatment of static final fields:

begin

```
// On CD-ROM in file classlife/ex3/Angry.java
interface Angry {
```

```
    String greeting = "Grrrr!";
```

```
    int angerLevel = Dog.getAngerLevel();
```

```
}
```

```
// On CD-ROM in file classlife/ex3/Dog.java
class Dog {
```

```
    static final String greeting = "Woof, woof, world!";
```

```
    static {
```

```
        System.out.println("Dog was initialized.");
```

```
}
```

```

static int getAngerLevel() {

    System.out.println("Anger was initialized");
    return 1;
}
}

// On CD-ROM in file classlife/ex3/Example3.java
class Example3 {

    // Invoking main() is an active use of Example4
    public static void main(String[] args) {

        // Using Angry.greeting is a passive use of Angry
        System.out.println(Angry.greeting);

        // Using Dog.greeting is a passive use of Dog
        System.out.println(Dog.greeting);
    }

    static {
        System.out.println("Example3 was initialized.");
    }
}

end

```

Running the Example4 application yields the following output:

begin

Example3 was initialized.

```
Grrrr!
```

```
Woof, woof, world!
```

```
end
```

Had `Angry` been initialized, the string `"Angry was initialized."` would have been written to the standard output. Likewise, had `Dog` been initialized, the string `"Dog was initialized."` would have been written to the standard output. As you can see from the above output, neither interface `Angry` or class `Dog` were ever initialized during the execution of the `Example3` application.

For more information about this special treatment of static final variables, see Chapter 8, "The Linking Model."

The Lifetime of an Object

Once a class has been loaded, linked, and initialized, it is ready for use. The program can access its static fields, invoke its static methods, or create instances of it. This section describes class instantiation and initialization, activities that take place at the beginning of an object's lifetime, and garbage collection and finalization, activities that mark the end of an object's lifetime.

Class Instantiation

In Java programs, classes can be instantiated explicitly or implicitly. The three ways a class can be instantiated explicitly are with the `new` operator, by invoking `newInstance()` on a `Class` object, or by invoking `clone()` on any existing object. Here is an example showing each way to create a new class instance:

```
begin
```

```
// On CD-ROM in file classlife/ex4/Example4.java
class Example4 implements Cloneable {
```

```
    Example4() {
        System.out.println("Created by invoking newInstance()");
    }
```

```
    Example4(String msg) {
        System.out.println(msg);
    }
```

```
    public static void main(String[] args)
        throws ClassNotFoundException, InstantiationException,
```

```

        IllegalAccessException, CloneNotSupportedException {

    // Create a new Example4 object with the new operator
    Example4 obj1 = new Example4("Created with new.");

    // Get a reference to the Class instance for Example4, then
    // invoke newInstance() on it to create a new Example4 object
    Class myClass = Class.forName("Example4");
    Example4 obj2 = (Example4) myClass.newInstance();

    // Make an identical copy of the the second Example4 object
    Example4 obj3 = (Example4) obj2.clone();
}

}

end

```

When executed, the `Example4` application prints this output:

```

begin

Created with new.
Created by invoking newInstance()

end

```

Besides the three ways listed above to explicitly instantiate objects in Java source code, there are several situations in which objects will be instantiated implicitly--without an explicit `new`, `newInstance()`, or `clone()` appearing in the source.

Possibly the first implicitly instantiated objects of any Java application are the `String` objects that hold the command line arguments. References to these objects, one for each command-line argument, are delivered in the `String` array passed as the sole parameter to the `main()` method of every application.

Two other ways a class can be instantiated implicitly involve the process of class loading. First, for every type a Java Virtual Machine loads, it implicitly instantiates a new `Class` object to represent that type. Second, when the Java Virtual Machine loads a class that contains `CONSTANT_String_info` entries in its constant pool, it may instantiate new `String` objects to represent those constant string literals. The process of transforming a `CONSTANT_String_info` entry in the method area to a

`String` instance on the heap is part of the process of constant pool resolution. This process is described in detail in Chapter 8, "The Linking Model."

Another way objects can be created implicitly is through the process of evaluating an expression that involves the string concatenation operator. If such an expression is not a compile-time constant, intermediate `String` and `StringBuffer` objects will be created in the process of evaluating the expression. Here is an example:

begin

```
// On CD-ROM in file classlife/ex5/Example5.java
class Example5 {
```

```
    public static void main(String[] args) {
```

```
        if (args.length < 2) {
```

```
            System.out.println("Must enter any two args.");
```

```
            return;
```

```
        }
```

```
        System.out.println(args[0] + args[1]);
```

```
    }
```

```
}
```

end

`javac` generates these bytecodes for `Example5`'s `main()` method:

begin

```
0 aload_0      // Push the objref from loc var 0 (args)
1 arraylength  // Pop arrayref, calc array length, push int length

2 iconst_2     // Push int constant 2
               // Pop 2 ints, compare, branch if (length < 2) to
3 if_icmpge 15 // offset 15.
               // Push objref from System.out
6 getstatic #11 <Field java.io.PrintStream out
               // Push objref of string literal
```

```

9 ldc #1 <String "Must enter any two args."
    // Pop objref to String param, objref to System.out,
    // invoke println()

11 invokevirtual #12 <Method void println(java.lang.String)

14 return    // Return void from main()
    // Push objref from System.out

15 getstatic #11 <Field java.io.PrintStream out

// The string concatenation operation begins here
    // Allocate mem for new StringBuffer object, and
    // initialize mem to default initial values, push
    // objref to new object

18 new #6 <Class java.lang.StringBuffer

21 dup    // Duplicate objref to StringBuffer object

22 aload_0    // Push ref from loc var 0 (args)

23 iconst_0    // Push int constant 0
    // Pop int, arrayref, push String at arrayref[int],

24 aaload    // which is args[0]
    // Pop objref, invoke String's class method
    // valueOf(), passing it the objref to the args[0]
    // String object. valueOf() calls toString() on the
    // ref, and returns (and pushes) the result, which
    // happens to be the original args[0] String. In this
    // case, the stack will look precisely the same
    // before and after this instruction is executed.
    // Thus here, the 1.1 javac compiler has
    // over-enthusiastically generated an unnecessary
    // instruction.

25 invokestatic #14 <Method java.lang.String valueOf(
    java.lang.Object)

```

```

        // Pop objref to args[0] String, objref of the
        // StringBuffer object, invoke <init() method on the
        // StringBuffer object passing the args[0] objref as
        // the only parameter.
28 invokespecial #9 <Method java.lang.StringBuffer(java.lang.String)
31 aload_0      // Push objref from loc var 0 (args)
32 iconst_1     // Push int constant 1
                // Pop int, arrayref, push String at arrayref[int],
33 aaload       // which is args[1]
                // Pop objref to args[1] String, objref of the
                // StringBuffer object (there's still another objref
                // to this same object on the stack because of the
                // dup instruction above), invoke append() method on
                // StringBuffer object, passing args[1] as the only
                // parameter. append() will return an objref to this
                // StringBuffer object, which will be pushed back
                // onto the stack.
34 invokevirtual #10 <Method java.lang.StringBuffer append(java.lang.String)
                // Pop objref to StringBuffer (pushed by append()),
                // invoke toString() on it, which returns the value
                // of the StringBuffer as a String object. Push
                // objref of String object.
37 invokevirtual #13 <Method java.lang.String toString()
// The string concatenation operation is now complete

        // Pop objref of concatenated String, objref of
        // System.out that was pushed by the getstatic
        // instruction at offset 15. Invoke println() on
        // System.out, passing the concatenated String as
        // the only parameter.

```



```

40 invokevirtual #12 <Method void println(java.lang.String)
43 return          // Return void from main()

end

```

The bytecodes for `Example5`'s `main()` method contain three implicitly generated `String` objects and one implicitly generated `StringBuffer` object. References to two of the `String` objects appear as arguments passed to `main()` in the `args` array, which are pushed onto the stack by the `aaload` instructions at offset 23 and 33. The `StringBuffer` is created with the `new` instruction at offset 18 and initialized with the `invokespecial` instruction at offset 28. The final `String`, which represents the concatenation of `args[0]` and `args[1]`, is created by calling `toString()` on the `StringBuffer` object via the `invokevirtual` instruction at offset 37.

When the Java Virtual Machine creates a new instance of a class, either implicitly or explicitly, it first allocates memory on the heap to hold the object's instance variables. Memory is allocated for all variables declared in the object's class and in all its superclasses, including instance variables that are hidden. As described in Chapter 5, "The Java Virtual Machine," memory for other implementation-dependent components of an object's image on the heap, such as a pointer to class data in the method area, are also likely allocated at this point. As soon as the virtual machine has set aside the heap memory for a new object, it immediately initializes the instance variables to default initial values. These are the same values shown above in Table 7-1 as default initial values for class variables.

Once the virtual machine has allocated memory for the new object and initialized the instance variables to default values, it is ready to give the instance variables their proper initial values. The Java Virtual Machine uses two techniques to do this, depending upon whether the object is being created because of a `clone()` invocation. If the object is being created because of a `clone()`, the virtual machine copies the values of the instance variables of the object being cloned into the new object. Otherwise, the virtual machine invokes an *instance initialization method* on the object. The instance initialization method initializes the object's instance variables to their proper initial values.

The Java compiler generates at least one instance initialization method for every class it compiles. In the Java class file, the instance initialization method is named "`<init.`" For each constructor in the source code of a class, the Java compiler generates one `<init()` method. If the class declares no constructors explicitly, the compiler generates a default no-arg constructor that just invokes the superclass's no-arg constructor. As with any other constructor, the compiler creates an `<init()` method in the class file that corresponds to this default constructor.

An `<init()` method can contain three kinds of code: an invocation of another `<init()` method, code that implements any instance variable initializers, and code for the body of the constructor. If a constructor begins with an explicit invocation of another constructor in the same class (a `this()` invocation) its corresponding `<init()` method will be composed of two parts:

- an invocation of the same-class `<init()` method
- the bytecodes that implement the body of the corresponding constructor

If a constructor does not begin with a `this()` invocation and the class is not `Object`, the `<init()` method will have three components:

- an invocation of a superclass `<init()` method
- the bytecodes for any instance variable initializers
- the bytecodes that implement the body of the corresponding constructor

If a constructor does not begin with a `this()` invocation and the class is `Object`, the first component in the above list is missing. Because `Object` has no superclass, its `<init()` method can't begin with a superclass `<init()` method invocation.

If a constructor begins with an explicit invocation of a superclass constructor (a `super()` invocation), its `<init()` method will invoke the corresponding superclass `<init()` method. For example, if a constructor begins with an explicit invocation of the "`super(int, String)` constructor," the corresponding `<init()` method will begin by invoking the superclass's "`<init(int, String)`" method. If a constructor does not begin with an explicit `this()` or `super()` invocation, the corresponding `<init()` method will invoke the superclass's no-arg `<init()` method by default.

Here's an example with three constructors, numbered one through three:

begin

```
// On CD-ROM in file classlife/ex6/Example6.java
class Example6 {

    private int width = 3;

    // Constructor one:
    // This constructor begins with a this() constructor invocation,
    // which gets compiled to a same-class <init() method
    // invocation.
    Example6() {
        this(1);
        System.out.println("Example6(), width = " + width);
    }

    // Constructor two:
    // This constructor begins with no explicit invocation of another
    // constructor, so it will get compiled to an <init() method
    // that begins with an invocation of the superclass's no-arg
    // <init() method.
```

```

Example6(int width) {
    this.width = width;
    System.out.println("Example6(int), width = " + width);
}

```

```

// Constructor three:
// This constructor begins with super(), an explicit invocation
// of the superclass's no-arg constructor. Its <init() method
// will begin with an invocation of the superclass's no-arg
// <init() method.

```

```

Example6(String msg) {
    super();
    System.out.println("Example6(String), width = " + width);
    System.out.println(msg);
}

```

```

public static void main(String[] args) {
    String msg
        = "The Agapanthus is also known as Lilly of the Nile.";
    Example6 one = new Example6();
    Example6 two = new Example6(2);
    Example6 three = new Example6(msg);
}

```

```

}

```

end

When executed, the Example6 application prints this output:

begin

```

Example6(int), width = 1
Example6(), width = 1

```

```

Example6(int), width = 2

```

```
Example6(String), width = 3
```

The Agapanthus is also known as Lilly of the Nile.

```
end
```

The bytecodes for Example6's no-arg <init() method (the <init() method that corresponds to constructor one) are:

```
begin
```

```
// The first component, the same-class <init() invocation, begins  
// here:
```

```
0 aload_0      // Push the objref from loc var 0 (this)  
1 iconst_1     // Push int constant 1  
               // Pop int and objref, invoke <init() method on  
               // objref (this), passing the int (a 1) as the  
               // only parameter.  
2 invokespecial #12 <Method Example6(int)
```

```
// The second component, the body of the constructor, begins
```

```
// here:
```

```
               // Push objref from System.out  
5 getstatic #16 <Field java.io.PrintStream out  
               // Allocate mem for new StringBuffer object, and  
               // initialize mem to default initial values, push  
               // objref to new object  
8 new #8 <Class java.lang.StringBuffer  
11 dup         // Duplicate objref to StringBuffer object  
               // Push objref to String literal from constant pool  
12 ldc #1 <String "Example6(), width = "  
               // Pop objref to literal String, pop objref of the  
               // StringBuffer object, invoke <init() method on the  
               // StringBuffer object passing the args[0] objref as  
               // the only parameter.
```

```

14 invokespecial #14 <Method java.lang.StringBuffer(
    java.lang.String)
17 aload_0      // Push objref from loc var 0 (this)
                // Pop this reference, Push int value of width field
18 getfield #19 <Field int width
                // Pop int (width), pop objref (StringBuffer object),
                // invoke append() on StringBuffer object passing the
                // width int as the only parameter. append() will add
                // the string representation of the int to the end of
                // the buffer, and return an objref to the same
                // StringBuffer object.
21 invokevirtual #15 <Method java.lang.StringBuffer append(int)
                // Pop objref to StringBuffer (pushed by append()),
                // invoke toString() on it, which returns the value
                // of the StringBuffer as a String object. Push
                // objref of String object.
24 invokevirtual #18 <Method java.lang.String toString()
                // Pop objref of String, pop objref of System.out
                // that was pushed by the getstatic instruction at
                // offset 5. Invoke println() on System.out,
                // passing the String as the only parameter:
                // System.out.println("Example6(), width = "
                //      + width);
27 invokevirtual #17 <Method void println(java.lang.String)
30 return      // Return void from <init()

end

```

Note that the `<init()` method for constructor one begins with an invocation of a same-class `<init()` method, then executes the body of the corresponding constructor. Because the constructor begins with a `this()` invocation, its corresponding `<init()` method doesn't contain bytecodes for the instance variable initializer.

The bytecodes for Example6's <init() method that takes an int parameter (the <init() method that corresponds to constructor two) is:

begin

// The first component, the superclass <init() invocation, begins
// here:

```
0 aload_0      // Push the objref from loc var 0 (this)
                // Pop objref (this), invoke the superclass's
                // no-arg<init() method on objref.
1 invokespecial #11 <Method java.lang.Object()
```

// The second component, the instance variable initializers, begins
// here:

```
4 aload_0      // Push the objref from loc var 0 (this)
5 iconst_3     // Push int constant 3
                // Pop int (3), pop objref (this), store 3 into
                // width instance variable of this object
6 putfield #19 <Field int width
```

// The third component, the body of the constructor, begins
// here:

```
9 aload_0      // Push the objref from loc var 0 (this)
10 iload_1     // Push int from loc var 1 (int param width)
                // Pop int (param width), pop objref (this), store
                // int param value into width field of this object:
                // this.width = width
11 putfield #19 <Field int width
                // Push objref from System.out
14 getstatic #16 <Field java.io.PrintStream out
                // Allocate mem for new StringBuffer object, and
                // initialize mem to default initial values, push
                // objref to new object
```

```

17 new #8 <Class java.lang.StringBuffer
20 dup          // Duplicate objref to StringBuffer object
                // Push objref to String literal from constant pool
21 ldc #3 <String "Example6(int), width = "
                // Pop objref to literal String, pop objref of the
                // StringBuffer object, invoke <init() method on the
                // StringBuffer object passing the args[0] objref as
                // the only parameter.
23 invokespecial #14 <Method java.lang.StringBuffer(
    java.lang.String)
26 iload_1      // Push int from loc var 1 (int param width)
                // Pop int (width), pop objref (StringBuffer object),
                // invoke append() on StringBuffer object passing the
                // width int as the only parameter. append() will add
                // the string representation of the int to the end of
                // the buffer, and return an objref to the same
                // StringBuffer object.
27 invokevirtual #15 <Method java.lang.StringBuffer append(int)
                // Pop objref to StringBuffer (pushed by append()),
                // invoke toString() on it, which returns the value
                // of the StringBuffer as a String object. Push
                // objref of String object.
30 invokevirtual #18 <Method java.lang.String toString()
                // Pop objref of String, pop objref of System.out
                // that was pushed by the getstatic instruction at
                // offset 14. Invoke println() on System.out,
                // passing the String as the only parameter:
                // System.out.println("Example6(int), width = "
                //      + width);
33 invokevirtual #17 <Method void println(java.lang.String)

```

```
36 return          // Return void from <init()
```

```
end
```

The `<init()` method for constructor two has three components. First it has an invocation of the superclass's (`Object`'s) no-arg `<init()` method. The compiler generated this invocation by default, because no explicit `super()` invocation appears as the first statement in the body of constructor two. Following the superclass `<init()` invocation is the second component: the bytecodes for width's instance variable initializer. Third, the `<init()` method contains the bytecodes for the body of constructor two.

The bytecodes for Example6's `<init()` method that takes a `String` parameter (the `<init()` method that corresponds to constructor three) are:

```
begin
```

```
// The first component, the superclass <init() invocation, begins
// here:
```

```
0 aload_0          // Push the objref from loc var 0 (this)
                   // Pop objref (this), invoke the superclass's
                   // no-arg<init() method on objref.
```

```
1 invokespecial #11 <Method java.lang.Object()
```

```
// The second component, the instance variable initializers, begins
```

```
// here:
```

```
4 aload_0          // Push the objref from loc var 0 (this)
5 iconst_3         // Push int constant 3
                   // Pop int (3), pop objref (this), store 3 into
                   // width instance variable of this object
```

```
6 putfield #19 <Field int width
```

```
// The third component, the body of the constructor, begins
```

```
// here:
```

```
           // Push objref from System.out
```

```
9 getstatic #16 <Field java.io.PrintStream out
```

```
           // Allocate mem for new StringBuffer object, and
```



```

        // initialize mem to default initial values, push
        // objref to new object
12 new #8 <Class java.lang.StringBuffer
15 dup        // Duplicate objref to StringBuffer object
               // Push objref to String literal from constant pool
16 ldc #2 <String "Example6(String), width = "
               // Pop objref to literal String, pop objref of the
               // StringBuffer object, invoke <init() method on the
               // StringBuffer object passing the args[0] objref as
               // the only parameter.
18 invokespecial #14 <Method java.lang.StringBuffer(
    java.lang.String)
21 aload_0    // Push objref from loc var 0 (this)
               // Pop this reference, Push int value of width field
22 getfield #19 <Field int width
               // Pop int (width), pop objref (StringBuffer object),
               // invoke append() on StringBuffer object passing the
               // width int as the only parameter. append() will add
               // the string representation of the int to the end of
               // the buffer, and return an objref to the same
               // StringBuffer object.
25 invokevirtual #15 <Method java.lang.StringBuffer append(int)
               // Pop objref to StringBuffer (pushed by append()),
               // invoke toString() on it, which returns the value
               // of the StringBuffer as a String object. Push
               // objref of String object.
28 invokevirtual #18 <Method java.lang.String toString()
               // Pop objref of String, pop objref of System.out
               // that was pushed by the getstatic instruction at
               // offset 9. Invoke println() on System.out,

```

```

        // passing the String as the only parameter:
        // System.out.println("Example6(String), width = "
        //      + width);

31 invokevirtual #17 <Method void println(java.lang.String)
        // Push objref from System.out

34 getstatic #16 <Field java.io.PrintStream out

37 aload_1      // Push objref from loc var 1 (param msg)
        // Pop objref of String, pop objref of System.out
        // that was pushed by the getstatic instruction at
        // offset 37. Invoke println() on System.out,
        // passing the String as the only parameter:
        // System.out.println(msg);

38 invokevirtual #17 <Method void println(java.lang.String)

41 return      // Return void from <init()

end

```

The `<init()` method for constructor three has the same three components as the `<init()` method for constructor two: a superclass `<init()` invocation, the bytecodes for width's initializer, and the bytecodes for the constructor body. One difference between constructor two and three is that constructor two does not begin with an explicit `this()` or `super()` invocation. As a result, the compiler places an invocation of the superclass's no-arg `<init()` method in constructor two's `<init()` method. By contrast, constructor three begins with an explicit `super()` invocation, which the compiler converts into the corresponding superclass `<init()` invocation in constructor three's `<init()` method.

For every class except `Object`, an `<init()` method must begin with an invocation of another `<init()` method belonging either to the same class or to the direct superclass. `<init()` methods are not allowed to catch exceptions thrown by the `<init()` method they invoke. If a subclass `<init()` method invokes a superclass `<init()` method that completes abnormally, for example, the subclass `<init()` method must also complete abnormally.

Garbage Collection and Finalization of Objects

As mentioned in earlier chapters, implementations of the Java Virtual Machine must have some kind of automatic storage management strategy for the heap, most likely a garbage collector. Applications can allocate memory for objects via the explicit and implicit ways described earlier in this chapter, but cannot explicitly free that memory. When an object becomes unreferenced by the application, the virtual machine may reclaim (garbage collect) that memory. Implementations can decide when to garbage collect unreferenced objects—even whether to garbage collect them at all. Java Virtual Machine implementations are not required to free memory occupied by unreferenced objects.

If a class declares a method named `finalize()` that returns `void`, the garbage collector will execute that method (called a "finalizer") once on an instance of that class, before it frees the memory space occupied by that instance. Here's an example of a class that declares a finalizer:

begin

```
// On CD-ROM in file classlife/ex7/Finale.java
class Finale {

    protected void finalize() {

        System.out.println("A Finale object was finalized.");

        //...

    }

    //...

}
```

end

Because a finalizer is a regular Java method, it may be invoked directly by the application. Such a direct invocation will not affect the automatic invocation of the finalizer by the garbage collector. The garbage collector may invoke an object's finalizer at most once, sometime after the object becomes unreferenced and before the memory it occupies is reused. If the object becomes referenced again (resurrected) as a result of executing the finalizer code, and then becomes unreferenced again later, the garbage collector must not automatically invoke the finalizer a second time.

Any exceptions thrown by the `finalize()` method during its automatic invocation by the garbage collector are ignored. The garbage collector may invoke `finalize()` methods in any order, using any thread, or even concurrently via multiple threads. Finalization is described in more detail in Chapter 9, "Garbage Collection."

Unloading and Finalization of Classes

In many ways, the lifetime of a class in the Java Virtual Machine is similar to the lifetime of an object. The virtual machine creates and initializes objects, allows the program to use the objects, and optionally garbage collects the objects after they are no longer referenced by the program. Similarly, the virtual machine loads, links, and initializes classes, allows the program to use the classes, and optionally unloads the classes after they are no longer referenced by the program. Just as the Java Virtual Machine will run an object's `finalize()` method (if its class declares one) before it garbage collects the object, the Java Virtual Machine will run a class's `classFinalize()` method (if the class declares one) before it unloads the class.

Garbage collection and unloading of classes is important in the Java Virtual Machine because Java programs can be dynamically extended at runtime by loading types through class loader objects. All loaded types occupy memory space in the method area. If a Java application continuously loads types through class loader objects, the memory footprint of the method area will continuously grow. If some of the dynamically loaded types are needed only temporarily, the memory space occupied by those types can be freed by unloading the types after they are no longer needed.

The way in which a Java Virtual Machine can tell whether a dynamically loaded type is still needed by the application is similar to the way it tells whether an object is still needed by the program. If the application has no references to the type, then the type can't affect the future course of computation. The type is unreachable and can be garbage collected.

Types loaded through the primordial class loader will always be reachable and never be unloaded. Only dynamically loaded types--those loaded through class loader objects--can become unreachable and be unloaded by the virtual machine. A dynamically loaded type is unreachable if its `Class` instance is found to be unreachable through the normal process of garbage collecting the heap.

There are two ways a `Class` instance of a dynamically loaded type can be reachable through the normal process of garbage collection. First and most obviously, a `Class` instance will be reachable if the application holds an explicit reference to it. Second, a `Class` instance will be reachable if there is a reachable object on the heap whose type data in the method area refers to the `Class` instance. As mentioned in Chapter 5, "The Java Virtual Machine," implementations must be able to locate the type data in the method area for an object's class, given only a reference to the object. For this reason, the image of an object on the heap likely includes some kind of pointer to its type data in the method area. From the type data, the virtual machine must be able to locate the `Class` instances for the object's class, all its superclasses, and all its superinterfaces. See Figure 7-2 for a graphical depiction of this way of "reaching" `Class` instances.




Figure 7-2. Reaching `Class` instances through reachable objects.

Figure 7-2 shows the paths a garbage collector must traverse from a reachable object of class `MyThread` through the type data in the method area to find reachable `Class` instances. In this figure, objects on the heap are shown as light gray circles; type data in the method area is shown as dark gray rectangles. The `MyThread` class has the following declaration:

```
begin
```

```
// On CD-ROM in file classlife/ex8/MyThread.java
class MyThread extends Thread implements Cloneable {

}
```

end

From the reachable `MyThread` object (shown in the bottom right hand corner of the figure), the garbage collector follows a pointer to `MyThread`'s type data, where it finds:

- A reference to `MyThread`'s `Class` instance on the heap
- A pointer to the type data for `MyThread`'s direct superinterface, `Cloneable`
- A pointer to the type data for `MyThread`'s direct superclass, `Thread`

From the type data for `Cloneable`, the garbage collector finds:

- A reference to `Cloneable`'s `Class` instance on the heap

From the type data for `Thread`, the garbage collector finds:

- A reference to `Thread`'s `Class` instance on the heap
- A pointer to the type data for `Thread`'s direct superinterface, `Runnable`
- A pointer to the type data for `Thread`'s direct superclass, `Object`

From the type data for `Runnable`, the garbage collector finds:

- A reference to `Runnable`'s `Class` instance on the heap

From the type data for `Object`, the garbage collector finds:

- A reference to `Object`'s `Class` instance on the heap

Thus, given only a reference to a reachable instance of class `MyThread`, the garbage collector is able to "reach" the `Class` instances for `MyThread` and all its supertypes: `Cloneable`, `Thread`, `Runnable`, and `Object`.

If a class declares a method, named `classFinalize()`, that is static, takes no parameters, and returns `void`, the Java Virtual Machine will execute that method (called a "class finalizer") before it unloads the class. The class finalizer may have any access permission. Here's an example:

begin

```
// On CD-ROM in file classlife/ex7/GrandFinale.java
class GrandFinale {

    static void classFinalize() {

        System.out.println(
```

```
        "Class GrandFinale was finally finalized.");  
        //...  
    }  
    //...  
}  
  
end
```

Similar to the `finalize()` method, the `classFinalize()` method will be called only once by the Java Virtual Machine. If a class's `classFinalize()` method resurrects the class (makes it reachable again), the next time the class is unloaded its `classFinalize()` method won't be invoked.

An example of dynamically loaded classes becoming unreachable and available for unloading is given at the end of Chapter 8, "The Linking Model."

On the CD-ROM

The CD-ROM contains the source code examples from this chapter in the `classlife` directory.

The Resources Page

For more information about the material presented in this chapter, visit the resources page:
<http://www.artima.com/insidejvm/classlife.html>

McGraw-Hill

A Division of The McGraw-Hill Companies



Chapter Eight

The Linking Model

From the programmer's perspective, one of the most important aspects of Java's architecture to understand is the linking model. As mentioned in earlier chapters, Java's linking model allows you to design class loader objects that extend your application in custom ways at run-time. Through class loader objects, your application can load and dynamically link to classes and interfaces that were unknown or did not even exist when your application was compiled.

The engine that drives Java's linking model is the process of resolution. The previous chapter described in detail all the various stages in the lifetime of a class, except for resolution. This chapter looks at resolution in depth, and shows how the process of resolution fits in with dynamic extension. It gives an overview of the linking model, explains constant pool resolution, describes method tables, shows how to write and use class loaders, and gives several examples.

Dynamic Linking and Resolution

When you compile a Java program, you get a separate class file for each class or interface in your program. Although the individual class files may appear to be independent, they actually harbor symbolic connections to one another and to the class files of the Java API. When you run your program, the Java Virtual Machine loads your program's classes and interfaces and hooks them together in a process of *dynamic linking*. As your program runs, the Java Virtual Machine builds an internal web of interconnected classes and interfaces.

A class file keeps all its symbolic references in one place, the constant pool. Each class file has a constant pool, and each class or interface loaded by the Java Virtual Machine has an internal version of its constant pool. The internal constant pool is an implementation-specific data structure that maps to the constant pool in the class file. Thus, after a type is initially loaded, all the symbolic references from the type reside in the type's internal constant pool.

At some point during the running of a program, if a particular symbolic reference is to be used, it must be *resolved*. Resolution is the process of finding the entity identified by the symbolic reference and replacing the symbolic reference with a direct reference. Because all symbolic references reside in the constant pool, this process is often called *constant pool resolution*.

As described in Chapter 6, "The Java Class File," the constant pool is organized as a sequence of items. Each item has a unique index, much like an array element. A symbolic reference is one kind of item that may appear in the constant pool. Java Virtual Machine instructions that use a symbolic reference specify the index in the constant pool where the symbolic reference resides. For example, the `getstatic`

opcode, which pushes the value of a static field onto the stack, is followed in the bytecode stream by an index into the constant pool. The constant pool entry at the specified index, a `CONSTANT_Fieldref_info` entry, reveals the fully qualified name of the class in which the field resides, and the name and type of the field.

Keep in mind that the Java Virtual Machine contains a separate internal constant pool for each class and interface it loads. When an instruction refers to the fifth item in the constant pool, it is referring to the fifth item in the constant pool for the current class, the class that defined the method the Java Virtual Machine is currently executing.

Several instructions, from the same or different methods, may refer to the same constant pool entry, but each constant pool entry is resolved only once. After a symbolic reference has been resolved for one instruction, subsequent attempts to resolve it by other instructions take advantage of the hard work already done, and use the same direct reference resulting from the original resolution.

Linking involves not only the replacement of symbolic references with direct ones, it also involves checking for correctness and permission. As mentioned in Chapter 7, "The Lifetime of a Class," the checking of symbolic references for existence and access permission (one aspect of the full verification phase) is likely performed during resolution. For example, when a Java Virtual Machine resolves a `getstatic` instruction to a field of another class, the Java Virtual Machine checks to make sure that:

- The other class exists.
- This class has permission to access the other class.
- The named field exists in the other class.
- The field has the expected type (symbolic references to fields include the field type).
- This class has permission to access the field.
- That the field is really static--a class variable and not an instance variable.

If any of these checks fail, an error is thrown and resolution fails. Otherwise, the symbolic reference is replaced by the direct reference and resolution succeeds.

As described in Chapter 7, "The Lifetime of a Class," different implementations of the Java Virtual Machine are permitted to perform resolution at different times during the execution of a program. An implementation may choose to link everything up front by following all symbolic references from the initial class, then all symbolic references from subsequent classes, until every symbolic reference has been resolved. In this case, the application would be completely linked before its `main()` method was ever invoked. This approach is called *early resolution*. Alternatively, an implementation may choose to wait until the very last minute to resolve each symbolic reference. In this case, the Java Virtual Machine would resolve a symbolic reference only when it is first used by the running program. This approach is called *late resolution*. Other implementations could choose a resolution strategy in-between these two extremes.

Although a Java Virtual Machine implementation has some freedom in choosing when to resolve symbolic references, every Java Virtual Machine must give the outward impression that it uses late resolution. No matter when a particular Java Virtual Machine performs its linking, it will always throw any error that results from attempting to resolve a symbolic reference at the point in the execution of the program where the symbolic reference was actually used for the first time. In this way, it will always appear to the user as if the linking were late. If a Java Virtual Machine does early linking, and during early linking discovers that a class file is missing, it won't report the class file missing by throwing the appropriate error until later in the program when something in that class file is actually used. If the class

is never used by the program, the error will never be thrown.

Resolution and Dynamic Extension

In addition to simply linking types at run-time, Java applications can decide at run-time which types to link. Java's architecture allows Java programs to be *dynamically extended*, the process of deciding at run-time other types to use, loading them, and using them. You can dynamically extend a Java application by creating a class loader object and using it to load types that are not part of your original application, but are determined by your running application. An example of dynamic extension is a Java-capable web browser loading class files for applets from across a network. When the browser starts, it doesn't know what class files it will be loading across the network.

To dynamically extend a Java application, you must include a class loader as part of your application. To create your own class loader, you write a subclass of `java.lang.ClassLoader` and implement the `loadClass()` method. When the `loadClass()` method of a class loader successfully loads a type, it returns a `java.lang.Class` object to represent the newly loaded (and optionally, verified, prepared, resolved, and initialized) type. Here's the declaration of the `loadClass()` method:

```
begin
// A method declared in class java.lang.ClassLoader:
protected abstract Class loadClass(String name, boolean resolve)
    throws ClassNotFoundException;
end
```

The class loader takes care of loading, linking, and initializing types. Admittedly the name `ClassLoader` might lead you to believe it is responsible only for loading, and not linking or initialization, but the name `ClassLoaderLinkerAndInitializer` is a bit unwieldy. Also, the purpose of creating your own class loader is to customize the load phase while still using the primordial link and initialization mechanisms. Class loaders differ not in how they link or initialize types, but in how they load types. In other words, class loaders distinguish themselves by the manner in which they produce a binary form for a type given a fully qualified name.

The `loadClass()` method encompasses all three steps of loading, linking, and initialization of a new type. You pass the fully qualified name of the requested type to `loadClass()` in the parameter `name`. If you want linking and initialization to be performed at this time, you pass `true` in the parameter `resolve`. If `loadClass()` is invoked with `resolve` set to `false`, it will load but not link and initialize the type. From your programs, you will likely invoke `loadClass()` with the `resolve` parameter always set to `true`. Normally, the `loadClass()` method is invoked with the `resolve` parameter set to `false` only by the virtual machine itself. (More on this later in this chapter.)

As mentioned in earlier chapters, each class loader--primordial or object--is awarded its own name space, which is populated by the names of the types it has loaded. Dynamic extension and name spaces are both supported by one aspect of the process of resolution: the way a virtual machine chooses a class loader when it resolves a symbolic reference to a type.

When the resolution of a constant pool entry requires loading a type, the virtual machine uses the same class loader that loaded the referencing type to load the referenced type. For example, imagine a `Cat`

class refers via a symbolic reference in its constant pool to a type named `Mouse`. Assume `Cat` was loaded by a class loader object. When the virtual machine resolves the reference to `Mouse`, it checks to see if `Mouse` has been loaded into the name space to which `Cat` belongs. (It checks to see if the class loader that loaded `Cat` has previously loaded a type named `Mouse`.) If not, the virtual machine requests `Mouse` from the same class loader that loaded `Cat`. This is true even if a class named `Mouse` had previously been loaded into a different name space. When a symbolic reference from a type loaded by the primordial class loader is resolved, the Java Virtual Machine uses the primordial class loader to load the referenced type. When a symbolic reference from a type loaded by a class loader object is resolved, the Java Virtual Machine uses the same class loader object to load the referenced type.

Constant Pool Resolution

This section describes the details of resolving each type of constant pool entry, including the errors that may be thrown during resolution. If an error is thrown during resolution, the error is seen as being thrown by the instruction that refers to the constant pool entry being resolved. Besides the errors described here, individual instructions that trigger the resolution of a constant pool entry may cause other errors to be thrown. For example, `getstatic` causes a `CONSTANT_Fieldref_info` entry to be resolved. If the entry is resolved successfully, the virtual machine performs one additional check: it makes sure the field is actually static (a class variable and not an instance variable). If the field is not static, the virtual machine throws an error. Any extra errors that may be thrown during resolution besides those described in this section are described for each individual instruction in Appendix A.

Resolution of `CONSTANT_Class_info` Entries

Of all the types of constant pool entries, the most complicated to resolve is `CONSTANT_Class_info`. This type of entry is used to represent symbolic references to classes (including array classes) and interfaces. Several instructions, such as `new` and `anewarray`, refer directly to `CONSTANT_Class_info` entries. Other instructions, such as `putfield` or `invokevirtual`, refer indirectly to `CONSTANT_Class_info` entries through other types of entry. For example, the `putfield` instruction refers to a `CONSTANT_Fieldref_info` entry. The `class_index` item of a `CONSTANT_Fieldref_info` gives the constant pool index of a `CONSTANT_Class_info` entry.

The details of resolving a `CONSTANT_Class_info` entry vary depending on whether or not the type is an array and whether the referencing type (the one that contains in its constant pool the `CONSTANT_Class_info` entry being resolved) was loaded via the primordial class loader or a class loader object.

Array Classes

A `CONSTANT_Class_info` entry refers to an array class if its `name_index` refers to a `CONSTANT_Utf8_info` string that begins with a left bracket, as in `"[I"`. As described in Chapter 6, "The Java Class File," internal array names contain one left bracket for each dimension, followed by an element type. If the element type begins with an `"L"`, as in `"Ljava.lang.Integer;"`, the array is an array of references. Otherwise, the element type is a primitive type, such as `"I"` for `int` or `"D"` for `double`, and the array is an array of primitive types.

In the case of an array of references, the virtual machine resolves the element type. For example, if resolving an array class with the name `"[[Ljava.lang.Integer;"`, the virtual machine would make certain class `java.lang.Integer` is loaded, linked, and initialized.

The end product of the resolution of a symbolic reference to an array class is a `Class` instance that represents the array class. If a `Class` instance has already been created for the array type being resolved, that same `Class` instance is used. Otherwise, the virtual machine creates a new `Class` instance to represent the newly resolved array type.

Non-Array Classes and Interfaces

A `CONSTANT_Class_info` entry whose `name_index` refers to a `CONSTANT_Utf8_info` string that doesn't begin with a left bracket is a symbolic reference to non-array class or an interface. Resolution of this kind of symbolic reference is a multiple step process.

The Java Virtual Machine performs the same basic steps, described below as Steps 1 and 2a through 2e, to resolve any symbolic reference (any `CONSTANT_Class_info` entry) to a non-array class or interface. In Step 1, the type is loaded. In Steps 2a through 2e, the type is linked and initialized. The precise way in which the virtual machine performs these steps depends on whether the referencing type was loaded via the primordial class loader or a class loader object.

Step 1. Load the Type and any Superclasses

Resolution of a non-array class or interface begins by making sure the type is loaded into the current name space. If the virtual machine determines that the referenced type hasn't yet been loaded into the current name space, it passes the fully qualified name of the type to a class loader. The class loader loads the type, parses the binary data into internal data structures, and creates a `Class` instance.

Once the referenced type is loaded in, the virtual machine can peer into its binary data. If the type is a class and not `java.lang.Object`, the virtual machine determines from the class's data the fully qualified name of the class's direct superclass. The virtual machine then checks to see if the superclass has been loaded into the current name space. If not, it loads the superclass. Once that class comes in, the virtual machine can again peer into its binary data to find its superclass. This process repeats all the way up to `Object`.

Through Step 1, the Java Virtual Machine makes sure a type is loaded, and if the type is a class, that all its superclasses are loaded. During this step, these types are not linked and initialized--just loaded.

If the referencing type (the one that contains in its constant pool the `CONSTANT_Class_info` entry being resolved) was loaded through the primordial class loader, the virtual machine uses the primordial class loader to load the referenced type (and its superclasses, if any). Otherwise, the referencing type was loaded through a class loader object, and the virtual machine uses the same class loader object to load the referenced type.

The Java Virtual Machine maintains a list of the names of all the types already loaded by each class loader. Each of these lists forms a name space inside the Java Virtual Machine. The virtual machine uses this list during resolution to determine whether a class has already been loaded by a particular class loader. If during resolution, the Java Virtual Machine encounters a symbolic reference to a class that was not previously loaded by the appropriate class loader, that class loader will be asked to load the class. If the appropriate class loader is a class loader object, the Java Virtual Machine will make the load request by invoking the class loader object's `loadClass()` method. On the other hand, if the Java Virtual Machine encounters a symbolic reference to a class previously loaded by the same class loader object, the class loader object will not be asked to load it again. Instead, the class previously loaded will automatically be used again. This helps ensure that only one class with a given name is loaded by a

particular class loader.

If the referencing type was loaded via the primordial class loader, the virtual machine will check the primordial class loader's name space to see if the class was already loaded. If not, the virtual machine will use the primordial class loader to load the referenced type in an implementation-dependent way. If the referenced type is a class, the virtual machine will make sure all the class's superclasses are loaded as well. If any of these classes haven't yet been loaded into the primordial loader's name space, the virtual machine will use the primordial class loader to load those classes.

If the referencing type was loaded via a class loader object, the virtual machine invokes `loadClass()` on the class loader object that loaded the referencing type, passing in parameter `name` the fully qualified name of the referenced type and setting the `resolve` flag to `true`. When `loadClass()` returns, it will have performed both Step 1 (loading) and Steps 2a through 2e (linking and initialization). (It will perform Steps 2a through 2e because the `resolve` flag is set to `true`.) Nonetheless, the `loadClass()` method will perform Step 1 completely before embarking on Steps 2a through 2e.

When a class loader object's `loadClass()` method is invoked, it has two ways to do its job. One way is to use the primordial class loader by invoking `findSystemClass()`. In this case, the primordial loader will load, link, and initialize the requested type. Alternatively, `loadClass()` can produce the binary data in its own custom way, then call `defineClass()` and--if `true` was passed in `loadClass()`'s `resolve` parameter--call `resolveClass()`. In this case, the class loader object would produce the binary form for the requested type in its own custom way, then used the primordial class loader's standard mechanism to link and initialize the type. Here are the declarations for `findSystemClass()`, `defineClass()`, and `resolveClass()`:

```
begin
```

```
// Three methods declared in class java.lang.ClassLoader:
```

```
protected final Class findSystemClass(String name)
```

```
    throws ClassNotFoundException;
```

```
protected final Class defineClass(String name, byte data[],
```

```
    int offset, int length);
```

```
protected final void resolveClass(Class c);
```

```
end
```

If the `loadClass()` method invokes `findSystemClass()`, the class loader object is in effect asking the virtual machine to use the primordial loader to resolve the reference completely. The `findSystemClass()` method accepts the fully qualified name of the type to resolve in its `name` parameter. If the primordial class loader is able to successfully resolve the reference, it returns a reference to a `Class` instance that represents the referenced type. The `loadClass()` method would then just return a reference to the same `Class` instance. If the primordial class loader is unsuccessful, however, `findSystemClass()` will complete abnormally by throwing a `ClassNotFoundException`. In this case, the `loadClass()` method can either throw `ClassNotFoundException` itself or attempt to load the type in its custom way.

Types brought in through the primordial class loader via `findSystemClass()` are marked as having no class loader object. Any types they refer to will be requested directly from the primordial loader.

If the `loadClass()` method produces the stream of binary data for the type in its own custom way (such as if it downloads the type's class file across a network), it must invoke `defineClass()`, passing a reference to a `byte` array containing the type's binary data in the Java class file format.

Invoking `defineClass()` will cause the virtual machine to do two things. First, the virtual machine will parse the binary data into internal data structures and create a `Class` instance to represent the type. Second, if the type is a class and not `java.lang.Object`, the virtual machine will retrieve the name of the class's direct superclass. It will then check to see if the direct superclass has been previously loaded by this class loader object (the one whose `loadClass()` method just invoked `defineClass()`). This is a check to see if the direct superclass already resides in the current name space. If the direct superclass has already been loaded by this class loader object, the `defineClass()` method returns the `Class` instance for the freshly loaded type. Otherwise, the virtual machine invokes `loadClass()` on the same class loader object, passing the fully qualified name of the direct superclass in the `name` parameter and setting the `resolve` flag to `false`. Because the `resolve` flag is `false`, the direct superclass (and recursively, all its superclasses) will be loaded, but not linked or initialized.

The reason the `resolve` flag is set to `false` when `loadClass()` is invoked by `defineClass()` is because this is Step 1: loading of the type and any superclasses. Linking and initialization, which would occur if `resolve` were set to `true`, is saved for Step 2.

When `defineClass()` has made certain the type and all its superclasses, if any, have been loaded into the current name space. It returns the new `Class` instance that represents the type. Step 1 is then complete.

During Step 1, the virtual machine may throw the following errors:

- If the binary data can't be produced (for example, if a class file of the appropriate name can't be found), the virtual machine throws `NoClassDefFoundError`.
- If the binary data is produced, but isn't of the proper structure or a recognized version (such as if the minor or major version number of a Java class file is too high), the virtual machine throws `NoClassDefFoundError`.
- If the binary data is produced and well formed, but doesn't contain the sought after class or interface (such as if file `CuteKitty.class` is discovered to contain class `HungryTiger` instead of `CuteKitty`), the virtual machine throws `NoClassDefFoundError`.
- If the class doesn't contain a superclass and isn't class `Object` itself, the virtual machine throws a `ClassFormatError`. (Note that this check has to be done here, during the loading step, because that one piece of information--the symbolic reference to the superclass--is needed by the virtual machine during this step. During Step 1, the virtual machine must load in all the superclasses recursively.)
- If a class appears as its own superclass, the virtual machine throws `ClassCircularityError`.

Step 2. Link and Initialize the Type and any Superclasses

At this point, the type being referred to by the `CONSTANT_Class_info` entry being resolved has

been loaded, but not linked or initialized. In addition, if the type being resolved is a class, all its superclasses have been loaded, but not necessarily linked or initialized. Some of the superclasses may be initialized at this point, because they may have been initialized during earlier resolutions.

As described in Chapter 7, "The Lifetime of a Class," superclasses must be initialized before subclasses. If the virtual machine is resolving a reference to a class (not an interface), it must make sure that the superclasses have been initialized, starting with `Object` and proceeding down the inheritance hierarchy to the referenced class. (Note that this is the opposite order in which they were loaded in Step 1.) If a type hasn't yet been linked, it must be linked before it is initialized.

Step 2a. Verify the Type

Step 2 begins with the official verification phase of linking, described in Chapter 7, "The Lifetime of a Class." As mentioned in Chapter 7, the process of verification may require that the virtual machine load new types. For example, when verifying a type, the Java Virtual Machine must check all a type's superclasses and superinterfaces to make sure they are all binary compatible with the type. If the type being verified has superinterfaces, they may not yet be loaded. During Step 1, the virtual machine makes sure all the type's superclasses are loaded, but not its superinterfaces. Thus, verification may require that a type's superinterfaces be loaded. At this point, the superinterfaces would be loaded and possibly linked, but definitely not initialized.

In addition, when verifying bytecodes, the Java Virtual Machine may need to load types to ensure the bytecodes are adhering to the semantics of the Java language. For example, if a reference to an instance of a particular class is assigned to a variable with a declared type of a different class, the virtual machine would have to load both types to make sure one is a subclass of the other. As with the superinterfaces described above, these classes would at this point be loaded and possibly linked, but definitely not initialized.

If during the verification process the Java Virtual Machine uncovers trouble, it throws `VerifyError`.

Step 2b. Prepare the Type

After the official verification phase is complete, the type must be prepared. As described in Chapter 7, "The Lifetime of a Class," during preparation the virtual machine allocates memory for class variables and implementation-dependent data structures such as method tables. If the virtual machine discovers during this step that a class that is not declared abstract contains an abstract method, the virtual machine throws `AbstractMethodError`.

Optional Step 2c. Resolve the Type

At this point, the type has been loaded, verified and prepared. As described in Chapter 7, "The Lifetime of a Class," a Java Virtual Machine implementation may optionally resolve the type at this point. Keep in mind that at this stage in the resolution process, Steps 1, 2a, and 2b have been performed on a referenced type to resolve a `CONSTANT_Class_info` entry in the constant pool of a referencing type. Step 2c is the resolution of symbolic references contained in the referenced type, not the referencing type.

For example, if the virtual machine is resolving a symbolic reference from class `Cat` to class `Mouse`, the virtual machine performs Steps 1, 2a, and 2b on class `Mouse`. At this stage of resolving the symbolic reference to `Mouse` contained in the constant pool of `Cat`, the virtual machine could

optionally (as Step 2c) resolve all the symbolic references contained in the constant pool for `Mouse`. If `Mouse`'s constant pool contains a symbolic reference to class `Cheese`, for example, the virtual machine could load and optionally link (but not initialize) `Cheese` at this time. The virtual machine mustn't attempt initialize `Cheese` here because `Cheese` is not being actively used. (Of course, `Cheese` may in fact have already been actively used elsewhere, so it have would been already be loaded into this name space, linked, and initialized.)

As mentioned earlier in this chapter, if an implementation does perform Step 2c at this point in the resolution process (early resolution), it must not report any errors until the symbolic references are actually used by the running program. For example, if during the resolution of `Mouse`'s constant pool, the virtual machine can't find class `Cheese`, it won't throw a `NoClassDefFound` error until (and unless) `Cheese` is actually used by the program.

Step 2d. Initialize the Type

At this point, the type has been loaded, verified, prepared and optionally resolved. At long last, the type is ready for initialization. As defined in Chapter 7, "The Lifetime of a Class," initialization consists of two steps. The initialization of the type's superclasses in top down order, if the type has any superclasses, and the execution of the type's class initialization method, if it has one. Step 2d just consists of executing the class initialization method, if one exists. Because Steps 2a through 2e are performed for all the referenced type's superclasses, from the top down, Step 2d will occur for superclasses before it occurs for subclasses.

If the class initialization method completes abruptly by throwing some exception that isn't a subclass of `Error`, the virtual machine throws `ExceptionInInitializerError` with the thrown exception as a parameter to the constructor. Otherwise, if the thrown exception is already a subclass of `Error`, that error is thrown. If the virtual machine can't create a new `ExceptionInInitializerError` because there isn't enough memory, it throws an `OutOfMemoryError`.

Step 2e. Check Access Permission

Lastly, after loading, linking, and initialization are complete, the virtual machine checks for access permission. If the referencing type does not have permission to access the referenced type, the virtual machine throws a `IllegalAccessError`. Step 2e is another activity that is logically part of verification, but that is performed at some other time than the official verification phase of Step 2a. In this case, the check for access permission is delayed to just after initialization. (Actually, the error can be detected at some other time, such as in Step 2a. But regardless of when or how it is detected, an access violation must be reported only after the Steps 1 and 2a through 2d have been successfully completed. As mentioned in Chapter 7, "The Lifetime of a Class," the Java Virtual Machine specification says when errors should be thrown, but doesn't dictate exactly when they should be detected.) Once this check is complete, Step 2e--and the entire process of resolving the `CONSTANT_Class_info` entry--is complete.

If an error occurred in any of Steps 1 through 2e, the resolution of the type fails. But if all went well up until the access permission check of Step 2e, the type is still usable in general, just not usable by the referencing type. If an error occurred before the access permission check, however, the type is unusable and must be marked as such or discarded.

The way a class loader object's `loadClass()` method accomplishes Steps 2a through 2e is by

invoking `resolveClass()`, passing the reference to the `Class` instance returned by `defineClass()`. Invoking this method causes all of Steps 2a through 2e to be performed on the type and all its superclasses, starting from the furthest up superclass that hasn't yet been linked and initialized and proceeding down the inheritance hierarchy.

Remember that when `defineClass()` was invoked, the virtual machine brought in the type's superclasses by calling `loadClass()` with the `resolve` flag set to `false`. So those superclasses were only loaded. They weren't linked or initialized. Only if the superclasses were already linked and initialized before `defineClass()` was invoked would they be already linked and initialized before `resolveClass()` is invoked.

Note that the `loadClass()` method should only invoke `resolveClass()` on types it imported via calling `defineClass()`, which loads a type but doesn't link or initialize it. (In addition, `loadClass()` should invoke `resolveClass()` only if the `resolve` parameter is `true`.) As mentioned above, `findSystemClass()` takes care of linking and initialization as well as loading. You needn't invoke `resolveClass()` on types returned by `findSystemClass()`.

Resolution of `CONSTANT_Fieldref_info` Entries

To resolve a constant pool entry of type `CONSTANT_Fieldref_info`, the virtual machine must first resolve the `CONSTANT_Class_info` entry specified in the `class_index` item. Therefore, any error that can be thrown because of the resolution of a `CONSTANT_Class_info` can be thrown during the resolution of a `CONSTANT_Fieldref_info`. If resolution of the `CONSTANT_Class_info` entry succeeds, the virtual machine checks to make sure the field exists and that the current class has permission to access it.

If the virtual machine discovers there is no field with the proper name and type in the referenced class, the virtual machine throws `NoSuchFieldError`. Otherwise, if the field exists, but the current class doesn't have permission to access the field, the virtual machine throws `IllegalAccessError`.

The virtual machine marks the entry as resolved and places a direct reference to the field in the data for the constant pool entry.

Resolution of `CONSTANT_Methodref_info` Entries

To resolve a constant pool entry of type `CONSTANT_Methodref_info`, the virtual machine must first resolve the `CONSTANT_Class_info` entry specified in the `class_index` item. Therefore, any error that can be thrown because of the resolution of a `CONSTANT_Class_info` can be thrown during the resolution of a `CONSTANT_Methodref_info`. If the resolution of the `CONSTANT_Class_info` entry succeeds, the virtual machine checks to make sure the method exists and that the current class has permission to access it.

If the virtual machine discovers there is no method with the proper name, return type, and number and types of parameters in the referenced class, the virtual machine throws `NoSuchMethodError`. Otherwise, if the method exists, but the current class doesn't have permission to access the method, the virtual machine throws `IllegalAccessError`.

The virtual machine marks the entry as resolved and places a direct reference to the method in the data

for the constant pool entry.

Resolution of `CONSTANT_InterfaceMethodref_info` Entries

To resolve a constant pool entry of type `CONSTANT_InterfaceMethodref_info`, the virtual machine must first resolve the `CONSTANT_Class_info` entry specified in the `class_index` item. Therefore, any error that can be thrown because of the resolution of a `CONSTANT_Class_info` can be thrown during the resolution of a `CONSTANT_InterfaceMethodref_info`. If the resolution of the `CONSTANT_Class_info` entry succeeds, the virtual machine checks to make sure the method exists. (The virtual machine need not check to make sure the current class has permission to access the method, because all methods declared in interfaces are implicitly public.)

If the virtual machine discovers there is no method with the proper name, return type, and number and types of parameters in the referenced interface, the virtual machine throws `NoSuchMethodError`.

The virtual machine marks the entry as resolved and places a direct reference to the method in the data for the constant pool entry.

Resolution of `CONSTANT_String_info` Entries

To resolve an entry of type `CONSTANT_String_info`, the virtual machine must place a reference to an *interned* `String` object in the data for the constant pool entry being resolved. The `String` object (an instance of class `java.lang.String`) must have the character sequence specified by the `CONSTANT_Utf8_info` entry identified by the `string_index` item of the `CONSTANT_String_info`.

Each Java Virtual Machine must maintain an internal list of references to `String` objects that have been "interned" during the course of running the application. Basically, a `String` object is said to be interned simply if it appears in the virtual machine's internal list of interned `String` objects. The point of maintaining this list is that any particular sequence of characters is guaranteed to appear in the list no more than once.

To intern a sequence of characters represented by a `CONSTANT_String_info` entry, the virtual machine checks to see if the sequence of characters is already in the list of interned strings. If so, the virtual machine uses the reference to the existing, previously-interned `String` object. Otherwise, the virtual machine creates a new `String` object with the proper character sequence and adds a reference to that `String` object to the list. To complete the resolution process for a `CONSTANT_String_info` entry, the virtual machine places the reference to the interned `String` object in the data of the constant pool entry being resolved.

In your Java programs, you can intern a string by invoking the `intern()` method of class `String`. All literal strings are interned via the process of resolving `CONSTANT_String_info` entries. If a string with the same sequence of Unicode characters has been previously interned, the `intern()` method returns a reference to the matching already-interned `String` object. If the `intern()` method is invoked on a `String` object that contains a sequence of characters that has not yet been interned, that object itself will be interned. The `intern()` method will return a reference to the same `String` object upon which it was invoked.

Here's an example:

begin

```
// On CD-ROM in file linking/ex1/Example1.java
class Example1 {

    // Assume this application is invoked with one command-line
    // argument, the string "Hi!".

    public static void main(String[] args) {

        // argZero, because it is assigned a String from the command
        // line, does not reference a string literal. This string
        // is not interned.
        String argZero = args[0];

        // literalString, however, does reference a string literal.
        // It will be assigned a reference to a String with the value
        // "Hi!" by an instruction that references a
        // CONSTANT_String_info entry in the constant pool. The
        // "Hi!" string will be interned by this process.
        String literalString = "Hi!";

        // At this point, there are two String objects on the heap
        // that have the value "Hi!". The one from arg[0], which
        // isn't interned, and the one from the literal, which
        // is interned.

        System.out.print("Before interning argZero: ");
        if (argZero == literalString) {
            System.out.println("they're the same string object!");
        }
        else {
```

```

        System.out.println("they're different string objects.");
    }

    // argZero.intern() returns the reference to the literal
    // string "Hi!" that is already interned. Now both argZero
    // and literalString have the same value. The non-interned
    // version of "Hi!" is now available for garbage collection.
    argZero = argZero.intern();
    System.out.print("After interning argZero: ");
    if (argZero == literalString) {
        System.out.println("they're the same string object!");
    }
    else {
        System.out.println("they're different string objects.");
    }
}

}

end

```

When executed with the string "Hi!" as the first command-line argument, the `Example1` application prints the following:

```

begin
Before interning argZero: they're different string objects.
After interning argZero: they're the same string object!
end

```

Resolution of Other Types of Entries

The `CONSTANT_Integer_info`, `CONSTANT_Long_info`, `CONSTANT_Float_info`, `CONSTANT_Double_info` entries contain the constant values they represent within the entry itself. These are straightforward to resolve. To resolve this kind of entry, many virtual machine implementations may not have to do anything but use the value as is. Other implementations, however, may choose to do some processing on it. For example, a virtual machine on a little-endian machine could choose to swap the byte order of the value at resolve time.

Entries of type `CONSTANT_Utf8_info` and `CONSTANT_NameAndType_info` are never referred to directly by instructions. They are only referred to via other types of entries, and resolved when those referring entries are resolved.

Compile-Time Resolution of Constants

As mentioned in Chapter 7, "The Lifetime of a Class," references to static final variables initialized to a compile-time constant are resolved at compile-time to a local copy of the constant value. This is true for constants of all the primitive types and of type `java.lang.String`.

This special treatment of constants facilitates two features of the Java language. First, local copies of constant values enable static final variables to be used as `case` expressions in `switch` statements. The two virtual machine instructions that implement `switch` statements in bytecodes, `tableswitch` and `lookupswitch`, require the `case` values in-line in the bytecode stream. These instructions do not support run-time resolution of `case` values. See Chapter 16, "Control Flow," for more information about these instructions.

The other motivation behind the special treatment of constants is conditional compilation. Java supports conditional compilation via `if` statements whose expressions resolve to a compile-time constant. Here is an example:

begin

```
// On CD-ROM in file linking/ex2/AntHill.java
class AntHill {
```

```
    static final boolean debug = true;
```

```
}
```

```
// On CD-ROM in file linking/ex2/Example2.java
class Example2 {
```

```
    public static void main(String[] args) {
```

```
        if (AntHill.debug) {
```

```
            System.out.println("Debug is true!");
```

```
        }
```

```
    }
```

```
}
```

end

Because of the special treatment of primitive constants, the Java compiler can decide whether or not to

include the body of the `if` statement in `Example2.main()` depending upon the value of `Anthill.debug`. Because `Anthill.debug` is `true` in this case, `javac` generates bytecodes for `Example2`'s `main()` method that include the body of the `if` statement, but not a check of `Anthill.debug`'s value. The constant pool of `Example2` has no symbolic reference to class `Anthill`. Here are the bytecodes for `main()`:

```
begin

    // Push objref from System.out
0  getstatic #8 <Field java.io.PrintStream out

    // Push objref to literal string "Debug is true!"
3  ldc #1 <String "Debug is true!"

    // Pop objref (to a String), pop objref(to
    // System.out), invoke println() on System.out
    // passing the string as the only parameter:
    // System.out.println("Debug is true!");
5  invokevirtual #9 <Method void println(java.lang.String)
8  return      // return void

end
```

If the reference to `Anthill.debug` were resolved at run-time, the compiler would always need to include a check of `Anthill.debug`'s value and the body of the `if` statement just in case value of `Anthill.debug` ever changed. The value of `Anthill.debug` can't change after it is compiled, of course, because it is declared as `final`. Still, you could change the source code of `Anthill` and recompile `Anthill`, but not recompile `Example2`.

Because the reference to `Anthill.debug` is resolved at compile-time the compiler can conditionally compile out the body of the `if` statement if `Anthill.debug` is discovered to be `false`. Note that this means you can't change the behavior of the `Example2` application just by setting `Anthill` to `false` and recompiling only `Anthill`. You have to recompile `Example2` as well.

`Example3`, shown below, is `Example2` with its name changed to `Example3` and compiled with an `Anthill` that has `debug` set to `false`:

```
begin

// On CD-ROM in file linking/ex3/Anthill.java
class Anthill {

    static final boolean debug = false;

}
```

```
// On CD-ROM in file linking/ex3/Example3.java
class Example3 {

    public static void main(String[] args) {

        if (AntHill.debug) {

            System.out.println("Debug is true!");

        }

    }

}

end
```

Here are the bytecodes generated by `javac` for `Example3`'s `main()` method:

```
begin
0 return    // return void
end
```

As you can see, the Java compiler has brazenly eliminated the entire `if` statement found in `Example3.main()`. There is not even a hint of the `println()` invocation in this very short bytecode sequence.

Direct References

The ultimate goal of constant pool resolution is to replace a symbolic reference with a direct reference. The form of symbolic references is well-defined in Chapter 6, "The Java Class File," but what form do direct references take? As you might expect, the form of direct references is yet another decision of the designers of individual Java Virtual Machine implementations. Nevertheless, there are some characteristics likely to be common among most implementations.

Direct references to types, class variables, and class methods are likely native pointers into the method area. A direct reference to a type can simply point to the implementation-specific data structure in the method area that holds the type data. A direct reference to a class variable can point to the class variable's value stored in the method area. A direct reference to a class method can point to a data structure in the method area that contains the data needed to invoke the method. For example, the data structure for a class method could include information such as whether or not the method is native. If the method is native, the data structure could include a function pointer to the dynamically linked native method implementation. If the method is not native, the data structure could include the method's bytecodes, `max_stack`, `max_locals`, and so on. If there is a just-in-time-compiled version of the method, the data structure could include a pointer to that just-in-time-compiled native code.

Direct references to instance variables and instance methods are offsets. A direct reference to an instance variable is likely the offset from the start of the object's image to the location of the instance variable. A direct reference to an instance method is likely an offset into a method table.

Using offset to represent direct references to instance variables and instance methods depends on a predictable ordering of the fields in a class's object image and the methods in a class's method table. Although implementation designers may choose any way of placing instance variables into an object image or methods into a method table, they will almost certainly use the same way for all types. Therefore, in any one implementation, the ordering of fields in an object and methods in a method table is defined and predictable.

As an example, consider this hierarchy of three classes and one interface:

begin

```
// On CD-ROM in file linking/ex4/Friendly.java
interface Friendly {
```

```
    void sayHello();
```

```
    void sayGoodbye();
```

```
}
```

```
// On CD-ROM in file linking/ex4/Dog.java
class Dog {
```

```
    // How many times this dog wags its tail when
```

```
    // saying hello.
```

```
    private int wagCount = ((int) (Math.random() * 5.0)) + 1;
```

```
    void sayHello() {
```

```
        System.out.print("Wag");
```

```
        for (int i = 0; i < wagCount; ++i) {
```

```
            System.out.print(", wag");
```

```
        }
```

```
        System.out.println(".");
```

```
}
```

```
public String toString() {  
  
    return "Woof!";  
  
}  
}
```

```
// On CD-ROM in file linking/ex4/CockerSpaniel.java  
class CockerSpaniel extends Dog implements Friendly {
```

```
    // How many times this Cocker Spaniel woofs when saying hello.
```

```
    private int woofCount = ((int) (Math.random() * 4.0)) + 1;
```

```
    // How many times this Cocker Spaniel wimpers when saying
```

```
    // goodbye.
```

```
    private int wimperCount = ((int) (Math.random() * 3.0)) + 1;
```

```
    public void sayHello() {
```

```
        // Wag that tail a few times.
```

```
        super.sayHello();
```

```
        System.out.print("Woof");
```

```
        for (int i = 0; i < woofCount; ++i) {
```

```
            System.out.print(", woof");
```

```
        }
```

```
        System.out.println("!");
```

```
    }
```

```
    public void sayGoodbye() {
```



```
        System.out.print("Wimper");  
        for (int i = 0; i < wimperCount; ++i) {  
            System.out.print(", wimper");  
        }  
        System.out.println(".");  
    }  
}
```

```
// On CD-ROM in file linking/ex4/Cat.java  
class Cat implements Friendly {
```

```
    public void eat() {  
  
        System.out.println("Chomp, chomp, chomp.");  
    }
```

```
    public void sayHello() {  
  
        System.out.println("Rub, rub, rub.");  
    }
```

```
    public void sayGoodbye() {  
  
        System.out.println("Scamper.");  
    }
```

```
    protected void finalize() {  
  
        System.out.println("Meow!");  
    }
```

```

    }
}

end

```

Assume these types are loaded into a Java Virtual Machine that organizes objects by placing the instance variables declared in superclasses into the object image before those declared in subclasses, and by placing the instance variables for each individual class in their order of appearance in the class file. Assuming there are no instance variables in class `Object`, the object images for `Dog`, `CockerSpaniel`, and `Cat` would appear as shown in Figure 8-1.




Figure 8-1. Some object images.

In this figure, the object image for `CockerSpaniel` best illustrates this particular virtual machine's approach to laying out objects. The instance variable for `Dog`, the superclass, appears before the instance variables for `CockerSpaniel`, the subclass. The instance variables of `CockerSpaniel` appear in order of declaration: `woofCount` first, then `wimperCount`.

Note that the `wagCount` instance variable appears at offset one in both `Dog` and `CockerSpaniel`. In this implementation of the Java Virtual Machine, a symbolic reference to the `wagCount` field of class `Dog` would be resolved to direct reference that is an offset of one. Regardless of whether the actual object being referred to was a `Dog`, a `CockerSpaniel`, or any other subclass of `Dog`, the `wagCount` instance variable would always appear at offset one in the object image.

A similar pattern emerges in method tables. A method table entry is associated in some way with data structures in the method area that contain sufficient data to enable the virtual machine to invoke the method. Assume that in the Java Virtual Machine implementation being described here, method tables are arrays of native pointers into the method area. The data structures that the method table entries point to are similar to the data structures described above for class methods. Assume that the particular Java Virtual Machine implementation that loads these types organizes its method tables by placing methods

for superclasses into the method table before those for subclasses, and by placing pointers for each class in the order the methods appear in the class file. The exception to the ordering is that methods that are overridden by a subclass appear in the slot where the overridden method first appears in a superclass.

The way this virtual machine would organize the method table for class `Dog` is shown in Figure 8-2. In this figure, the method table entries that point to methods defined in class `Object` are shown in dark gray. Entries that point to methods defined in `Dog` are shown in light gray.




Figure 8-2. The method table for class `Dog`.

Note that only non-private instance methods appear in this method table. Class methods, which are invoked via the `invokestatic` instruction, need not appear here, because they are statically bound and don't need the extra indirection of a method table. Private methods and instance initialization methods need not appear here because they are invoked via the `invokespecial` instruction and are therefore statically bound. Only methods that are invoked with `invokevirtual` or `invokeinterface` appear in this method table. See Chapter 19, "Method Invocation and Return," for a discussion of the different invocation instructions.

By looking at the source code, you can see that `Dog` overrides the `toString()` method defined in class `Object`. In `Dog`'s method table, the `toString()` method appears only once, in the same slot (offset seven) in which it appears in the method table for `Object`. The pointer residing at offset seven in `Dog`'s method table points to the data for `Dog`'s implementation of `toString()`. In this implementation of the Java Virtual Machine, the pointer to the method data for `toString()` will appear at offset seven for every method table of every class.

Below the methods declared in `Object`, which appear first in this method table, come the methods declared in `Dog` that don't override any method in `Object`. There is only one such method,

`sayHello()`, which has the method table offset 11. All of `Dog`'s subclasses will either inherit or override this implementation of `sayHello()`, and some version of `sayHello()` will always appear at offset 11 of any subclass of `Dog`.

Figure 8-3 shows the method table for `CockerSpaniel`. Note that because `CockerSpaniel` declares `sayHello()` and `sayGoodbye()`, the pointers for those methods point to the data for `CockerSpaniel`'s implementation of those methods. Because `CockerSpaniel` inherits `Dog`'s implementation of `toString()`, the pointer for that method (which is still at offset seven) points the data for `Dog`'s implementation of that method. `CockerSpaniel` inherits all other methods from `Object`, so the pointers for those methods point directly into `Object`'s type data. Note also that `sayHello()` is sitting at offset eleven, the same offset it has in `Dog`'s method table.




Figure 8-3. The method table for class `CockerSpaniel`.

When the virtual machine resolves a symbolic reference (a `CONSTANT_Methodref_info` entry) to the `toString()` method of any class, the direct reference is method table offset seven. When the virtual machine resolves a symbolic reference to the `sayHello()` method of `Dog` or any of its subclasses, the direct reference is method table offset eleven. When the virtual machine resolves a symbolic reference to the `sayGoodbye()` method of `CockerSpaniel` or any of its subclasses, the direct reference is the method table offset twelve.

Once a symbolic reference to an instance method is resolved to a method table offset, the virtual machine must still actually invoke the method. To invoke an instance method, the virtual machine goes through the object to get at the method table for the object's class. As mentioned in Chapter 5, "The Java Virtual Machine," given a reference to an object, every virtual machine implementation must have some way to get at the type data for that object's class. In addition, given a reference to an object, the method table (a part of the type data for the object's class) is usually very quickly accessible. (One potential scheme is shown in Figure 5-7.) Once the virtual machine has the method table for the object's class, it uses the offset to find the actual method to invoke. Voila!

The virtual machine can always depend on method table offsets when it has a reference of a class type (a `CONSTANT_Methodref_info` entry). If the `sayHello()` method appears in offset seven in class `Dog`, it will appear in offset seven in any subclass of `Dog`. The same is not true, however, if the reference is of an interface type (a `CONSTANT_InterfaceMethodref_info` entry). With direct references to instance methods accessed through an interface reference there is no guaranteed method table offset. Consider the method table for class `Cat`, shown in Figure 8-4.




Figure 8-4. The method table for class `Cat`.

Note that both `Cat` and `CockerSpaniel` implement the `Friendly` interface. A variable of type `Friendly` could hold a reference to a `Cat` object or a `CockerSpaniel` object. With that reference, your program could invoke `sayHello()` or `sayGoodbye()` on a `Cat`, a `CockerSpaniel`, or any other object whose class implements the `Friendly` interface. The `Example4` application demonstrates this:

```
begin

// On CD-ROM in file linking/ex4/Example4.java
class Example4 {

    public static void main(String[] args) {

        Dog dog = new CockerSpaniel();
```

```

    dog.sayHello();

    Friendly fr = (Friendly) dog;

    // Invoke sayGoodbye() on a CockerSpaniel object through a
    // reference of type Friendly.
    fr.sayGoodbye();

    fr = new Cat();

    // Invoke sayGoodbye() on a Cat object through a reference
    // of type Friendly.
    fr.sayGoodbye();
}
}

end

```

In Example4, local variable `fr` invokes `sayGoodbye()` on both a `CockerSpaniel` object and a `Cat` object. The same constant pool entry, a `CONSTANT_InterfaceMethodref_info` entry, is used to invoke this method on both objects. But when the virtual machine resolves the symbolic reference to `sayHello()`, it can't just save a method table offset and expect that offset to always work in future uses of the constant pool entry.

The trouble is that classes that implement the `Friendly` interface aren't guaranteed to have a common superclass that also implements `Friendly`. As a result, the methods declared in `Friendly` aren't guaranteed to be in the same place in all method tables. If you compare the method table for `CockerSpaniel` against the method table for `Cat`, for example, you'll see that in `CockerSpaniel`, `sayHello()`'s pointer occupies offset 11. But in `Cat`, `sayHello()` occupies offset 12. Likewise, `CockerSpaniel`'s `sayGoodbye()` method pointer resides in offset 12, but `Cat`'s `sayGoodbye()` method pointer resides at offset 13.

Thus, whenever the Java Virtual Machine invokes a method from an interface reference, it must search the method table of the object's class until it finds the appropriate method. This is why invoking instance methods on interface references can be significantly slower than invoking instance methods on class references. Virtual machine implementations can attempt to be smart, of course, about how they search through a method table. For example, an implementation could save the last index at which they found

the method and try there first the next time. Or an implementation could build data structures during preparation that help them search through method tables given an interface reference. Nevertheless, invoking a method given an interface reference will almost certainly be to some extent slower than invoking a method given a class reference.

_quick Instructions

The Java Virtual Machine specification describes a technique used by Sun's Java Virtual Machine implementation to speed up the interpretation of bytecodes. This technique is likely common to many Java Virtual Machine implementations, because many implementations are based on Sun's code. In this scheme, opcodes that refer to constant pool entries are replaced by a "_quick" opcode when the constant pool entry is resolved. When the virtual machine encounters a _quick instruction, it knows the constant pool entry is already resolved and can therefore execute the instruction faster.

The core instruction set of the Java Virtual Machine consists of 200 single-byte opcodes, all of which are described in Appendix A, "Instruction Set by Opcode Mnemonic." These 200 opcodes are the only opcodes you will ever see in class files. Sun's virtual machine implementation uses another 25 single-byte opcodes internally, the "_quick" opcodes.

For example, when Sun's virtual machine resolves a constant pool entry referred to by an `ldc` instruction (opcode value 0x12), it replaces the `ldc` opcode byte in the bytecode stream with an `ldc_quick` instruction (opcode value 0xcb). This technique is part of the process of replacing a symbolic reference with a direct reference in Sun's virtual machine.

For some instructions, in addition to overwriting the normal opcode with a _quick opcode, the virtual machine overwrites the operands of the instruction with data that represents the direct reference. For example, in addition to replacing an `invokevirtual` opcode with an `invokevirtual_quick`, the virtual machine also puts the method table offset and the number of arguments into the two operand bytes that follow every `invokevirtual` instruction. Placing the method table offset in the bytecode stream following the `invokevirtual_quick` opcode saves the virtual machine the time it would take to look up the offset in the resolved constant pool entry.

Example: The Linking of the Salutation Application

As an example of Java's linking model, consider the `Salutation` application shown below:

```
begin
```

```
// On CD-ROM in file linking/ex5/Salutation.java
class Salutation {
```

```
    private static final String hello = "Hello, world!";
```

```
    private static final String greeting = "Greetings, planet!";
```

```
    private static final String salutation = "Salutations, orb!";
```

```
    private static int choice = (int) (Math.random() * 2.99);
```

```

public static void main(String[] args) {

    String s = hello;

    if (choice == 1) {
        s = greeting;
    }

    else if (choice == 2) {
        s = salutation;
    }

    System.out.println(s);
}
}

end

```

Assume that you have asked a Java Virtual Machine to run `Salutation`. When the virtual machine starts, it attempts to invoke the `main()` method of `Salutation`. It quickly realizes, however, that it can't invoke `main()`. The invocation of a method declared in a class is an active use of that class, which is not allowed until the class is initialized. Thus, before the virtual machine can invoke `main()`, it must initialize `Salutation`. And before it can initialize `Salutation`, it must load and link `Salutation`. So, the virtual machine hands the fully qualified name of `Salutation` to the primordial class loader, which retrieves the binary form of the class, parses the binary data into internal data structures, and creates an instance of `java.lang.Class`. The constant pool for `Salutation` is shown in Table 8-1.

Table 8-1. Class `Salutation`'s constant pool

Index	Type	Value
1	CONSTANT_String_info	30
2	CONSTANT_String_info	31
3	CONSTANT_String_info	39
4	CONSTANT_Class_info	37
5	CONSTANT_Class_info	44
6	CONSTANT_Class_info	45
7	CONSTANT_Class_info	46

8	CONSTANT_Class_info	47
9	CONSTANT_Methodref_info	7, 16
10	CONSTANT_Fieldref_info	4, 17
11	CONSTANT_Fieldref_info	8, 18
12	CONSTANT_Methodref_info	5, 19
13	CONSTANT_Methodref_info	6, 20
14	CONSTANT_Double_info	2.99
16	CONSTANT_NameAndType_info	26, 22
17	CONSTANT_NameAndType_info	41, 32
18	CONSTANT_NameAndType_info	49, 34
19	CONSTANT_NameAndType_info	50, 23
20	CONSTANT_NameAndType_info	51, 21
21	CONSTANT_Utf8_info	"()D"
22	CONSTANT_Utf8_info	"()V"
23	CONSTANT_Utf8_info	"(Ljava/lang/String;)V"
24	CONSTANT_Utf8_info	"([Ljava/lang/String;)V"
25	CONSTANT_Utf8_info	"<clinit"
26	CONSTANT_Utf8_info	"<init"
27	CONSTANT_Utf8_info	"Code"
28	CONSTANT_Utf8_info	"ConstantValue"
29	CONSTANT_Utf8_info	"Exceptions"
30	CONSTANT_Utf8_info	"Greetings, planet!"
31	CONSTANT_Utf8_info	"Hello, world!"
32	CONSTANT_Utf8_info	"I"
33	CONSTANT_Utf8_info	"LineNumberTable"
34	CONSTANT_Utf8_info	"Ljava/io/PrintStream;"
35	CONSTANT_Utf8_info	"Ljava/lang/String;"
36	CONSTANT_Utf8_info	"LocalVariables"
37	CONSTANT_Utf8_info	"Salutation"
38	CONSTANT_Utf8_info	"Salutation.java"
39	CONSTANT_Utf8_info	"Salutations, orb!"
40	CONSTANT_Utf8_info	"SourceFile"
41	CONSTANT_Utf8_info	"choice"
42	CONSTANT_Utf8_info	"greeting"
43	CONSTANT_Utf8_info	"hello"
44	CONSTANT_Utf8_info	"java/io/PrintStream"
45	CONSTANT_Utf8_info	"java/lang/Math"
46	CONSTANT_Utf8_info	"java/lang/Object"
47	CONSTANT_Utf8_info	"java/lang/System"
48	CONSTANT_Utf8_info	"main"

```

49  CONSTANT_Utf8_info      "out"
50  CONSTANT_Utf8_info      "println"
51  CONSTANT_Utf8_info      "random"
52  CONSTANT_Utf8_info      "salutation"

```

As part of the loading process for `Salutation`, the Java Virtual Machine must make sure all of `Salutation`'s superclasses have been loaded. To start this process, the virtual machine looks into `Salutation`'s type data at the `super_class` item, which is a seven. The virtual machine looks up entry seven in the constant pool, and finds a `CONSTANT_Class_info` entry that serves as a symbolic reference to class `java.lang.Object`. See Figure 8-5 for a graphical depiction of this symbolic reference. The virtual machine resolves this symbolic reference, which causes it to load class `Object`. Because `Object` is the top of `Salutation`'s inheritance hierarchy, the virtual machine and links and initializes `Object` as well.




Figure 8-5 The symbolic reference from `Salutation` to `Object`

Now that the Java Virtual Machine has loaded the `Salutation` class and loaded, linked and initialized all its superclasses, the virtual machine is ready to link `Salutation`. As the first step in the linking process, the virtual machine verifies the integrity of the binary representation of class `Salutation`. Assume this implementation of the Java Virtual Machine performs all verification up front, except for the verification of symbolic references. So by the time this official verification phase of linking is completed, the virtual machine will have verified:

1. that `Salutation`'s binary data is structurally correct

2. that `Salutation` correctly implements the semantics of the Java language
3. that `Salutation`'s bytecodes won't crash the virtual machine

After the Java Virtual Machine has verified `Salutation`, it must prepare for `Salutation`'s use by allocating any memory needed by the class. At this stage, the virtual machine allocates memory for `Salutation`'s class variable, `choice`, and gives it a default initial value. Because the `choice` class variable is an `int`, it receives the default initial value of zero.

The three literal Strings--`hello`, `greeting`, and `salutation`--are constants, not class variables. They do not occupy memory space as class variables in the method area. They don't receive default initial values. Because they are declared static and final, they appear as `CONSTANT_String_info` entries in `Salutation`'s constant pool. The constant pool for `Salutation` that was generated by `javac` is shown in Table 8-1. The entries that represent `Salutation`'s constant strings are: for `greeting`, entry one; for `hello`, entry two; and for `salutation`, entry three.

After the processes of verification and preparation have successfully completed, the class is ready for resolution. As mentioned above, different implementations of the Java Virtual Machine may perform the resolution phase of linking at different times. Resolution of `Salutation` is optional at this point in its lifetime. Java Virtual Machines are not required to perform resolution until each symbolic reference is actually used by the program. If a symbolic reference is never actually used by a program, the virtual machine is not required to resolve it.

A Java Virtual Machine implementation could perform the recursive resolution process, described above for `Salutation`, at this point in the lifetime of a program. If so, the program would be completely linked before `main()` is ever invoked. A different Java Virtual Machine implementation could perform none of the resolution process at this point. Instead, it could resolve each symbolic reference the first time it is actually used by the running program. Other implementations could choose a resolution strategy between these two extremes. Although different implementations may perform resolution at different times, all implementations will ensure that a type is loaded, verified, prepared, and initialized before it is used.

Assume this implementation of the Java Virtual Machine uses late resolution. As each symbolic reference is used for the first time by the program, it will be checked for accuracy and converted into a direct reference. Assume also that this implementation uses the technique of replacing the opcode that refers to the constant pool with `_quick` equivalents.

Once this Java Virtual Machine implementation has loaded, verified, and prepared `Salutation`, it is ready to initialize it. As mentioned above, the Java Virtual Machine must initialize all superclasses of a class before it can initialize the class. In this case, the virtual machine has already initialized `Object`, the superclass of `Salutation`.

After the virtual machine has made sure all of `Salutation`'s superclasses have been initialized (in this case, just class `Object`), it is ready to invoke `Salutation`'s `<clinit()` method. Because `Salutation` contains a class variable, `width`, that has an initializer that doesn't resolve at compile-time to a constant, the compiler does place a `<clinit()` method into `Salutation`'s class file.

Here's the `<clinit()` method for `Salutation`:

```

begin

    // Invoke class method Math.random(), passing no
    // parameters. Push double result.

0 invokestatic #13 <Method double random()

    // Push double constant 2.99 from constant pool.

3 ldc2_w #14 <Double 2.99

6 dmul      // Pop two doubles, multiple, push double result.

7 d2i      // Pop double, convert to int, push int result.

    // Pop int, store int Salutation.choice

8 putstatic #10 <Field int choice

11 return   // Return void from <clinit()

end

```

The Java Virtual Machine executes `Salutation.clinit()` method to set the `choice` field to its proper initial value. Before executing `<clinit()`, `choice` has its default initial value of zero. After executing `<clinit()`, it has one of three values chosen pseudo-randomly: zero, one, or two.

The first instruction of the `<clinit()` method, `invokestatic #13`, refers to constant pool entry 13, a `CONSTANT_Methodref_info` that represents a symbolic reference to the `random()` method of class `java.lang.Math`. You can see a graphical depiction of this symbolic reference in Figure 8-6. The Java Virtual Machine resolves this symbolic reference, which causes it to load, link, and initialize class `java.lang.Math`. It places a direct reference to the `random()` method into constant pool entry 13, marks the entry as resolved, and replaces the `invokestatic` opcode with `invokestatic_quick`.




Figure 8-6. The symbolic reference from `Salutation` to `Math.random`

Having completed the resolution process for constant pool entry 13, the Java Virtual Machine is ready to invoke the method. When the virtual machine actually invokes the `random()` method, it will load, link, and initialize any types referenced symbolically from `Math`'s constant pool and `random()`'s code. When this method returns, the virtual machine will push the returned `double` value onto the `main()` method's operand stack.

To execute the next instruction, `ldc2_w #14`, the virtual machine looks into constant pool entry 14 and finds an unresolved `CONSTANT_Double_info` entry. The virtual machine resolves this entry to the double value 2.99, marks the entry as resolved, and replaces the `ldc2_w` opcode with `ldc2_w_quick`. Once the virtual machine has resolved constant pool entry 14, it pushes the constant double value, 2.99, onto the operand stack.

Note that this entry, a `CONSTANT_Double_info`, does not refer to any other constant pool entry or item outside this class. The eight bytes of the double value 2.99 are specified within the entry itself.

Note also that in this constant pool, there is no entry with an index of 15. As mentioned in Chapter 6, "The Java Class File," entries of type `CONSTANT_Double_info` and `CONSTANT_Long_info` occupy two slots in the constant pool. Thus, the `CONSTANT_Double_info` at index 14 is considered to occupy both indices 14 and 15.

To execute the next instruction, `dmul`, the virtual machine pops two `doubles`, multiplies them, and pushes the `double` result. For the next instruction, the virtual machine pops the `double`, converts it to `int`, and pushes the `int` result. Assume that for this particular execution of `Salutation`, the result of this operation is the `int` value two.

The next instruction, `putstatic #10`, uses another symbolic reference from the constant pool, this one to the `choice` variable of `Salutation` itself. This instruction illustrates that a class's bytecodes

use symbolic references to refer not only to fields and methods of other types, but also to its own fields and methods. When the virtual machine executes this instruction, it looks up constant pool entry 10 and finds an as yet unresolved `CONSTANT_Fieldref_info` item. See Figure 8-7 For a graphical depiction of this symbolic reference. The virtual machine resolves the reference by locating the `choice` class variable in `Salutation`'s type data in the method area, and placing a pointer to the actual variable data in constant pool entry 10. It marks the entry as resolved and replaces the `putstatic` opcode with `putstatic_quick`.




Figure 8-7. The symbolic reference from `Salutation` to its own `choice` field.

Once it has resolved the `CONSTANT_Fieldref_info` entry for `choice`, the virtual machine pops an `int` (in this case a two) from the operand stack and places it into the `choice` variable. The execution of the `putstatic` instruction is now complete.

Lastly, the virtual machine executes the `return` instruction, which signals to the virtual machine that the `<clinit()` method, and hence the initialization of class `Salutation`, is complete.

Now that class `Salutation` has been initialized, it is finally ready for use. The Java Virtual Machine invokes `main()`, and the program begins. Here's the bytecode sequence for `Salutation`'s `main()` method:

begin

```

// Push objref to literal string from constant pool
// entry 2

```

```

0 ldc #2 <String "Hello, world!"

```

```

2 astore_1 // Pop objref into loc var 1: String s = hello;

```

```

// Push int from static field Salutation.choice. Note

```

```

        // that by this time, choice has definitely been
        // given its proper initial value.

3  getstatic #10 <Field int choice

6  iconst_1      // Push int constant 1
                // Pop two ints, compare, if not equal branch to 16:

7  if_icmpne 16 // if (choice == 1) {
                // Here, choice does equal 1. Push objref to string
                // literal from constant pool:

10 ldc #1 <String "Greetings, planet!"

12 astore_1      // Pop objref into loc var 1: s = greeting;

13 goto 26       // Branch unconditionally to offset 26
                // Push int from static field Salutation.choice

16 getstatic #10 <Field int choice

19 iconst_2      // Push int constant 2
                // Pop two ints, compare, if not equal branch to 26:

20 if_icmpne 26 // if (choice == 2) {
                // Here, choice does equal 2. Push objref to string
                // literal from constant pool:

23 ldc #3 <String "Salutations, orb!"

25 astore_1      // Pop objref into loc var 1: String s = salutation;
                // Push objref from System.out

26 getstatic #11 <Field java.io.PrintStream out

29 aload_1      // Push objref (to a String) from loc var 1
                // Pop objref (to a String), pop objref(to
                // System.out), invoke println() on System.out
                // passing the string as the only parameter:
                // System.out.println(s);

30 invokevirtual #12 <Method void println(java.lang.String)

33 return       // Return void from main()

```

end

The first instruction in `main()`, `ldc #2`, uses a symbolic reference to the string literal "Hello, world!". When the virtual machine executes this instruction, it looks up constant pool entry two and finds a `CONSTANT_String_info` item that hasn't yet been resolved. See Figure 8-8 For a graphical depiction of the symbolic reference to this string literal.




Figure 8-8. The symbolic reference from `Salutation` to "Hello, world!".

As part of executing the `ldc` instruction, the virtual machine resolves the constant pool entry. It creates and interns a new `String` object with the value "Hello, world!", places a reference to the string object in the constant pool entry, marks the entry as resolved, and replaces the `ldc` opcode with an `ldc_quick`.

Now that the virtual machine has resolved the "Hello, world!" string literal, it pushes the reference to that `String` object onto the stack. The next instruction, `astore_1`, pops the reference and stores it into local variable position one, the `S` variable.

To execute the next instruction, `getstatic #10`, the virtual machine looks up constant pool entry 10 and discovers a `CONSTANT_Fieldref_info` entry that has already been resolved. This entry, a symbolic reference to `Salutation`'s own `choice` field, was resolved by the `putstatic #10` instruction in the `<clinit()` method. The virtual machine simply replaces the `getstatic` opcode with `getstatic_quick`, and pushes the `int` value of `choice` onto the stack.

To execute `main()`'s next instruction, `iconst_1`, the virtual machine simply pushes `int` one onto the operand stack. For the next instruction, `ificmpne 16`, the virtual machine pops the top two `ints` and subtracts one from the other. In this case, since the value of `choice` was set by the `<clinit()` method to be two, the result of the subtraction is not zero. As a consequence, the virtual machine takes

the branch. It updates the pc register so that the next instruction it executes is the `getstatic` instruction at offset 16.

The `getstatic` instruction at offset 16 refers to the same constant pool entry referred to by the `getstatic` instruction at offset three: constant pool entry 10. When the virtual machine executes the `getstatic` at offset 16, it looks up constant pool entry 10 and finds a `CONSTANT_Fieldref_info` entry that is already resolved. It replaces the `getstatic` opcode with `getstatic_quick`, and pushes the `int` value of `Salutation's choice` class variable (a two) onto the operand stack.

To execute the next instruction, `iconst_2`, the virtual machine pushes an `int` two onto the stack. For the next instruction, another `ificmpne` 26, the virtual machine again pops two `ints` and subtracts one from the other. This time, however, both `ints` equal two, so the result of the subtraction is zero. As a consequence, the virtual machine does not take the branch and continues on to execute the next instruction in the bytecode array, another `ldc`.

The `ldc` instruction at offset 23 refers to constant pool entry three, a `CONSTANT_String_info` entry that serves as a symbolic reference to the string literal `"Salutations, orb!"`. The virtual machine looks up this entry in the constant pool and discovers it is as yet unresolved. To resolve the entry, the virtual machine creates and interns a new `String` object with the value `"Salutations, orb!"`, places a reference to the new object in the data for constant pool entry three, and replaces the `ldc` opcode with `ldc_quick`. Having resolved the string literal, the virtual machine pushes the reference to the `String` object onto the stack.

To execute the next instruction, `astore_1`, the virtual machine pops the object reference to the `"Salutations, orb!"` string literal off the stack and stores it into local variable slot one, overwriting the reference to `"Hello, world!"` written there by the `astore_1` instruction at offset two.

The next instruction, `getstatic #11`, uses a symbolic reference to a public static class variable of `java.lang.System` with the name `out` and the type `java.io.PrintStream`. This symbolic reference occupies the `CONSTANT_Fieldref_info` entry at index 11 in the constant pool. See Figure 8-9 For a graphical depiction of this symbolic reference.




Figure 8-9. The symbolic reference from `Salutation` to `System.out`

To resolve the reference to `System.out`, the Java Virtual Machine must load, link, and initialize `java.lang.System` to make sure it has a public static field, named `out`, of type `java.io.PrintStream`. Then, the virtual machine will replace the symbolic reference with a direct reference, such as a native pointer, so that any future uses of `System.out` by `Salutation` won't require resolution and will be faster. Lastly, the virtual machine will replace the `getstatic` opcode with `getstatic_quick`.

Once the virtual machine has successfully resolved the symbolic reference, it will push the reference to `System.out` onto the stack. To execute the next instruction, `aload_1`, the virtual machine simply pushes onto the stack the object reference from local variable one, which is the reference to the `"Salutations, orb!"` string literal.

To execute the next instruction, `invokevirtual #12`, the Java Virtual Machine looks up constant pool entry 12 and finds an unresolved `CONSTANT_Methodref_info` entry, a symbolic reference to the `println()` method of `java.io.PrintStream`. See Figure 8-10 for a graphical depiction of this symbolic reference. The virtual machine loads, links, and initializes `java.io.PrintStream`, and makes sure it has a `println()` method that is `public`, returns `void`, and takes a `String` argument. It marks the entry as resolved and puts a direct reference (an index into `PrintStream`'s method table) into the data for the resolved constant pool entry. Lastly, the virtual machine replaces the `invokevirtual` opcode with `invokevirtual_quick`, and places the method table index and the number of arguments accepted by the method as operands to the `invokevirtual_quick` opcode.




Figure 8-10. The symbolic reference from `Salutation` to `PrintStream.println()`

When the virtual machine actually invokes the `println()` method, it will load, link, and initialize any types referenced symbolically from `PrintStream`'s constant pool and `println()`'s code.

The next instruction is the last instruction the `main()` method: `return`. Because `main()` was being executed by the only non-daemon thread running in the `Salutation` application, executing the `return` instruction will cause the virtual machine to exit. Note that constant pool entry one, which contained a symbolic reference to the "Greetings, planet!" string literal, was never resolved during this execution of the `Salutation` application. Because `choice` happened to be initialized with a value of two, the instruction that referred to constant pool entry one, the `ldc #1` instruction at offset 10, was never executed. As a result, the virtual machine never created a `String` object with the value "Greetings, planet!".

Example: The Dynamic Extension of the Greet Application

As an example of an application that performs dynamic extension through class loader objects, consider the following class:

```

begin

// On CD-ROM in file linking/ex6/Greet.java
import COM.artima.greeter.*;

public class Greet {

    // Arguments to this application:

```

```

//      args[0] - path name of directory in which class files
//
//          for greeters are stored
//
//      args[1], args[2], ... - class names of greeters to load
//
//          and invoke the greet() method on.
//
//
// All greeters must implement the COM.artima.greeter.Greeter
// interface.
//
static public void main(String[] args) {

    if (args.length <= 1) {
        System.out.println(
            "Enter base path and greeter class names as args.");
        return;
    }

    GreeterClassLoader gcl = new GreeterClassLoader(args[0]);

    for (int i = 1; i < args.length; ++i) {
        try {

            // Load the greeter specified on the command line
            Class c = gcl.loadClass(args[i], true);

            // Instantiate it into a greeter object
            Object o = c.newInstance();

            // Cast the Object ref to the Greeter interface type
            // so greet() can be invoked on it
            Greeter greeter = (Greeter) o;

```

```

        // Greet the world in this greeter's special way
        greeter.greet();
    }
    catch (Exception e) {
        e.printStackTrace();
    }
}
}
}

end

```

The **Greet** application is a fancy incarnation of the typical "Hello, world!" program. **Greet** uses a class loader object to dynamically extend itself with classes--called "greeters"--that do the actual work of telling the world hello.

A greeter is any class that implements the **COM.artima.greeter.Greeter** interface:

```

begin

// On CD-ROM in file linking/ex6/COM/artima/greeter/Greeter.java
package COM.artima.greeter;

public interface Greeter {

    void greet();

}

end

```

As you can see from the code above, the **Greeter** interface declares only one method: **greet()**. When a greeter object's **greet()** method is invoked, the object should say hello to the world in its own unique way. Here are a few examples of greeters:

```

begin

// On CD-ROM in file linking/ex6/greeters/Hello.java
import COM.artima.greeter.Greeter;

```

```
public class Hello implements Greeter {
```

```
    public void greet() {
```

```
        System.out.println("Hello, world!");
```

```
    }
```

```
}
```

```
// On CD-ROM in file linking/ex6/greeters/Greetings.java  
import COM.artima.greeter.Greeter;
```

```
public class Greetings implements Greeter {
```

```
    public void greet() {
```

```
        System.out.println("Greetings, planet!");
```

```
    }
```

```
}
```

```
// On CD-ROM in file linking/ex6/greeters/Salutations.java  
import COM.artima.greeter.Greeter;
```

```
public class Salutations implements Greeter {
```

```
    public void greet() {
```

```
        System.out.println("Salutations, orb!");
```

```
    }
```

```
}
```

```
// On CD-ROM in file linking/ex6/greeters/HowDoYouDo.java  
import COM.artima.greeter.Greeter;
```

```
public class HowDoYouDo implements Greeter {
```

```

    public void greet() {
        System.out.println("How do you do, globe!");
    }
}

end

```

Greeters can be more complex than the above four examples. Hereís an example of a greeter that chooses a greeting based on the time of day:

```

begin

// On CD-ROM in file linking/ex6/greeters/HiTime.java
import COM.artima.greeter.Greeter;

import java.util.Date;

public class HiTime implements Greeter {

    public void greet() {

        // Date's no-arg constructor initializes itself to the
        // current date and time
        Date date = new Date();
        int hours = date.getHours();

        // Some hours: midnight, 0; noon, 12; 11PM, 23;
        if (hours = 4 && hours <= 11) {
            System.out.println("Good morning, world!");
        }
        else if (hours = 12 && hours <= 16) {
            System.out.println("Good afternoon, world!");
        }
        else if (hours = 17 && hours <= 21) {

```

```

        System.out.println("Good evening, world!");
    }
    else if (hours = 22 || hours <= 3) {
        System.out.println("Good night, world!");
    }
    else {
        // This should never happen.
        System.out.println("Oh oh, the clock is broken, world!");
    }
}
}

end

```

The **Greet** application doesn't know at compile-time what greeter classes it will load and where those classes will be stored. At run-time it takes a directory path as its first command-line argument and greeter class names as subsequent arguments. It attempts to load the greeters using the path name as a base directory.

For example, imagine you invoke the **Greet** application with the following command line:

```

begin
java Greet greeters Hello
end

```

In this command line, **java** is the name of the Java Virtual Machine executable. **Greet** is the class name of the **Greet** application. **greeters** is the name of a directory relative to the current directory in which the **Greet** application should look for greeters. **Hello** is the name of the greeter.

When the **Greet** application is invoked with the above command line, it attempts to load **greeters/Hello.class** and invoke **Hello**'s **greet()** method. If the **Hello.class** file is indeed sitting in a directory named **greeters**, the application will print:

```

begin
Hello, world!
end

```

The **Greet** application can handle more than one greeter. If you invoke the it with the following command line:


```
begin
```

```
java Greet greeters Hello Greetings Salutations HowDoYouDo
```

```
end
```

The `Greet` application will load each of the four greeters listed and invoke their `greet()` methods, yielding the following output:

```
begin
```

```
Hello, world!
```

```
Greetings, planet!
```

```
Salutations, orb!
```

```
How do you do, globe!
```

```
end
```

The `Greet` application works by first checking to make sure there are at least two command-line arguments: a directory path and at least one greeter class name. It then instantiates a new `GreeterClassLoader` object, which will be responsible for loading the greeters. (The inner workings of class `GreeterClassLoader`, a subclass of `java.lang.ClassLoader`, will be described later in this section.) The constructor for `GreeterClassLoader` accepts a `String` that it uses as a directory path in which to look for greeters.

After it has created the `GreeterClassLoader` object, the `Greet` application invokes its `loadClass()` method for each greeter name that appears on the command line. When it invokes `loadClass()`, it passes the greeter class name in `name` and sets the `resolve` flag to `true`:

```
begin
```

```
// Load the greeter specified on the command line
```

```
Class c = gcl.loadClass(args[i], true);
```

```
end
```

If the `loadClass()` method is unsuccessful, it throws an exception or error. If the `loadClass()` method is successful, it returns the `Class` instance for the newly loaded, linked, and initialized type.

Once `loadClass()` has returned a `Class` instance, the `Greet` application's `main()` method instantiates a new instance of the greeter by calling `newInstance()` on the `Class` instance:

```
begin
```

```
// Instantiate it into a greeter object
```

```
Object o = c.newInstance();
```

```
end
```

When the `newInstance()` method is invoked on a `Class` object, the virtual machine creates and initializes a new instance of the class represented by the `Class` object. To initialize the new instance, the the virtual machine invokes its no-arg constructor.

The `Greet` application then casts the `Object` reference that points to the greeter object to a `Greeter` reference:

```
begin
// Cast the Object ref to the Greeter interface type
// so greet() can be invoked on it
Greeter greeter = (Greeter) o;
end
```

Finally, armed with a `Greeter` reference, the `main()` method invokes the `greet()` method on the greeter object:

```
begin
// Greet the world in this greeter's special way
greeter.greet();
end
```

Here's the code for `GreeterClassLoader`:

```
begin
// On CD-ROM in file
// linking/ex6/COM/artima/greeter/GreeterClassLoader.java
package COM.artima.greeter;

import java.io.*;
import java.util.Hashtable;

public class GreeterClassLoader extends ClassLoader {

    // types is this class loader object's private
    // cache of type names it has loaded. For each type
    // name, the Hashtable stores a reference to the
    // Class instance for that type
```

```

private Hashtable types = new Hashtable();

// basePath gives the path to which this class
// loader appends "<typename.class" to get the
// full path name of the class file to load
private String basePath;

public GreeterClassLoader(String basePath) {

    this.basePath = basePath;
}

public synchronized Class loadClass(String className,
    boolean resolveIt) throws ClassNotFoundException {

    Class result;
    byte classData[];

    // Check the local cache
    result = (Class) types.get(className);
    if (result != null) {
        // Return a cached class
        return result;
    }

    // Check with the primordial class loader
    try {
        result = super.findSystemClass(className);
        // Return a system class
        return result;
    }

```

```

    }
    catch (ClassNotFoundException e) {
    }

    // Don't attempt to load a system file except through
    // the primordial class loader
    if (className.startsWith("java.")) {
        throw new ClassNotFoundException();
    }

    // Try to load it from the basePath directory.
    classData = getTypeFromBasePath(className);
    if (classData == null) {
        System.out.println("GCL - Can't load class: "
            + className);
        throw new ClassNotFoundException();
    }

    // Parse it
    result = defineClass(className, classData, 0,
        classData.length);
    if (result == null) {
        System.out.println("GCL - Class format error: "
            + className);
        throw new ClassFormatError();
    }

    if (resolveIt) {
        resolveClass(result);
    }

```

```

        // Add newly loaded type to private name cache
        types.put(className, result);

        // Return class from basePath directory
        return result;
    }

private byte[] getTypeFromBasePath(String className) {

    InputStream in = null;
    String fileName = basePath + "/" + className + ".class";

    try {
        in = new FileInputStream(fileName);
    } catch (FileNotFoundException e) {
        return null;
    }

    ByteArrayOutputStream out = new ByteArrayOutputStream();

    try {
        int c = in.read();
        while (c != -1) {
            out.write(c);
            c = in.read();
        }
    }
    catch (IOException e) {
        System.out.println("GCL: Error reading bytes of:"

```

```

        + className);
    e.printStackTrace();
    return null;
}

return out.toByteArray();
}
}

end

```

The `GreeterClassLoader` declares two instance variables, `classes` and `basePath`. The `classes` variable, a `Hashtable`, is used to store for each type this class loader object loads the type's name and a reference to its `Class` instance. The `basePath` variable, a `String`, is used to store the directory path (passed to `GreeterClassLoader`'s constructor) in which the `loadClass()` method should look for the class file of the type it has been requested to load.

The `loadClass()` method begins by checking to see if the requested type has already been loaded by this class loader. It does this by looking up the type name passed in parameter `name` in the `Hashtable`. If the name exists in the `Hashtable`, `loadClass()` retrieves and returns the `Class` instance for that name:

```

begin
// Check the local cache
result = (Class) classes.get(className);
if (result != null) {
    // Return a cached class
    return result;
}
end

```

Why does `GreeterClassLoader` maintain a list of the types it has loaded? As mentioned earlier in this chapter, the virtual machine maintains its own list of type names that have already been requested of each class loader. When loading superclasses in Step 1 of resolution, the virtual machine always checks its own internal list before automatically invoking `loadClass()`. As a result, the virtual machine will never automatically invoke `loadClass()` on a class loader object with the name of a type already loaded by that class loader object. Nevertheless, `GreeterClassLoader` maintains its own list of the names of the types it has already loaded. Why? Because even though the virtual machine will never ask a class loader object to load the same type twice, the application just might.

As an example, imagine the `Greet` application were invoked with this command line:

```
begin
java Greet greeters Hello Hello Hello Hello Hello
end
```

Given this command line, the `Greet` application would invoke `loadClass()` with the name `Hello` five times on the same `GreeterClassLoader` object. The first time, the `GreeterClassLoader` would load the class and add its name and `Class` instance to the `Hashtable`. The next four times, however, the `GreeterClassLoader` would simply extract the `Class` instance for `Hello` from its `Hashtable` and return that. It would only load class `Hello` once.

If the `loadClass()` method determines that the requested type has not been loaded into its name space, it next passes the name of the requested type to `findSystemClass()`:

```
begin
// Check with the primordial class loader
try {
    result = super.findSystemClass(className);
    // Return a system class
    return result;
}
catch (ClassNotFoundException e) {
}
end
```

When the `findSystemClass()` method is invoked, the primordial class loader attempts to load the type. If it successfully loads the type, the primordial class loader also links and initializes the type. (As mentioned earlier in this chapter, if the type is a class, the primordial class loader makes sure all the class's superclasses are loaded, linked, and initialized as well.) `findSystemClass()` returns the `Class` instance created by the primordial class loader when it successfully loaded the type, and `loadClass()` returns that same `Class` instance. Note that classes loaded via `findSystemClass()` are marked as having been loaded by the primordial class loader. The `loadClass()` method of `GreeterClassLoader` does not add the name of such types to its `Hashtable`.

If the primordial class loader is unable to load the type, `findSystemClass()` throws `ClassNotFoundException`. In this case, the `loadClass()` method next checks to make sure the requested class is not part of the `java` package:

```
begin
// Don't attempt to load a system file except through
```

```
// the primordial class loader

if (className.startsWith("java.")) {
    throw new ClassNotFoundException();
}

end
```

This check prevents members of the standard java packages (java.lang, java.io, etc.) from being loaded by anything but the primordial class loader. As mentioned in Chapter 3, "Security," class loader objects should not attempt to custom load types that declare themselves to be part of the Java API (or any other "restricted" packages) but which can't be loaded by the primordial class loader.

If the type name doesn't begin with "java," the loadClass() method next invokes getTypeFromBasePath(), which attempts to import the binary data in the class loader object's custom way:

```
begin

// Try to load it from the basePath directory.

classData = getTypeFromBasePath(className);

if (classData == null) {
    System.out.println("GCL - Can't load class: "
        + className);
    throw new ClassNotFoundException();
}

end
```

The getTypeFromBasePath() method looks for a file with the type name plus a ".class" extension in the base directory passed to the GreeterClassLoader's constructor. If the getTypeFromBasePath() method is unable to find the file, it returns a null result and the loadClass() method throws ClassNotFoundException. Otherwise, loadClass() invokes defineClass(), passing the byte array returned by getTypeFromBasePath():

```
begin

// Parse it

result = defineClass(className, classData, 0,
    classData.length);

if (result == null) {
    System.out.println("GCL - Class format error: "
```



```

        + className);
    throw new ClassFormatError();
}
end

```

The `defineClass()` method uses the primordial class loader to complete the loading process: to parse the binary data into internal data structures and create a `Class` instance. The `defineClass()` method does not link and initialize the type. (As mentioned earlier in this chapter, if the type is a class, the `defineClass()` method also make sure all the class's superclass are loaded. It does this by invoking `loadClass()` on this class loader object for each superclass in the hierarchy that hasn't yet been loaded by this class loader, with the `resolve` flag set to `false`.)

If `defineClass()` is successful, it returns the newly created `Class` instance. The `loadClass()` method checks to see if `resolve` were set to `true`. If so, it invokes `resolveClass()`, passing the `Class` instance returned by `defineClass()`. The `resolveClass()` method links and initializes the class. (As mentioned earlier in this chapter, `resolveClass()` also makes sure that, if the type is a class, all its superclasses are linked and initialized as well.)

```

begin
if (resolveIt) {
    resolveClass(result);
}
end

```

Finally, the `loadClass()` method adds the name and `Class` instance for the newly loaded type into the class loader's `Hashtable`, and returns the `Class` instance:

```

begin
// Add newly loaded type to private name cache
classes.put(className, result);

// Return class from basePath directory
return result;
end

```

The `Greet` application demonstrates the flexibility inherent in Java's linking model. The `Greet` application does not know at compile time what greeters it will be loading and dynamically linking to at run-time. In the examples above, class `Greet` invokes the `greet()` method in classes `Hello`, `Greetings`, `Salutations`, and `HowDoYouDo`. But if you look at `Greet`'s constant pool, there is no symbolic reference to any of these classes. There is only a symbolic reference to their shared superinterface, `COM.artima.greeter.Greeter`. Greeters themselves, so long as they implement

the `COM.artima.greeter.Greeter` interface, can be anything and can be written and compiled anytime, even after the `Greet` application itself is compiled.

Example: Unloading Unreachable Greeters

As an example of dynamically loaded types becoming unreachable and getting unloaded by the virtual machine, consider the following application:

```
begin
```

```
// On CD-ROM in file linking/ex6/GreetAndForget.java
import COM.artima.greeter.*;
```

```
public class GreetAndForget {
```

```
    // Arguments to this application:
```

```
    //     args[0] - path name of directory in which class files
```

```
    //             for greeters are stored
```

```
    //     args[1], args[2], ... - class names of greeters to load
```

```
    //             and invoke the greet() method on.
```

```
    //
```

```
    // All greeters must implement the COM.artima.greeter.Greeter
```

```
    // interface.
```

```
    //
```

```
    static public void main(String[] args) {
```

```
        if (args.length <= 1) {
```

```
            System.out.println(
```

```
                "Enter base path and greeter class names as args.");
```

```
            return;
```

```
        }
```

```
        for (int i = 1; i < args.length; ++i) {
```

```
            try {
```

```

GreeterClassLoader gcl =
    new GreeterClassLoader(args[0]);

// Load the greeter specified on the command line
Class c = gcl.loadClass(args[i], true);

// Instantiate it into a greeter object
Object o = c.newInstance();

// Cast the Object ref to the Greeter interface type
// so greet() can be invoked on it
Greeter greeter = (Greeter) o;

// Greet the world in this greeter's special way
greeter.greet();

// Forget the class loader object, Class
// instance, and greeter object
gcl = null;
c = null;
o = null;
greeter = null;

// At this point, the types loaded through the
// GreeterClassLoader object created at the top of
// this for loop are unreferenced and can be unloaded
// by the virtual machine.
}

catch (Exception e) {
    e.printStackTrace();
}

```

```

        }
    }
}

```

end

The **GreetAndForget** application accepts the same command line arguments as the **Greet** application of the previous example. The first argument is a base directory path name where the **GreetAndForget** application will look for greeters. Subsequent arguments are greeter names. To understand this example you should be familiar with the **Greet** application of the previous example.

Imagine you invoke the **GreetAndForget** application with the following command line:

```

begin
java GreetAndForget greeters Surprise HiTime Surprise
end

```

The code for the **HiTime** greeter, which selects a different greeting based on the time of day, is shown above in the previous section of this chapter. The code for the **Surprise** greeter, which pseudo-randomly selects one of four helper greeters--**Hello**, **Greetings**, **Salutations**, or **HowDoYouDo**--and invokes its **greet()** method, is shown here:

```

begin
// On CD-ROM in file linking/ex6/greeters/Surprise.java
import COM.artima.greeter.Greeter;

public class Surprise implements Greeter {

    public void greet() {

        // Choose one of four greeters pseudo-randomly and
        // invoke its greet() method.
        int choice = (int) (Math.random() * 3.99);

        Greeter g;

```

```

switch(choice) {

case 0:
    g = new Hello();
    g.greet();
    break;

case 1:
    g = new Greetings();
    g.greet();
    break;

case 2:
    g = new Salutations();
    g.greet();
    break;

case 3:
    g = new HowDoYouDo();
    g.greet();
    break;

}

}

end

```

Given the command line shown above, the `GreetAndForget` application invokes the `greet()` method of the `Surprise` greeter first, then the `HiTime` greeter, then the `Surprise` greeter again. `GreetAndForget`'s actual output would vary depending on the time of day and `Surprise`'s pseudo-random mood. For the purposes of this example, assume that you typed in the above command, hit return, and got the following output:

```
begin
How do you do, globe!
Good afternoon, world!
Greetings, planet!
end
```

This output indicates `Surprise` chose to execute `HowDoYouDo`'s `greet()` method the first time around and `Greetings`'s `greet()` method the second time around.

The first pass through `GreetAndForget`'s for loop, the virtual machine loads the `Surprise` class and invokes its `greet()` method. The constant pool for `Surprise` includes a symbolic reference to each the four helper greeters that it may choose: `Hello`, `Greetings`, `Salutations`, and `HowDoYouDo`. Assuming the Java Virtual Machine that you used to run the `GreetAndForget` application uses late resolution, only one of these four symbolic references will be resolved during the first pass of `GreetAndForget`'s for loop: the symbolic reference to `HowDoYouDo`. The virtual machine resolves this symbolic reference when it executes the bytecodes that correspond to the following statement in `Surprise`'s `greet()` method:

```
begin
g = new HowDoYouDo();
end
```

To resolve the symbolic reference from `Surprise`'s constant pool to `HowDoYouDo`, the virtual machine invokes the `GreeterClassLoader` object's `loadClass()` method, passing the string "HowDoYouDo" in the `name` parameter and setting the `resolve` flag to `true`. The virtual machine uses the `GreeterClassLoader` object to load `HowDoYouDo` because `Surprise` was loaded through the `GreeterClassLoader` object. As mentioned earlier in this chapter, when the Java Virtual Machine resolves a symbolic reference, it uses the same class loader that loaded the referencing type (in this case, `Surprise`) to load the referenced type (in this case, `HowDoYouDo`).

Once `Surprise`'s `greet()` method has created a new `HowDoYouDo` instance, it invokes its `greet()` method:

```
begin
g.greet();
end
```

As the virtual machine executes `HowDoYouDo`'s `greet()` method, it must resolve two symbolic references from `HowDoYouDo`'s constant pool--one to class `java.lang.System` and another to class `java.io.PrintStream`. To resolve these symbolic references, the virtual machine invokes the `GreeterClassLoader` object's `loadClass()` method, once with the name `java.lang.System` and once with the name `java.io.PrintStream`. As before, the virtual machine uses the `GreeterClassLoader` object to load these classes because the referencing class--in this case, `HowDoYouDo`--was loaded through the `GreeterClassLoader` object. But these two

classes, both members of the Java API, will end up being loaded by the primordial class loader anyway.

Remember that before the `loadClass()` method of `GreeterClassLoader` attempts to look for a requested type in the base directory (in this case, directory `greeters`), it invokes `findSystemClass()`. Because `findSystemClass()`, which uses the primordial class loader to try and load the requested type, is able to load both `java.lang.System` and `java.io.PrintStream`, the `loadClass()` method will simply return the `Class` instance returned by `findSystemClass()`. These classes will be marked not as having been loaded by the `GreeterClassLoader` object, but as having been loaded by the primordial class loader. To resolve any references from `java.lang.System` or `java.io.PrintStream`, the virtual machine will not invoke the `loadClass()` method of the `GreeterClassLoader` object. It will just use the primordial class loader directly.

As a result, after `Surprise`'s `greet()` method has returned, there will be two types marked as having been loaded by the `GreeterClassLoader` object: class `Surprise` and class `HowDoYouDo`. These two types will be in the virtual machine's internal list of the types loaded by the `GreeterClassLoader` object and will appear in the `Hashtable` pointed to by the `GreeterClassLoader` object's `classes` instance variable.

Just after `Surprise`'s `greet()` method returns, the `Class` instances for `Surprise` and `HowDoYouDo` are reachable by the application. The garbage collector will not reclaim the space occupied by these `Class` instances, because there are ways for the application's code to access and use them. See Figure 8-11 for a graphical depiction of the reachability of these two `Class` instances.




Figure 8-11. The reachability of the `Class` instances for `Surprise` and `HowDoYouDo`.

The `Class` instance for `Surprise` can be reached in three ways. First, it can be reached directly from local variable `c` of `GreetAndForget`'s `main()` method. Second, it can be reached from local

variables `o` and `greeter`, which both point to the same `Surprise` object. From the `Surprise` object, the virtual machine can get at `Surprise`'s type data, which includes a reference to `Surprise`'s `Class` object. The third way the `Class` instance for `Surprise` can be reached is through the `gcl` local variable of `GreetAndForget`'s `main()` method. This local variable points to the `GreeterClassLoader` object, which includes a reference to a `HashTable` object in which a reference to `Surprise`'s `Class` instance is stored.

The `Class` instance for `HowDoYouDo` can be reached in two ways. One way is identical to the one of the paths to the `Class` instance for `Surprise`: the `gcl` local variable of `GreetAndForget`'s `main()` method points to the `GreeterClassLoader` object, which includes a reference to a `HashTable` object. The `Hashtable` contains a reference to `HowDoYouDo`'s `Class` instance. The other way to reach `HowDoYouDo`'s class instance is through `Surprise`'s constant pool.

When the virtual machine resolved the symbolic reference from `Surprise`'s constant pool to `HowDoYouDo`, it replaced the symbolic reference with a direct reference. The direct reference points to `HowDoYouDo`'s type data, which includes a reference to `HowDoYouDo`'s `Class` instance.

Thus, starting from `Surprise`'s constant pool, the `Class` instance to `HowDoYouDo` is reachable. But why would the garbage collector look at direct references emanating from `Surprise`'s constant pool in the first place? Because `Surprise`'s `Class` instance is reachable. When the garbage collector finds that it can reach `Surprise`'s `Class` instance, it makes sure it marks the `Class` instances for any types that are directly referenced from `Surprise`'s constant pool as reachable. If `Surprise` is still live, the virtual machine can't unload any types `Surprise` may need to use.

Note that of the three ways, described above, that `Surprise`'s `Class` instance can be reached, none of them involve a constant pool of another type. `Surprise` does not appear as a symbolic reference in the constant pool for `GreetAndForget`. `Class GreetAndForget` did not know about `Surprise` at compile-time. Instead, the `GreetAndForget` application decided at run-time to load and link to class `Surprise`. Thus, the `Class` instance for class `Surprise` is only reachable by starting from the local variables of `GreetAndForget`'s `main()` method. Unfortunately for `Surprise` (and ultimately for `HowDoYouDo`), this does not constitute a very firm grasp on life.

The next four statements in `GreetAndForget`'s `main()` method, will change the reachability situation completely:

```
begin
// Forget the class loader object, Class
// instance, and greeter object
gcl = null;
c = null;
o = null;
greeter = null;
end
```


These statements null out all four starting places from which `Surprise` `Class` instance is reachable. As a result, after these statements have been executed, the `Class` instance for `Surprise` is no longer reachable. These statements also render unreachable the `Class` instance for `HowDoYouDo`, the `Surprise` instance that was formerly pointed to by the `o` and `greeter` variables, the `GreeterClassLoader` instance that was formerly pointed to by the `gcl` variable, and the `Hashtable` instance that was pointed to by the `classes` variable of the `GreeterClassLoader` object. All five of these objects are now available for garbage collection.

When (and if) the garbage collector gets around to freeing the unreferenced `Class` instances for `Surprise` and `HowDoYouDo`, it can also free up all the associated type data in the method area for `Surprise` and `HowDoYouDo`. Because these `Class` instances are unreachable, the types themselves are unreachable and can be unloaded by the virtual machine.

Note that two iterations of the for loop later (given the command line shown above), the `GreetAndForget` application will again load class `Surprise`. Keep in mind that the virtual machine will not reuse the type data for `Surprise` that was loaded during the first pass of the for loop. Granted, that type data became available for unloading at the end of the first pass. But even if the `Class` instance for `Surprise` hadn't become unreferenced at the end of the first pass, the type data from the first pass wouldn't be reused during the third pass.

With each pass of the for loop, the `main()` method of `GreetAndForget` creates a new `GreeterClassLoader` object. Thus, every greeter that `GreetAndForget` loads is loaded through a different class loader object. For example, if you invoke the `GreetAndForget` application with the `Hello` greeter listed five times on the command line, the application will create five instances of class `GreeterClassLoader`. The `Hello` greeter will be loaded five times by five different class loader objects. The method area will contain five different copies of the type data for `Hello`. The heap will contain five `Class` instances that represent the `Hello` class--one for each name space into which `Hello` is loaded. When one of the `Class` instances for `Hello` becomes unreferenced, only the `Hello` type data associated with that particular `Class` instance would be available for unloading.

On the CD-ROM

The CD-ROM contains the source code examples from this chapter in the `linking` directory.

The Resources Page

For more information about the material presented in this chapter, visit the resources page: <http://www.artima.com/insidejvm/linking.html>.

Chapter Nine

Garbage Collection

The Java Virtual Machine's heap stores all objects created by a running Java application. Objects are created by the `new`, `newarray`, `anewarray`, and `multianewarray` instructions, but never freed explicitly by the code. Garbage collection is the process of automatically freeing objects that are no longer referenced by the program.

This chapter does not describe an official Java garbage-collected heap, because none exists. As mentioned in earlier chapters, the Java Virtual Machine specification does not require any particular garbage collection technique. It doesn't even require garbage collection at all. But until infinite memory is invented, most Java Virtual Machine implementations will likely come with garbage-collected heaps. This chapter describes various garbage collection techniques and explains how garbage collection works in Java Virtual Machines.

Accompanying this chapter on the CD-ROM is an applet that interactively illustrates the material presented in the chapter. The applet, named *Heap of Fish*, simulates a garbage-collected heap in a Java Virtual Machine. The simulation--which demonstrates a compacting, mark-and-sweep collector--allows you to interact with the heap as if you were a Java program: you can allocate objects and assign references to variables. The simulation also allows you to interact with the heap as if you were the Java Virtual Machine: you can drive the processes of garbage collection and heap compaction. At the end of this chapter, you will find a description of this applet and instructions on how to use it.

Why Garbage Collection?

The name "garbage collection" implies that objects no longer needed by the program are "garbage" and can be thrown away. A more accurate and up-to-date metaphor might be "memory recycling." When an object is no longer referenced by the program, the heap space it occupies can be recycled so that the space is made available for subsequent new objects. The garbage collector must somehow determine which objects are no longer referenced by the program and make available the heap space occupied by such unreferenced objects. In the process of freeing unreferenced objects, the garbage collector must run any finalizers of objects being freed.

In addition to freeing unreferenced objects, a garbage collector may also combat heap fragmentation. Heap fragmentation occurs through the course of normal program execution. New objects are allocated, and unreferenced objects are freed such that free portions of heap memory are left in between portions occupied by live objects. Requests to allocate new objects may have to be filled by extending the size of the heap even though there is enough total unused space in the existing heap. This will happen if there is not enough contiguous free heap space available into which the new object will fit. On a virtual memory

system, the extra paging (or swapping) required to service an ever growing heap can degrade the performance of the executing program. On an embedded system with low memory, fragmentation could cause the virtual machine to "run out of memory" unnecessarily.

Garbage collection relieves you from the burden of freeing allocated memory. Knowing when to explicitly free allocated memory can be very tricky. Giving this job to the Java Virtual Machine has several advantages. First, it can make you more productive. When programming in non-garbage-collected languages you can spend many late hours (or days or weeks) chasing down an elusive memory problem. When programming in Java you can use that time more advantageously by getting ahead of schedule or simply going home to have a life.

A second advantage of garbage collection is that it helps ensure program integrity. Garbage collection is an important part of Java's security strategy. Java programmers are unable to accidentally (or purposely) crash the Java Virtual Machine by incorrectly freeing memory.

A potential disadvantage of a garbage-collected heap is that it adds an overhead that can affect program performance. The Java Virtual Machine has to keep track of which objects are being referenced by the executing program, and finalize and free unreferenced objects on the fly. This activity will likely require more CPU time than would have been required if the program explicitly freed unnecessary memory. In addition, programmers in a garbage-collected environment have less control over the scheduling of CPU time devoted to freeing objects that are no longer needed.

Garbage Collection Algorithms

Any garbage collection algorithm must do two basic things. First, it must detect garbage objects. Second, it must reclaim the heap space used by the garbage objects and make the space available again to the program.

Garbage detection is ordinarily accomplished by defining a set of roots and determining *reachability* from the roots. An object is reachable if there is some path of references from the roots by which the executing program can access the object. The roots are always accessible to the program. Any objects that are reachable from the roots are considered "live." Objects that are not reachable are considered garbage, because they can no longer affect the future course of program execution.

The root set in a Java Virtual Machine is implementation dependent, but would always include any object references in the local variables and operand stack of any stack frame and any object references in any class variables. Another source of roots are any object references, such as strings, in the constant pool of loaded classes. The constant pool of a loaded class may refer to strings stored on the heap, such as the class name, superclass name, superinterface names, field names, field signatures, method names, and method signatures. Another source of roots may be any object references that were passed to native methods that either haven't been "released" by the native method. (Depending upon the native method interface, a native method may be able to release references by simply returning, by explicitly invoking a call back that releases passed references, or some combination of both.) Another potential source of roots is any part of the Java Virtual Machine's runtime data areas that are allocated from the garbage-collected heap. For example, the class data in the method area itself could be placed on the garbage-collected heap in some implementations, allowing the same garbage collection algorithm that frees objects to detect and unload unreferenced classes.

Any object referred to by a root is reachable and is therefore a live object. Additionally, any objects referred to by a live object are also reachable. The program is able to access any reachable objects, so

these objects must remain on the heap. Any objects that are not reachable can be garbage collected because there is no way for the program to access them.

The Java Virtual Machine can be implemented such that the garbage collector knows the difference between a genuine object reference and a primitive type (for example, an `int`) that appears to be a valid object reference. (For example, an `int` that would point to an object on the heap if it were interpreted as a native pointer.) Some garbage collectors, however, may choose not to distinguish between genuine object references and look-alikes. Such garbage collectors are called *conservative* because they may not always free every unreferenced object. Sometimes a garbage object will be wrongly considered to be live by a conservative collector, because an object reference look-alike referred to it. Conservative collectors trade off an increase in garbage collection speed for occasionally not freeing some actual garbage.

Two basic approaches to distinguishing live objects from garbage are *reference counting* and *tracing*. Reference counting garbage collectors distinguish live objects from garbage objects by keeping a count for each object on the heap. The count keeps track of the number of references to that object. Tracing garbage collectors actually trace out the graph of references starting with the root nodes. Objects that are encountered during the trace are marked in some way. After the trace is complete, unmarked objects are known to be unreachable and can be garbage collected.

Reference Counting Collectors

Reference counting was an early garbage collection strategy. In this approach, a reference count is maintained for each object on the heap. When an object is first created and a reference to it is assigned to a variable, the object's reference count is set to one. When any other variable is assigned a reference to that object, the object's count is incremented. When a reference to an object goes out of scope or is assigned a new value, the object's count is decremented. Any object with a reference count of zero can be garbage collected. When an object is garbage collected, any objects that it refers to have their reference counts decremented. In this way the garbage collection of one object may lead to the subsequent garbage collection of other objects.

An advantage of this approach is that a reference counting collector can run in small chunks of time closely interwoven with the execution of the program. This characteristic makes it particularly suitable for real-time environments where the program can't be interrupted for very long. A disadvantage is that reference counting does not detect *cycles*: two or more objects that refer to one another. An example of a cycle is a parent object that has a reference to a child object that has a reference back to the parent. These objects will never have a reference count of zero even though they may be unreachable by the roots of the executing program. Another disadvantage of reference counting is the overhead of incrementing and decrementing the reference count each time.

Because of the disadvantages inherent in the reference counting approach, this technique is currently out of favor. It is more likely that the Java Virtual Machines you encounter in the real world will use a tracing algorithm in their garbage-collected heaps.

Tracing Collectors

Tracing garbage collectors trace out the graph of object references starting with the root nodes. Objects that are encountered during the trace are marked in some way. Marking is generally done by either setting flags in the objects themselves or by setting flags in a separate bitmap. After the trace is complete, unmarked objects are known to be unreachable and can be garbage collected.

The basic tracing algorithm is called "mark and sweep." This name refers to the two phases of the garbage collection process. In the mark phase, the garbage collector traverses the tree of references and marks each object it encounters. In the sweep phase, unmarked objects are freed, and the resulting memory is made available to the executing program. In the Java Virtual Machine, the sweep phase must include finalization of objects.

Compacting Collectors

Garbage collectors of Java Virtual Machines will likely have a strategy to combat heap fragmentation. Two strategies commonly used by mark and sweep collectors are compacting and copying. Both of these approaches move objects on the fly to reduce heap fragmentation. Compacting collectors slide live objects over free memory space toward one end of the heap. In the process the other end of the heap becomes one large contiguous free area. All references to the moved objects are updated to refer to the new location.

Updating references to moved objects is sometimes made simpler by adding a level of indirection to object references. Instead of referring directly to objects on the heap, object references refer to a table of object handles. The object handles refer to the actual objects on the heap. When an object is moved, only the object handle must be updated with the new location. All references to the object in the executing program will still refer to the updated handle, which did not move. While this approach simplifies the job of heap defragmentation, it adds a performance overhead to every object access.

Copying Collectors

Copying garbage collectors move all live objects to a new area. As the objects are moved to the new area, they are placed side by side, thus eliminating any free space that may have separated them in the old area. The old area is then known to be all free space. The advantage of this approach is that objects can be copied as they are discovered by the traversal from the root nodes. There are no separate mark and sweep phases. Objects are copied to the new area on the fly, and forwarding pointers are left in their old locations. The forwarding pointers allow the garbage collector to detect references to objects that have already been moved. The garbage collector can then assign the value of the forwarding pointer to the references so they point to the object's new location.

A common copying collector algorithm is called "stop and copy." In this scheme, the heap is divided into two regions. Only one of the two regions is used at any time. Objects are allocated from one of the regions until all the space in that region has been exhausted. At that point program execution is stopped and the heap is traversed. Live objects are copied to the other region as they are encountered by the traversal. When the stop and copy procedure is finished, program execution resumes. Memory will be allocated from the new heap region until it too runs out of space. At that point the program will once again be stopped. The heap will be traversed and live objects will be copied back to the original region. The cost associated with this approach is that twice as much memory is needed for a given amount of heap space because only half of the available memory is used at any time.

You can see a graphical depiction of a garbage-collected heap that uses a stop and copy algorithm in Figure 9-1. This figure shows nine snapshots of the heap over time. In the first snapshot, the lower half of the heap is unused space. The upper half of the heap is partially filled by objects. That portion of the heap that contains objects is painted with diagonal gray lines. The second snapshot shows that the top half of the heap is gradually being filled up with objects, until it becomes full as shown in the third snapshot.

At that point, the garbage collector stops the program and traces out the graph of live objects starting with the root nodes. It copies each live object it encounters down to the bottom half of the heap, placing each object next to the previously copied object. This process is shown in snapshot four.

Snapshot five shows the heap after the garbage collection has finished. Now the top half of the heap is unused, and the bottom half is partially filled with live objects. The sixth snapshot shows the bottom half is now becoming gradually filled with objects, until it too becomes full in snapshot seven.

Once again, the garbage collector stops the program and traces out the graph of live objects. This time, it copies each live object it encounters up to the top half of the heap, as shown in snapshot eight. Snapshot nine shows the result of the garbage collection: the bottom half is once again unused space and the top half is partially filled with objects. This process repeats again and again as the program executes.




Figure 9-1. A "stop and copy" garbage-collected heap.

Generational Collectors

One disadvantage of simple stop and copy collectors is that *all* live objects must be copied at every collection. This facet of copying algorithms can be improved upon by taking into account two facts that have been empirically observed in most programs in a variety of languages:

1. Most objects created by most programs have very short lives.
2. Most programs create some objects that have very long lifetimes.

A major source of inefficiency in simple copying collectors is that they spend much of their time copying the same long-lived objects again and again.

Generational collectors address this inefficiency by grouping objects by age and garbage collecting younger objects more often than older objects. In this approach, the heap is divided into two or more sub-heaps, each of which serves one "generation" of objects. The youngest generation is garbage collected most often. As most objects are short-lived, only a small percentage of young objects are likely

to survive their first collection. Once an object has survived a few garbage collections as a member of the youngest generation, the object is promoted to the next generation: it is moved to another sub-heap. Each progressively older generation is garbage collected less often than the next younger generation. As objects "mature" (survive multiple garbage collections) in their current generation, they are moved to the next older generation.

The generational collection technique can be applied to mark and sweep algorithms as well as copying algorithms. In either case, dividing the heap into generations of objects can help improve the efficiency of the basic underlying garbage collection algorithm.

Adaptive Collectors

An adaptive garbage collection algorithm takes advantage of the fact that some garbage collection algorithms work better in some situations, while others work better in other situations. An adaptive algorithm monitors the current situation on the heap and adjusts its garbage collection technique accordingly. It may tweak the parameters of a single garbage collection algorithm as the program runs. It may switch from one algorithm to another on the fly. Or it may divide the heap into sub-heaps and use different algorithms on different sub-heaps simultaneously.

With an adaptive approach, designers of Java Virtual Machine implementations need not choose just one garbage collection technique. They can employ many techniques, giving each algorithm work for which it is best suited.

Finalization

In Java, an object may have a finalizer: a method that the garbage collector must run on the object prior to freeing the object. The potential existence of finalizers complicates the job of any garbage collector in a Java Virtual Machine.

To add a finalizer to a class, you simply declare a method in that class as follows:

begin

```
// On CD-ROM in file gc/ex2/Example2.java
class Example2 {

    protected void finalize() throws Throwable {

        //...

        super.finalize();

    }

    //...

}
```

end

A garbage collector must examine all objects it has discovered to be unreferenced to see if any include a `finalize()` method.

Because of finalizers, a garbage collector in the Java Virtual Machine must perform some extra steps each time it garbage collects. First, the garbage collector must in some way detect unreferenced objects (call this Pass I). Then, it must examine the unreferenced objects it has detected to see if any declare a finalizer. If it has enough time, it may at this point in the garbage collection process finalize all unreferenced objects that declare finalizers.

After executing all finalizers, the garbage collector must once again detect unreferenced objects starting with the root nodes (call this Pass II). This step is needed because finalizers can "resurrect" unreferenced objects and make them referenced again. Finally, the garbage collector can free all objects that were found to be unreferenced in both Passes I and II.

To reduce the time it takes to free up some memory, a garbage collector can optionally insert a step between the detection of unreferenced objects that have finalizers and the running of those finalizers. Once the garbage collector has performed Pass I and found the unreferenced objects that need to be finalized, it can run a miniature trace starting not with the root nodes but with the objects waiting to be finalized. Any objects that are (1) not reachable from the root nodes (those detected during Pass I) and (2) not reachable from the objects waiting to be finalized cannot be resurrected by any finalizer. These objects can be freed immediately.

If an object with a finalizer becomes unreferenced, and its finalizer is run, the garbage collector must in some way ensure that it never runs the finalizer on that object again. If that object is resurrected by its own finalizer or some other object's finalizer and later becomes unreferenced again, the garbage collector must treat it as an object that has no finalizer.

As you program in Java, you must keep in mind that it is the garbage collector that runs finalizers on objects. Because it is not generally possible to predict exactly when unreferenced objects will be garbage collected, it is not possible to predict when object finalizers will be run. As mentioned in Chapter 2, "Platform Independence," you should avoid writing programs for which correctness depends upon the timely finalization of objects. For example, if a finalizer of an unreferenced object releases a resource that is needed again later by the program, the resource will not be made available until after the garbage collector has run the object finalizer. If the program needs the resource before the garbage collector has gotten around to finalizing the unreferenced object, the program is out of luck.

Heap of Fish: A Simulation

The *Heap of Fish* applet, shown in Figures 9-2 through 9-5, demonstrates a compacting, mark and sweep, garbage-collected heap. To facilitate compaction, this heap uses indirect handles to objects instead of direct references. It is called *Heap of Fish* because the only type of objects stored on the heap for this demonstration are fish objects, defined as follows:

```
insert

// On CD-ROM in file gc/ex1/YellowFish.java

class YellowFish {
```



```
YellowFish myFriend;  
  
}
```

```
// On CD-ROM in file gc/ex1/BlueFish.java  
  
class BlueFish {  
  
    BlueFish myFriend;  
    YellowFish myLunch;  
  
}
```

```
// On CD-ROM in file gc/ex1/RedFish.java  
  
class RedFish {  
  
    RedFish myFriend;  
    BlueFish myLunch;  
    YellowFish mySnack;  
  
}
```

end

As you can see, there are three classes of fish: red, yellow, and blue. The red fish is the largest as it has three instance variables. The yellow fish, with only one instance variable, is the smallest fish. The blue fish has two instance variables and is therefore medium-sized.

The instance variables of fish objects are references to other fish objects. `BlueFish.myLunch`, for example, is a reference to a `YellowFish` object. In this implementation of a garbage-collected heap, a reference to an object occupies four bytes. Therefore, the size of the instance data of a `RedFish` object is twelve bytes, a `BlueFish` object is eight bytes, and a `YellowFish` object is four bytes.

Heap of Fish has five modes, which can be selected via radio buttons at the bottom left of the applet. When the applet starts it is in swim mode. Swim mode is just a gratuitous animation, vaguely

reminiscent of the familiar image of a big fish about to eat a medium-sized fish, which is about to eat a small fish. The other four modes--allocate fish, assign references, garbage collect, and compact heap--allow you to interact with the heap. In the allocate fish mode, you can instantiate new fish objects. In the assign references mode, you can build a network of local variables and fish that refer to other fish. In garbage collect mode, a mark and sweep operation will free any unreferenced fish. The compact heap mode allows you to slide heap objects so that they are side by side at one end of the heap, leaving all free memory as one large contiguous block at the other end of the heap.

Allocate Fish

The allocate fish mode, shown in Figure 9-2, allows you to allocate new fish objects on the heap. In this mode you can see the two parts that make up the heap: the object pool and handle pool. The object pool is a contiguous block of memory from which space is taken for new objects. It is structured as a series of memory blocks. Each memory block has a four-byte header that indicates the length of the memory block and whether or not it is free. The headers are shown in the applet as black horizontal lines in the object pool.




Figure 9-2. The allocate fish mode of the *Heap of Fish* applet.

The object pool in *Heap of Fish* is implemented as an array of `ints`. The first header is always at `objectPool[0]`. The object pool's series of memory blocks can be traversed by hopping from header to header. Each header gives the length of its memory block, which also reveals where the next header is going to be. The header of the next memory block will be the first `int` immediately following the current memory block.

When a new object is allocated, the object pool is traversed until a memory block is encountered with

enough space to accommodate the new object. Allocated objects in the object pool are shown as colored bars. `YellowFish` objects are shown in yellow, `BlueFish` in blue, and in red. Free memory blocks, those that currently contain no fish, are shown in white.

The handle pool in *Heap of Fish* is implemented as an array of objects of a class named `ObjectHandle`. An `ObjectHandle` contains information about an object, including the vital index into the object pool array. The object pool index functions as a reference to the object's instance data in the object pool. The `ObjectHandle` also reveals information about the class of the fish object. As mentioned in Chapter 5, "The Java Virtual Machine," every object on the heap must in some way be associated with its class information stored in the method area. In *Heap of Fish*, the `ObjectHandle` associates each allocated object with information such as its class--whether it is a `RedFish`, `BlueFish`, or `YellowFish`--and some data used in displaying the fish in the applet user interface.

The handle pool exists to make it easier to defragment the object pool through compaction. References to objects, which can be stored in local variables of a stack or the instance variables of other objects, are not direct indexes into the object pool array. They are instead indexes into the handle pool array. When objects in the object pool are moved for compaction, only the corresponding `ObjectHandle` must be updated with the object's new object pool array index.

Each handle in the handle pool that refers to a fish object is shown as a horizontal bar painted the same color as the fish to which it refers. A line connects each handle to its fish instance variables in the object pool. Those handles that are not currently in use are drawn in white.

Assign References

The assign references mode, shown in Figure 9-3, allows you to build a network of references between local variables and allocated fish objects. A reference is merely a local or instance variable that contains a valid object reference. There are three local variables which serve as the roots of garbage collection, one for each class of fish. If you do not link any fish to local variables, then all fish will be considered unreachable and freed by the garbage collector.




Figure 9-3. The assign references mode of the *Heap of Fish* applet.

The assign references mode has three sub-modes: move fish, link fish, and unlink fish. You can select the sub-mode via radio buttons at the bottom of the canvas upon which the fish appear. In move fish mode, you can click on a fish and drag it to a new position. You might want to do this to make your links more visible or just because you feel like rearranging fish in the sea.

In link fish mode, you can click on a fish or local variable and drag a link to another fish. The fish or local variable you initially drag from will be assigned a reference to the fish you ultimately drop upon. A line will be shown connecting the two items. A line connecting two fish will be drawn between the nose of the fish with the reference to the tail of the referenced fish.

Class `YellowFish` has only one instance variable, `myFriend`, which is a reference to a `YellowFish` object. Therefore, a yellow fish can only be linked to one other yellow fish. When you link two yellow fish, the `myFriend` variable of the "dragged from" fish will be assigned the reference to the "dropped upon" fish. If this action were implemented in Java code, it might look like:

```
insert
```

```
// Fish are allocated somewhere
```

```
YellowFish draggedFromFish = new YellowFish();
```

```
YellowFish droppedUponFish = new YellowFish();
```

```
// Sometime later the assignment takes place

draggedFromFish.myFriend = droppedUponFish;

end
```

Class **BlueFish** has two instance variables, **BlueFish myFriend** and **YellowFish myLunch**. Therefore, a blue fish can be linked to one blue fish and one yellow fish. Class **RedFish** has three instance variables, **RedFish myFriend**, **BlueFish myLunch**, and **YellowFish mySnack**. Red fish can therefore link to one instance of each class of fish.

In unlink fish mode, you can disconnect fish by moving the cursor over the line connecting two fish. When the cursor is over the line, the line will turn black. If you click a black line the reference will be set to null and the line will disappear.

Garbage Collect

The garbage collect mode, shown in Figure 9-4, allows you to drive the mark and sweep algorithm. The Step button at the bottom of the canvas takes you through the garbage collection process one step at a time. You can reset the garbage collector at any time by clicking the Reset button. However, once the garbage collector has swept, the freed fish are gone forever. No manner of frantic clicking of the Reset button will bring them back.




Figure 9-4. The garbage collect mode of the *Heap of Fish* applet.

The garbage collection process is divided into a mark phase and a sweep phase. During the mark phase,

the fish objects on the heap are traversed depth-first starting from the local variables. During the sweep phase, all unmarked fish objects are freed.

At the start of the mark phase, all local variables, fish, and links are shown in white. Each press of the Step button advances the depth-first traversal one more node. The current node of the traversal, either a local variable or a fish, is shown in magenta. As the garbage collector traverses down a branch, fish along the branch are changed from white to gray. Gray indicates the fish has been reached by the traversal, but there may yet be fish further down the branch that have not been reached. Once the terminal node of a branch is reached, the color of the terminal fish is changed to black and the traversal retreats back up the branch. Once all links below a fish have been marked black, that fish is marked black and the traversal returns back the way it came.

At the end of the mark phase, all reachable fish are colored black and any unreachable fish are colored white. The sweep phase then frees the memory occupied by the white fish.

Compact Heap

The compact heap mode, shown in Figure 9-5, allows you to move one object at a time to one end of the object pool. Each press of the Slide button will move one object. You can see that only the object instance data in the object pool moves; the handle in the handle pool does not move.




Figure 9-5. The compact heap mode of the *Heap of Fish* applet.

The *Heap of Fish* applet allows you to allocate new fish objects, link fish, garbage collect, and compact the heap. These activities can be done in any order as much as you please. By playing around with this applet you should be able to get a good idea how a mark and sweep garbage-collected heap works. There is some text at the bottom of the applet that should help you as you go along. Happy clicking.

On the CD-ROM

The CD-ROM contains the source code examples from this chapter in the `gc` directory. The *Heap of Fish* applet is contained in a web page on the CD-ROM in file `applets/HeapOfFish.html`. The source code for this applet is found alongside its class files, in the `applets/HeapOfFish` directory.

The Resources Page

For more information about the material presented in this chapter, visit the resources page:
<http://www.artima.com/insidejvm/gc.html>

Chapter Ten

Stack and Local Variable Operations

As mentioned in Chapter 5, "The Java Virtual Machine," the abstract specification of the execution engine is defined in terms of an instruction set. The remaining chapters of this book (Chapters 10 through 20) are a tutorial of that instruction set. They describe the instructions in functional groups and give relevant background information for each group.

This chapter covers instructions that deal with the operand stack and local variables. As the Java Virtual Machine is a stack-based machine, almost all of its instructions involve the operand stack in some way. Most instructions push values, pop values, or both as they perform their functions. This chapter describes the instructions that focus most exclusively on the operand stack--those that push constants onto the operand stack, perform generic stack operations, and transfer values back and forth between the operand stack and local variables.

Accompanying this chapter on the CD-ROM is an applet that interactively illustrates the material presented in the chapter. The applet, named *Fibonacci Forever*, simulates the Java Virtual Machine executing a method that generates the Fibonacci sequence. In the process, it demonstrates how the virtual machine pushes constants, pops values to local variables, and pushes values from local variables. At the end of this chapter, you will find a description of this applet and the bytecodes it executes.

Pushing Constants Onto the Stack

Many opcodes push constants onto the stack. Opcodes indicate the constant value to push in three different ways. The constant value is either implicit in the opcode itself, follows the opcode in the bytecode stream as an operand, or is taken from the constant pool.

Some opcodes by themselves indicate a type and constant value to push. For example, the `iconst_1` opcode tells the Java Virtual Machine to push integer value one. Such opcodes are defined for some commonly pushed numbers of various types. All of these instructions are redundant to the instructions that take operands from the bytecode stream or refer to the constant pool, but these are more efficient. Because these instructions occupy only one byte in the bytecode stream, they increase the efficiency of bytecode execution and reduce the size of bytecode streams. The opcodes that push `ints` and `floats` are shown in Table 10-1.

Table 10-1. Pushing single-word constants onto the stack

Opcode	Operand(s)	Description
--------	------------	-------------

<code>iconst_m1</code>	(none)	pushes <code>int</code> -1 onto the stack
<code>iconst_0</code>	(none)	pushes <code>int</code> 0 onto the stack
<code>iconst_1</code>	(none)	pushes <code>int</code> 1 onto the stack
<code>iconst_2</code>	(none)	pushes <code>int</code> 2 onto the stack
<code>iconst_3</code>	(none)	pushes <code>int</code> 3 onto the stack
<code>iconst_4</code>	(none)	pushes <code>int</code> 4 onto the stack
<code>iconst_5</code>	(none)	pushes <code>int</code> 5 onto the stack
<code>fconst_0</code>	(none)	pushes <code>float</code> 0 onto the stack
<code>fconst_1</code>	(none)	pushes <code>float</code> 1 onto the stack
<code>fconst_2</code>	(none)	pushes <code>float</code> 2 onto the stack

The opcodes shown in the previous table push `ints` and `floats`, which are single-word values. Each slot on the Java stack is one word in size (at least 32 bits wide). Therefore each time an `int` or `float` is pushed onto the stack, it occupies one slot.

The opcodes shown in Table 10-2 push `longs` and `doubles`. `long` and `double` values occupy 64 bits. Each time a `long` or `double` is pushed onto the stack, its value occupies two slots on the stack. Opcodes that indicate a specific `long` or `double` value to push are shown in Table 10-2.

Table 10-2. Pushing dual-word constants onto the stack

Opcode	Operand(s)	Description
<code>lconst_0</code>	(none)	pushes <code>long</code> 0 onto the stack
<code>lconst_1</code>	(none)	pushes <code>long</code> 1 onto the stack
<code>dconst_0</code>	(none)	pushes <code>double</code> 0 onto the stack
<code>dconst_1</code>	(none)	pushes <code>double</code> 1 onto the stack

One other opcode pushes an implicit constant value onto the stack. The `aconst_null` opcode, shown in Table 10-3, pushes a `null` object reference onto the stack.

As mentioned in earlier chapters, the format of an object reference depends upon the Java Virtual Machine implementation. An object reference will somehow refer to a Java object on the garbage-collected heap. A `null` object reference indicates an object reference variable does not currently refer to any valid object. The `aconst_null` opcode is used in the process of assigning `null` to an object reference variable.

Table 10-3. Pushing a `null` reference onto the stack

Opcode	Operand(s)	Description
<code>aconst_null</code>	(none)	pushes a <code>null</code> object reference onto the stack

Two opcodes indicate the constant to push with an operand that immediately follows the opcode. These opcodes, shown in Table 10-4, are used to push integer constants that are within the valid range for `byte` or `short` types. The `byte` or `short` that follows the opcode is expanded to an `int` before it is

pushed onto the stack. Operations on `byte`s and `short`s that have been pushed onto the stack are actually done on their `int` equivalents.

Table 10-4. Pushing `byte` and `short` constants onto the stack

Opcode	Operand(s)	Description
<code>bipush</code>	<code>byte1</code>	expands <code>byte1</code> (a <code>byte</code> type) to an <code>int</code> and pushes it onto the stack
<code>sipush</code>	<code>byte1</code> , <code>byte2</code>	expands <code>byte1</code> , <code>byte2</code> (a <code>short</code> type) to an <code>int</code> and pushes it onto the stack

Three opcodes push constants from the constant pool. These opcodes take operands that specify a constant pool index. The Java Virtual Machine looks up the constant pool entry given the index, determines the constant's type and value, and pushes it onto the stack.

The constant pool index is an unsigned value that immediately follows the opcode in the bytecode stream. Opcodes `ldc` and `ldc_w` push a single-word item onto the stack, either an `int`, `float`, or an object reference to a `String`. The difference between `ldc` and `ldc_w` is that `ldc` can only refer to constant pool locations one through 255 because its index is just 1 byte. (Constant pool location zero is unused.) `ldc_w` has a 2-byte index, so it can refer to any constant pool location. `ldc2_w` also has a 2-byte index, and it is used to refer to any constant pool location containing a `long` or `double`, which occupy two words. The opcodes that push constants from the constant pool are shown in Table 10-5.

Table 10-5. Pushing constant pool entries onto the stack

Opcode	Operand(s)	Description
<code>ldc</code>	<code>indexbyte1</code>	pushes single-word value from constant pool entry specified by <code>indexbyte1</code> onto the stack
<code>ldc_w</code>	<code>indexbyte1</code> , <code>indexbyte2</code>	pushes single-word value from constant pool entry specified by <code>indexbyte1</code> , <code>indexbyte2</code> onto the stack
<code>ldc2_w</code>	<code>indexbyte1</code> , <code>indexbyte2</code>	pushes dual-word value from constant pool entry specified by <code>indexbyte1</code> , <code>indexbyte2</code> onto the stack

All string literals from Java source code end up as entries in a constant pool. If multiple classes of the same application use the same string literal, that string literal will appear in the class file of every class that uses it. For example, if three classes use the string literal `"Harumph!"`, that string will appear in the constant pool of each of three class files. Methods of those classes can use the `ldc` or `ldc_w` instructions to push onto the operand stack a reference to a `String` object that has the value `"Harumph!"`.

As mentioned in Chapter 8, "The Linking Model," the Java Virtual Machine resolves all string literals that have the same sequence of characters into the same `String` object. In other words, if multiple classes use the same literal string, say `"Harumph!"`, the Java Virtual Machine will only create one `String` object with the value `"Harumph!"` to represent all of those string literals.

When the virtual machine resolves the constant pool entry for a literal string, it "interns" the string. First, it checks to see if the string's sequence of characters have already been interned. If so, it just uses the same reference as the already-interned string. Otherwise, it creates a new `String` object, adds a reference to the new `String` object to its set of interned strings, and uses the reference to the newly-interned string.

Generic Stack Operations

Although most instructions in the Java Virtual Machine's instruction set operate on a particular type, some instructions manipulate the stack independent of type. As mentioned in Chapter 5, "The Java Virtual Machine," these generic (typeless) instructions cannot be used to break up dual-word values. These instructions are shown in Table 10-6.

Table 10-6. Stack manipulation

Opcode	Operand (s)	Description
<code>nop</code>	(none)	do nothing
<code>pop</code>	(none)	pop the top word from the operand stack
<code>pop2</code>	(none)	pop the top two words from the operand stack
<code>swap</code>	(none)	swap the top operand stack two words
<code>dup</code>	(none)	duplicate top operand stack word
<code>dup2</code>	(none)	duplicate top two operand stack words
<code>dup_x1</code>	(none)	duplicate top operand stack word and put two down
<code>dup_x2</code>	(none)	duplicate top operand stack word and put three down
<code>dup2_x1</code>	(none)	duplicate top two operand stack words and put three down
<code>dup2_x2</code>	(none)	duplicate top two operand stack words and put four down

The last four instructions shown in Table 10-6 can be a bit difficult to understand. Consult the description of these instructions in Appendix A for a picture of the stack before and after these instructions have been executed.

Pushing Local Variables Onto the Stack

Several opcodes exist that push `int` and `float` local variables onto the operand stack. Some opcodes are defined that implicitly refer to a commonly used local variable position. For example, `iload_0` loads the `int` local variable at position zero. Other local variables are pushed onto the stack by an opcode that takes the local variable index from the first byte following the opcode. The `iload` instruction is an example of this type of opcode. The first byte following `iload` is interpreted as an unsigned 8-bit index that refers to a local variable.

The opcodes that push `int` and `float` local variables onto the stack are shown in Table 10-7.

Table 10-7. Pushing single-word local variables onto the stack

Opcode	Operand(s)	Description
<code>iload</code>	<code>vindex</code>	pushes <code>int</code> from local variable position <code>vindex</code>
<code>iload_0</code>	(none)	pushes <code>int</code> from local variable position zero
<code>iload_1</code>	(none)	pushes <code>int</code> from local variable position one
<code>iload_2</code>	(none)	pushes <code>int</code> from local variable position two
<code>iload_3</code>	(none)	pushes <code>int</code> from local variable position three
<code>fload</code>	<code>vindex</code>	pushes <code>float</code> from local variable position <code>vindex</code>
<code>fload_0</code>	(none)	pushes <code>float</code> from local variable position zero
<code>fload_1</code>	(none)	pushes <code>float</code> from local variable position one
<code>fload_2</code>	(none)	pushes <code>float</code> from local variable position two
<code>fload_3</code>	(none)	pushes <code>float</code> from local variable position three

Table 10-8 shows the instructions that push local variables of type `long` and `double` onto the stack. These instructions move two words from the local variable section of the stack frame to the operand stack section.

Table 10-8. Pushing dual-word local variables onto the stack

Opcode	Operand(s)	Description
<code>lload</code>	<code>vindex</code>	pushes <code>long</code> from local variable positions <code>vindex</code> and (<code>vindex + 1</code>)
<code>lload_0</code>	(none)	pushes <code>long</code> from local variable positions zero and one
<code>lload_1</code>	(none)	pushes <code>long</code> from local variable positions one and two
<code>lload_2</code>	(none)	pushes <code>long</code> from local variable positions two and three
<code>lload_3</code>	(none)	pushes <code>long</code> from local variable positions three and four
<code>dload</code>	<code>vindex</code>	pushes <code>double</code> from local variable positions <code>vindex</code> and (<code>vindex + 1</code>)
<code>dload_0</code>	(none)	pushes <code>double</code> from local variable positions zero and one
<code>dload_1</code>	(none)	pushes <code>double</code> from local variable positions one and two
<code>dload_2</code>	(none)	pushes <code>double</code> from local variable positions two and three
<code>dload_3</code>	(none)	pushes <code>double</code> from local variable positions three and four

The final group of opcodes that push local variables move object references (which occupy one word) from the local variables section of the stack frame to the operand section. These opcodes are shown in Table 10-9.

Table 10-9. Pushing object reference local variables onto the stack

Opcode	Operand(s)	Description
<code>aload</code>	<code>vindex</code>	pushes object reference from local variable position <code>vindex</code>
<code>aload_0</code>	(none)	pushes object reference from local variable position zero
<code>aload_1</code>	(none)	pushes object reference from local variable position one
<code>aload_2</code>	(none)	pushes object reference from local variable position two
<code>aload_3</code>	(none)	pushes object reference from local variable position three

Popping to Local Variables

For each opcode that pushes a local variable onto the stack there exists a corresponding opcode that pops the top of the stack back into the local variable. The mnemonics of the pop opcodes can be formed from the mnemonics of the push opcodes by replacing "load" with "store." The opcodes that pop `ints` and `floats` from the top of the operand stack to a local variable are listed in Table 10-10. Each of these opcodes moves one single-word value from the top of the stack to a local variable.

Table 10-10. Popping single-word values into local variables

Opcode	Operand(s)	Description
<code>istore</code>	<code>vindex</code>	pops <code>int</code> to local variable position <code>vindex</code>
<code>istore_0</code>	(none)	pops <code>int</code> to local variable position zero
<code>istore_1</code>	(none)	pops <code>int</code> to local variable position one
<code>istore_2</code>	(none)	pops <code>int</code> to local variable position two
<code>istore_3</code>	(none)	pops <code>int</code> to local variable position three
<code>fstore</code>	<code>vindex</code>	pops <code>float</code> to local variable position <code>vindex</code>
<code>fstore_0</code>	(none)	pops <code>float</code> to local variable position zero
<code>fstore_1</code>	(none)	pops <code>float</code> to local variable position one
<code>fstore_2</code>	(none)	pops <code>float</code> to local variable position two
<code>fstore_3</code>	(none)	pops <code>float</code> to local variable position three

Table 10-11 shows the instructions that pop values of type `long` and `double` into a local variable. These instructions move a dual-word value from the top of the operand stack to a local variable.

Table 10-11. Popping dual-word values into local variables

Opcode	Operand(s)	Description
<code>lstore</code>	<code>vindex</code>	pops <code>long</code> to local variable positions <code>vindex</code> and

	(vindex + 1)
lstore_0 (none)	pops long to local variable positions zero and one
lstore_1 (none)	pops long to local variable positions one and two
lstore_2 (none)	pops long to local variable positions two and three
lstore_3 (none)	pops long to local variable positions three and four
dstore vindex	pops double to local variable positions vindex and (vindex + 1)
dstore_0 (none)	pops double to local variable positions zero and one
dstore_1 (none)	pops double to local variable positions one and two
dstore_2 (none)	pops double to local variable positions two and three
dstore_3 (none)	pops double to local variable positions three and four

The final group of opcodes that pops to local variables are shown in Table 10-12. These opcodes pop an object reference from the top of the operand stack to a local variable.

Table 10-12. Popping object references into local variables

Opcode	Operand (s)	Description
astore vindex		pops object reference to local variable position vindex
astore_0 (none)		pops object reference to local variable position zero
astore_1 (none)		pops object reference to local variable position one
astore_2 (none)		pops object reference to local variable position two
astore_3 (none)		pops object reference to local variable position three

The wide Instruction

Unsigned 8-bit local variable indexes, such as the one that follows the `iload` instruction, limit the number of local variables in a method to 256. A separate instruction, `wide`, can extend an 8-bit index by another 8 bits, which raises the local variable limit to 65,536. The `wide` opcode modifies other opcodes. `wide` can precede an instruction, such as `iload`, that takes an 8-bit unsigned local variable index. Two bytes that form a 16-bit unsigned index into the local variables follows the `wide` opcode and the modified opcode.

Table 10-13 lists all but two of the opcodes that can be modified by `wide`. The other two opcodes, `inc` and `ret`, are described in later chapters. The `inc` instruction and its `wide` variant are described in Chapter 12, "Integer Arithmetic." The `ret` instruction and its `wide` variant are described in Chapter 18, "Finally Clauses."

Table 10-13. Popping object references into local variables

Opcode	Operand(s)	Description
wide	iload, indexbyte1, indexbyte2	pushes <code>int</code> from local variable position index
wide	lload, indexbyte1, indexbyte2	pushes <code>long</code> from local variable position index
wide	float, indexbyte1, indexbyte2	pushes <code>float</code> from local variable position index
wide	dload, indexbyte1, indexbyte2	pushes <code>double</code> from local variable position index
wide	aload, indexbyte1, indexbyte2	pushes object reference from local variable position index
wide	istore, indexbyte1, indexbyte2	pops <code>int</code> to local variable position index
wide	lstore, indexbyte1, indexbyte2	pops <code>long</code> to local variable position index
wide	fstore, indexbyte1, indexbyte2	pops <code>float</code> to local variable position index
wide	dstore, indexbyte1, indexbyte2	pops <code>double</code> to local variable position index
wide	astore, indexbyte1, indexbyte2	pops object reference to local variable position index

When verifying bytecode sequences that include `wide` instructions, the opcode modified by `wide` is seen as an operand to `wide`. Jump instructions are not allowed to jump directly to an opcode modified by `wide`. For example, if a bytecode sequence include the instruction:

```
begin
wide iload 257
end
```

No other opcode of that method's bytecode sequence would be allowed to jump directly to the `iload` opcode. In this case, the `iload` opcode must always be executed as an operand to the `wide` opcode.

Fibonacci Forever: A Simulation

The *Fibonacci Forever* applet, shown in Figure 10-1, demonstrates a Java Virtual Machine executing a sequence of bytecodes that generate the Fibonacci series. The applet is embedded in a web page on the CD-ROM in file `applets/FibonacciForever.html`. The bytecode sequence in the simulation was generated by the `javac` compiler for the `calcSequence()` method of the class shown below:

```
begin
// On CD-ROM in file stackops/ex1/Fibonacci.java
class Fibonacci {
```

```

static void calcSequence() {
    long fiboNum = 1;
    long a = 1;
    long b = 1;

    for (;;) {
        fiboNum = a + b;
        a = b;
        b = fiboNum;
    }
}
}

```

end

The `calcSequence()` method produces the Fibonacci series and places each Fibonacci number successively in the `fiboNum` variable. The first two numbers of the Fibonacci series are both ones. Each subsequent number is calculated by summing the previous two numbers, as in: 1, 1, 2, 3, 5, 8, 13, 21, 34, 55, and so on.

The bytecodes generated by `javac` for the `calcSequence()` method are shown below:

begin

```

0 lconst_1    // Push long constant 1
1 lstore_0    // Pop long into local vars 0 & 1: long a = 1;
2 lconst_1    // Push long constant 1
3 lstore_2    // Pop long into local vars 2 & 3: long b = 1;
4 lconst_1    // Push long constant 1
5 lstore 4    // Pop long into local vars 4 & 5: long fiboNum = 1;
7 lload_0     // Push long from local vars 0 & 1
8 lload_2     // Push long from local vars 2 & 3
9 ladd        // Pop two longs, add them, push result
10 lstore 4   // Pop long into local vars 4 & 5: fiboNum = a + b;

```



```

12 lload_2    // Push long from local vars 2 & 3
13 lstore_0   // Pop long into local vars 0 & 1: a = b;
14 lload 4    // Push long from local vars 4 & 5
16 lstore_2   // Pop long into local vars 2 & 3: b = fiboNum;
17 goto 7     // Jump back to offset 7: for (;;) {}

end

```


The `javac` compiler placed local variable `a` from the source into local variable slots 0 and 1 on the stack frame. It put `b` into slots 2 and 3 and `fiboNum` into slots 4 and 5. As this method calculates each successive Fibonacci number, it places the number into the `fiboNum` variable. As you run the simulation, therefore, you will see the Fibonacci series appear in the `long` value stored in local variable slots 4 and 5.

You may notice that `long` values are split across the two words they occupy in the local variables by placing the lower half (bits 0 through 31) in the first slot and the upper half (bits 32 through 63) in the second slot. For example, the lower half of the `fiboNum` variable is stored in local variable slot 4. The upper half of `fiboNum` is stored in local variable slot 5. On the operand stack, a similar representation is used. When a `long` value is pushed onto the operand stack, the lower half of the word is pushed, then the upper half.

Keep in mind that this manner of representing `long` values in the local variables and on the operand stack is an artifact of this particular (simulated) implementation of the Java Virtual Machine. As mentioned in Chapter 5, "The Java Virtual Machine," the specification does not dictate any particular way to layout `long`s and `doubles` across the two words they occupy on the stack frame.

Although according to the best mathematical minds, the Fibonacci series does indeed go on forever, the `calcSequence()` method is able to generate Fibonacci numbers only for a while. Unfortunately for `calcSequence()`, the `long` type has a finite range. The highest Fibonacci number this simulation can calculate, therefore, is the highest Fibonacci number that can be represented in a `long`: 7540113804746346429L. After the simulation arrives at this point in the Fibonacci series, the next addition will overflow.

To drive the *Fibonacci Forever* simulation, use the Step, Reset, Run, and Stop buttons. Each time you press the Step button, the simulator will execute the instruction pointed to by the pc register. If you press the Run button, the simulation will continue with no further coaxing on your part until you press the Stop button. To start the simulation over, press the Reset button. For each step of the simulation, a panel at the bottom of the applet contains an explanation of what the next instruction will do. Happy clicking.



On the CD-ROM

The CD-ROM contains the source code examples from this chapter in the `stackops` directory. The *Fibonacci Forever* applet is contained in a web page on the CD-ROM in file `applets/FibonacciForever.html`. The source code for this applet is found alongside its class files, in the `applets/JVMSimulators` and `applets/JVMSimulators`.

The Resources Page

For more information about the material presented in this chapter, visit the resources page: <http://www.artima.com/insidejvm/stackops.html>

Chapter Eleven

Type Conversion

This chapter describes the instructions that convert values from one primitive type to another inside the Java Virtual Machine. It also looks at some of the idiosyncrasies of type conversion that arise from the virtual machine's limited support of `bytes`, `shorts`, and `chars`.

Accompanying this chapter on the CD-ROM is an applet that interactively illustrates the material presented in the chapter. The applet, named *Conversion Diversion*, simulates the Java Virtual Machine executing a method that performs type conversion. At the end of this chapter, you will find a description of this applet and the bytecodes it executes.

The Conversion Opcodes

The Java Virtual Machine has many opcodes that convert from one primitive type to another. No operands follow the conversion opcodes in the bytecode stream. The value to convert is taken from the top of the stack. The Java Virtual Machine pops the value at the top of the stack, converts it, and pushes the result back onto the stack. Opcodes that convert between `int`, `long`, `float`, and `double` are shown in Table 11-1. There is an opcode for each possible from-to combination of these four types:

Table 11-1. Converting ints, longs, floats, and doubles

Opcode	Operand(s)	Description
i2l	(none)	converts <code>int</code> to <code>long</code>
i2f	(none)	converts <code>int</code> to <code>float</code>
i2d	(none)	converts <code>int</code> to <code>double</code>
l2i	(none)	converts <code>long</code> to <code>int</code>
l2f	(none)	converts <code>long</code> to <code>float</code>
l2d	(none)	converts <code>long</code> to <code>double</code>
f2i	(none)	converts <code>float</code> to <code>int</code>
f2l	(none)	converts <code>float</code> to <code>long</code>
f2d	(none)	converts <code>float</code> to <code>double</code>
d2i	(none)	converts <code>double</code> to <code>int</code>
d2l	(none)	converts <code>double</code> to <code>long</code>
d2f	(none)	converts <code>double</code> to <code>float</code>

Opcodes that convert from an `int` to a type smaller than `int` are shown in Table 11-2. These opcodes pop one `int` off the operand stack, convert it to an `int` value that fits into a `byte`, `short`, or `char`, and push the resulting `int` back onto the operand stack. The `i2b` instruction truncates the popped `int` value to a `byte`, then sign-extends back out to an `int`. The `i2s` instruction truncates the popped `int` value to a `short`, then sign-extends back out to an `int`. The `i2c` instruction truncates the popped `int` value to a `char`, then zero-extends back out to an `int`.

Table 11-2. Converting ints, bytes, chars, and shorts.

Opcode	Operand(s)	Description
<code>i2b</code>	(none)	converts <code>int</code> to <code>byte</code>
<code>i2c</code>	(none)	converts <code>int</code> to <code>char</code>
<code>i2s</code>	(none)	converts <code>int</code> to <code>short</code>

No opcodes exist that convert directly from a `long`, `float`, or `double` to the types smaller than `int`. Therefore converting from a `float` to a `byte`, for example, requires two steps. First the `float` must be converted to an `int` with `f2i`, then the resulting `int` can be converted to a `byte` with `i2b`.

Although opcodes exist that convert an `int` to primitive types smaller than `int` (`byte`, `short`, and `char`), no opcodes exist that convert in the opposite direction. This is because any `bytes`, `shorts`, or `chars` are effectively converted to `int` before being pushed onto the stack. The instructions that retrieve `byte`, `short`, and `char` values from arrays or objects on the heap and push them onto the stack convert the values to `int`. These instructions are described in Chapter 15, "Objects and Arrays."

Arithmetic operations upon `bytes`, `shorts`, and `chars` are done by first converting the values to `int`, performing the arithmetic operations on the `ints`, and being happy with an `int` result. This means that if you add 2 `bytes` you get an `int`, and if you want a `byte` result you must explicitly convert the `int` result back to a `byte`. For example, the following code won't compile:

```
begin
```

```
// On the CD-ROM in file opcodes/ex1/BadArithmetic.java
class BadArithmetic {
```

```
    static byte addOneAndOne() {
        byte a = 1;
        byte b = 1;
        byte c = a + b;
        return c;
    }
```

```
}
```

end

When presented with the above code, `javac` objects with the following remark:

begin

BadArithmetic.java(7): Incompatible type for declaration. Explicit cast needed to convert `int` to `byte`.

```
byte c = a + b;
```

^

end

To remedy the situation, the you must explicitly convert the `int` result of the addition of `a + b` back to a `byte`, as in the following code:

begin

```
// On the CD-ROM in file opcodes/ex1/GoodArithmetic.java
class GoodArithmetic {
```

```
    static byte addOneAndOne() {
```

```
        byte a = 1;
```

```
        byte b = 1;
```

```
        byte c = (byte) (a + b);
```

```
        return c;
```

```
    }
```

```
}
```

end

This makes `javac` so happy it drops a `GoodArithmetic.class` file, which contains the following bytecode sequence for the `addOneAndOne()` method:

begin

```
0 iconst_1    // Push int constant 1.
```

```
1 istore_0    // Pop into local variable 0, which is a: byte a = 1;
```

```
2 iconst_1    // Push int constant 1 again.
```

```
3 istore_1    // Pop into local variable 1, which is b: byte b = 1;
```

```
4 iload_0     // Push a (a is already stored as an int in local
```

```

        // variable 0).
5 iload_1    // Push b (b is already stored as an int in local
            // variable 1).
6 iadd       // Perform addition. Top of stack is
            // now (a + b), an int.
7 i2b        // Convert int result to byte (result still
            // occupies 32 bits).
8 istore_2    // Pop into local variable 3, which is
            // byte c: byte c = (byte) (a + b);
9 iload_2     // Push the value of c so it can be returned.
10 ireturn   // Proudly return the result of the addition: return c;

end

```

Conversion Diversion: A Simulation

The *Conversion Diversion* applet, shown in Figure 11-1, demonstrates a Java Virtual Machine executing a sequence of bytecodes. The applet is embedded in a web page on the CD-ROM in file `applets/ConversionDiversion.html`. The bytecode sequence in the simulation was generated by `javac` for the `Convert()` method of the class shown below:

```

begin

// On CD-ROM in file opcodes/ex1/Diversion.java
class Diversion {

    static void Convert() {

        byte imByte = 0;

        int imInt = 125;

        for (;;) {

            ++imInt;

            imByte = (byte) imInt;

            imInt *= -1;

```

```

        imByte = (byte) imInt;
        imInt *= -1;
    }
}
}

```

end

The bytecodes generated by javac for `Convert()` are shown below:

begin

```

0  iconst_0    // Push int constant 0.
1  istore_0    // Pop to local variable 0, which is
                // imByte: byte imByte = 0;

2  bipush 125  // Expand byte constant 125 to int and push.

4  istore_1    // Pop to local variable 1, which
                // is imInt: int imInt = 125;

5  iinc 1 1    // Increment local variable 1 (imInt) by 1: ++imInt;

8  iload_1     // Push local variable 1 (imInt).

9  i2b         // Truncate and sign extend top of stack so it
                // has a valid byte value.

10 istore_0    // Pop to local variable 0 (imByte):
                // imByte = (byte) imInt;

11 iload_1     // Push local variable 1 (imInt) again.

12 iconst_m1   // Push integer -1.

13 imul        // Pop top two ints, multiply, push result.

14 istore_1    // Pop result of multiply to local variable 1 (imInt):
                // imInt *= -1;

15 iload_1     // Push local variable 1 (imInt).

16 i2b         // Truncate and sign extend top of stack so it has
                // a valid byte value.

17 istore_0    // Pop to local variable 0 (imByte):

```

```

        // imByte = (byte) imInt;

18 iload_1    // Push local variable 1 (imInt) again.
19 iconst_m1  // Push integer -1.
20 imul       // Pop top two ints, multiply, push result.
21 istore_1   // Pop result of multiply to local variable 1 (imInt):
        // imInt *= -1;

22 goto 5     // Jump back to the iinc instruction: for (;;) {}

end

```

The `Convert()` method demonstrates the manner in which the Java Virtual Machine converts from `int` to `byte`. `imInt` starts out as 125. Each pass through the while loop, it is incremented and converted to a `byte`. Then it is multiplied by -1 and again converted to a `byte`. The simulation quickly shows what happens at the edges of the valid range for the `byte` type.

The maximum value for a `byte` is 127. The minimum value is -128. Values of type `int` that are within this range convert directly to `byte`. As soon as the `int` gets beyond the valid range for `byte`, however, things get interesting.

The Java Virtual Machine converts an `int` to a `byte` by truncating and sign extending. The highest order bit, the "sign bit," of `longs`, `ints`, `shorts`, and `bytes` indicate whether or not the integer value is positive or negative. If the sign bit is zero, the value is positive. If the sign bit is one, the value is negative. Bit 7 of a `byte` value is its sign bit. To convert an `int` to a `byte`, bit 7 of the `int` is copied to bits 8 through 31. This produces an `int` that has the same numerical value that the `int`'s lowest order byte would have if it were interpreted as a `byte` type. After the truncation and sign extension, the `int` will contain a valid `byte` value.

The simulation applet shows what happens when an `int` that is just beyond the valid range for `byte` types gets converted to a `byte`. For example, when the `imInt` variable has a value of 128 (0x00000080) and is converted to `byte`, the resulting `byte` value is -128 (0xffffffff80). Later, when the `imInt` variable has a value of -129 (0xffffffff7f) and is converted to `byte`, the resulting `byte` value is 127 (0x0000007f).

To drive the *Conversion Diversion* simulation, use the Step, Reset, Run, and Stop buttons. Each time you press the Step button, the simulator will execute the instruction pointed to by the pc register. If you press the Run button, the simulation will continue with no further coaxing on your part until you press the Stop button. To start the simulation over, press the Reset button. For each step of the simulation, a panel at the bottom of the applet contains an explanation of what the next instruction will do. Happy clicking.

Applet Viewer: ConversionDiversi.on.class

Applet

CONVERSION DIVERSION

Local Variables

index	hex value	value
0	ffffff80	-128
1	00000080	128

optop
2

Operand Stack

offset	hex value	value
0	00000080	128
1	ffffff	-1

optop >

pc
13

The Method

offset	bytecode	mnemonic
5	84	iinc 1, 1
6	01	
7	01	
8	1b	iload_1
9	91	i2b
10	3b	istore_0
11	1b	iload_1
12	02	iconst_m1
pc >	13	imul
14	3c	istore_1
15	1b	iload_1
16	91	i2b

Step Run

Reset Stop

imul will pop two integers, multiply them, and push the result.

On the CD-ROM

The CD-ROM contains the source code examples from this chapter in the `opcodes` directory. The *Conversion Diversion* applet is contained in a web page on the CD-ROM in file `applets/ConversionDiversi.on.html`. The source code for this applet is found alongside its class files, in the `applets/JVMSimulators` and `applets/JVMSimulators/COM/artima/jvmsim` directories.

The Resources Page

For more information about the material presented in this chapter, visit the resources page: <http://www.artima.com/insidejvm/typeconv.html>

Chapter Twelve

Integer Arithmetic

This chapter describes integer arithmetic in the Java Virtual Machine. It explains two's complement arithmetic (the mechanism used by the Java Virtual Machine to perform integer arithmetic) and describes the instructions that perform integer arithmetic.

Accompanying this chapter on the CD-ROM are two applets that interactively illustrate the material presented in the chapter. One applet, named *Inner Int*, allows you to view and manipulate the individual components that make up a two's complement number. The other applet, named *Prime Time*, simulates the Java Virtual Machine executing a method that generates prime numbers. At the end of this chapter, you will find a description of this applet and the bytecodes it executes.

Two's Complement Arithmetic

All integer types supported by the Java Virtual Machine--bytes, shorts, ints, and longs--are signed two's-complement numbers. The two's-complement scheme allows both positive and negative integers to be represented. The most significant bit of a two's-complement number is its sign bit. The sign bit is one for negative numbers and zero for positive numbers and for the number zero.

The number of unique values that can be represented by the two's-complement scheme is two raised to the power of the total number of bits. For example, the `short` type in Java is a 16-bit signed two's-complement integer. The number of unique integers that can be represented by this scheme is 2^{16} , or 65,536. Half of the `short` type's range of values are used to represent zero and positive numbers; the other half of the `short` type's range are used to represent negative numbers. The range of negative values for a 16-bit two's-complement number is -32,768 (0x8000) to -1 (0xffff). Zero is 0x0000. The range of positive values is one (0x0001) to 32,767 (0x7fff).

Positive numbers are intuitive in that they are merely the base two representation of the number. Negative numbers can be calculated by adding the negative number to two raised to the power of the total number of bits. For example, the total number of bits in a `short` is 16, so the two's-complement representation of a negative number in the valid range for a `short` (-32,768 to -1) can be calculated by adding the negative number to 2^{16} , or 65,536. The two's-complement representation for -1 is $65,536 + (-1)$ or 65,535 (0xffff). The two's-complement representation for -2 is $65,536 + (-2)$ or 65,534 (0xfffe).

Addition is performed on two's-complement signed numbers in the same way it would be performed on unsigned binary numbers. The two numbers are added, overflow is ignored, and the result is interpreted as a signed two's-complement number. This will work as long as the result is actually within the range of


valid values for the type. For example, to add $4 + (-2)$, just add `0x00000004` and `0xffffffff`. The result is actually `0x100000002`, but because there are only 32 bits in an `int`, the overflow is ignored and the result becomes `0x00000002`.

Overflow in integer operations does not throw any exception in the Java Virtual Machine. The result is merely truncated to fit into the result type (either `int` or `long`). For example, adding `ints` `0x7fffffff` and `1` yields `0x80000000`. This means that the Java Virtual Machine will report that $2,147,483,647 + 1 = -2,147,483,648$, if the type of the values being added are `ints` and not `longs`. As you program in Java, you must keep in mind that overflow can happen and make sure you choose the appropriate type, `int` or `long`, in each situation. Integer division by zero does throw an `ArithmeticException`, so you should also keep in mind that this exception could be thrown and catch it if necessary.

If you encounter a situation in which `long` just isn't long enough, you can use the `BigInteger` class of the `java.math` package. Instances of this class are arbitrary-length integers. The `BigInteger` class supports all arithmetic operations on arbitrary-length integers that are provided for the primitive types by the Java Virtual Machine and the `java.lang.Math` package.

Inner Int: A Java int Reveals its Inner Nature

The *Inner Int* applet, shown in Figure 12-1, lets you play around with the two's-complement format of integers in the Java Virtual Machine. The Max and Min buttons will give you the maximum and minimum values of the `int` type. By clicking Max followed by `++` you can increment beyond the maximum integer and see what happens. Clicking Min followed by `--` lets you decrement beyond the minimum integer. Both of these result in overflow, but no exceptions are thrown by the Java Virtual Machine. The applet is embedded in a web page on the CD-ROM in file `applets/InnerInt.html`.



Arithmetic Opcodes

The Java Virtual Machine offers several opcodes that perform integer arithmetic operations on `ints` and `longs`. As mentioned in earlier chapters, values of type `byte`, `short`, and `char` are converted to `int` before they take part in arithmetic operations. For each opcode that performs arithmetic on `ints`, there is a corresponding opcode that performs the same operation on `longs`.

Integer addition can be performed on `ints` and `longs`. Table 12-1 shows the opcodes that pop the top two values on the stack, add them, and push the resulting sum. The values to be added must have been pushed onto the stack by previous instructions. The type of the values is indicated by the opcode itself,

and the result always has the same type as the numbers being added. No exceptions are thrown for any of these opcodes. Overflow is just ignored.

Table 12-1. Integer addition

Opcode	Operand(s)	Description
iadd	(none)	pops two <code>ints</code> , adds them, and pushes the <code>int</code> result
ladd	(none)	pops two <code>longs</code> , adds them, and pushes the <code>long</code> result

Table 12-2 shows the exception to the rule that arithmetic opcodes take their operands from the stack. The `inc` opcode performs an addition on a local variable of type `int`. The local variable to increment is specified by the first byte that follows the `inc` instruction in the bytecode stream. The amount to add to the local variable is taken from the second byte following the `inc` instruction. The second byte is interpreted as an 8-bit signed two's-complement number. The local variable and 8-bit signed value are added, and the result is written back to the local variable. This opcode can be used to change a local variable value by any number between and including -128 through 127. This opcode makes for more efficient incrementing and decrementing of variables that are used to control execution of loops, such as `for` or `while`. As with the `add` instruction, no exceptions are thrown. Overflow is ignored.

The second row in Table 12-2 shows the `wide` variant of the `inc` instruction. As mentioned in Chapter 10, "Stack and Local Variable Operations," the `wide` instruction is used to extend unsigned local variable indexes from eight bits to sixteen. Sixteen bits allows instructions to address up to 65,536 local variable slots. In the `inc` case, the `wide` instruction also extends the size of the signed increment value from eight bits to sixteen. This means the `wide` variant of the `inc` opcode can be used to change a local variable value by any number between and including -32,768 through 32,767.

Table 12-2. Increment a local variable by a constant

Opcode	Operand(s)	Description
inc	vindex, const	adds const to an <code>int</code> at local variable position vindex
wide	inc, indexbyte1, indexbyte2, constbyte1, constbyte2	adds const to an <code>int</code> at local variable position index end table

Integer subtraction is performed on `ints` and `longs` via the opcodes shown in Table 12-3. Each opcode causes the top two values of the appropriate type to be popped off the stack. The topmost value is subtracted from the value just beneath it. The result is pushed back onto the stack. No exceptions are thrown by these opcodes.

Table 12-3. Integer subtraction

Opcode	Operand(s)	Description
isub	(none)	pops two <code>ints</code> , subtracts them, and pushes the <code>int</code>

		result
<code>lsub</code>	(none)	pops two <code>longs</code> , subtracts them, and pushes the <code>long</code> result

Integer multiplication of `ints` and `longs` is accomplished via the opcodes shown in Table 12-4. Each opcode causes two values of the same type to be popped off the stack and multiplied. The result, of the same type as the numbers being multiplied, is pushed back onto the stack. No exceptions are thrown.

Table 12-4. Integer multiplication

Opcode	Operand (s)	Description
<code>imul</code>	(none)	pops two <code>ints</code> , multiplies them, and pushes the <code>int</code> result
<code>lmul</code>	(none)	pops two <code>longs</code> , multiplies them, and pushes the <code>long</code> result

The opcodes that perform division on `ints` and `longs` are shown in Table 12-5. The division opcodes cause the top two values of the appropriate type to be popped off the stack. The value just beneath the topmost value is divided by the topmost value. (In other words, the first value pushed onto the stack is used as the dividend or numerator. The second value pushed--the top of the stack--is used as the divisor or denominator.) The result of the division is pushed back onto the stack. Integer division yields a result that is truncated down to the nearest integer value between it and zero. Integer division by zero throws an `ArithmeticException`.

Table 12-5. Integer division

Opcode	Operand (s)	Description
<code>idiv</code>	(none)	pops two <code>ints</code> , divides them, and pushes the <code>int</code> result
<code>ldiv</code>	(none)	pops two <code>longs</code> , divides them, and pushes the <code>long</code> result

The remainder operation is performed on `ints` and `longs` via the opcodes shown in Table 12-6. These opcodes cause the top two values to be popped from the stack. The value just beneath the topmost value is divided by the topmost value, and the remainder of that division is pushed back onto the stack. As with the division opcodes, integer remainder by zero throws an `ArithmeticException`.

Table 12-6. Integer remainder

Opcode	Operand (s)	Description
<code>irem</code>	(none)	pops two <code>ints</code> , divides them, and pushes the <code>int</code> remainder
<code>lrem</code>	(none)	pops two <code>longs</code> , divides them, and pushes the <code>long</code> remainder

remainder

The opcodes shown in Table 12-7 perform arithmetic negation on `ints` and `longs`. The negation opcodes pop the top value from the stack, negate it, and push the result.

Table 12-7. Integer negation

Opcode	Operand(s)	Description
<code>ineg</code>	(none)	pops an <code>int</code> , negates it, and pushes the result
<code>lneg</code>	(none)	pops a <code>long</code> , negates it, and pushes the result

Prime Time

: A Simulation

The *Prime Time* applet, shown in Figure 12-2, demonstrates a Java Virtual Machine executing a sequence of bytecodes that generates prime numbers. The applet is embedded in a web page on the CD-ROM in file `applets/PrimeTime.html`. The bytecode sequence in the simulation was generated by the `javac` compiler for the `findPrimes()` method of the class shown below:

begin

```
// On CD-ROM in file integer/ex1/PrimeFinder.java
class PrimeFinder {
```

```
    static void findPrimes() {
```

```
        int primeNum = 1;
```

```
        int numToCheck = 2;
```

```
        for (;;) {
```

```
            boolean foundPrime = true;
```

```
            for (int divisor = numToCheck / 2; divisor > 1;
                --divisor) {
```

```
                if (numToCheck % divisor == 0) {
```

```
                    foundPrime = false;
```

```

        break;
    }
}

if (foundPrime) {
    primeNum = numToCheck;
}

++numToCheck;
}
}
}

end

```

The `findPrimes()` method places prime numbers, one at a time and in increasing numerical order, into the `primeNum` variable. To find the primes, it checks each positive integer in increasing numerical order starting with integer value two. It keeps the current number it is checking in the `numToCheck` variable. The outer for loop, a "forever" loop, keeps this process going indefinitely. To check a number, it divides the number by smaller integers looking for a zero remainder. If it encounters a zero remainder, then the number has integral factors other than one and itself and therefore isn't prime.

For each number to check, the `findPrimes()` method divides the number by two. The result of this integer division is the first value checked as a possible integral divisor for the number.

In the inner for loop, the `findPrimes()` method tries each number as a divisor between the result of the division by two and the divisor two. If the remainder of any of these divisions is zero, it breaks out of the inner for loop and skips to the next number to check. If it reaches divisor two and has found no divisor that yields a zero remainder, it has found the next prime number. It exits the inner for loop and sets `primeNum` equal to the number to check.

For example, when `numToCheck` is 10, `findPrimes()` first divides 10 by 2 to get the first divisor, 5. It then performs the remainder operation on 10 and 5 and discovers a zero remainder. So it breaks out of the inner for loop and sets `numToCheck` to 11. (It doesn't ever set `primeNum` to 10.) It divides 11 by 2 to get the first divisor to check, once again 5. It performs the remainder operation on 11 and integers 5, 4, 3, and 2, none of which yield a zero remainder. It completes the inner for loop, sets `primeNum` equal to 11, and continues on to check 12.

The bytecodes generated by `javac` for the `findPrimes()` method are shown below:

`begin`

```

0 iconst_1      // Push int constant 1
1 istore_0      // Pop into local var 0: int primeNum = 1;

2 iconst_2      // Push int constant 2
3 istore_1      // Pop into local var 1: int numToCheck = 2;

// The outer for loop (the "forever" loop) begins here:

4 iconst_1      // Push int constant 1
5 istore_2      // Pop into local var 2: boolean foundPrime = true;

// The inner for loop begins here. First, initialize divisor.
6 iload_1       // Push int in local var 1 (numToCheck)
7 iconst_2      // Push int constant 2
8 idiv          // Pop two ints, divide them, push int result
                // Pop int into local var 3:
9 istore_3      // int divisor = numToCheck / 2;

// Next, test the inner for loop's termination condition
10 goto 27      // Jump to for loop condition check

// The body of the inner for loop begins here.
13 iload_1      // Push the int in local var 1 (numToCheck)
14 iload_3      // Push the int in local var 3 (divisor)
15 irem         // Pop two ints, remainder them, push result
                // Pop int, jump if equal to zero:
16 ifne 24      // if (numToCheck % divisor == 0)
19 iconst_0     // Push int constant 0
20 istore_2     // Pop into local var 2: foundPrime = false;
21 goto 32      // Jump out of inner for loop

// At this point, the body of the inner for loop is done. Now just

```



```

// perform the third statement of the for expression: decrement
// divisor.
24 iinc 3 -1    // Increment local var 3 by -1: --divisor

// The test for the inner for loop's termination condition
// begins here. This loop will keep on looping while (divisor != 1).
27 iload_3      // Push int from local var 3 (divisor)
28 iconst_1     // Push int constant 1
29 if_icmpgt 13  // Pop top two ints, jump if greater than

// At this point, the inner for loop has completed. Next check
// to see if a prime number was found.
32 iload_2      // Push int from local var 2 (foundPrime)
33 ifeq 38      // Pop top int, jump if zero: if (foundPrime) {
36 iload_1      // Push int from local var 1 (numToCheck)
37 istore_0     // Pop into local var 0: primeNum = numToCheck;
38 iinc 1 1     // Increment local var 1 by 1: ++numToCheck;
41 goto 4       // Jump back to top of outer for loop.

end

```

The **javac** compiler placed local variable **primeNum** from the source into local variable slots 0 on the stack frame. It put **numToCheck** into slot 1, **foundPrime** into slot 2, and **divisor** into slot 3. As mentioned above, as this method finds each successive prime number, it places the number into the **primeNum** variable. As you run the simulation, therefore, you will see the prime numbers appear sequentially in the **int** value stored in local variable slot 0.

One thing to note about this bytecode sequence is that it demonstrates the way in which **booleans** from Java source code are treated on the stack frame by Java bytecodes. The value stored in local variable slot 2, which represents the **boolean foundPrime** variable from the source, is an **int**. It is set to **true** or **false** by instructions that push a constant **int** zero or one. Its boolean value is checked by instructions that compare an **int** against zero.

Another thing to note about this simulation is that eventually the **numToCheck** value will overflow. When it does, the virtual machine will throw no exceptions. It will just continue executing the **findPrimes()** method with **int** values that no longer hold any relationship to the prime numbers.

To drive the Prime Time simulation, use the Step, Reset, Run, and Stop buttons. Each time you press the

Step button, the simulator will execute the instruction pointed to by the pc register. If you press the Run button, the simulation will continue with no further coaxing on your part until you press the Stop button. To start the simulation over, press the Reset button. For each step of the simulation, a panel at the bottom of the applet contains an explanation of what the next instruction will do. Happy clicking.

Applet Viewer: PrimeTime.class

Applet

PRIME TIME

Local Variables				The Method			
index	hex	value	value	offset	bytecode	mnemonic	
0	00000007	7		13	1b	iload_1	
1	00000009	9		14	1d	iload_3	
2	00000001	1		pc > 15	70	irem	
3	00000004	4		16	9a	ifne 24	
				17	00		
				18	08		
				19	03	iconst_0	
				20	3d	istore_2	
				21	a7	goto 32	
				22	00		
				23	0b		
				24	84	iinc 3 -1	
				25	03		

Operand Stack			
offset	hex	value	value
0	00000009	9	
1	00000004	4	

optop >

Step Run

Reset Stop

irem will pop two integers, remainder them, and push the result.

On the CD-ROM

The CD-ROM contains the source code examples from this chapter in the *integer* directory. The Prime Time applet is contained in a web page on the CD-ROM in file *applets/PrimeTime.html*. The source code for this applet is found alongside its class files, in the *applets/JVMSimulators* and *applets/JVMSimulators*.

The Resources Page

For more information about the material presented in this chapter, visit the resources page: <http://www.artima.com/insidejvm/integer.html>.

Chapter Thirteen

Logic

This chapter describes the instructions that perform bitwise logical operations inside the Java Virtual Machine. These instructions include opcodes to perform shifting and boolean operations on integers. The boolean operations are performed on individual bits of integer values.

Accompanying this chapter on the CD-ROM is an applet that interactively illustrates the material presented in the chapter. The applet, named *Logical Results*, simulates the Java Virtual Machine executing a method that includes uses several of the logic opcodes. At the end of this chapter, you will find a description of this applet and the bytecodes it executes.

The Logic Opcodes

The Java Virtual Machine's logic capabilities operate on `ints` and `longs`. These operations treat `ints` and `longs` not so much as signed two's-complement numbers, necessarily, but as generic bit patterns. Integer shifting is accomplished via the `ishl`, `ishr`, and `iushr` opcodes. Java's `operator` is implemented by `ishl`. The `operator` is implemented by `ishr`, and the `operator` is implemented by `iushr`. The difference between `ishr` and `iushr` is that only `ishr` does sign extension. Table 13-1 shows the instructions that shift `ints` left and right.

Table 13-1. Shifting `ints`

Opcode	Operand(s)	Description
<code>ishl</code>	(none)	shifts <code>int</code> left
<code>ishr</code>	(none)	arithmetic shifts <code>int</code> right
<code>iushr</code>	(none)	logical shifts <code>int</code> right

Table 13-2 shows the instructions that shift `longs` left and right.

Table 13-2. Shifting `longs`

Opcode	Operand(s)	Description
<code>lshl</code>	(none)	shifts <code>long</code> left
<code>lshr</code>	(none)	arithmetic shifts <code>long</code> right

lshr (none) logical shifts long right

The opcodes shown in Table 13-3 perform bitwise logical operations on `ints`. The opcodes implement Java's `&`, `|`, and `^` operators.

Table 13-3. Bitwise logical operations on `ints`

Opcode	Operand(s)	Description
<code>iand</code>	(none)	boolean ands two <code>ints</code>
<code>ior</code>	(none)	boolean ors two <code>ints</code>
<code>ixor</code>	(none)	boolean xors two <code>ints</code>

Table 13-4 shows the opcodes that perform bitwise logical operations on `longs`.

Table 13-4. Bitwise logical operations on `longs`

Opcode	Operand(s)	Description
<code>land</code>	(none)	boolean ands two <code>longs</code>
<code>lor</code>	(none)	boolean ors two <code>longs</code>
<code>lxor</code>	(none)	boolean xors two <code>longs</code>

As mentioned in previous chapters, there is no native `boolean` type in the Java Virtual Machine. The Java Virtual Machine uses `ints` to represent `booleans`. The instruction set includes many instructions that decide whether or not to jump based on an `int` value interpreted as a `boolean`. This instructions are covered in Chapter 16, "Control Flow."

Logical Results: A Simulation

The *Logical Results* applet, shown in Figure 13-1, demonstrates a Java Virtual Machine executing a sequence of bytecodes. This applet is in a web page on the CD-ROM in file `applets/LogicalResults.html`. The bytecode sequence in the simulation was generated by `javac` for the `incrementLogically()` method of the `VulcanCounter` class:

```
begin

// On CD-ROM in file opcodes/ex1/VulcanCounter.java
class VulcanCounter {

    static void incrementLogically() {

        int spock = 0;

        for (;;) {

            int tempSpock = spock;

            for (int i = 0; i < 32; ++i) {
```

```

        int mask = 0x1 i;
        if ((tempSpock & mask) == 0) {
            tempSpock |= mask; // Change 0 to 1
            break;
        }
        else {
            tempSpock &= ~mask; // Change 1 to 0
        }
    }
    spock = tempSpock;
}
}
}

```

end

The bytecodes generated by javac for incrementLogically() are:

begin

```

0 iconst_0    // Push int constant 0.
1 istore_0    // Pop to local variable 0: int spock = 0;
2 iload_0     // Push local variable 0 (spock).
3 istore_1    // Pop to local variable 1: int tempSpock = spock;
4 iconst_0    // Push int constant 0.
5 istore_2    // Pop to local variable 2: int i = 0;
6 goto 35     // Jump unconditionally ()
9 iconst_1    // Push int constant 1.
10 iload_2    // Push local variable 2 (i).
11 ishl       // Arithmetic shift left top int (i) by next
              // to top int (1).
12 istore_3    // Pop to local variable 3: int mask = i 0x1;
13 iload_1    // Push local variable 1 (tempSpock).

```

```

14 iload_3      // Push local variable 3 (mask).
15 iand         // Bitwise AND top two ints: (spock & mask)
16 ifne 26      // Jump if top of stack is not equal to zero:
                // if ((spock & mask) == 0) {
19 iload_1      // Push local variable 1 (tempSpock).
20 iload_3      // Push local variable 3 (mask).
21 ior          // Bitwise OR top two ints (tempSpock | mask)
22 istore_1     // Pop to local variable 1: tempSpock |= mask;
23 goto 41      // Jump unconditionally (to just after
                // inner for): break;
26 iload_1      // Push local variable 1 (tempSpock).
27 iload_3      // Push local variable 3 (mask).
28 iconst_m1    // Push -1.
29 ixor         // Bitwise EXCLUSIVE-OR top two ints: ~mask
30 iand         // Bitwise AND top two ints: tempSpock & (~mask)
31 istore_1     // Pop to local variable 1: tempSpock &= ~mask;
32 iinc 2 1     // Increment local variable 2 by 1: ++i
35 iload_2      // Push local variable 2 (i).
36 bipush 32    // Push integer constant 32.
38 if_icmplt 9   // Jump (to top of inner for) if "next to top"
                // integer is less than "top" integer: i < 32
41 iload_1      // Push local variable 1 (tempSpock)
42 istore_0     // Pop to local variable 0: spock = tempSpock;
43 goto 2       // Jump unconditionally (to top of outer for).

end

```

The `incrementLogically()` method repeatedly increments an int without using the `+` or `++` operators. It uses only the logical operators: `&`, `|`, and `~`. An increment is accomplished by searching through the bits of the current int, starting with the lowest order bit, and turning ones to zeros. As soon as a one is encountered in the current int, it is changed to zero and the search stops. The resultant int now represents the old int incremented by one. The process is started

again on the new int. Each incremented number is stored in the spock variable. spock is local variable zero in the compiled bytecodes, so you can watch local variable zero count 0, 1, 2, 3, and so on.

To drive the *Logical Results* simulation, use the Step, Reset, Run, and Stop buttons. Each time you press the Step button, the simulator will execute the instruction pointed to by the pc register. If you press the Run button, the simulation will continue with no further coaxing on your part until you press the Stop button. To start the simulation over, press the Reset button. For each step of the simulation, a panel at the bottom of the applet contains an explanation of what the next instruction will do. Happy clicking.

Applet Viewer: LogicalResults.class

Applet

LOGICAL RESULTS

Local Variables			
index	hex value	value	
0	00000003	3	
1	00000003	3	
2	00000000	0	
3	00000001	1	

optop 3

Operand Stack			
offset	hex value	value	
0	00000003	3	
1	00000001	1	
2	ffffff	-1	

optop >

The Method			
pc	offset	bytecode	mnemonic
	26	1b	iload_1
	27	1d	iload_3
	28	02	iconst_m1
pc >	29	82	ixor
	30	7e	iand
	31	3c	istore_1
	32	84	iinc 2 1
	33	02	
	34	01	
	35	1c	iload_2
	36	10	bipush 32
	37	20	
	38	a1	if_icmplt 9

Step Run

Reset Stop

ixor will pop the top two ints off the stack, bitwise-xor them, and push the result back onto the stack.

On the CD-ROM

The CD-ROM contains the source code examples from this chapter in the opcodes directory. The *Logical Results* applet is in a web page on the CD-ROM in file applets/LogicalResults.html. The source code for this applet is found alongside its class files, in the applets/JVMSimulators and applets/JVMSimulators/COM/artima/jvmsim directories.

The Resources Page

For more information about the material presented in this chapter, visit the resources page: <http://www.artima.com/insidejvm/logic.html>

Chapter Fourteen

Floating-Point Arithmetic

This chapter describes the floating point numbers and the instructions that perform floating point arithmetic inside the Java Virtual Machine. The floating point numbers described here conform to the IEEE 754 floating point standard, which is the standard to which all Java Virtual Machine implementations must adhere.

Accompanying this chapter on the CD-ROM are two applets that interactively illustrate the material presented in the chapter. One applet, named *Inner Float*, allows you to view and manipulate the individual components that make up a floating point number. The other applet, named *Circle of Squares*, simulates the Java Virtual Machine executing a method that uses several of the floating point opcodes. Within this chapter, you will find descriptions of both of these applets.

Floating Point Numbers

The Java Virtual Machine's floating-point support adheres to the IEEE-754 1985 floating-point standard. This standard defines the format of 32-bit and 64-bit floating-point numbers and defines the operations upon those numbers. In the Java Virtual Machine, floating-point arithmetic is performed on 32-bit `floats` and 64-bit `doubles`. For each opcode that performs arithmetic on `floats`, there is a corresponding opcode that performs the same operation on `doubles`.

A floating-point number has four parts--a sign, a mantissa, a radix, and an exponent. The sign is either a 1 or -1. The mantissa, always a positive number, holds the significant digits of the floating-point number. The exponent indicates the positive or negative power of the radix that the mantissa and sign should be multiplied by. The four components are combined as follows to get the floating-point value:

begin math formula

$$\text{sign} * \text{mantissa} * \text{radix}^{\text{exponent}}$$

end math formula

Floating-point numbers have multiple representations, because one can always multiply the mantissa of any floating-point number by some power of the radix and change the exponent to get the original number. For example, as shown in Table 14-1, the number -5 can be represented equally by any of the following forms in radix 10.

Table 14-1. Forms of -5

Sign	Mantissa	Radix ^{exponent}
-1	50	10 ⁻¹
-1	5	10 ⁰
-1	0.5	10 ¹
-1	0.05	10 ²

For each floating-point number there is one representation that is said to be *normalized*. A floating-point number is normalized if its mantissa is within the range defined by the following relation:

begin math formula

$$1/\text{radix} \leq \text{mantissa} < 1$$

end math formula

A normalized radix 10 floating-point number has its decimal point just to the left of the first non-zero digit in the mantissa. The normalized floating-point representation of -5 is $-1 * 0.5 * 10^1$. In other words, a normalized floating-point number's mantissa has no non-zero digits to the left of the decimal point and a non-zero digit just to the right of the decimal point. Any floating-point number that doesn't fit into this category is said to be *denormalized*. Note that the number zero has no normalized representation, because it has no non-zero digit to put just to the right of the decimal point. "Why be normalized?" is a common exclamation among zeros.

Floating-point numbers in the Java Virtual Machine use a radix of two, so they have the following form:

begin math formula

$$\text{sign} * \text{mantissa} * 2^{\text{exponent}}$$

end math formula

The mantissa of a floating-point number in the Java Virtual Machine is expressed as a binary number. A normalized mantissa has its binary point (the base-two equivalent of a decimal point) just to the left of the most significant non-zero digit. Because the binary number system has just two digits--zero and one--the most significant digit of a normalized mantissa is always a one.

The most significant bit of a `float` or `double` is its sign bit. The mantissa occupies the 23 least significant bits of a `float` and the 52 least significant bits of a `double`. The exponent, 8 bits in a `float` and 11 bits in a `double`, sits between the sign and mantissa. The format of a `float` is shown below. The sign bit is shown as an "s," the exponent bits are shown as "e," and the mantissa bits are shown as "m":

begin math formula

end math formula

Exponents that are neither all ones nor all zeros indicate the power of two by which to multiply the normalized mantissa. The power of two can be determined by interpreting the exponent bits as a positive number, and then subtracting a bias from the positive number. For a `float`, the bias is 126. For a `double`, the bias is 1023. For example, an exponent field in a `float` of 00000001 yields a power of two by subtracting the bias (126) from the exponent field interpreted as a positive integer (1). The power of two, therefore, is $(1 - 126)$, which is -125. This is the smallest possible power of two for a `float`. At the other extreme, an exponent field of 11111110 yields a power of two of $(254 - 126)$ or 128. The number 128 is the largest power of two available to a `float`. Several examples of normalized `floats` are shown in Table 14-3.

Table 14-3. Normalized float values

Value	float bits (sign exponent mantissa)	Unbiased exponent
Largest positive (finite) float	0 11111110 111111111111111111111111	128
Largest negative (finite) float	1 11111110 111111111111111111111111	128
Smallest normalized float	1 00000001 000000000000000000000000	-125
Pi	0 10000000 10010010000111111011011	2

An exponent of all zeros indicates the mantissa is denormalized, which means the unstated leading bit is a zero instead of a one. The power of two in this case is the same as the lowest power of two available to a normalized mantissa. For the `float`, this is -125. This means that normalized mantissas multiplied by two raised to the power of -125 have an exponent field of 00000001, while denormalized mantissas multiplied by two raised to the power of -125 have an exponent field of 00000000.

The allowance for denormalized numbers at the bottom end of the range of exponents supports gradual underflow. If the lowest exponent was instead used to represent a normalized number, underflow to zero would occur for larger numbers. In other words, leaving the lowest exponent for denormalized numbers allows smaller numbers to be represented. The smaller denormalized numbers have fewer bits of precision than normalized numbers, but this is preferable to underflowing to zero as soon as the exponent reaches its minimum normalized value. Table 14-4 shows several denormalized floating-point values.


Table 14-4. Denormalized float values

Value	float bits (sign exponent mantissa)
Smallest positive (non-zero) float	0 00000000 000000000000000000000001
Smallest negative (non-zero) float	1 00000000 000000000000000000000001
Largest denormalized float	1 00000000 111111111111111111111111
Positive zero	0 00000000 000000000000000000000000
Negative zero	1 00000000 000000000000000000000000

Inner Float: A Java float Reveals its Inner Nature

The applet shown in Figure 14-1 lets you play around with the floating-point format. The value of a `float` is displayed in several formats. The radix two scientific notation format shows the mantissa and exponent in base ten. Before being displayed, the actual mantissa is multiplied by 2

²⁴, which yields an integral number, and the unbiased exponent is decremented by 24. Both the integral mantissa and exponent are then easily converted to base ten and displayed. The applet is embedded in a web page on the CD-ROM in file `applets/InnerFloat.html`. The text of the web page includes several button-press sequences that demonstrate various properties of floating point numbers.



The Floating Point Opcodes

Table 14-5 shows the opcodes that pop two floating-point values from the top of the stack, add them, and push the result. The type of the values is indicated by the opcode itself, and the result always has the same type as the numbers being added. No exceptions are thrown by these opcodes. Overflow results in a positive or negative infinity, and underflow results in a positive or negative zero.

Table 14-5. Floating-point addition

Opcode	Operand(s)	Description
fadd	(none)	pops two floats, adds them, and pushes the float result
dadd	(none)	pops two doubles, adds them, and pushes the double result

Subtraction is performed on floats and doubles via the opcodes shown in Table 14-6. Each opcode causes the top two values of the appropriate type to be popped off the stack. The topmost value is subtracted from the value just beneath the topmost value. The result is pushed back onto the stack. No exceptions are thrown by either of these opcodes.

Table 14-6. Floating-point subtraction

Opcode	Operand(s)	Description
fsub	(none)	pops two floats, subtracts them, and pushes the

		float result
dsub	(none)	pops two doubles, subtracts them, and pushes the double result

Multiplication of floats and doubles is accomplished via the opcodes shown in Table 14-7. Each opcode causes two values of the same type to be popped off the stack and multiplied. The result, of the same type as the numbers being multiplied, is pushed back onto the stack. No exceptions are thrown.

Table 14-7. Floating-point multiplication

Opcode	Operand (s)	Description
fmul	(none)	pops two floats, multiplies them, and pushes the float result
dmul	(none)	pops two doubles, multiplies them, and pushes the double result

The division is performed on floats and doubles by the opcodes shown in Table 14-8. The division opcodes cause the top two values of the appropriate type to be popped off the stack. The value immediately beneath the topmost value is divided by the topmost value. (In other words, the value pushed first is the dividend or numerator. The value pushed second--the top of the stack--is the divisor or denominator.) The result of the division is pushed back onto the stack.

Table 14-8. Floating-point division

Opcode	Operand (s)	Description
fdiv	(none)	pops two floats, divides them, and pushes the float result
ddiv	(none)	pops two doubles, divides them, and pushes the double result

No exception is thrown as a result of any floating-point division. Floating-point division of a finite value by zero yields a positive or negative infinity. Floating-point division of zero by zero yields NaN. A summary of the result of dividing various combinations of infinity, zero, NaN, and finite values is given in Table 14-10.

The remainder operation is performed on floats and doubles via the opcodes shown in Table 14-9. The following opcodes cause the top two values to be popped from the stack. The value just beneath the topmost value is divided by the topmost value, and the remainder of that division is pushed back onto the stack.

Table 14-9. Floating-point remainder

Opcode	Operand (s)	Description
--------	----------------	-------------

<code>frem</code>	(none)	pops two <code>float</code> s, divides them, and pushes the <code>float</code> remainder
<code>drem</code>	(none)	pops two <code>double</code> s, divides them, and pushes the <code>double</code> remainder

No exception is thrown as a result of any floating-point remainder operation. Floating-point remainder of any value divided by zero yields a NaN result. A summary of the result of computing the remainder various combinations of infinity, zero, NaN, and finite values is given in Table 14-10.

Table 14-10. Results of various floating-point divisions

a	b	a/b	a%b
Finite	+0.0	+Infinity	NaN
Finite	+Infinity	+0.0	a
+0.0	+0.0	NaN	NaN
+Infinity	Finite	+Infinity	NaN
+Infinity	+Infinity	NaN	NaN

The remainder operation provided by `frem` and `drem` follows the same rule as the integer remainder operation provided by `irem` and `lrem`:

```
begin
(a/b)*b + a%b == a
end
```

In the floating point case, the division `(a/b)` must be converted to an `int` or `long` to remove the fraction, as in:

```
begin
((long)(a/b))*b + a%b == a
end
```

This remainder does not follow to the IEEE 754 standard. To obtain a floating point remainder as defined by IEEE 754, use the `IEEEremainder()` method of `java.lang.Math`.

The opcodes shown in Table 14-11 perform arithmetic negation on `float`s and `double`s. Negation opcodes pop the top value from the stack, negates it, and pushes the result.

Table 14-11. Floating-point negation

Opcode	Operand(s)	Description
<code>fneg</code>	(none)	pops a <code>float</code> , negates it, and pushes the result
<code>dneg</code>	(none)	pops a <code>double</code> , negates it, and pushes the result

Circle of Squares: A Simulation

The *Circle of Squares* applet, shown in Figure 14-2, demonstrates a Java Virtual Machine executing a sequence of bytecodes that perform floating-point arithmetic. This applet is part of a web page on the CD-ROM in file `applets/CircleOfSquares.html`. The bytecode sequence in the simulation was generated by `javac` for the `squareItForever()` method of the class shown below:

begin

```
// On CD-ROM in file opcodes/ex1/SquareCircle.java
class SquareCircle {
```

```
    static void squareItForever() {
```

```
        float f = 2;
```

```
        for (;;) {
```

```
            f *= f;
```

```
            f = 0 - f;
```

```
        }
```

```
    }
```

```
}
```

end

The bytecodes generated by `javac` for `squareItForever()` are shown below:

begin

```
0 fconst_2  // Push float constant 2.
```

```
1 fstore_0  // Pop to local variable 0 (float f): float f = 2;
```

```
2 fload_0   // Push local variable 0 (float f).
```

```
3 fload_0   // Push local variable 0 (float f).
```

```
4 fmul      // Pop top two floats, multiply, push float result.
```

```
5 fstore_0  // Pop to local variable 0 (float f): f *= f;
```

```
6 fconst_0  // Push float constant 0.
```

```
7 fload_0   // Push local variable 0 (float f).
```

```
8 fsub      // Subtract top float from next to top float:
```

```
    // imByte = (byte) imInt;
```

```

9 fstore_0 // Pop result to local variable 0 (float f): f = 0 - f;
10 goto 2   // Jump back to the first fload_0 instruction:
           // for (;;) {}

end

```

The `squareItForever()` method repeatedly squares a `float` value until it hits infinity. Each time the `float` is squared it is also negated. The `float` starts out as 2. It only takes seven iterations before infinity is reached, which isn't nearly as long as it takes in real life. The hex representation of the bits that make up the `float` are shown in the "hex value" column in the applet. The "value" column shows the number as humans are used to seeing it. This human-friendly value is generated by the `Float.toString()` method.

To drive the *Circle of Squares* simulation, use the Step, Reset, Run, and Stop buttons. Each time you press the Step button, the simulator will execute the instruction pointed to by the pc register. If you press the Run button, the simulation will continue with no further coaxing on your part until you press the Stop button. To start the simulation over, press the Reset button. For each step of the simulation, a panel at the bottom of the applet contains an explanation of what the next instruction will do. Happy clicking.

Applet Viewer: CircleOfSquares.class

Applet

CIRCLE OF SQUARES

Local Variables			
index	hex value	value	
0	c7800000	-65536.0	

optop: 1

Operand Stack			
offset	hex value	value	
0	4f800000	4.2949673E9	

optop > 1

pc: 5

The Method			
offset	bytecode	mnemonic	
0	0d	fconst_2	
1	43	fstore_0	
2	22	fload_0	
3	22	fload_0	
4	6a	fmul	
pc > 5	43	fstore_0	
6	0b	fconst_0	
7	22	fload_0	
8	66	fsub	
9	43	fstore_0	
10	a7	goto 2	
11	ff		
12	f8		

Step Run Reset Stop

fstore_0 will pop the float off the top of the stack and store it in local variable 0.

On the CD-ROM

The CD-ROM contains the source code examples from this chapter in the `opcodes` directory.

The *Inner Float* applet is embedded in a web page on the CD-ROM in file `applets/InnerFloat.html`. The text of the web page includes several button-press sequences that demonstrate various properties of floating point numbers. The source code for this applet is found alongside its class files in the `applets/InnerFloat` directory.

The *Circle of Squares* applet is in a web page on the CD-ROM in file `applets/CircleOfSquares.html`. The source code for this applet is found alongside its class files, in the `applets/JVMSimulators` and `applets/JVMSimulators/COM/artima/jvmsim` directories.

The Resources Page

For more information about the material presented in this chapter, visit the resources page:
<http://www.artima.com/insidejvm/float.html>

Chapter Fifteen

Objects and Arrays

This chapter describes the instructions that create and manipulate objects and arrays inside the Java Virtual Machine. All of these instructions involve the heap. Most of them refer to entries in the constant pool.

Accompanying this chapter on the CD-ROM is an applet that interactively illustrates the material presented in the chapter. The applet, named *Three-Dimensional Array*, simulates the Java Virtual Machine executing a method that allocates and initializes a three dimensional array. At the end of this chapter, you will find a description of this applet and the bytecodes it executes.

A Refresher on Objects and Arrays

As mentioned in earlier chapters, in the Java Virtual Machine, memory is allocated on the garbage-collected heap only as objects. There is no way to allocate memory for a primitive type on the heap, except as part of an object. If you want to use a primitive type where an `Object` reference is needed, you can allocate a wrapper object for the type from the `java.lang` package. For example, there is an `Integer` class that wraps an `int` type with an object. Only object references and primitive types can reside on the Java stack as local variables. Objects can never reside on the Java stack.

The architectural separation of objects and primitive types in the Java Virtual Machine is reflected in the Java programming language, in which objects cannot be declared as local variables--only object references and primitive types. Upon declaration, an object reference refers to nothing. Only after the reference has been explicitly initialized--either with a reference to an existing object or with a call to `new`--does the reference refer to an actual object.

In the Java Virtual Machine instruction set, all objects are instantiated and accessed with the same set of opcodes, except for arrays. As mentioned in earlier chapters, arrays in Java are full-fledged objects. Like any other Java object, arrays are created dynamically. Array references can be used anywhere a reference to type `Object` is called for, and any method of `Object` can be invoked on an array. Still, in the Java Virtual Machine, array operations are handled by special bytecodes.

As with any other object, arrays cannot be declared as local variables; only array references can. Array objects themselves always contain either an array of primitive types or an array of object references. If you declare an array of objects, you get an array of object references. The objects themselves must be explicitly created with `new` and assigned to the elements of the array.

Opcodes for Objects

Instantiation of new objects is accomplished via the `new` opcode, shown in Table 15-1. Two one-byte operands follow the `new` opcode. These two bytes are combined to form an unsigned 16-bit index into the constant pool. The constant pool entry at the specified offset gives information about the class of the new object. If it hasn't already, the virtual machine resolves the constant pool entry. It creates a new instance of the object on the heap, initializes the object's instance variables to default initial values, and pushes the reference to the new object onto the stack.

Table 15-1. Object creation

Opcode	Operand(s)	Description
<code>new</code>	<code>indexbyte1</code> , <code>indexbyte2</code>	creates a new object on the heap, pushes reference

Table 15-2 shows the opcodes that put and get object fields. These opcodes, `putfield` and `getfield`, operate only on fields that are instance variables. Static variables are accessed by `putstatic` and `getstatic`, which are described later. The `putfield` and `getfield` instructions each take two one-byte operands, which are combined to form an unsigned 16-bit index into the constant pool. The constant pool entry at that index contains information about the type, size, and offset of the field. If it hasn't already, the virtual machine resolves the constant pool entry. The object reference is taken from the stack in both the `putfield` and `getfield` instructions. The `putfield` instruction takes the instance variable value from the stack, and the `getfield` instruction pushes the retrieved instance variable value onto the stack.

Table 15-2. Accessing instance variables

Opcode	Operand(s)	Description
<code>putfield</code>	<code>indexbyte1</code> , <code>indexbyte2</code>	set field, indicated by index, of object to value (both taken from stack)
<code>getfield</code>	<code>indexbyte1</code> , <code>indexbyte2</code>	pushes field, indicated by index, of object (taken from stack)

Class variables are accessed via the `getstatic` and `putstatic` opcodes, as shown in Table 15-3. Both `getstatic` and `putstatic` take two one-byte operands, which are combined by the Java Virtual Machine to form a 16-bit unsigned offset into the constant pool. The constant pool item at that location gives information about one static field of a class. If it hasn't already, the virtual machine resolves the constant pool entry. Because there is no particular object associated with a static field, there is no object reference used by either `getstatic` or `putstatic`. The `putstatic` instruction takes the value to assign from the stack. The `getstatic` instruction pushes the retrieved value onto the stack.

Table 15-3. Accessing class variables

Opcode	Operand(s)	Description
<code>putstatic</code>	<code>indexbyte1</code> , <code>indexbyte2</code>	set field, indicated by index, of object to value (both taken from stack)

<code>getstatic</code>	<code>indexbyte1,</code> <code>indexbyte2</code>	pushes field, indicated by index, of object (taken from stack)
------------------------	---	---

The following opcodes check to see whether the object reference on the top of the stack refers to an instance of the class or interface indexed by the operands following the opcode. In both cases, the virtual machine forms an unsigned 16-bit index into the constant pool from the two bytes that follow the opcode. If it hasn't already, the virtual machine resolves the constant pool entry.

The `checkcast` instruction throws `CheckCastException` if the object is not an instance of the specified class or interface. Otherwise, `checkcast` does nothing. The object reference remains on the stack and execution is continued at the next instruction. This instruction ensures that casts are safe at run time and forms part of the Java Virtual Machine's security blanket.

The `instanceof` instruction, shown in Table 15-4, pops the object reference from the top of the stack and pushes true or false. If the object is indeed an instance of the specified class or interface, then true is pushed onto the stack, otherwise, false is pushed onto the stack. The `instanceof` instruction is used to implement the `instanceof` keyword of the Java language, which allows you to test whether an object is an instance of a particular class or interface.

Table 15-4. Type checking

Opcode	Operand(s)	Description
<code>checkcast</code>	<code>indexbyte1,</code> <code>indexbyte2</code>	Throws <code>ClassCastException</code> if objectref on stack cannot be cast to class at index
<code>instanceof</code>	<code>indexbyte1,</code> <code>indexbyte2</code>	Pushes <code>true</code> if objectref on stack is an instance of class at index, else pushes <code>false</code>

Opcodes for Arrays

Instantiation of new arrays is accomplished via the `newarray`, `anewarray`, and `multianewarray` opcodes, shown in Table 15-6. The `newarray` opcode is used to create arrays of primitive types other than object references. The particular primitive type is specified by a single one-byte operand, the "atype," following the `newarray` opcode. The `newarray` instruction can create arrays for `byte`, `short`, `char`, `int`, `long`, `float`, `double`, or `boolean`. Table 15-5 shows the legal values for atype and the corresponding array types.

Table 15-5. Values for atype

Array Type	atype
<code>T_BOOLEAN</code>	4
<code>T_CHAR</code>	5
<code>T_FLOAT</code>	6
<code>T_DOUBLE</code>	7
<code>T_BYTE</code>	8

T_SHORT	9
T_INT	10
T_LONG	11

Note that the instructions that create arrays is one place in the Java Virtual Machine where **boolean** values are referred to explicitly. This enables implementations, especially those that must operate under constrained memory requirements, to compress arrays of **boolean** values into bit maps. In such a representation, each **boolean** element of an array would be represented by one bit. On implementations that have more memory available, **boolean** arrays may instead be implemented as arrays of **byte**s. This representation would likely yield better performance than a bit map approach, though it would consume more memory. No matter what internal implementation a particular virtual machine uses for **boolean** arrays, the elements of a **boolean** array are accessed with the same opcodes used to access **byte** arrays. These opcodes are described later in this chapter.

The **anewarray** instruction creates an array of object references. Two one-byte operands follow the **anewarray** opcode and are combined to form an unsigned 16-bit index into the constant pool. A description of the class of object for which the array is to be created is found in the constant pool at the specified index. If it hasn't already, the virtual machine resolves the constant pool entry. This instruction allocates space for the array of object references and initializes the references to null.

The **multianewarray** instruction is used to allocate multidimensional arrays, which are simply arrays of arrays. A multidimensional array could alternatively be allocated with repeated use of the **anewarray** and **newarray** instructions. The **multianewarray** instruction simply compresses the bytecodes needed to create multidimensional arrays into one instruction. Two one-byte operands follow the **multianewarray** opcode and are combined to form an unsigned 16-bit index into the constant pool. A description of the class of object for which the array is to be created is found in the constant pool at the specified index. If it hasn't already, the virtual machine resolves the symbolic reference. Immediately following the two one-byte operands that form the constant pool index is an unsigned one-byte operand that specifies the number of dimensions in this multidimensional array. The sizes for each dimension are popped off the stack. This instruction allocates space for all arrays that are needed to implement the multidimensional arrays.

The constant pool entry referred to by a **multianewarray** instruction contains a **Constant_CLASS** entry with an array class name. For example, the constant pool entry for a four-dimensional array of float would have the name "[[[[F". The class name in the constant pool entry may have more (but not less) left-brackets than indicated by the dimensions byte. The virtual machine always creates the number of dimensions specified by the dimensions byte.

Table 15-6. Creating new arrays

Opcode	Operand(s)	Description
newarray	atype	pops length, allocates new array of primitive types of type indicated by atype , pushes objectref of new array
anewarray	indexbyte1 , indexbyte2	pops length, allocates a new array of objects of class indicated by indexbyte1 and indexbyte2 , pushes objectref of new array

<code>multianewarray</code>	<code>indexbyte1,</code> <code>indexbyte2,</code> <code>dimensions</code>	pops dimensions number of array lengths, allocates a new multidimensional array of class indicated by <code>indexbyte1</code> and <code>indexbyte2</code> , pushes objectref of new array
-----------------------------	---	---

Table 15-7 shows the `arraylength` instruction, which pops an array reference off the top of the stack and pushes the length of that array.

Table 15-7. Getting an array length

Opcode	Operand (s)	Description
<code>arraylength</code>	(none)	pops objectref of an array, pushes length of that array

The opcodes shown in Table 15-8 retrieve an element from an array. The array index and array reference are popped from the stack, and the value at the specified index of the specified array is pushed back onto the stack. The `baload` opcode converts the `byte` or `boolean` value to `int` by sign extending, then pushes the `int`. Likewise, the `saload` opcode converts the short value to `int` by sign extending, then pushes the `int`. The `caload` instruction converts the `char` to an `int` by zero-extending, then pushes the `int`.

Table 15-8. Retrieving an array element

Opcode	Operand (s)	Description
<code>baload</code>	(none)	pops index and arrayref of an array of <code>byte</code> s or <code>boolean</code> s, pushes <code>arrayref[index]</code>
<code>caload</code>	(none)	pops index and arrayref of an array of <code>char</code> s, pushes <code>arrayref[index]</code>
<code>saload</code>	(none)	pops index and arrayref of an array of <code>short</code> s, pushes <code>arrayref[index]</code>
<code>iaload</code>	(none)	pops index and arrayref of an array of <code>int</code> s, pushes <code>arrayref[index]</code>
<code>laload</code>	(none)	pops index and arrayref of an array of <code>long</code> s, pushes <code>arrayref[index]</code>
<code>faload</code>	(none)	pops index and arrayref of an array of <code>float</code> s, pushes <code>arrayref[index]</code>
<code>daload</code>	(none)	pops index and arrayref of an array of <code>double</code> s, pushes <code>arrayref[index]</code>
<code>aaload</code>	(none)	pops index and arrayref of an array of objectrefs, pushes <code>arrayref[index]</code>

Table 15-9 shows the opcodes that store a value into an array element. The value, index, and array

reference are popped from the top of the stack. The **bastore** instruction just stores the lower eight bits of the popped **int** value. The **sastore** and **castore** instructions just store the lower 16 bits of the popped **int** value.

Table 15-9. Storing to an array element

Opcode	Operand (s)	Description
bastore	(none)	pops value, index, and arrayref of an array of bytes or booleans , assigns <code>arrayref[index] = value</code>
castore	(none)	pops value, index, and arrayref of an array of chars , assigns <code>arrayref[index] = value</code>
sastore	(none)	pops value, index, and arrayref of an array of shorts , assigns <code>arrayref[index] = value</code>
iastore	(none)	pops value, index, and arrayref of an array of ints , assigns <code>arrayref[index] = value</code>
lastore	(none)	pops value, index, and arrayref of an array of longs , assigns <code>arrayref[index] = value</code>
fastore	(none)	pops value, index, and arrayref of an array of floats , assigns <code>arrayref[index] = value</code>
dastore	(none)	pops value, index, and arrayref of an array of doubles , assigns <code>arrayref[index] = value</code>
aastore	(none)	pops value, index, and arrayref of an array of objectrefs , assigns <code>arrayref[index] = value</code>

Three-Dimensional Array: A Simulation

The *Three-Dimensional Array* applet, shown in Figure 15-1, demonstrates a Java Virtual Machine executing a sequence of bytecodes. This applet is embedded in a web page on the CD-ROM in file `applets/ThreeDArray.html`.

The bytecode sequence in the simulation was generated by `javac` for the `initAnArray()` method of the class shown below:

```
begin

// On CD-ROM in file opcodes/ex1/ThreeDTree.java
class ThreeDTree {

    static void initAnArray() {

        int[][][] threeD = new int[5][4][3];

        for (int i = 0; i < 5; ++i) {
```

```

        for (int j = 0; j < 4; ++j) {
            for (int k = 0; k < 3; ++k) {
                threeD[i][j][k] = i + j + k;
            }
        }
    }
}

```

end

The bytecodes generated by javac for `initAnArray()` are shown below:

begin

```

0 iconst_5      // Push constant int 5.
1 iconst_4      // Push constant int 4.

2 iconst_3      // Push constant int 3.
3 multianewarray #2 dim #3 <Class [[[I
                // Create a new multi-dimensional array using
                // constant pool entry #2 as the class (which
                // is [[[I, an 3D array of ints) with a
                // dimension of 3.

7 astore_0      // Pop object ref into local variable 0:
                // int threeD[][][] = new int[5][4][3];

8 iconst_0      // Push constant int 0.
9 istore_1      // Pop int into local variable 1: int i = 0;
10 goto 54      // Go to section of code that tests outer loop.
13 iconst_0     // Push constant int 0.
14 istore_2     // Pop int into local variable 2: int j = 0;
15 goto 46      // Go to section of code that tests middle loop.
18 iconst_0     // Push constant int 0.
19 istore_3     // Pop int into local variable 3: int k = 0;

```



```

20 goto 38          // Go to section of code that tests inner loop.
23 aload_0          // Push object ref from local variable 0.
24 iload_1           // Push int from local variable 1 (i).
25 aload            // Pop index and arrayref, push object ref
                    // at arrayref[index] (gets threeD[i]).
26 iload_2           // Push int from local variable 2 (j).
27 aload            // Pop index and arrayref, push object ref
                    // at arrayref[index] (gets threeD[i][j]).
28 iload_3           // Push int from local variable 3 (k).
                    // Now calculate the int that will be assigned
                    // to threeD[i][j][k]
29 iload_1           // Push int from local variable 1 (i).
30 iload_2           // Push int from local variable 2 (j).
31 iadd              // Pop two ints, add them, push int result (i + j).
32 iload_3           // Push int from local variable 3 (k).
33 iadd              // Pop two ints, add them, push int
                    // result (i + j + k).
34 iastore           // Pop value, index, and arrayref; assign
                    // arrayref[index] = value:
                    // threeD[i][j][k] = i + j + k;
35 iinc 3 1          // Increment by 1 the int in local variable 3: ++k;
38 iload_3           // Push int from local variable 3 (k).
39 iconst_3          // Push constant int 3.
40 if_icmplt 23      // Pop right and left ints, jump if left < right:
                    // for (...; k < 3;...)
43 iinc 2 1          // Increment by 1 the int in local variable 2: ++j;
46 iload_2           // Push int from local variable 2 (j).
47 iconst_4          // Push constant int 4.
48 if_icmplt 18      // Pop right and left ints, jump if left < right:
                    // for (...; j < 4;...)

```

```

51 iinc 1 1      // Increment by 1 the int in local variable 1: ++i;
54 iload_1      // Push int from local variable 1 (i).
55 iconst_5     // Push constant int 5.
56 if_icmplt 13  // Pop right and left ints, jump if left < right:
                // for (...; i < 5;...)

59 return

end

```

The `initAnArray()` method merely allocates and initializes a three-dimensional array. This simulation demonstrates how the Java Virtual Machine handles multidimensional arrays. In response to the `multianewarray` instruction, which in this example requests the allocation of a three-dimensional array, the Java Virtual Machine creates a tree of one-dimensional arrays. The reference returned by the `multianewarray` instruction refers to the base one-dimensional array in the tree. In the `initAnArray()` method, the base array has five components--`threeD[0]` through `threeD[4]`. Each component of the base array is itself a reference to a one-dimensional array of four components, accessed by `threeD[0][0]` through `threeD[4][3]`. The components of these five arrays are also references to arrays, each of which has three components. These components are `ints`, the elements of this multidimensional array, and they are accessed by `threeD[0][0][0]` through `threeD[4][3][2]`.

In response to the `multianewarray` instruction in the `initAnArray()` method, the Java Virtual Machine creates one five-dimensional array of arrays, five four-dimensional arrays of arrays, and twenty three-dimensional arrays of `ints`. The Java Virtual Machine allocates these 26 arrays on the heap, initializes their components such that they form a tree, and returns the reference to the base array.

To assign an `int` value to an element of the three-dimensional array, the Java Virtual Machine uses `aaload` to get a component of the base array. Then the Java Virtual Machine uses `aaload` again on this component--which is itself an array of arrays--to get a component of the branch array. This component is a reference to a leaf array of `ints`. Finally the Java Virtual Machine uses `iastore` to assign an `int` value to the element of the leaf array. The Java Virtual Machine uses multiple one-dimensional array accesses to accomplish operations on multidimensional arrays.

To drive the *Three-Dimensional Array* simulation, use the Step, Reset, Run, and Stop buttons. Each time you press the Step button, the simulator will execute the instruction pointed to by the pc register. If you press the Run button, the simulation will continue with no further coaxing on your part until you press the Stop button. To start the simulation over, press the Reset button. For each step of the simulation, a panel at the bottom of the applet contains an explanation of what the next instruction will do. Happy clicking.

Applet Viewer: ThreeDArray.class

Applet

THREE DIMENSIONAL ARRAY

Local Variables				The Method			
index	hex value	value	offset	bytecode	mnemonic		
0	OBJREF	[Ljava.lang.Object;	23	2a	aload_0		
1	00000000	0	24	1b	iload_1		
2	00000001	1	25	32	aaload		
3	00000002	2	26	1c	iload_2		
			27	32	aaload		
			28	1d	iload_3		
			29	1b	iload_1		
			30	1c	iload_2		
			31	60	iadd		
			32	1d	iload_3		
			33	60	iadd		
			34	4f	iastore		
			35	84	iinc 3 1		

Operand Stack			
offset	hex value	value	
0	OBJREF	[]	
1	00000002	2	
2	00000000	0	
3	00000001	1	

pc: 31

op: 4

pc >

op: >

Step Run

Reset Stop

iadd will pop the top two ints off the stack, add them, and push the result back onto the stack.

On the CD-ROM

The CD-ROM contains the source code examples from this chapter in the `opcodes` directory. The *Three-Dimensional Array* applet is part of a web page on the CD-ROM in file `applets/ThreeDArray.html`. The source code for this applet is found alongside its class files, in the `applets/JVMSimulators` and `applets/JVMSimulators/COM/artima/jvmsim` directories.

The Resources Page

For more information about the material presented in this chapter, visit the resources page: <http://www.artima.com/insidejvm/objects.html>

Chapter Sixteen

Control Flow

This chapter describes the instructions that cause the Java Virtual Machine to conditionally or unconditionally branch to a different location within the same method. This includes the instructions that implement the `if`, `if-else`, `while`, `do-while`, `for`, and `switch` statements of Java source code.

Accompanying this chapter on the CD-ROM is an applet that interactively illustrates the material presented in the chapter. The applet, named *Saying Tomato*, simulates the Java Virtual Machine executing a method that includes bytecodes that perform table jumps (the compiled version of a `switch` statement in Java source code). At the end of this chapter, you will find a description of this applet and the bytecodes it executes.

Conditional Branching

In Java source code, you specify basic control-flow within a method with `if`, `if-else`, `while`, `do-while`, `for`, and `switch` statements. With the exception of `switch`, Java compilers draw from the same set of opcodes when translating all of these source code constructs to bytecodes. For example, the simplest control-flow construct Java offers is the `if` statement. When a Java program is compiled, the `if` statement may be translated to any one of a variety of opcodes, depending upon the nature of the `if`'s expression. Each opcode either pops one or two values from the top of the stack and does a comparison. The opcodes that pop only one value compare that value with zero. The opcodes that pop two values compare one with the other. If the comparison succeeds (success is defined differently by each individual opcode), the Java Virtual Machine branches--or jumps--to the offset given as an operand to the comparison opcode.

For each of the conditional branch opcodes, the Java Virtual Machine follows the same procedure to determine the next instruction to execute. The virtual machine first performs the comparison specified by the opcode. If the comparison fails, the virtual machine continues execution with the instruction that immediately follows the conditional branch instruction. If the comparison succeeds, the virtual machine forms a signed 16-bit offset from two operand bytes that trail the opcode. It adds the offset to the contents of the current pc register (the address of the conditional branch opcode) to get the target address. The target address must point to an opcode of an instruction in the same method. Execution continues as the target address.

One family of `if` opcodes, shown in Table 16-1, performs integer comparisons against zero. When the Java Virtual Machine encounters one of these opcodes, it pops one `int` off the stack and compares it with zero.

Table 16-1. Conditional branch: integer comparison with zero

Opcode	Operand(s)	Description
<code>ifeq</code>	<code>branchbyte1,</code> <code>branchbyte2</code>	pop <code>int</code> value, if <code>value == 0</code> , branch to offset
<code>ifne</code>	<code>branchbyte1,</code> <code>branchbyte2</code>	pop <code>int</code> value, if <code>value != 0</code> , branch to offset
<code>iflt</code>	<code>branchbyte1,</code> <code>branchbyte2</code>	pop <code>int</code> value, if <code>value < 0</code> , branch to offset
<code>ifle</code>	<code>branchbyte1,</code> <code>branchbyte2</code>	pop <code>int</code> value, if <code>value <= 0</code> , branch to offset
<code>ifgt</code>	<code>branchbyte1,</code> <code>branchbyte2</code>	pop <code>int</code> value, if <code>value > 0</code> , branch to offset
<code>ifge</code>	<code>branchbyte1,</code> <code>branchbyte2</code>	pop <code>int</code> value, if <code>value >= 0</code> , branch to offset

Another family of `if` opcodes, shown in Table 16-2, pops two integers off the top of the stack and compares them against one another. The virtual machine branches if the comparison succeeds. Just before these opcodes are executed, `value2` is on the top of the stack; `value1` is just beneath `value2`.

Table 16-2. Conditional branch: comparison of two integers

Opcode	Operand(s)	Description
<code>if_icmpeq</code>	<code>branchbyte1,</code> <code>branchbyte2</code>	pop <code>int</code> <code>value2</code> and <code>value1</code> , if <code>value1 == value2</code> , branch to offset
<code>if_icmpne</code>	<code>branchbyte1,</code> <code>branchbyte2</code>	pop <code>int</code> <code>value2</code> and <code>value1</code> , if <code>value1 != value2</code> , branch to offset
<code>if_icmplt</code>	<code>branchbyte1,</code> <code>branchbyte2</code>	pop <code>int</code> <code>value2</code> and <code>value1</code> , if <code>value1 < value2</code> , branch to offset
<code>if_icmple</code>	<code>branchbyte1,</code> <code>branchbyte2</code>	pop <code>int</code> <code>value2</code> and <code>value1</code> , if <code>value1 <= value2</code> , branch to offset
<code>if_icmpgt</code>	<code>branchbyte1,</code> <code>branchbyte2</code>	pop <code>int</code> <code>value2</code> and <code>value1</code> , if <code>value1 > value2</code> , branch to offset
<code>if_icmpge</code>	<code>branchbyte1,</code> <code>branchbyte2</code>	pop <code>int</code> <code>value2</code> and <code>value1</code> , if <code>value1 >= value2</code> , branch to offset

The opcodes shown above in Table 16-2 operate on `ints`. These opcodes also are used for comparisons of types `short`, `byte`, and `char`; the Java Virtual Machine always manipulates types smaller than `int` by first converting them to `ints` and then manipulating the `ints`.

A third family of opcodes, shown in Table 16-3, takes care of comparisons of the other primitive types: `long`, `float`, and `double`. These opcodes don't cause a branch by themselves. Instead, they push the `int` value that represents the result of the comparison--0 for equal to, 1 for greater than, and -1 for less than--and then use one of the `int` compare opcodes introduced above to force the actual branch.

Table 16-3. Comparison of `longs`, `floats`, and `doubles`

Opcode	Operand(s)	Description
<code>lcmp</code>	(none)	pop long value2 and value1, compare, push int result
<code>fcmpg</code>	(none)	pop float value2 and value1, compare, push int result
<code>fcmpl</code>	(none)	pop float value2 and value1, compare, push int result
<code>dcmpg</code>	(none)	pop double value2 and value1, compare, push int result
<code>dcmpl</code>	(none)	pop double value2 and value1, compare, push int result

The two opcodes for float comparisons (*fcmpg* and *fcmpl*) differ only in how they handle NaN ("not a number"). In the Java Virtual Machine, comparisons of floating-point numbers always fail if one of the values being compared is NaN. If neither value being compared is NaN, both *fcmpg* and *fcmpl* instructions push a 0 if the values are equal, a 1 if value1 is greater than value2, and a -1 if value1 is less than value2. But if one or both of the values is NaN, the *fcmpg* instruction pushes a 1, whereas the *fcmpl* instruction pushes a -1. Because both of these operands are available, any comparison between two float values can push the same result onto the stack independent of whether the comparison failed because of a NaN. This is also true for the two opcodes that compare double values: *dcmpg* and *dcmpl*.

A fourth family of if opcodes, shown in Table 16-4, pops one object reference off the top of the stack and compares it with null. If the comparison succeeds, the Java Virtual Machine branches.

Table 16-4. **Conditional branch: object reference comparison with null**

Opcode	Operand(s)	Description
<code>ifnull</code>	branchbyte1, branchbyte2	pop reference value, if value == null, branches to offset
<code>ifnonnull</code>	branchbyte1, branchbyte2	pop reference value, if value != null, branches to offset

The last family of if opcodes, which is shown in Table 16-5, pops two object references off the stack and compares them with each other. In this case, there are only two comparisons that make sense: "equals" and "not equals." If the references are equal, then they refer to the exact same object on the heap. If not, they refer to two different objects. As with all the other if opcodes, if the comparison succeeds, the Java Virtual Machine branches.

Table 16-5. **Conditional branch: comparison of two object references**

Opcode	Operand(s)	Description
<code>if_acmpeq</code>	branchbyte1, branchbyte2	pop reference value2 and value1, if value1 == value2, branch to offset
<code>if_acmpne</code>	branchbyte1,	pop reference value2 and value1, if

branchbyte2 value1 != value2, branch to offset

Unconditional Branching

That's all of the opcodes that cause the Java Virtual Machine to branch conditionally. One other family of opcodes, however, causes the virtual machine to branch unconditionally. Not surprisingly, these opcodes, shown in Table 16-6, are called "*goto*." To execute a *goto* instruction, the virtual machine forms a signed 16-bit offset from two operand bytes that follow the *goto* opcode. (To execute a *goto_w* instruction, the virtual machine forms a signed 32-bit offset from four operand bytes that follow the *goto_w* opcode.) The virtual machine adds this offset to the current contents of the pc register. The resulting address must contain an opcode of an instruction in the current method. The virtual machine continues execution at this instruction.

Table 16-6. **Unconditional branch**

Opcode	Operand(s)	Description
<i>goto</i>	branchbyte1, branchbyte2	branch to offset
<i>goto_w</i>	branchbyte1, branchbyte2, branchbyte3, branchbyte4	branch to offset

The opcodes shown in Tables 16-1 through 16-6 are sufficient to express in bytecodes the any control flow specified in Java source code with an *if*, *if-else*, *while*, *do-while*, or *for* statement. The above opcodes also could be used to express a *switch* statement, but the Java Virtual Machine's instruction set includes two opcodes specially designed for the *switch* statement: *tableswitch* and *lookupswitch*.

Conditional Branching with Tables

The *tableswitch* and *lookupswitch* instructions, shown in Table 16-7, both include one default branch offset and a variable-length set of case value/branch offset pairs. Both instructions pop the key (the value of the expression in the parentheses immediately following the *switch* keyword) from the stack. The key is compared with all the case values. If a match is found, the branch offset associated with the case value is taken. If no match is found, the default branch offset is taken.

The difference between *tableswitch* and *lookupswitch* is in how they indicate the case values. The *lookupswitch* instruction is more general-purpose than *tableswitch*, but *tableswitch* is usually more efficient. Both instructions are followed by zero to three bytes of padding--enough so that the byte immediately following the padding starts at an address that is a multiple of four bytes from the beginning of the method. (These two instructions, by the way, are the only ones in the entire Java Virtual Machine instruction set that involve alignment on a greater than one-byte boundary.) For both instructions, the next four bytes after the padding is the default branch offset.

After the zero- to three-byte padding and the four-byte default branch offset, the *lookupswitch* opcode is followed by a four-byte value, *npairs*, which indicates the number of case value/branch offset pairs that will follow. The case value is an *int*, which highlights the fact that *switch* statements in Java require a key expression that is an *int*, *short*, *char*, or *byte*. If you attempt to use a *long*, *float*, or *double* as a *switch* key, your program won't compile. The branch offset associated with each case value

is another four-byte offset. The value/branch offset pairs must appear in increasing numerical order of case value.

In the `tableswitch` instruction, the zero- to three-byte padding and the four-byte default branch offset are followed by low and high int values. The low and high values indicate the endpoints of a range of case values included in this `tableswitch` instruction. Following the low and high values are $high - low + 1$ branch offsets--one branch offset for high, one for low, and one for each integer case value in between high and low. The branch offset for low immediately follows the high value.

Thus, when the Java Virtual Machine encounters a `lookupswitch` instruction, it must check the key against each case value until it finds a match, encounters a case value greater than the key (the values/branch offset pairs are sorted in increasing numerical order of case value), or runs out of case values. If it doesn't find a match, it uses the default branch offset. On the other hand, when the Java Virtual Machine encounters a `tableswitch` instruction, it can simply check to see if the key is within the range defined by low and high. If not, it takes the default branch offset. If so, it just subtracts low from key to get an offset into the list of branch offsets. This way, it can determine the appropriate branch offset without having to check each case value.

Table 16-7. **Table jumping**

Opcode	Operand(s)	Description
<code>lookupswitch</code>	<0-3 byte paddefaultbyte1, defaultbyte2, defaultbyte3, defaultbyte4, npairs1, npairs2, pairs3, npairs4, case value/branch offset pairs...	pop key, match key with case values, if match found jump to associated branch offset, else jump to default branch offset
<code>tableswitch</code>	<0-3 byte paddefaultbyte1, defaultbyte2, defaultbyte3, defaultbyte4, lowbyte1, lowbyte2, lowbyte3, lowbyte4, highbyte1, highbyte2, highbyte3, highbyte4, branch offsets...	pop key, if not in low/high range jump to default branch offset, else get the (key - low) branch offset and jump

Other than the opcodes described above, the only Java Virtual Machine instructions that affect control flow are those that deal with throwing and catching exceptions, finally clauses, and invoking and returning from methods. These opcodes are discussed in later chapters.

Saying Tomato: A Simulation

The ***Saying Tomato*** applet, shown in Figure 16-1, demonstrates a Java Virtual Machine executing a sequence of bytecodes. The applet is part of a web page on the CD-ROM in file `applets/SayingTomato.html`. The bytecode sequence in the simulation was generated by the `javac` compiler for the `argue()` method of the class shown below:

```
begin

// On CD-ROM in file opcodes/ex1/Struggle.java
class Struggle {
```



```

public final static int TOMAYTO = 0;
public final static int TOMAHTO = 1;

static void argue() {

    int say = TOMAYTO;

    for (;;) {

        switch (say) {

            case TOMAYTO:

                say = TOMAHTO;
                break;

            case TOMAHTO:

                say = TOMAYTO;
                break;
        }
    }
}

end

```

The bytecodes generated by javac for the argue() method are shown below:

```

begin
0 iconst_0      // Push constant 0 (TOMAYTO)
1 istore_0      // Pop into local var 0: int say = TOMAYTO;

```

```

2 iload_0      // Push key for switch from local var 0
               // Perform switch statement: switch (say) {...
               // Low case value is 0, high case value is 1
               // Default branch offset will goto 2

3 tableswitch 0 to 1: default=2

    0: 24  // case 0 (TOMAYTO): goto 24
    1: 29  // case 1 (TOMAHTO): goto 29

               // Note that the next instruction starts at address
               // 24, which means the tableswitch took up 21 bytes

24 iconst_1    // Push constant 1 (TOMAHTO)
25 istore_0    // Pop into local var 0: say = TOMAHTO
26 goto 2      // Branch unconditionally to 2, top of while loop
29 iconst_0    // Push constant 1 (TOMAYTO)
30 istore_0    // Pop into local var 0: say = TOMAYTO
31 goto 2      // Branch unconditionally to 2, top of while loop

end

```


The `argue()` method merely switches the value of `say` back and forth between `TOMAYTO` and `TOMAHTO`. Because the values of `TOMAYTO` and `TOMAHTO` were consecutive (`TOMAYTO` was a 0 and `TOMAHTO` was a 1), the `javac` compiler used a `tableswitch`. The `tableswitch` is a more efficient instruction than a `lookupswitch`, and the equivalent `lookupswitch` instruction would occupy 28 bytes--4 bytes more than `tableswitch`.

It turns out that even if `TOMAYTO` were a 0 and `TOMAHTO` were a 2, the `javac` compiler still would have used a `tableswitch`, because even with the extra default branch offset in there for a 1, the `tableswitch` instruction would occupy only 28 bytes--the same number of bytes as the equivalent `lookupswitch`. Both instructions occupy the same number of bytes, but `tableswitch` is more efficient, so it is used. As soon as you make `TOMAHTO` a 3, however, `javac` starts using a `lookupswitch`. This is because a `tableswitch` would now need two default branch offsets in its list (for 1 and 2), which would push its size up to 32 bytes. Thus, a `lookupswitch` now would require fewer bytes than a `tableswitch`--so `javac` chooses the `lookupswitch`.

The branch offsets for the case values cause the Java Virtual Machine to hop down to code that will change the value of the `say` local variable. The value of `say` will alternate between `TOMAYTO` and `TOMAHTO` indefinitely, until the user aborts the program, thereby calling the whole thing off.

To drive the *Saying Tomato* simulation, use the Step, Reset, Run, and Stop buttons. Each time you press the Step button, the simulator will execute the instruction pointed to by the pc register. If you press the

Run button, the simulation will continue with no further coaxing on your part until you press the Stop button. To start the simulation over, press the Reset button. For each step of the simulation, a panel at the bottom of the applet contains an explanation of what the next instruction will do. Happy clicking.



On the CD-ROM

The CD-ROM contains the source code examples from this chapter in the `opcodes` directory. The *Saying Tomato* applet is contained in a web page on the CD-ROM in file `applets/SayingTomato.html`. The source code for this applet is found alongside its class files, in the `applets/JVMSimulators` and `applets/JVMSimulators/COM/artima/jvmsim` directories.

The Resources Page

For more information about the material presented in this chapter, visit the resources page: <http://www.artima.com/insidejvm/flow.html>.

Chapter Seventeen

Exceptions

This chapter shows how exceptions are implemented in bytecodes. It describes the instruction for throwing an exception explicitly, explains exception tables, and shows how catch clauses work.

Accompanying this chapter on the CD-ROM is an applet that interactively illustrates the material presented in the chapter. The applet, named *Play Ball!*, simulates the Java Virtual Machine executing a method that throws and catches exceptions. At the end of this chapter, you will find a description of this applet and the bytecodes it executes.

Throwing and Catching Exceptions

Exceptions allow you to smoothly handle unexpected conditions that occur as your programs run. To demonstrate the way the Java Virtual Machine handles exceptions, consider a class named `NitPickyMath` that provides methods that perform addition, subtraction, multiplication, division, and remainder on integers. `NitPickyMath` performs these mathematical operations the same as the normal operations offered by Java's `+`, `-`, `*`, `/`, and `%` operators, except the methods in `NitPickyMath` throw checked exceptions on overflow, underflow, and divide-by-zero conditions. The Java Virtual Machine will throw an `ArithmeticException` on an integer divide-by-zero, but will not throw any exceptions on overflow and underflow. The exceptions thrown by the methods of `NitPickyMath` are defined as follows:

```
begin
```

```
// On CD-ROM in file except/ex1/OverflowException.java
class OverflowException extends Exception {

}
```

```
// On CD-ROM in file except/ex1/UnderflowException.java
class UnderflowException extends Exception {

}
```

```
// On CD-ROM in file except/ex1/DivideByZeroException.java
class DivideByZeroException extends Exception {
```

```
}
```

```
end
```

A simple method that catches and throws exceptions is the `remainder()` method of class `NitPickyMath`:

```
begin
```

```
// On CD-ROM in file except/ex1/NitpickyMath.java  
class NitpickyMath {
```

```
    static int add(int a, int b)  
        throws OverflowException, UnderflowException {  
  
        long longA = (long) a;  
        long longB = (long) b;  
        long result = a + b;  
        if (result > Integer.MAX_VALUE) {  
            throw new OverflowException();  
        }  
        if (result < Integer.MIN_VALUE) {  
            throw new UnderflowException();  
        }  
        return (int) result;  
    }  
}
```

```
    static int subtract(int minuend, int subtrahend)  
        throws OverflowException, UnderflowException {  
  
        long longMinuend = (long) minuend;  
        long longSubtrahend = (long) subtrahend;  
        long result = longMinuend - longSubtrahend;  
        if (result > Integer.MAX_VALUE) {
```

```

        throw new OverflowException();
    }
    if (result < Integer.MIN_VALUE) {
        throw new UnderflowException();
    }
    return (int) result;
}

```

```

static int multiply(int a, int b)
    throws OverflowException, UnderflowException {

    long longA = (long) a;
    long longB = (long) b;
    long result = a * b;
    if (result > Integer.MAX_VALUE) {
        throw new OverflowException();
    }
    if (result < Integer.MIN_VALUE) {
        throw new UnderflowException();
    }
    return (int) result;
}

```

```

static int divide(int dividend, int divisor)
    throws OverflowException, DivideByZeroException {

    // Overflow can occur in division when dividing the
    // negative integer of the largest possible magnitude
    // (Integer.MIN_VALUE) by -1, because this would just flip
    // the sign, but there is no way to represent that number

```

```

// in an int.
if ((dividend == Integer.MIN_VALUE) && (divisor == -1)) {
    throw new OverflowException();
}
try {
    return dividend / divisor;
}
catch (ArithmeticException e) {
    throw new DivideByZeroException();
}
}

static int remainder(int dividend, int divisor)
    throws OverflowException, DivideByZeroException {

    // Overflow can occur in division when dividing the
    // negative integer of the largest possible magnitude
    // (Integer.MIN_VALUE) by -1, because this would just flip
    // the sign, but there is no way to represent that number
    // in an int.
    if ((dividend == Integer.MIN_VALUE) && (divisor == -1)) {
        throw new OverflowException();
    }
    try {
        return dividend % divisor;
    }
    catch (ArithmeticException e) {
        throw new DivideByZeroException();
    }
}
}

```

```
}
```

```
end
```

The `remainder()` method simply performs the remainder operation (%) upon the two *ints* passed as arguments. The remainder operation throws an `ArithmeticException` if the divisor of the remainder operation is a zero. This method catches this `ArithmeticException` and throws a `DivideByZeroException`.

The difference between these two exceptions is that `DivideByZeroException` is *checked* and `ArithmeticException` is *unchecked*. Because the `ArithmeticException` is unchecked, a method need not declare this exception in a `throws` clause even though it might throw it. Any exceptions that are subclasses of either `Error` or `RuntimeException` are unchecked. (`ArithmeticException` is a subclass of `RuntimeException`.) By catching `ArithmeticException` and then throwing `DivideByZeroException`, the `remainder()` method forces its clients to deal with the possibility of a divide-by-zero exception, either by catching it or declaring `DivideByZeroException` in their own `throws` clauses.

`javac` generates the following bytecode sequence for the `remainder` method:

```
begin
```

```
// The main bytecode sequence for remainder():
```

```
0 iload_0 // Push local variable 0 (arg passed as dividend)
          // Push the minimum integer value
1 ldc #1 <Integer -2147483648
          // If the dividend isn't equal to the minimum integer,
          // jump to the remainder calculation
3 if_icmpne 19

6 iload_1 // Push local variable 1 (arg passed as divisor)
          // Push -1
7 iconst_m1
          // If the divisor isn't equal to -1, jump
          // to the remainder calculation
8 if_icmpne 19
          // This is an overflow case, so throw an exception
          // Create a new OverflowException, push reference to
```



```

        // it onto the stack
11 new #4 <Class OverflowException
14 dup      // Make a copy of the reference
           // Pop one copy of the reference and invoke the <init
           // method of new OverflowException object
15 invokespecial #10 <Method OverflowException()
           // Pop the other reference to the OverflowException
           // and throw it
18 athrow

// Calculate the remainder

19 iload_0  // Push local variable 0 (arg passed as dividend)
20 iload_1  // Push local variable 1 (arg passed as divisor)
21 irem     // Pop divisor; pop dividend; calculate, push remainder
22 ireturn  // Return int on top of stack (the remainder)

// The bytecode sequence for the catch (ArithmeticException) clause:

23 pop      // Pop the reference to the ArithmeticException
           // because it isn't used by this catch clause.
24 new #2 <Class DivideByZeroException
           // Create and push reference to new object of
           // class DivideByZeroException.
27 dup      // Duplicate the reference to the new object on the top
           // of the stack because it must be both initialized
           // and thrown. The initialization will consume
           // the copy of the reference created by the dup.
28 invokespecial #9 <Method DivideByZeroException()

```

```

        // Call the no-arg <init method for the
        // DivideByZeroException to initialize it. This
        // instruction will pop the top reference to the object.
31 athrow // Pop the reference to a Throwable object, in this
        // case the DivideByZeroException,
        // and throw the exception.

end

```

The bytecode sequence of the `remainder()` method has two separate parts. The first part is the normal path of execution for the method. This part goes from pc offset zero through 22. The second part is the catch clause, which goes from pc offset 23 through 31.

The Exception Table

The *irem* instruction in the main bytecode sequence may throw an `ArithmeticException`. If this occurs, the Java Virtual Machine knows to jump to the bytecode sequence that implements the catch clause by looking up and finding the exception in a table. Each method that catches exceptions is associated with an *exception table* that is delivered in the class file along with the bytecode sequence of the method. The exception table has one entry for each exception that is caught by each try block. Each entry has four pieces of information:

- the start point
- the end point
- the pc offset within the bytecode sequence to jump to
- a constant pool index of the exception class that is being caught.

The exception table for the `remainder()` method of class `NitPickyMath` is shown below:

```

begin

Exception table:
  from    to  target type
    19     23    23   <Class java.lang.ArithmeticException

end

```

The above exception table indicates that from pc offset 19 through 22, inclusive, `ArithmeticException` is caught. The try block's endpoint value, listed in the table under the label "to", is always one more than the last pc offset for which the exception is caught. In this case the endpoint value is listed as 23, but the last pc offset for which the exception is caught is 22. This range, 19 to 22 inclusive, corresponds to the bytecode sequence that implements the code inside the try block of `remainder()`. The target listed in the table is the pc offset to jump to if an `ArithmeticException` is thrown between the pc offsets 19 and

22, inclusive.

If an exception is thrown during the execution of a method, the Java Virtual Machine searches through the exception table for a matching entry. An exception table entry matches if the current program counter is within the range specified by the entry, and if the exception class thrown is the exception class specified by the entry (or is a subclass of the specified exception class). The Java Virtual Machine searches through the exception table in the order in which the entries appear in the table. When the first match is found, the virtual machine sets the program counter to the new pc offset location and continues execution there. If no match is found, the virtual machine pops the current stack frame and rethrows the same exception. When the Java Virtual Machine pops the current stack frame, it effectively aborts execution of the current method and returns to the method that called this method. But instead of continuing execution normally in the previous method, it throws the same exception in that method, which causes it to go through the same process of searching through the exception table of that method.

A Java programmer can throw an exception with a throw statement such as the one in the catch (ArithmeticException) clause of remainder(), where a DivideByZeroException is created and thrown. The bytecode that does the throwing is shown in Table 17-1:

Table 17-1. **Throwing Exceptions**

Opcode	Operand (s)	Description
athrow	(none)	pops Throwable object reference, throws the exception

The *athrow* instruction pops the top word from the stack and expects it to be a reference to an object that is a subclass of Throwable (or Throwable itself). The exception thrown is of the type defined by the popped object reference.

Play Ball!: A Simulation

The *Play Ball!* Applet, shown in Figure 17-1, demonstrates a Java Virtual Machine executing a sequence of bytecodes. This applet is embedded in a web page on the CD-ROM in file `applets/PlayBall.html`. The bytecode sequence in the simulation was generated by `javac` for the `playBall` method of the class shown below:

```
begin

// On CD-ROM in file except/ex2/Ball.java
class Ball extends Exception {

}

// On CD-ROM in file except/ex2/Pitcher.java
class Pitcher {

    private static Ball ball = new Ball();
```

```

static void playBall() {
    int i = 0;
    for (;;) {
        try {
            if (i % 4 == 3) {
                throw ball;
            }
            ++i;
        }
        catch (Ball b) {
            i = 0;
        }
    }
}
}

```

end

The bytecodes generated by javac for the playBall method are shown below:

begin

// The main bytecode sequence for playBall():

```

0 iconst_0      // Push constant 0
1 istore_0      // Pop into local var 0: int i = 0;
                // The try block starts here (see the
                // exception table, below).
2 iload_0       // Push local var 0
3 iconst_4      // Push constant 4
4 irem          // Calc remainder of top two operands
5 iconst_3      // Push constant 3

```

```

6 if_icmpne 13 // Jump if remainder not equal to 3:
    // if (i % 4 == 3) {
    // Push the static field at constant pool
    // location #6, which is the Ball exception eager
    // to be thrown

9 getstatic #6 <Field Ball ball

12 athrow      // Heave it home: throw ball;

13 iinc 0 1    // Increment the int at local var 0 by 1: ++i;
    // The try block ends here (see the
    // exception table, below).

16 goto 2      // jump always back to 2: for (;;) {}

// The bytecode sequence for the catch (Ball) clause:

19 pop         // Pop the exception reference because it is unused
20 iconst_0    // Push constant 0
21 istore_0    // Pop into local var 0: i = 0;
22 goto 2      // Jump always back to 2: for (;;) {}

```

Exception table:

from	to	target	type
2	16	19	<Class Ball


end

The `playball()` method loops forever. Every fourth pass through the loop, `playball()` throws a `Ball` and catches it, just because it's fun. Since the try block and the catch clause are both within the endless while loop, the fun never stops. The local variable `i` starts at 0 and increments each pass through the loop. When the `if` statement is true, which happens every time `i` is equal to 3, the `Ball` exception is thrown.

The Java Virtual Machine checks the exception table and discovers that there is indeed an applicable entry. The entry's valid range is from 2 to 15, inclusive, and the exception is thrown at pc offset 12. The exception caught by the entry is of class `Ball`, and the exception thrown is of class `Ball`. Given this perfect match, the virtual machine pushes the thrown exception object onto the stack, and continues

execution at pc offset 19. The catch clause, which starts at offset 19, merely resets *int i* to 0, and the loop starts over.

To drive the *Play Ball!* simulation, use the Step, Reset, Run, and Stop buttons. Each time you press the Step button, the simulator will execute the instruction pointed to by the pc register. If you press the Run button, the simulation will continue with no further coaxing on your part until you press the Stop button. To start the simulation over, press the Reset button. For each step of the simulation, a panel at the bottom of the applet contains an explanation of what the next instruction will do. Happy clicking.



On the CD-ROM

The CD-ROM contains the source code examples from this chapter in the `except` directory. The *Play Ball!* applet is contained in a web page on the CD-ROM in file `applets/PlayBall.html`. The source code for this applet is found alongside its class files, in the `applets/JVMSimulators` and `applets/JVMSimulators/COM/artima/jvmsim` directories.

The Resources Page

For more information about the material presented in this chapter, visit the resources page: <http://www.artima.com/insidejvm/except.html>.

Chapter Eighteen

Finally Clauses

This chapter describes the way finally clauses are implemented in bytecodes. It describes the relevant instructions and shows how they are used. The chapter also describes some surprising behaviors exhibited by finally clauses in Java source code and explains this behavior at the bytecode level.

Accompanying this chapter on the CD-ROM is an applet that interactively illustrates the material presented in the chapter. The applet, named *Hop Around*, simulates the Java Virtual Machine executing a method that includes finally clauses. At the end of this chapter, you will find a description of this applet and the bytecodes it executes.

Miniature Subroutines

In bytecodes, finally clauses act as "miniature subroutines" within a method. The subroutine for a finally clause is "invoked" at each exit point inside a try block and its associated catch. After the finally clause completes--as long as it completes by executing past the last statement in the finally clause, not by throwing an exception or executing a return, continue, or break--the miniature subroutine itself "returns." Execution continues just past the point where the miniature subroutine was called in the first place.

The opcode that causes a Java Virtual Machine to jump to a miniature subroutine is the *jsr* instruction. The *jsr* opcode takes a two-byte operand, the 16-bit signed offset from the location of the *jsr* instruction where the miniature subroutine begins. A second variant of the *jsr* instruction is *jsr_w*, which performs the same function as *jsr* but takes a wide (four-byte) operand. When the Java Virtual Machine encounters a *jsr* or *jsr_w* instruction, it pushes a return address onto the stack, then continues execution at the start of the miniature subroutine. The return address is the address (an offset or native pointer) of the bytecode immediately following the *jsr* or *jsr_w* opcode and its operands. The type of the address is `returnAddress`.

After a miniature subroutine completes, it invokes the *ret* instruction, which returns from the subroutine. The *ret* instruction takes one operand, an index into the local variables where the return address is stored. The opcodes that deal with finally clauses are summarized in Table 18-1:

Table 18-1. **Finally clauses**

Opcode	Operand(s)	Description
--------	------------	-------------

<code>jsr</code>	<code>branchbyte1, branchbyte2</code>	pushes the return address, branches to offset
<code>jsr_w</code>	<code>branchbyte1, branchbyte2, branchbyte3, branchbyte4</code>	pushes the return address, branches to wide offset
<code>ret</code>	<code>index</code>	returns to the address stored in local variable index
<code>wide ret</code>	<code>indexbyte1, indexbyte2</code>	returns to the address stored in local variable index

Don't confuse a miniature subroutine with a Java method. Java methods use a different set of instructions. Instructions such as *invokevirtual* or *invokespecial* cause a Java method to be invoked, and instructions such as *return*, *areturn*, or *ireturn* cause a Java method to return. The *jsr* instruction does not cause a Java method to be invoked. It merely causes a jump to a different opcode within the same method. Likewise, the *ret* instruction doesn't return from a method. It just causes the virtual machine to jump back to the opcode in the same method that immediately follows the calling *jsr* opcode and its operands. In this book, the bytecodes that implement a `finally` clause are called a "miniature subroutine" because they act like a small subroutine within the bytecode stream of a single method.

Asymmetrical Invocation and Return

You might think that the *ret* instruction should pop the return address off the stack, because that is where it was pushed by the *jsr* instruction. But it doesn't. At the start of each subroutine, the return address is popped off the top of the stack and stored in a local variable--the same local variable from which the *ret* instruction later gets it. This asymmetrical manner of working with the return address is necessary because `finally` clauses (and therefore, miniature subroutines) themselves can throw exceptions or include `return`, `break`, or `continue` statements. Because of this possibility, the extra return address that was pushed onto the stack by the *jsr* instruction must be removed from the stack right away, so it won't still be there if the `finally` clause exits with a `break`, `continue`, `return`, or thrown exception.

As an illustration, consider the following code, which includes a `finally` clause that exits with a `break` statement. The result of this code is that, irrespective of the parameter `bVal` passed to method `surpriseTheProgrammer()`, the method returns `false`:

begin

```
// On CD-ROM in file opcodes/ex3/Surprise.java
class Surprise {
```

```
    static boolean surpriseTheProgrammer(boolean bVal) {
        while (bVal) {
            try {
                return true;
            }
        }
    }
}
```



```

        finally {
            break;
        }
    }
    return false;
}
}

```

end

The example above shows why the return address must be stored into a local variable at the beginning of the `finally` clause. Because the `finally` clause exits with a `break`, it never executes the ***ret*** instruction. As a result, the Java Virtual Machine never goes back to finish up the "return true" statement. Instead, it just goes ahead with the `break` and drops down past the closing curly brace of the `while` statement. The next statement is "return false," which is precisely what the Java Virtual Machine does.

The behavior shown by this `finally` clause, which exits with a `break`, is also shown by `finally` clauses that exit with a `return` or `continue`, or by throwing an exception. In any of these cases, the ***ret*** instruction at the end of the `finally` clause is never executed. Because the ***ret*** instruction is not guaranteed to be executed, it can't be relied on to remove the return address from the stack. Therefore, the return address is stored into a local variable at the beginning of the `finally` clause's miniature subroutine.

For a complete example of a `finally` clause, consider the following method, which contains a `try` block with two exit points. In this example, both exit points are `return` statements:

```

begin

// On CD-ROM in file opcodes/ex1/Nostalgia.java
class Nostalgia {

    static int giveMeThatOldFashionedBoolean(boolean bVal) {
        try {
            if (bVal) {
                return 1;
            }
            return 0;
        }
    }
}

```

```

        finally {
            System.out.println("Got old fashioned.");
        }
    }
}

end

```

The `giveMeThatOldFashionedBoolean()` method compiles to the following bytecodes:

```

begin
// The bytecode sequence for the try block:
0 iload_0    // Push local variable 0 (bval parameter)
1 ifeq 11    // Pop int, if equal to 0, jump to 11 (just past
              // the if statement): if (bval) {}
4 iconst_1   // Push int 1
5 istore_1   // Pop an int (the 1), store into local variable 1
6 jsr 24     // Jump to the mini-subroutine for the finally clause
9 iload_1    // Push local variable 1 (the 1)
10 ireturn   // Return int on top of the stack (the 1): return 1;
11 iconst_0   // Push int 0
12 istore_1   // Pop an int (the 0), store into local variable 1
13 jsr 24     // Jump to the mini-subroutine for the finally clause
16 iload_1    // Push local variable 1 (the 0)
17 ireturn   // Return int on top of the stack (the 0): return 0;

// The bytecode sequence for a catch clause that catches any kind
// of exception thrown from within the try block.
18 astore_2   // Pop the reference to the thrown exception, store
              // into local variable 2
19 jsr 24     // Jump to the mini-subroutine for the finally clause
22 aload_2    // Push the reference (to the thrown exception) from
              // local variable 2

```

```

23 athrow    // Rethrow the same exception

// The miniature subroutine that implements the finally block.
24 astore_3  // Pop the return address, store it in local variable 3
25 getstatic #7 <Field java.io.PrintStream out
           // Get a reference to java.lang.System.out
28 ldc #1 <String "Got old fashioned."
           // Push reference to "Got old fashioned." String from
           // the constant pool
30 invokevirtual #8 <Method void println(java.lang.String)
           // Invoke System.out.println()
33 ret 3     // Return to return address stored in local variable 3

end

```

The bytecodes for the try block include two *jsr* instructions. Another *jsr* instruction is contained in the catch clause. The catch clause is added by the compiler because if an exception is thrown during the execution of the try block, the finally block must still be executed. Therefore, the catch clause merely invokes the miniature subroutine that represents the finally clause, then throws the same exception again. The exception table for the `giveMeThatOldFashionedBoolean()` method, shown below, indicates that any exception thrown between and including addresses 0 and 17 (all the bytecodes that implement the try block) are handled by the catch clause that starts at address 18.

begin

Exception table:

from	to	target	type
0	18	18	any

end

The bytecodes of the finally clause begin by popping the return address off the stack and storing it into local variable three. At the end of the finally clause, the *ret* instruction takes its return address from the proper place, local variable three.

Hop Around: A Simulation

The *Hop Around* applet, shown in Figure 18-1, demonstrates a Java Virtual Machine executing a sequence of bytecodes. The applet is embedded in a web page on the CD-ROM in file `applets/HopAround.html`. The bytecode sequence in the simulation was generated by the `javac`

compiler for the hopAround() method of the class shown below:

begin

```
// On CD-ROM in file opcodes/ex1/Clown.java
class Clown {
```

```
    static int hopAround() {
        int i = 0;
        while (true) {
            try {
                try {
                    i = 1;
                }
                finally { // The first finally clause
                    i = 2;
                }
                i = 3;
                // This return never completes, because of
                // the continue in the second finally clause
                return i;
            }
            finally { // The second finally clause
                if (i == 3) {
                    // This continue overrides the return statement
                    continue;
                }
            }
        }
    }
}
```

end

The bytecodes generated by javac for the hopAround() method are shown below:

begin

```
0  iconst_0      // Push constant 0
1  istore_0      // Pop into local var 0: int i = 0;

// Both try blocks start here (see exception table, below):
2  iconst_1      // Push constant 1
3  istore_0      // Pop into local var 0: i = 1;
4  jsr 18        // Jump to mini-subroutine at offset 18 (the
                // first finally clause)
7  goto 24       // Jump to offset 24 (to just below first
                // finally clause)

// Catch clause for the first finally clause:
10 astore 4      // Pop the reference to thrown exception, store
                // in local variable 4
12 jsr 18        // Jump to mini-subroutine at offset 18 (the
                // first finally clause)
15 aload 4       // Push the reference (to thrown exception)
                // from local variable 4
17 athrow        // Rethrow the same exception

// The first finally clause:
18 astore 5      // Store the return address in local variable 5
20 iconst_2      // Push constant 2
21 istore_0      // Pop into local var 0: i = 2;
22 ret 5         // Jump to return address stored in local variable 5

// Bytecodes for the code just after the first finally clause:
24 iconst_3      // Push constant 3
```

```

25 istore_0      // Pop into local var 0: int i = 3;

// Bytecodes for the return statment:
26 iload_0       // Push the int from local
                  // variable 0 (i, which is 3)
27 istore_1      // Pop and store the int into local
                  // variable 1 (the return value, i)
28 jsr 39        // Jump to mini-subroutine at offset 39 (the
                  // second finally clause)
31 iload_1       // Push the int from local variable 1 (the
                  // return value)
32 ireturn       // Return the int on the top of the stack

// Catch clause for the second finally clause:
33 astore_2      // Pop the reference to thrown exception, store
                  // in local variable 2
34 jsr 39        // Jump to mini-subroutine at offset 39 (the
                  // second finally clause)
37 aload_2       // Push the reference (to thrown exception)
                  // from local variable 2
38 athrow        // Rethrow the same exception

// The second finally clause:
39 astore_3      // Store the return address in local variable 3
40 iload_0       // Push the int from local variable 0 (i)
41 iconst_3      // Push constant 3
42 if_icmpeq 47   // If the top two ints on the stack are equal, jump
                  // to offset 47: if (i == 3) {
45 ret 3         // Jump to return address stored in local variable 3
47 goto 2        // Jump to offset 2 (the top of the while

```

```
// block): continue;
```

Exception table:

from	to	target	type
2	4	10	any
2	31	31	any


end

The `hopAround()` method returns from the first `finally` clause by executing past the closing curly brace, but returns from the second `finally` clause by executing a `continue` statement. The first `finally` clause, therefore, exits via its *ret* instruction. But because the second `finally` clause exits via a `continue`, its *ret* instruction is never executed. The `continue` statement causes the Java Virtual Machine to jump to the top of the `while` loop again. This results in an endless loop, even though it is a `return` statement that originally causes the second `finally` clause to be executed in the first place. The `continue` statement in the `finally` clause supersedes the `return` statement, so the method never returns.

Note that the bytecodes that implement the `return` statement store a copy of the return value into local variable one before jumping to the miniature subroutine that represents the second `finally` clause. Then, after the miniature subroutine returns (in this case it never does, because the `continue` is always executed), the return value is retrieved from local variable one and returned.

This highlights the way the Java Virtual Machine returns values when `finally` clauses are also executed. Rather than returning the value of `i` after the `finally` clause is executed, the Java Virtual Machine will return the value that `i` had just *before* the `finally` clause was executed. This means that even if the `finally` clause changes the value of `i`, the method will still return the value that `i` had when the `return` statement was first reached, before the `finally` clause was invoked. If you wanted the `finally` clause to be able to change the return value of the method, you would have to put an actual `return` statement with the new return value into the `finally` clause itself.

To drive the *Hop Around* simulation, use the Step, Reset, Run, and Stop buttons. Each time you press the Step button, the simulator will execute the instruction pointed to by the pc register. If you press the Run button, the simulation will continue with no further coaxing on your part until you press the Stop button. To start the simulation over, press the Reset button. For each step of the simulation, a panel at the bottom of the applet contains an explanation of what the next instruction will do. Happy clicking.



On the CD-ROM

The CD-ROM contains the source code examples from this chapter in the `opcodes` directory. The *Hop Around* applet is contained in a web page on the CD-ROM in file `applets/HopAround.html`. The source code for this applet is found alongside its class files, in the `applets/JVMSimulators` and `applets/JVMSimulators/COM/artima/jvmsim` directories.

The Resources Page

For more information about the material presented in this chapter, visit the resources page: <http://www.artima.com/insidejvm/finally.html>

Chapter Nineteen

Method Invocation and Return

The Java Virtual Machine's instruction set has four different instructions that invoke methods. This chapter describes these four instructions and the situations in which each is used.

Method Invocation

The Java programming language provides two basic kinds of methods: instance methods and class (or static) methods. The difference between these two kinds of methods are:

1. instance methods require an instance before they can be invoked; class methods do not,
2. instance methods use dynamic (late) binding; class methods use static (early) binding.

When the Java Virtual Machine invokes a class method, it selects the method to invoke based on the type of the object reference, which is always known at compile-time. When the virtual machine invokes an instance method, by contrast, it selects the method to invoke based on the actual class of the object, which may be known only at run-time.

The Java Virtual Machine uses two different instructions to invoke these two different kinds of methods: `invokevirtual` for instance methods and `invokestatic` for class methods. These instructions are shown in Table 19-1.

Table 19-1. Method invocation

Opcode	Operand(s)	Description
<code>invokevirtual</code>	<code>indexbyte1</code> , <code>indexbyte2</code>	pop objectref and args, invoke instance method at constant pool index
<code>invokestatic</code>	<code>indexbyte1</code> , <code>indexbyte2</code>	pop args, invoke class method at constant pool index

As mentioned in earlier chapters, references to methods are initially symbolic. All invoke instructions, such as `invokevirtual` and `invokestatic`, refer to a constant pool entry that initially contains a symbolic reference. When the Java Virtual Machine encounters an invoke instruction, if the symbolic reference has not yet been resolved, the virtual machine resolves it as part of the execution of the invoke instruction.

To resolve a symbolic reference, the Java Virtual Machine locates the method being referred to symbolically and replaces the symbolic reference with a direct reference. A direct reference, such as a pointer or offset, allows the virtual machine to invoke the method more quickly if the reference is ever used again in the future.

For example, upon encountering an `invokevirtual` instruction, the Java Virtual Machine forms an unsigned 16-bit index into the constant pool of the current class from the `indexbyte1` and `indexbyte2` operands that follow the `invokevirtual` opcode. The constant pool entry contains a symbolic reference to the method to invoke.

During resolution, the Java Virtual Machine also performs several verification checks. These checks ensure that Java language rules are followed and that the `invoke` instruction is safe to execute. For example, the virtual machine first makes sure the symbolically referenced method even exists. If it exists, it checks to make sure the current class can legally access the method. For instance, if the method is private, it must be a member of the current class. If any of these checks fail, the Java Virtual Machine throws an exception. Chapter 5, "The Java Virtual Machine," gives an overview of the resolution process. Chapter 8, "The Linking Model," describes the process in detail.

Once a method has been resolved, the Java Virtual Machine is ready to invoke it. If the method is an instance method, it must be invoked on an object. For every instance method invocation, the virtual machine expects a reference to the object (`objectref`) to be on the stack. In addition to `objectref`, the virtual machine expects the arguments (`args`) required by the method, if any, to be on the stack. If the method is a class method, which doesn't require an `objectref` because class methods aren't invoked on an object, only the `args` are on the stack.

The `objectref` and `args` (or just `args`, in the case of a class method) must be pushed onto the calling method's operand stack by the instructions that precede the `invoke` instruction.

Invoking a Java Method

As mentioned in Chapter 5, "The Java Virtual Machine," the virtual machine creates a new stack frame for each Java (not native) method it invokes. The stack frame contains space for the method's local variables, its operand stack, and any other information required by a particular virtual machine implementation. The size of the local variables and operand stack are calculated at compile-time and placed into the class file, so the virtual machine knows just how much memory will be needed by the method's stack frame. When it invokes a method, it creates a stack frame of the proper size for that method. The virtual machine pushes the new stack frame onto the Java stack.

For an instance method, the virtual machine pops the `objectref` and `args` from the operand stack of the calling method's stack frame. It places the `objectref` on the new stack frame as local variable 0, and all the `args` as local variable 1, 2, and so on. The `objectref` is the implicit `this` pointer that is passed to any instance method.

For a class method, the virtual machine just pops the `args` from the operand stack of the calling method's frame and places them onto the new stack frame as local variable 0, 1, 2, and so on.

Once the `objectref` and `args` (or just the `args`, for a class method) have been placed into the local variables of the new frame, the virtual machine makes the new stack frame current and sets the program counter to point to the first instruction in the new method.

Invoking a Native Method

As mentioned in Chapter 5, "The Java Virtual Machine," the virtual machine invokes native methods in an implementation dependent manner. When invoking a native method, the virtual machine does not push a new stack frame onto the Java stack for the native method. At the point in which the thread enters the native method, it leaves the Java stack behind. When the native method returns, the Java stack will once again be used.

Other Forms of Method Invocation

Although instance methods are normally invoked with `invokevirtual`, in certain situations, two other opcodes may be used: `invokespecial` and `invokeinterface`. These opcodes are shown in Table 19-2.

The `invokespecial` instruction is used for three situations in which an instance method must be invoked based on the type of the reference, not on the class of the object. The three situations are:

1. instance initialization (`<init()`) methods
2. private methods
3. methods invoked with the `super` keyword

`Invokeinterface` is used to invoke an instance method given a reference to an interface.

Table 19-2. Method invocation

Opcode	Operand(s)	Description
<code>invokespecial</code>	<code>indexbyte1</code> , <code>indexbyte2</code>	pop objectref and args, invoke instance method at constant pool index
<code>invokeinterface</code>	<code>indexbyte1</code> , <code>indexbyte2</code>	pop objectref and args, invoke instance method at constant pool index

As mentioned in Chapter 7, "The Lifetime of a Class," class initialization (or `<clinit()`) methods are always invoked directly by the Java Virtual Machine itself, never by any bytecodes. There is no instruction in the Java Virtual Machine's instruction set that will invoke a method named `<clinit()`. If some class file attempts to invoke a `<clinit()` method with any of the four instructions described in this chapter, the virtual machine will throw an exception.

The invokespecial instruction

`Invokespecial` differs from `invokevirtual` primarily in that `invokespecial` normally (with one special exception) selects a method based on the type of the reference rather than the class of the object. In other words, it does static binding instead of dynamic binding. In each of the three situations where `invokespecial` is used, dynamic binding wouldn't yield the desired result.

invokespecial and <init()

As mentioned in Chapter 7, "The Lifetime of a Class," the <init() method, or instance initialization method, is where the compiler places code for constructors and instance variable initializers. A class gets one <init() method in the class file for each constructor in the source. If you don't explicitly declare a constructor in the source, the compiler will generate a default no-arg constructor for you. This default constructor also ends up as an <init() method in the class file. So just as every class will have at least one constructor, every class will have at least one <init() method. These methods are always invoked with `invokespecial`.

The <init() methods are called only when a new instance is created. At least one <init() method will be invoked for each class along the inheritance path of the newly created object, and multiple <init() methods could be invoked for any one class along that path.

The reason `invokespecial` is used to invoke <init() methods is that subclass <init() methods need to be able to invoke superclass <init() methods. This is how multiple <init() methods get invoked when an object is instantiated. The virtual machine invokes an <init() method declared in the object's class. That <init() method first invokes either another <init() method in the same class, or an <init() method in its superclass. This process continues all the way up to `Object`.

For example, consider this code:

begin

```
// On CD-ROM in file invoke/ex1/Dog.java
class Dog {

}
```

```
// On CD-ROM in file invoke/ex1/CockerSpaniel.java
class CockerSpaniel extends Dog {
```

```
    public static void main(String args[]) {

        CockerSpaniel bootsie = new CockerSpaniel();

    }
}
```

end

When you invoke `main()`, the virtual machine will allocate space for a new `CockerSpaniel` object, then invoke `CockerSpaniel`'s default no-arg <init() method to initialize that space. That method will invoke `Dog`'s <init() method, which will invoke `Object`'s <init() method. Here are the bytecodes for the `main()` method of class `CockerSpaniel`:

begin

```

0 new #1 <Class CockerSpaniel
3 invokespecial #3 <Method <init() of class CockerSpaniel

6 return

end

```

`CockerSpaniel`'s `main()` method allocates memory for the new `CockerSpaniel` object and initializes that memory to default initial values with the `new` instruction. (The "#1" specifies the constant pool entry that refers to the class to instantiate, in this case, class `CockerSpaniel`.) The `new` instruction pushes a reference to the newly created `CockerSpaniel` object onto the stack. The `main()` method then calls the `<init()` method of class `CockerSpaniel` using `invokespecial` on that object reference. (The "#3" specifies the constant pool entry that contains the reference to `CockerSpaniel`'s `<init()` method.) The Java Virtual Machine pushes a new frame onto the Java Stack and places the object reference into local variable 0 of the new frame. Here are the bytecodes for the `<init()` method of class `CockerSpaniel`:

```

begin

0 aload_0
1 invokespecial #4 <Method <init() of class Dog

4 return

end

```

As mentioned above, this `<init()` method corresponds to the default no-arg constructor generated automatically by the compiler for class `CockerSpaniel`. It first pushes the reference to the object being initialized onto the stack from local variable 0, then it invokes the `Dog`'s `<init()` method on that reference. (The "#4" specifies the constant pool entry that contains the reference to `Dog`'s `<init()` method.) Here are the bytecodes for `Dog`'s `<init()` method:

```

begin

0 aload_0
1 invokespecial #3 <Method <init() of class java.lang.Object

4 return

end

```

This `<init()` method corresponds to the default no-arg constructor generated automatically by the compiler for class `Dog`. It first pushes the reference to the object being initialized onto the stack from local variable 0, then it invokes the `Object`'s `<init()` method on that reference. (The "#3" specifies the constant pool entry that contains the reference to `Object`'s `<init()` method. Remember that this is not the same constant pool as that referred to by the methods of class `CockerSpaniel`. Each class has its own constant pool.) When all three of these `<init()` methods have returned, the new `CockerSpaniel` object created by `main()` is fully initialized and ready for use.

Because every class has at least one `<init()` method, it is common that classes have `<init()` methods

with the same signature. (A method's *signature* is its name and the number and types of its arguments.) The `<init()` methods for the three classes in the inheritance path for `CockerSpaniel`, for example, all have the same signature. `CockerSpaniel`, `Dog`, and `Object` all contain a method named `<init()` that takes no arguments.

It would be impossible to invoke a `Dog`'s `<init()` method from `CockerSpaniel`'s `<init()` method using `invokevirtual`, because `invokevirtual` would perform dynamic binding and invoke `CockerSpaniel`'s `<init()` method. With `invokespecial`, however, `Dog`'s `<init` method can be invoked from `CockerSpaniel`'s `<init()` method, because the type of the reference placed in `CockerSpaniel`'s class file (in constant pool entry "#4") is `Dog`.

invokespecial and Private Methods

In the case of private instance methods, it must be possible for a subclass to declare an instance method with the same signature as a private instance method in a superclass. (`invokespecial` is not used to invoke private class methods, just private instance methods. Private class methods are invoked with `invokestatic`.) For example, consider the following code in which `interestingMethod()` is declared as private in a superclass and with package access in a subclass:

begin

```
// On CD-ROM in file invoke/ex2/Superclass.java
class Superclass {
```

```
    private void interestingMethod() {
        System.out.println("Superclass's interesting method.");
    }
```

```
    void exampleMethod() {
        interestingMethod();
    }
```

```
}
```

```
// On CD-ROM in file invoke/ex2/Subclass.java
class Subclass extends Superclass {
```

```
    void interestingMethod() {
        System.out.println("Subclass's interesting method.");
    }
```

```

    public static void main(String args[]) {
        Subclass me = new Subclass();
        me.exampleMethod();
    }
}

```

end

When you invoke `main()` in `Subclass` as defined above, it instantiates a new `Subclass` object and invokes `exampleMethod()` on it. Here are the bytecodes for the `main()` method of class `Subclass`:

begin

```

0 new #2 <Class Subclass
3 dup
4 invokespecial #6 <Method <init() of class Subclass
7 astore_1
8 aload_1
9 invokevirtual #8 <Method void exampleMethod() of class Superclass
12 return

```

end

`Subclass` inherits `exampleMethod()` from `Superclass`. When `main()` invokes `exampleMethod()` on a `Subclass` object, it uses `invokevirtual`. The Java Virtual Machine will create and push a new frame onto the stack and begin executing the bytecodes of `exampleMethod()` as defined in class `Superclass`. Here are the bytecodes for `exampleMethod()`:

begin

```

0 aload_0
1 invokespecial #7 <Method void interestingMethod() of Superclass
4 return

```

end

Note that `exampleMethod()` first pushes the reference passed in local variable 0 (the hidden `this` passed to all instance methods) onto the stack. It then invokes `interestingMethod()` on that reference with the `invokespecial` instruction. The class identified in the constant pool entry

specified by the `invokespecial` instruction, entry #7, is `Superclass`. The Java Virtual Machine invokes the `interestingMethod()` defined in `Superclass` *even though* the object is an instance of class `Subclass` and there is an accessible `interestingMethod()` defined in `Subclass`.

The program correctly prints "Superclass's interesting method". If `interestingMethod()` had been invoked with an `invokevirtual` instruction instead of `invokespecial`, the program would have printed "Subclass's interesting method". Why? Because the virtual machine would choose the `interestingMethod()` to call based on the actual class of the object, which is `Subclass`. So it would use `Subclass`'s `interestingMethod()`. With `invokespecial` however, the virtual machine selects the method based on the type of the reference, so `Superclass`'s version of `interestingMethod()` is invoked.

`invokespecial` and `super`

In the case of invocation of a method with the `super` keyword, as in `super.someMethod()`, you want the superclass's version of a method to be invoked, even if the current class overrides the method. Once again, `invokevirtual` would invoke the current class's version, so it can't be used. For an example, consider this code:

begin

```
// On CD-ROM in file invoke/ex3/Cat.java
class Cat {
```

```
    void someMethod() {
    }
```

```
}
```

```
// On CD-ROM in file invoke/ex3/TabbyCat.java
class TabbyCat extends Cat{
```

```
    void someMethod() {
        super.someMethod();
    }
```

```
}
```

end

Here are the bytecodes for `TabbyCat`'s `someMethod()`:

begin

```
0 aload_0
```



```
1 invokespecial #4 <Method void someMethod() of class Cat
4 return

end
```

Had `invokevirtual` been used in this case, the `someMethod()` defined in class `TabbyCat` would have been invoked. Because `invokespecial` is used and the constant pool entry (#4, in this case) indicates the `someMethod()` declared in class `Cat` should be invoked, the Java Virtual Machine correctly invokes the superclass's (`Cat`'s) version of `someMethod()`.

Whether or not the Java Virtual Machine actually uses static binding to execute an `invokespecial` instruction, or instead uses a special kind of dynamic binding, depends on whether or not the referring class's `ACC_SUPER` flag is set. Prior to JDK version 1.0.2, the `invokespecial` instruction was called `invokenonvirtual` and always resulted in static binding. It turned out, however, that `invokenonvirtual`'s stubborn insistence on static binding did not yield a correct implementation of the semantics of the Java language in all cases. (In other words, there was a "bug" in the instruction set.) In JDK 1.0.2, the `invokenonvirtual` instruction was renamed `invokespecial` and its semantics were changed. In addition, a new flag, `ACC_SUPER`, was added to the `access_flags` item of the Java class file.

The `ACC_SUPER` flag item of a class file indicates to the Java Virtual Machine which semantics it should use to execute `invokespecial` instructions it encounters in the bytecodes of that class file. If the `ACC_SUPER` flag is not set, the virtual machine uses the old (`invokenonvirtual`) semantics. Else, if the `ACC_SUPER` flag is set, the virtual machine uses the new semantics. All new Java compilers are supposed to set the `ACC_SUPER` flag in every class file they generate.

According to the old semantics, the virtual machine will perform static binding in all cases when executing `invokespecial`. By contrast, the new semantics requires static binding in all cases except one: the invocation of superclass methods.

According to the new semantics, when the Java Virtual Machine resolves an `invokespecial` instruction's symbolic reference to a superclass method, it dynamically searches the current class's superclasses to find the nearest superclass implementation of the method. (As used here, "nearest" means the implementation of the method declared in the superclass that is closest to the current class in its inheritance hierarchy.) In most cases, the virtual machine will likely discover that the nearest implementation of the method is in the superclass listed in the symbolic reference. But it is possible that the virtual machine could find the nearest implementation in a different superclass.

For example, imagine you create an inheritance hierarchy of three classes: `Animal`, `Dog`, and `CockerSpaniel`. Assume class `Dog` extends class `Animal`, class `CockerSpaniel` extends class `Dog`, and that a method defined in `CockerSpaniel` uses `invokespecial` to invoke a non-private superclass method named `walk()`. Assume also that when you compiled `CockerSpaniel`, the compiler set the `ACC_SUPER` flag. In addition, assume that when you compiled `CockerSpaniel`, class `Animal` defined a `walk()` method, but `Dog` didn't. In that case, the symbolic reference from `CockerSpaniel` to the `walk()` method would give `Animal` as its class. When the `invokespecial` instruction in `CockerSpaniel`'s method is executed, the virtual machine would dynamically select and invoke `Animal`'s `walk()` method.

Now imagine that later, you added a `walk()` method to `Dog`, and recompiled `Dog`, but didn't recompile `CockerSpaniel`. `CockerSpaniel`'s symbolic reference to the superclass `walk()` method still claims `Animal` as its class, even though there is now an implementation of `walk()` in `Dog`'s class file. Nevertheless, when the `invokespecial` instruction in `CockerSpaniel`'s method is executed, the virtual machine would dynamically select and invoke `Dog`'s implementation of the `walk()` method.

The invokeinterface Instruction

The `invokeinterface` opcode performs the same function as `invokevirtual`, it invokes instance methods and uses dynamic binding. The difference between these two instructions is that `invokevirtual` is used when the type of the reference is a *class*, whereas `invokeinterface` is used when the type of the reference is an *interface*.

The Java Virtual Machine uses a different opcode to invoke a method on an interface reference because it can't make as many assumptions about the method table offset given an interface reference as it can given a class reference. (Method tables are described in Chapter 8, "The Linking Model.") Given a class reference, a method will always occupy the same position in the method table, independent of the actual class of the object. This is not true given an interface reference. The method could occupy different locations for different classes that implement the same interface.

For examples of the use of `invokeinterface` in bytecodes, see the "Examples of Method Invocation" section later in this chapter.

Invocation Instructions and Speed

As you might imagine, invoking a method given an interface reference is likely slower than invoking a method given a class reference. When the Java Virtual Machine encounters an `invokevirtual` instruction and resolves the symbolic reference to a direct reference to an instance method, that direct reference is likely an offset into a method table. From that point forward, the same offset can be used. For an `invokeinterface` instruction, however, the virtual machine will have to search through the method table every single time the instruction is encountered, because it can't assume the offset is the same as the previous time.

The fastest instructions will likely be `invokespecial` and `invokestatic`, because methods invoked by these instructions are statically bound. When the Java Virtual Machine resolves the symbolic reference for these instructions and replaces it with a direct reference, that direct reference will likely include a pointer to the actual bytecodes.

Examples of Method Invocation

The following code illustrates the various ways the Java Virtual Machine invokes methods and shows which invocation opcode is used in each situation:

```
begin

// On CD-ROM in file invoke/ex4/InYourFace.java
interface InYourFace {

    void interfaceMethod ();
```

```
}
```

```
// On CD-ROM in file invoke/ex4/ItsABirdItsAPlaneItsSuperclass.java  
class ItsABirdItsAPlaneItsSuperclass implements InYourFace {
```

```
    ItsABirdItsAPlaneItsSuperclass(int i) {  
        super();                // invokespecial (of an <init)  
    }
```

```
    static void classMethod() {  
    }
```

```
    void instanceMethod() {  
    }
```

```
    final void finalInstanceMethod() {  
    }
```

```
    public void interfaceMethod() {  
    }  
}
```

```
// On CD-ROM in file invoke/ex4/Subclass.java  
class Subclass extends ItsABirdItsAPlaneItsSuperclass {
```

```
    Subclass() {  
        this(0);                // invokespecial (of an <init)  
    }
```

```
    Subclass(int i) {  
        super(i);                // invokespecial (of an <init)
```

```

    }

    private void privateMethod() {
    }

    void instanceMethod() {
    }

    final void anotherFinalInstanceMethod() {
    }

    void exampleInstanceMethod() {

        instanceMethod();           // invokevirtual
        super.instanceMethod();     // invokespecial

        privateMethod();           // invokespecial

        finalInstanceMethod();     // invokevirtual
        anotherFinalInstanceMethod(); // invokevirtual

        interfaceMethod();         // invokevirtual

        classMethod();             // invokestatic
    }
}

// On CD-ROM in file invoke/ex4/UnrelatedClass.java
class UnrelatedClass {

```

```

public static void main(String args[]) {

    Subclass sc = new Subclass(); // invokespecial (of an <init)

    Subclass.classMethod();        // invokestatic
    sc.classMethod();              // invokestatic
    sc.instanceMethod();           // invokevirtual
    sc.finalInstanceMethod();      // invokevirtual
    sc.interfaceMethod();          // invokevirtual

    InYourFace iyf = sc;

    iyf.interfaceMethod();         // invokeinterface
}
}

end

```

Here are the bytecodes generated by `javac` for each of these classes (The compiler generates no bytecodes for the `InYourFace` interface.):

```

begin
// Methods of class ItsABirdItsAPlaneItsSuperclass

// Method <init(int)
0 aload_0
1 invokespecial #4 <Method <init() of class java.lang.Object
4 return

// Method void classMethod()
0 return

// Method void instanceMethod()
0 return

```

```

// Method void finalInstanceMethod()

    0 return

// Method void interfaceMethod()

    0 return

// -----

// Methods of class Subclass

// Method <init()

    0 aload_0
    1 iconst_0
    2 invokespecial #4 <Method <init(int) of class Subclass
    5 return

// Method <init(int)

    0 aload_0
    1 iload_1
    2 invokespecial #3 <Method <init(int) of class
        ItsABirdItsAPlaneItsSuperclass
    5 return

// Method void privateMethod()

    0 return

// Method void instanceMethod()

    0 return

// Method void anotherFinalInstanceMethod()

    0 return

```

```

// Method void exampleInstanceMethod()
0 aload_0
1 invokevirtual #9 <Method void instanceMethod() of class Subclass
4 aload_0
5 invokespecial #8 <Method void instanceMethod() of class
    ItsABirdItsAPlaneItsSuperclass
8 aload_0
9 invokespecial #11 <Method void privateMethod() of class Subclass
12 aload_0
13 invokevirtual #7 <Method void finalInstanceMethod() of class
    ItsABirdItsAPlaneItsSuperclass
16 aload_0
17 invokevirtual #5 <Method void anotherFinalInstanceMethod() of
    class Subclass
20 aload_0
21 invokevirtual #10 <Method void interfaceMethod() of class
    ItsABirdItsAPlaneItsSuperclass
24 invokestatic #6 <Method void classMethod() of class
    ItsABirdItsAPlaneItsSuperclass
27 return

// -----
// Methods of class UnrelatedClass

// Method <init()
0 aload_0
1 invokespecial #7 <Method java.lang.Object()
4 return

```

```

// Method void main(java.lang.String[])
0 new #3 <Class Subclass
3 dup
4 invokespecial #6 <Method <init() of class Subclass
7 astore_1
8 invokestatic #8 <Method void classMethod() of class
    ItsABirdItsAPlaneItsSuperclass
11 invokestatic #8 <Method void classMethod() of class
    ItsABirdItsAPlaneItsSuperclass
14 aload_1
15 invokevirtual #10 <Method void instanceMethod() of class Subclass
18 aload_1
19 invokevirtual #9 <Method void finalInstanceMethod() of class
    ItsABirdItsAPlaneItsSuperclass
22 aload_1
23 invokevirtual #12 <Method void interfaceMethod() of class
    ItsABirdItsAPlaneItsSuperclass
26 aload_1
27 astore_2
28 aload_2
29 invokeinterface (args 1) #11 <Method void interfaceMethod() of
    interface InYourFace
34 return

end

```

Returning from Methods

There are several opcodes that return from a method, one for each type of return value. These opcodes, which are shown in Table 19-3, take no operands. If there is a return value, it must be on the operand stack. The return value is popped off the operand stack and pushed onto the operand stack of the calling method's stack frame. The current stack frame is popped, and the calling method's stack frame becomes current. The program counter is reset to the instruction in the calling method just following the instruction that invoked the returning method.

Table 19-3. Returning from methods

Opcode	Operand (s)	Description
<code>ireturn</code>	<code>none</code>	pop <code>int</code> , push onto stack of calling method and return
<code>lreturn</code>	<code>none</code>	pop <code>long</code> , push onto stack of calling method and return
<code>freturn</code>	<code>none</code>	pop <code>float</code> , push onto stack of calling method and return
<code>dreturn</code>	<code>none</code>	pop <code>double</code> , push onto stack of calling method and return
<code>areturn</code>	<code>none</code>	pop object reference, push onto stack of calling method and return
<code>return</code>	<code>none</code>	return void

The `ireturn` instruction is used for methods that return `int`, `char`, `byte`, or `short`.

On the CD-ROM

The CD-ROM contains the source code examples from this chapter in subdirectories of the `invoke` directory.

The Resources Page

For more information about the material presented in this chapter, visit the resources page:
<http://www.artima.com/insidejvm/invoke.html>

Chapter Twenty

Thread Synchronization

One of the strengths of the Java programming language is its support for multithreading at the language level. Much of this support centers on *synchronization*: coordinating activities and data access among multiple threads. The mechanism that Java uses to support synchronization is the *monitor*. This chapter describes monitors and shows how they are used by the Java Virtual Machine. It describes how one aspect of monitors, the locking and unlocking of data, is supported in the instruction set.

Monitors

A monitor supports two kinds of thread synchronization: *mutual exclusion* and *cooperation*. Mutual exclusion, which is supported in the Java Virtual Machine via object locks, enables multiple threads to independently work on shared data without interfering with each other. Cooperation, which is supported in the Java Virtual Machine via the wait and notify methods of class `Object`, enables threads to work together towards a common goal.

A monitor is like a building that contains one special room that can be occupied by only one thread at a time. The room usually contains some data. From the time a thread enters this room to the time it leaves, it has exclusive access to any data in the room. Entering the monitor building is called "entering the monitor." Entering the special room inside the building is called "acquiring the monitor." Occupying the room is called "owning the monitor," and leaving the room is called "releasing the monitor." Leaving the entire building is called "exiting the monitor."

In addition to being associated with a bit of data, a monitor is associated with one or more bits of code called *critical sections*. A critical section is code that needs to be executed as one indivisible operation. In other words, one thread must be able to execute a critical section from beginning to end without another thread concurrently executing a critical section of the same monitor. A monitor enforces this one-thread-at-a-time execution of its critical sections. The only way a thread can enter a monitor is by arriving at the beginning of one of the critical sections associated with that monitor. The only way a thread can move forward and execute the critical section is by acquiring the monitor.

When a thread arrives at the beginning of a critical section, it is placed into an *entry set* for the associated monitor. The entry set is like the front hallway of the monitor building. If no other thread is waiting in the entry set and no other thread currently owns the monitor, the thread acquires the monitor and continues executing the critical section. When the thread finishes executing the critical section, it exits (and releases) the monitor.

If a thread arrives at the beginning of a critical section of a monitor that is already owned by another

thread, the newly arrived thread must wait in the entry set. When the current owner exits the monitor, the newly arrived thread must compete with any other threads also waiting in the entry set. Only one thread will win the competition and acquire the monitor.

The first kind of synchronization listed above, mutual exclusion, refers to the mutually exclusive execution of critical sections by multiple threads. At any one time, only one thread can be executing a critical section of a particular monitor. In general, mutual exclusion is important only when multiple threads are sharing data. If two threads are not working with any common data, they usually can't interfere with each other and needn't execute in a mutually exclusive way. On a Java Virtual Machine implementation that doesn't time slice, however, a higher priority thread that is never blocked will interfere with any lower priority threads, even if none of the threads share data. The higher priority thread will monopolize the CPU at the expense of the lower priority threads. Lower priority threads will never get any CPU time. In such a case, a monitor that protects no data may be used to orchestrate these threads to ensure all threads get some CPU time. Nevertheless, in most cases a monitor protects data that is accessed through the critical section code. In cases where the data can be accessed only through the critical sections, the monitor enforces mutually exclusive access to that data.

The other kind of synchronization listed above as supported by monitors is cooperation. Whereas mutual exclusion helps keep threads from interfering with one another while sharing data, cooperation helps threads to work together towards some common goal.

Cooperation is important when one thread needs some data to be in a particular state and another thread is responsible for getting the data into that state. For example, one thread, a "read thread," may be reading data from a buffer that another thread, a "write thread," is filling. The read thread needs the buffer to be in a "not empty" state before it can read any data out of the buffer. If the read thread discovers that the buffer is empty, it must wait. The write thread is responsible for filling the buffer with data. Once the write thread has done some more writing, the read thread can do some more reading.

The form of monitor used by the Java Virtual Machine is called a "Wait and Notify" monitor. (It is also sometimes called a "Signal and Continue" monitor.) In this kind of monitor, a thread that currently owns the monitor can suspend itself inside the monitor by executing a *wait command*. When a thread executes a wait, it releases the monitor and enters a *wait set*. The thread will stay suspended in the wait set until some time after another thread executes a *notify command* inside the monitor. When a thread executes a notify, it continues to own the monitor until it releases the monitor of its own accord, either by executing a wait or by completing the critical section. After the notifying thread has released the monitor, the waiting thread will be resurrected and will reacquire the monitor.

The kind of monitor used in the Java Virtual Machine is sometimes called a Signal and Continue monitor because after a thread does a notify (the signal) it retains ownership of the monitor and continues executing the critical section (the continue). At some later time, the notifying thread releases the monitor and a waiting thread is resurrected. Presumably, the waiting thread suspended itself because the data protected by the monitor wasn't in a state that would allow the thread to continue doing useful work. Also, the notifying thread presumably executed the notify command after it had placed the data protected by the monitor into the state desired by the waiting thread. But because the notifying thread continued, it may have altered the state after the notify such that the waiting thread still can't do useful work. Alternatively, a third thread may have acquired the monitor after the notifying thread released it but before the waiting thread acquired it, and the third thread may have changed the state of the protected data. As a result, a notify must often be considered by waiting threads merely as a hint that the desired state *may* exist. Each time a waiting thread is resurrected, it may need to check the state again to determine whether it can move forward and do useful work. If it finds the data still isn't in the desired state, the thread could execute another wait or give up and exit the monitor.

As an example, consider once again the scenario described above that involves a buffer, a read thread, and a write thread. Assume the buffer is protected by a monitor. When a read thread enters the monitor that protects the buffer, it checks to see if the buffer is empty. If the buffer is not empty, the read thread reads (and removes) some data from the buffer. Satisfied, it exits the monitor. On the other hand, if the buffer is empty, the read thread executes a wait command. As soon as it executes the wait, the read thread is suspended and placed into the monitor's wait set. In the process, the read thread releases the monitor, which becomes available to other threads. At some later time, the write thread enters the monitor, writes some data into the buffer, executes a notify, and exits the monitor. When the write thread executes the notify, the read thread is marked for eventual resurrection. After the write thread has exited the monitor, the read thread is resurrected as the owner of the monitor. If there is any chance that some other thread has come along and consumed the data left by the write thread, the read thread must explicitly check to make sure the buffer is not empty. If there is no chance that any other thread has consumed the data, then the read thread can just assume the data exists. The read thread reads some data from the buffer and exits the monitor.

A graphical depiction of the kind of monitor used by a Java Virtual Machine is shown in Figure 20-1. This figure shows the monitor as three rectangles. In the center, a large rectangle contains a single thread, the monitor's owner. On the left, a small rectangle contains the entry set. On the right, another small rectangle contains the wait set. Active threads are shown as dark gray circles. Suspended threads are shown as light gray circles.




Figure 20-1. A Java monitor.

Figure 20-1 also shows several numbered doors that threads must "pass through" to interact with the monitor. When a thread arrives at the start of a critical section, it enters the monitor via the leftmost door, door number one, and finds itself in the rectangle that houses the entry set. If no thread currently owns the monitor and no other threads are waiting in the entry set, the thread passes immediately through the next door, door number two, and becomes the owner of the monitor. As the monitor owner, the thread continues executing the critical section. If, on the other hand, there is another thread currently claiming ownership of the monitor, the newly arrived thread must wait in the entry set, possibly along with other threads already waiting there. The newly arrived thread is blocked and therefore doesn't

execute any further into the critical section.

Figure 20-1 shows three threads suspended in the entry set and four threads suspended in the wait set. These threads will remain where they are until the current owner of the monitor--the active thread--releases the monitor. The active thread can release the monitor in either of two ways: it can complete the critical section it is executing or it can execute a wait command. If it completes the critical section, it exits the monitor via the door at the bottom of the central rectangle, door number five. If it executes a wait command, it releases the monitor as it travels through door number three, the door to the wait set.

If the former owner did not execute a notify before it released the monitor, then only the threads in the entry set will compete to acquire the monitor. If the former owner did execute a notify, then the entry set threads will have to compete with one or more threads from the wait set. If a thread from the entry set wins the competition, it passes through the door number two and becomes the new owner of the monitor. If a thread from the wait set wins the competition, it exits the wait set and reacquires the monitor as it passes through door number four. Note that doors three and four are the only ways a thread can enter or exit the wait set. A thread can only execute a wait command if it currently owns the monitor, and it can't leave the wait set without automatically becoming again the owner of the monitor.

In the Java Virtual Machine, threads can optionally specify a timeout when they execute a wait command. If a thread does specify a timeout, and no other thread executes a notify before the timeout expires, the waiting thread in effect receives an automatic notify from the virtual machine. After the timeout expires, the waiting thread will be resurrected even if no other thread has executed an explicit notify.

The Java Virtual Machine offers two kinds of notify commands: "notify" and "notify all." A notify command selects one thread arbitrarily from the wait set and marks it for eventual resurrection. A notify all command marks all threads currently in the wait set for eventual resurrection.

To a great extent, the manner in which a Java Virtual Machine implementation selects the next thread from the wait or entry sets is a decision of individual implementation designers. For example, implementation designers can decide how to select:

- a thread from the wait set given a notify command
- the order to resurrect threads from the wait set given a notify all command
- the order to allow threads from the entry set to acquire the monitor
- how to choose between threads suspended in the wait set versus the entry set after a notify command

You might think it would make sense to implement entry set and wait sets as first-in-first-out (FIFO) queues, so that the thread that waits the longest will be the first chosen to acquire the monitor. Alternatively, it might make sense to have ten FIFO queues, one for each priority a thread can have inside the Java Virtual Machine. The virtual machine could then choose the thread that has been waiting the longest in the highest priority queue that contains any waiting threads. Implementations may take approaches such as these, but you can't depend on it. Implementations are free to implement the entry and wait sets as last-in-first-out (LIFO) queues, to select lower priority threads before higher priority threads, or to do anything else that may not seem to make sense. Implementations are free to select threads in an arbitrary manner that defies analysis and yields surprising orderings.

As a programmer, you must not rely on any particular selection algorithm or treatment of priorities, at least if you are trying to write a Java program that is platform independent. For example, because you

don't know what order threads in the wait set will be chosen for resurrection by the notify command, you should use notify (as opposed to notify all) only when you are absolutely certain there will only be one thread suspended in the wait set. If there is a chance more than one thread will be suspended in the wait set at any one time, you should probably use notify all. Otherwise, on some Java Virtual Machine implementations a particular thread may be stuck in the wait set for a very long time. If a notify always selects the most recent arrival from the wait set and the wait set always contains multiple threads, some threads that have been waiting the longest may never be resurrected.

Object Locking

As mentioned in earlier chapters, some of the Java Virtual Machine's runtime data areas are shared by all threads, others are private to individual threads. Because the heap and method area are shared by all threads, Java programs need to coordinate multi-threaded access to two kinds of data:

- instance variables, which are stored on the heap
- class variables, which are stored in the method area

Programs never need to coordinate access to local variables, which reside on Java stacks, because data on the Java stack is private to the thread to which the Java stack belongs.

In the Java Virtual Machine, every object and class is logically associated with a monitor. For objects, the associated monitor protects the object's instance variables. For classes, the monitor protects the class's class variables. If an object has no instance variables, or a class has no class variables, the associated monitor protects no data.

To implement the mutual exclusion capability of monitors, the Java Virtual Machine associates a *lock* (sometimes called a *mutex*) with each object and class. A lock is like a privilege that only one thread can "own" at any one time. Threads need not obtain a lock to access instance or class variables. If a thread does obtain a lock, however, no other thread can access the locked data until the thread that owns the lock releases it. (To "lock an object" is to acquire the monitor associated with that object.)

Class locks are actually implemented as object locks. As mentioned in earlier chapters, when the Java Virtual Machine loads a class file, it creates an instance of class `java.lang.Class`. When you lock a class, you are actually locking that class's `Class` object.

A single thread is allowed to lock the same object multiple times. For each object, the Java Virtual Machine maintains a count of the number of times the object has been locked. An unlocked object has a count of zero. When a thread acquires the lock for the first time, the count is again incremented to one. Each time the thread acquires a lock on the same object, the count is again incremented. (Only the thread that already owns an object's lock is allowed to lock it again. No other thread can lock the object until the owning thread releases the lock.) Each time the thread releases the lock, the count is decremented. When the count reaches zero, the lock is released and made available to other threads.

A thread in the Java Virtual Machine requests a lock when it arrives at the beginning of a critical section. In Java, there are two kinds of critical sections: synchronized statements and synchronized methods. (These are described in detail later in this chapter.) Each critical section in a Java program is associated with an object reference. When a thread arrives at the first instruction in a critical section, the thread must obtain a lock on the referenced object. The thread is not allowed to execute the code until it obtains the lock. Once it has obtained the lock, the thread enters the block of protected code. When the thread leaves the block, no matter how it leaves the block, it releases the lock on the associated object.

Note that as a Java programmer, you never explicitly lock an object. Object locks are internal to the Java Virtual Machine. In your Java programs, you identify the critical sections of your program by writing synchronized statements and methods. As the Java Virtual Machine runs your program, it automatically locks an object or class every time it encounters a critical section.

Synchronization Support in the Instruction Set

As mentioned above, the language provides two built-in ways to identify critical sections in your programs: synchronized statements and synchronized methods. These two mechanisms, which implement the mutual exclusion aspect of synchronization, are supported by the Java Virtual Machine's instruction set.

Synchronized Statements

To create a synchronized statement, you use the synchronized keyword with an expression that evaluates to an object reference, as in the `reverseOrder()` method below:

begin

```
// On CD-ROM in file threads/ex1/KitchenSync.java
class KitchenSync {
```

```
    private int[] intArray = new int[10];
```

```
    void reverseOrder() {
```

```
        synchronized (this) {
```

```
            int halfWay = intArray.length / 2;
```

```
            for (int i = 0; i < halfWay; ++i) {
```

```
                int upperIndex = intArray.length - 1 - i;
```

```
                int save = intArray[upperIndex];
```

```
                intArray[upperIndex] = intArray[i];
```

```
                intArray[i] = save;
```

```
            }
```

```
        }
```

```
    }
```

```
    // ...
```

```
}
```

end

In the above case, the statements contained within the synchronized block will not be executed until a lock is acquired on the current object (`this`). If instead of a `this` reference, the expression yielded a reference to another object, the lock associated with that object would be acquired before the thread continued. If the expression yields a reference to an instance of class `CLASS`, the lock associated with the class is acquired.

Two opcodes, `monitorenter` and `monitorexit`, are used for synchronization blocks within methods. These opcodes are shown in the Table 20-1.

Table 20-1. Monitors

3 columns

Opcode	Operand(s)	Description
<code>monitorenter</code>	<code>none</code>	<code>pop objectref</code> , acquire the lock associated with <code>objectref</code>
<code>monitorexit</code>	<code>none</code>	<code>pop objectref</code> , release the lock associated with <code>objectref</code>

end table

When `monitorenter` is encountered by the Java Virtual Machine, it acquires the lock for the object referred to by `objectref` on the stack. If the thread already owns the lock for that object, the count that is associated with the lock is incremented. Each time `monitorexit` is executed for the thread on the object, the count is decremented. When the count reaches zero, the monitor is released.

Here is the bytecode sequence generated by the `reverseOrder()` method of the `KitchenSync` class:

begin

```
// First place the reference to the object to lock into local
// variable 1. This local variable will be used by both the
```

```
// monitorenter and monitorexit instructions.
```

```
0 aload_0          // Push local var 0 (the this reference)
```

```
1 astore_1         // Store into local var 1
```

```
// Now acquire the lock on the referenced object
```

```
                // Push local var 1 (the this reference; the
```

```
2 aload_1          // object to lock)
```

```
                // Pop reference, acquire the lock
```

```
3 monitorenter     // on referenced object
```



```

// The code of the synchronized block begins here. A thread will not
// execute the next instruction, aload_0, until a lock has been
// successfully acquired on the this reference above.

4  aload_0          // Push the object ref at loc var 0 (the this ref)
                        // Pop object ref, push ref to instance variable
                        // intArray

5  getfield #4 <Field int intArray[]

8  arraylength      // Pop array ref, push int array length

9  iconst_2         // Push constant int 2

10 idiv             // Pop two ints, divide, push int result
                        // Pop int into local var 3:

11 istore_3         // int halfway = intArray.length/2;

// This is the start of the code for the for loop

12 iconst_0         // Push constant int 0

13 istore 4         // Pop into local var 2: int i = 0;

15 goto 65          // Jump to for loop condition check

// This is the start of the body of the for loop

18 aload_0          // Push the object ref at loc var 0 (the this ref)
                        // Pop object ref, push ref to instance variable
                        // intArray

19 getfield #4 <Field int intArray[]

22 arraylength      // Pop array ref, push int array length

23 iconst_1         // Push constant int 1

24 isub             // Pop two ints, subtract, push int result

25 iload 4          // Push int at local var 4 (i)

27 isub             // Pop two ints, subtract, push int result
                        // Pop int into local var 5:

28 istore 5         // int upperindex = intArray.length - 1 - i;

```

```

30 aload_0      // Push the object ref at loc var 0 (the this ref)
                // Pop object ref, push ref to instance variable
                // intArray
31 getfield #4 <Field int intArray[]
34 iload 5      // Push int at local var 5 (upperIndex)
36 iaload       // Pop index, arrayref, push int at arrayref[index]
                // Pop into local var 6:
37 istore 6     // int save = intArray[upperIndex];
39 aload_0      // Push the object ref at loc var 0 (the this ref)
                // Pop object ref, push ref to instance variable
                // intArray
40 getfield #4 <Field int intArray[]
43 iload 5      // Push int at local var 5 (upperIndex)
45 aload_0      // Push the object ref at loc var 0 (the this ref)
                // Pop object ref, push ref to instance variable
                // intArray
46 getfield #4 <Field int intArray[]
49 iload 4      // Push int at local var 4 (i)
51 iaload       // Pop index, arrayref, push int at arrayref[index]
                // Set arrayref[index] = value:
52 iastore      // intArray[upperIndex] = intArray[i];
53 aload_0      // Push the object ref at loc var 0 (the this ref)
                // Pop object ref, push ref to instance variable
                // intArray
54 getfield #4 <Field int intArray[]
57 iload 4      // Push int at local var 4 (i)
59 iload 6      // Push int at local var 6 (save)
                // Set arrayref[index] = value:
61 iastore      // intArray[i] = save;

```

```

// The body of the for loop is now done, this instruction does
// the incrementing of the loop variable i
62 iinc 4 1      // Increment by 1 int at local var 4: ++i;

// This is the for loop condition check:
65 iload 4       // Push int at local var 4 (i)
67 iload_3       // Push int at local var 3 (halfway)
                // Pop two ints, compare, jump if less than to
68 if_icmplt 18   // top of for loop body: for (; i < halfway;)
// The code of the synchronized block ends here

// The next two instructions unlock the object, making it available
// for other threads. The reference to the locked object was stored
// in local variable 1 above.
71 aload_1       // Push local var 1 (the this reference)
72 monitorexit   // Pop ref, unlock object
73 return        // return normally from method

// This is a catch clause for any exception thrown (and not caught
// from within the synchronized block. If an exception is thrown,
// the locked object is unlocked, making it available for other
// threads.
74 aload_1       // Push local var 1 (the this reference)
75 monitorexit   // Pop ref, unlock object
76 athrow        // rethrow the same exception

// The exception table shows the "catch all" clause covers the
// entire synchronized block, from just after the lock is acquired
// to just before the lock is released.
Exception table:

```

from to target type

4 71 74 any

end

Note that a catch clause ensures the locked object will be unlocked even if an exception is thrown from within the synchronized block. No matter how the synchronized block is exited, the object lock acquired when the thread entered the block will definitely be released.

Synchronized Methods

To synchronize an entire method, you just include the synchronized keyword as one of the method qualifiers, as in:

begin

```
// On CD-ROM in file threads/ex1/HeatSync.java
class HeatSync {
```

```
    private int[] intArray = new int[10];
```

```
    synchronized void reverseOrder() {
```

```
        int halfWay = intArray.length / 2;
```

```
        for (int i = 0; i < halfWay; ++i) {
```

```
            int upperIndex = intArray.length - 1 - i;
```

```
            int save = intArray[upperIndex];
```

```
            intArray[upperIndex] = intArray[i];
```

```
            intArray[i] = save;
```

```
        }
```

```
    }
```

```
    // ...
```

```
}
```

end

The Java Virtual Machine does not use any special opcodes to invoke or return from synchronized methods. When the virtual machine resolves the symbolic reference to a method, it determines whether the method is synchronized. If so, the virtual machine acquires a lock before invoking the method. For

an instance method, the virtual machine acquires the lock associated with the object upon which the method is being invoked. For a class method, it acquires the lock associated with the class to which the method belongs (it locks a `CLASS` object). After a synchronized method completes, whether it completes by returning or by throwing an exception, the virtual machine releases the lock.

Here are the bytecodes that `javac` generates for `HeatSync's reverseOrder()` method:

begin

```
0 aload_0          // Push the object ref at loc var 0 (the this ref)
                   // Pop object ref, push ref to instance variable
                   // intArray

1 getfield #4 <Field int intArray[]

4 arraylength      // Pop array ref, push int array length

5 iconst_2         // Push constant int 2

6 idiv             // Pop two ints, divide, push int result
                   // Pop int into local var 1:

7 istore_1         // int halfway = intArray.length/2;

// This is the start of the code for the for loop

8 iconst_0         // Push int constant 0

9 istore_2         // Pop into local var 2: int i = 0;

10 goto 54         // Jump to for loop condition check

// This is the start of the body of the for loop

13 aload_0         // Push the object ref at loc var 0 (the this ref)
                   // Pop object ref, push ref to instance variable
                   // intArray

14 getfield #4 <Field int intArray[]

17 arraylength     // Pop array ref, push int array length

18 iconst_1        // Push constant int 1

19 isub            // Pop two ints, subtract, push int result

20 iload_2         // Push int at local var 2 (i)

21 isub            // Pop two ints, subtract, push int result
```

```

                // Pop int into local var 3:
22 istore_3      // int upperindex = intArray.length - 1 - i;
23 aload_0      // Push the object ref at loc var 0 (the this ref)
                // Pop object ref, push ref to instance variable
                // intArray
24 getfield #4 <Field int intArray[]
27 iload_3      // Push int at local var 3 (upperIndex)
28 iaload       // Pop index, arrayref, push int at arrayref[index]
                // Pop into local var 4:
29 istore 4     // int save = intArray[upperIndex];
31 aload_0      // Push the object ref at loc var 0 (the this ref)
                // Pop object ref, push ref to instance variable
                // intArray
32 getfield #4 <Field int intArray[]
35 iload_3      // Push int at local var 3 (upperIndex)
36 aload_0      // Push the object ref at loc var 0 (the this ref)
                // Pop object ref, push ref to instance variable
                // intArray
37 getfield #4 <Field int intArray[]
40 iload_2      // Push int at local var 2 (i)
41 iaload       // Pop index, arrayref, push int at arrayref[index]
                // Pop value, index, and arrayref,
                // Set arrayref[index] = value:
42 iastore      // intArray[upperIndex] = intArray[i];
43 aload_0      // Push the object ref at loc var 0 (the this ref)
                // Pop object ref, push ref to instance variable
                // intArray
44 getfield #4 <Field int intArray[]
47 iload_2      // Push int at local var 2 (i)
48 iload 4      // Push int at local var 4 (save)

```

```

        // Pop value, index, and arrayref,
        // Set arrayref[index] = value:
50 iastore      // intArray[i] = save;

// The body of the for loop is now done, this instruction does
// the incrementing of the loop variable i
51 iinc 2 1      // Increment by 1 int at local var 2: ++i;

// This is the for loop condition check:
54 iload_2      // Push int at local var 2 (i)
55 iload_1      // Push int at local var 1 (halfway)
                // Pop two ints, compare, jump if less than to
56 if_icmplt 13  // top of for loop body: for (; i < halfway;)

59 return      // return (void) from method

end

```

If you compare these bytecodes with the ones shown earlier for `KitchenSync`'s `reverseOrder()` method, you will see that these bytecodes are in effect those of `KitchenSync` with the support for entering and exiting the monitor removed. Instructions at offset 0 through 56 of `HeatSync`'s bytecodes correspond to instructions at offset 4 through 68 of `KitchenSync`'s bytecodes. Because `HeatSync`'s `reverseOrder()` method doesn't need a local variable slot to store the reference to the locked object, the local variable positions used by each method are different. The function of the instructions themselves, however, match up exactly.

Another difference between the two `reverseOrder()` methods is that the compiler doesn't create an exception table for `HeatSync`'s `reverseOrder()` method. In `HeatSync`'s case, an exception table isn't necessary. When this method is invoked, the Java Virtual Machine automatically acquires the lock on the `this` object. If this method completes abruptly, just as if it completes normally, the virtual machine will release the lock on the `this` object automatically.

A synchronized class (static) method operates in the same way as the synchronized instance method shown in the example above. The one difference is that instead of acquiring a lock on `this` (as there is no `this` in a class method), the thread must acquire a lock on the appropriate `Class` instance.

Coordination Support in Class Object

Class `Object` declares five methods that enable programmers to access the Java Virtual Machine's

support for the coordination aspect of synchronization. These methods are declared public and final, so they are inherited by all classes. They can only be invoked from within a synchronized method or statement. In other words, the lock associated with an object must already be acquired before any of these methods are invoked. The methods are listed in Table 20-2.

Table 20-2. The wait and notify methods of class `Object`

Method	Description
<code>void wait();</code>	Enter a monitor's wait set until notified by another thread
<code>void wait(long timeout);</code>	Enter a monitor's wait set until notified by another thread or <code>timeout</code> milliseconds elapses
<code>void wait(long timeout, int nanos);</code>	Enter a monitor's wait set until notified by another thread or <code>timeout</code> milliseconds plus <code>nanos</code> nanoseconds elapses
<code>void notify();</code>	Wake up one thread waiting in the monitor's wait set. (If no threads are waiting, do nothing.)
<code>void notifyAll();</code>	Wake up all threads waiting in the monitor's wait set. (If no threads are waiting, do nothing.)

On the CD-ROM

The CD-ROM contains the source code examples from this chapter in a subdirectory of the `threads` directory.

The Resources Page

For more information about the material presented in this chapter, visit the resources

page: <http://www.artima.com/insidejvm/threads.html>

Appendix A

Instruction Set by Opcode Mnemonic

Appendix A lists Java Virtual Machine instructions in alphabetical order by opcode mnemonic. All 201 instructions that may legally appear in the bytecode streams stored in Java class files are described in detail in Appendix A.

Besides the opcodes for the 201 instructions that may appear in class files, the Java Virtual Machine specification describes two other families of opcodes: the reserved opcodes and the "`_quick`" opcodes. None of these opcodes can legally appear in the bytecode streams of Java class files.

The Java Virtual Machine specification lists three reserved opcodes, which are shown in Table A-1. These opcodes are reserved for internal use by Java Virtual Machine implementations and tools. The specification guarantees that these three opcodes will not be part of any future extension of the Java Virtual Machine's instruction set. As you may have guessed, the intended purpose of the `breakpoint` opcode is to provide a way for debuggers to implement breakpoints. The intended purpose of the other two reserved opcodes, `impdep1` and `impdep2`, is to serve as "back doors" to implementation-dependent software functionality or "traps" to implementation-dependent hardware functionality.

Table A-1. The reserved opcodes

mnemonic	byte value
<code>breakpoint</code>	202 (0xca)
<code>impdep1</code>	254 (0xfe)
<code>impdep2</code>	255 (0xff)

The Java Virtual Machine specification also lists 25 "`_quick`" opcodes, which Sun's virtual machine implementation uses internally to speed up the interpreted execution of bytecodes. This optimization technique is described in general in this book in Chapter 8, "The Linking Model," but the individual "`_quick`" instructions are not described in detail in Appendix A. (The "`_quick`" opcodes do appear in Appendix C, which lists their mnemonics and corresponding opcode byte values.) Like the reserved opcodes, the "`_quick`" opcodes may not legally appear in Java class files. But unlike the reserved opcodes, the Java Virtual Machine specification leaves open the possibility that these opcodes will take on new meanings in future extensions of the instruction set.

For each instruction listed in Appendix A, you will find the following information:

- the mnemonic
- a short description
- the opcode's byte value in decimal and hex
- the format of the instruction
- the operand stack before and after the instruction is executed
- a description of any constraints associated with the instruction
- a description of the execution of the instruction
- a description of any exceptions or errors that may be thrown by the instruction

The format of each instruction appears as a list of comma-separated items next to the label "Instruction Format," with each item representing one byte in the bytecode stream. The opcode mnemonic appears first (in fixed-width font), followed by any operands (shown in *italics*).

For each instruction, Appendix A shows two snapshots of the operand stack. One snapshot, labelled "Before," shows the contents of the current method's operand stack just before the instruction is executed. The other snapshot, labelled "After," shows the operand stack immediately after execution of the instruction. In both snapshots, the operand stack appears as a list of comma-separated items, with each item representing one word. Although everywhere else in this book the operand stack is shown growing downwards (the top of the stack appears at the bottom of the picture), in Appendix A the operand stack is shown growing sideways, from left to right. In each snapshot, the top of the stack is the rightmost item shown. Unaffected portions of the operand stack are shown as an ellipsis, "...".

For each instruction, the section labelled "Description" combines three kinds of information: constraints, the process of execution, and exceptions and errors. As mentioned in Chapter 3, "Security," Java Virtual Machine implementations are required to enforce at run-time certain constraints on the bytecodes of every method they execute. Whenever you see the word "must" in an instruction's description in Appendix A, you are reading about a constraint every implementation must enforce when executing the instruction. For example, the `goto` instruction, which causes an unconditional jump, may only cause a jump to another opcode of the same method. In Appendix A's description for `goto`, this constraint is stated as, "The target address *must* be the address of an opcode within the same method as the `goto` opcode."

The designers of each individual Java Virtual Machine implementation can decide how and when to detect violations of the bytecode constraints. If any implementation detects a constraint violation, it must report the violation by throwing a `VerifyError` when (and if) the running program attempts to execute the instruction.

In addition to describing the constraints placed on each instruction, Appendix A describes the process of executing each instruction and lists any errors or exceptions that may be thrown during the course of executing the instruction. Besides the errors and exceptions explicitly listed in the instruction descriptions, one other family of errors, subclasses of `VirtualMachineError`, can be thrown at any time as the result of executing any instruction. Four subclasses of `VirtualMachineError`, and the circumstances under which they will be thrown, are:

- `OutOfMemoryError` - the virtual machine has run out of real or virtual memory, and the garbage collector can't reclaim enough space to enable the thread to continue.
- `StackOverflowError` - a thread has exhausted its supply of memory for stack space (usually because the application has an unbounded recursion).
- `InternalError` - the virtual machine has encountered a bug in its own implementation that prevents it from properly implementing the semantics of the Java language.

- **UnknownError** - the virtual machine has encountered some error condition, but is unable to report the actual condition by throwing the appropriate exception or error.

aaload

- Load reference from array

Opcode:

50 (0x32)

Instruction Format:

aaload

Stack:

Before:..., arrayref, index

After:..., value

Description:

To execute the **aaload** instruction, the Java Virtual Machine first pops two words from the operand stack. The arrayref word must be a **reference** that refers to an array of **references**. The index word must be an **int**. The virtual machine retrieves from the arrayref array the **reference** value specified by index and pushes it onto the operand stack.

If arrayref is **null**, the Java Virtual Machine throws **NullPointerException**. Otherwise, if index is not a legal index into the arrayref array, the virtual machine throws **ArrayIndexOutOfBoundsException**.

For more information about the **aaload** instruction, see Chapter 15, "Objects and Arrays."

aastore

- Store **reference** into array

Opcode:

83 (0x53)

Instruction Format:

aastore

Stack:

Before: ..., arrayref, index, value

After: ...

Description:

To execute the **aaload** instruction, the Java Virtual Machine first pops three words from the operand stack. The arrayref word must be a **reference** that refers to an array of **floats**. The index word must be an **int**, and the value word must be a **reference**. The type of value must be assignment compatible with the component type of the arrayref array. The virtual machine stores **reference** value into the arrayref array location specified by index.

If arrayref is **null**, the Java Virtual Machine throws **NullPointerException**. Else, if index is not a legal index into the arrayref array, the virtual machine throws **ArrayIndexOutOfBoundsException**. Otherwise, if the actual type of value is not assignment compatible with the actual type of the components of the arrayref array, the virtual machine throws **ArrayStoreException**.

For more information about the **aastore** instruction, see Chapter 15, "Objects and Arrays."

aconst_null

- Push **null** object reference

Opcode:

1 (0x1)

Instruction Format:

aconst_null

Stack:

Before: ...

After: ..., null

Description:

To execute the **aconst_null** instruction, the Java Virtual Machine pushes a **null** object reference onto the operand stack. (Note that the Java Virtual Machine specification does not dictate any actual value for **null**. The specification leaves that decision to the designers of each individual implementation.)

For more information about the **aconst_null** instruction, see Chapter 10, "Stack Operations."

aload

- Load *reference* from local variable

Opcode:

25 (0x19)

Instruction Format:

a load, index

Stack:

Before:...

After:..., value

Description:

The index operand, which serves as an 8-bit unsigned index into the local variables of the current frame, must specify a local variable word that contains a *reference*. To execute the *a load* instruction, the Java Virtual Machine pushes onto the operand stack the *reference* contained in the local variable word specified by index.

Note that the *wide* instruction can precede the *a load* instruction, to allow a local variable to be accessed with a 16-bit unsigned offset.

Note also that even though the *astore* instruction may be used to pop a *returnAddress* value off the operand stack and into a local variable, the *a load* instruction cannot be used to push a *returnAddress* value back onto the operand stack. For more information about the use of *returnAddress*, see Chapter 18, "Finally Clauses."

For more information about the *a load* instruction, see Chapter 10, "Stack Operations."

a load_0

- Load *reference* from local variable 0

Opcode:

42 (0x2a)

Instruction Format:

a load_0

Stack:

Before:...

After:..., value

Description:

The local variable word at index zero must contain a *reference*. To execute the **aload_0** instruction, the Java Virtual Machine pushes onto the operand stack the *reference* value contained in the local variable word zero.

Note that even though the **astore_0** instruction may be used to pop a *returnAddress* value off the operand stack and into a local variable, the **aload_0** instruction cannot be used to push a *returnAddress* value back onto the operand stack. For more information about the use of *returnAddress*, see Chapter 18, "Finally Clauses."

For more information about the **aload_0** instruction, see Chapter 10, "Stack Operations."

aload_1

- Load *reference* from local variable 1

Opcode:

43 (0x2b)

Instruction Format:

aload_1

Stack:

Before:...

After:..., value

Description:

The local variable word at index one must contain a *reference*. To execute the **aload_1** instruction, the Java Virtual Machine pushes onto the operand stack the *reference* value contained in the local variable word one.

Note that even though the **astore_1** instruction may be used to pop a *returnAddress* value off the operand stack and into a local variable, the **aload_1** instruction cannot be used to push a *returnAddress* value back onto the operand stack. For more information about the use of *returnAddress*, see Chapter 18, "Finally Clauses."

For more information about the **aload_1** instruction, see Chapter 10, "Stack Operations."

aload_2

- Load *reference* from local variable 2

Opcode:

44 (0x2c)

Instruction Format:

aload_2

Stack:

Before:...

After:..., value

Description:

The local variable word at index two must contain a *reference*. To execute the **a**load_2 instruction, the Java Virtual Machine pushes onto the operand stack the *reference* value contained in the local variable word two.

Note that even though the **a**store_2 instruction may be used to pop a *returnAddress* value off the operand stack and into a local variable, the **a**load_2 instruction cannot be used to push a *returnAddress* value back onto the operand stack. For more information about the use of *returnAddress*, see Chapter 18, "Finally Clauses."

For more information about the **a**load_2 instruction, see Chapter 10, "Stack Operations."

aload_3

- Load *reference* from local variable 3

Opcode:

45 (0x2d)

Instruction Format:

aload_3

Stack:

Before:...

After:..., value

Description:

The local variable word at index three must contain a **reference**. To execute the **aload_3** instruction, the Java Virtual Machine pushes onto the operand stack the **reference** value contained in the local variable word three.

Note that even though the **astore_3** instruction may be used to pop a **returnAddress** value off the operand stack and into a local variable, the **aload_3** instruction cannot be used to push a **returnAddress** value back onto the operand stack. For more information about the use of **returnAddress**, see Chapter 18, "Finally Clauses."

For more information about the **aload_3** instruction, see Chapter 10, "Stack Operations."

anewarray

- Allocate new array of reference type components

Opcode:

189 (0xbd)

Instruction Format:

anewarray, indexbyte1, indexbyte2

Stack:

Before: ..., count

After: arrayref

Description:

The top word of the operand stack, count, must be an **int**. To execute the **anewarray** instruction, the Java Virtual Machine first forms an unsigned 16-bit index into the constant pool by calculating $(\text{indexbyte1} \ll 8) \mid \text{indexbyte2}$. The virtual machine then looks up the constant pool entry specified by the calculated index. The constant pool entry at that index must be a **CONSTANT_Class_info** entry. If it hasn't already, the virtual machine resolves the entry. The entry may be a class, interface, or array type.

If the resolution is successful, the Java Virtual Machine pops count and creates on the heap an array of size count of the reference type specified by the resolved **CONSTANT_Class_info** entry. The virtual machine initializes each array element to its default initial value (null) and pushes arrayref, a reference to the new array, onto the operand stack.

As a result of executing this instruction, the Java Virtual Machine may throw any of the linking errors listed in Chapter 8, "The Linking Model," as possible during resolution of a `CONSTANT_Class_info` entry. If resolution succeeds, but count is less than zero, the virtual machine throws `NegativeArraySizeException`.

For more information about the `anewarray` instruction, see Chapter 15, "Objects and Arrays."

`areturn`

- Return reference from method

Opcode:

176 (0xb0)

Instruction Format:

`areturn`

Stack:

Before:..., `objectref`

After:[empty]

Description:

The return type of the returning method must be reference. The top word of the operand stack, `objectref`, must be a reference that is assignment compatible with the type represented by the returning method's descriptor. To execute the `areturn` instruction, the Java Virtual Machine pops `objectref` from the operand stack of the current frame and pushes it onto the operand stack of the invoking method's frame. The virtual machine discards any other words that may still be on the returning method's frame. If the returning method is synchronized, the monitor that was acquired when the method was invoked is released. The invoking method's frame is made current, and the virtual machine continues execution in the invoking method.

For more information about monitors, see Chapter 20, "Thread Synchronization." For more information about the `areturn` instruction, see Chapter 19, "Method Invocation and Return."

`arraylength`

- Get length of array

Opcode:

190 (0xbe)

Instruction Format:

arraylength

Stack:

Before:..., arrayref

After:..., length

Description:

The top word of the operand stack, arrayref, must be a reference that points to an array. To execute the arraylength instruction, the Java Virtual Machine pops arrayref and pushes the length of the array pointed to by arrayref.

If arrayref is null, the Java Virtual Machine throws NullPointerException.

For more information about the arraylength instruction, see Chapter 15, "Objects and Arrays."

astore

- Store reference or returnAddress into local variable

Opcode:

58 (0x3a)

Instruction Format:

astore, index

Stack:

Before:..., value

After:...

Description:

The index operand must specify a valid 8-bit unsigned index into the

local variables of the current frame. The value on the top of the operand stack must be a reference or a returnAddress. To execute the astore instruction, the Java Virtual Machine pops the reference or returnAddress value from the top of the operand stack and stores it into the local variable word specified by index.

Note that the wide instruction can precede the astore instruction, to enable a value to be stored into a local variable specified by a 16-bit unsigned offset.

For more information about the astore instruction, see Chapter 10, "Stack Operations."

astore_0

- Store reference or returnAddress into local variable 0

Opcode:

75 (0x4b)

Instruction Format:

astore_0

Stack:

Before:..., value

After:...

Description:

The index zero must be a valid index into the local variables of the current stack frame, and the value word on the top of the operand stack must be a reference or a returnAddress. To execute the astore_0 instruction, the Java Virtual Machine pops the reference or returnAddress value from the top of the operand stack and stores it into the local variable word at index zero.

For more information about the astore_0 instruction, see Chapter 10, "Stack Operations."

astore_1

- Store reference or returnAddress into local variable 1

Opcode:

76 (0x4c)

Instruction Format:

astore_1

Stack:

Before:..., value

After:...

Description:

*The index one must be a valid index into the local variables of the current stack frame, and the value word on the top of the operand stack must be a reference or a returnAddress. To execute the *astore_1* instruction, the Java Virtual Machine pops the reference or returnAddress value from the top of the operand stack and stores it into the local variable word at index one.*

*For more information about the *astore_1* instruction, see Chapter 10, "Stack Operations."*

astore_2

- Store reference or returnAddress into local variable 2

Opcode:

77 (0x4d)

Instruction Format:

astore_2

Stack:

Before:..., value

After:...

Description:

*The index two must be a valid index into the local variables of the current stack frame, and the value word on the top of the operand stack must be a reference or a returnAddress. To execute the *astore_2* instruction, the Java Virtual Machine pops the reference or*

returnAddress value from the top of the operand stack and stores it into the local variable word at index two.

For more information about the *astore_2* instruction, see Chapter 10, "Stack Operations."

astore_3

- Store reference or *returnAddress* into local variable 3

Opcode:

78 (0x4e)

Instruction Format:

astore_3

Stack:

Before:..., value

After:...

Description:

The index three must be a valid index into the local variables of the current stack frame, and the value word on the top of the operand stack must be a reference or a *returnAddress*. To execute the *astore_3* instruction, the Java Virtual Machine pops the reference or *returnAddress* value from the top of the operand stack and stores it into the local variable word at index three.

For more information about the *astore_3* instruction, see Chapter 10, "Stack Operations."

athrow

- Throw exception or error

Opcode:

191 (0xbf)

Instruction Format:

athrow

Stack:

Before:..., objectref

Afterobjectref

Note that "Before" shows the operand stack of the frame belonging to the method containing the *athrow* instruction being executed. "After" shows the operand stack of the frame belonging to the method in which the catch clause is found, if a catch clause is found. If no catch clause is found, the thread exits and there are no more operand stacks for that thread.

Description:

The top word of the operand stack, *objectref*, must be a reference that points either to an instance of class *java.lang.Throwable* or to an instance of some subclass of *java.lang.Throwable*. To execute the *athrow* instruction, the Java Virtual Machine pops *objectref* from the operand stack. The virtual machine "throws" the exception by searching through the current method's exception table for the most recent catch clause that catches either the class of the throwable object pointed to by *objectref*, or a subclass of the throwable object's class. If the current method's exception table contains a matching entry, the virtual machine extracts the address of the handler to jump to from the matching exception table entry. The virtual machine pops any words remaining on the operand stack, pushes the *objectref*, sets the program counter to the handler address, and continues execution there. If the current method's exception table doesn't have a matching catch clause, the virtual machine pops the current method's entire frame and rethrows the exception in the previous method. This process repeats until either a matching catch clause is found or the stack frames for all the methods along the current thread's call stack have been popped. If no catch clause is found by this process, the current thread exits.

If the *objectref* word is null, the virtual machine throws *NullPointerException*.

For more information about the *athrow* instruction, see Chapter 17, "Exceptions."

baload

- Load byte or boolean from array

Opcode:

51 (0x33)

Instruction Format:

baload

Stack:

Before:..., arrayref, index

After:..., value

Description:

To execute the `baload` instruction, the Java Virtual Machine first pops two words from the operand stack. The `arrayref` word must be a reference that refers to an array of bytes or booleans. The `index` word must be an `int`. The virtual machine retrieves from the `arrayref` array the byte or boolean value specified by `index`, sign-extends it to an `int`, and pushes it onto the operand stack.

If `arrayref` is null, the Java Virtual Machine throws `NullPointerException`. Otherwise, if `index` is not a legal index into the `arrayref` array, the virtual machine throws `ArrayIndexOutOfBoundsException`.

For more information about the `baload` instruction, see Chapter 15, "Objects and Arrays."

`bastore`

- Store into byte or boolean array

Opcode:

84 (0x54)

Instruction Format:

bastore

Stack:

Before:..., arrayref, index, value

After:...

Description:

To execute the `bastore` instruction, the Java Virtual Machine first

pops three words from the operand stack. The arrayref word must be a reference that refers to an array of bytes or booleans. The index and value words must be ints. The virtual machine truncates the int value to a byte and stores it into the arrayref array location specified by index.

If arrayref is null, the Java Virtual Machine throws NullPointerException. Otherwise, if index is not a legal index into the arrayref array, the virtual machine throws ArrayIndexOutOfBoundsException.

For more information about the bastore instruction, see Chapter 15, "Objects and Arrays."

bipush

- Push 8-bit signed integer

Opcode:

16 (0x10)

Instruction Format:

bipush, byte

Stack:

Before:...

After:..., value

Description:

To execute the bipush instruction, the Java Virtual Machine first sign-extends operand byte, an 8-bit signed integer, to an int. The virtual machine then pushes the resulting int value onto the operand stack.

For more information about the bipush instruction, see Chapter 10, "Stack Operations."

caload

- Load char from array

Opcode:

52 (0x34)

Instruction Format:

caload

Stack:

Before:..., arrayref, index

After:..., value

Description:

To execute the caload instruction, the Java Virtual Machine first pops two words from the operand stack. The arrayref word must be a reference that refers to an array of chars. The index word must be an int. The virtual machine retrieves from the arrayref array the char value specified by index, zero-extends it to an int, and pushes it onto the operand stack.

If arrayref is null, the Java Virtual Machine throws NullPointerException. Otherwise, if index is not a legal index into the arrayref array, the virtual machine throws ArrayIndexOutOfBoundsException.

For more information about the caload instruction, see Chapter 15, "Objects and Arrays."

castore

- Store into char array

Opcode:

85 (0x55)

Instruction Format:

castore

Stack:

Before:..., arrayref, index, value

After:...

Description:

To execute the `castore` instruction, the Java Virtual Machine first pops three words from the operand stack. The `arrayref` word must be a reference that refers to an array of `chars`. The `index` and `value` words must be `ints`. The virtual machine truncates the `int` value to a `char` and stores it into the `arrayref` array location specified by `index`.

If `arrayref` is `null`, the Java Virtual Machine throws `NullPointerException`. Otherwise, if `index` is not a legal index into the `arrayref` array, the virtual machine throws `ArrayIndexOutOfBoundsException`.

For more information about the `castore` instruction, see Chapter 15, "Objects and Arrays."

checkcast

- Make sure object is of given type

Opcode:

192 (0xc0)

Instruction Format:

`checkcast, indexbyte1, indexbyte2`

Stack:

Before:..., `objectref`

After:..., `objectref`

Description:

The top word of the stack, `objectref`, must be a reference. To execute the `checkcast` instruction, the Java Virtual Machine first forms an unsigned 16-bit index into the constant pool by calculating $(\text{indexbyte1} \ll 8) \mid \text{indexbyte2}$. The virtual machine then looks up the constant pool entry specified by the calculated index. The constant pool entry at that index must be a `CONSTANT_Class_info` entry. If it hasn't already, the virtual machine resolves the entry. The entry may be a class, interface, or array type. If `objectref` is `null` or if `objectref` can be cast to the resolved type, the stack remains unchanged. Otherwise, the virtual machine throws `ClassCastException`.

To determine whether the object pointed to by `objectref` can be cast to the resolved type, the virtual machine first determines whether the object is a class instance or array. (It can't be an interface

instance, because interfaces can't be instantiated.) If it is a class instance, and the resolved type is a class, not an interface, the object can be cast to the resolved class if the object's class is the resolved class or a subclass of the resolved class. Else, if it is a class instance, and the resolved type is an interface, not a class, the object can be cast to the resolved interface if the object's class implements the resolved interface. Otherwise, the object is an array. If the resolved type is a class, it must be `java.lang.Object`. Else, if the resolved type is an array of primitive types, the object must be an array of the same primitive type. Otherwise, the resolved type must be an array with a component type of some reference type, and the object must be an array with a component type that can be cast to the component type of the resolved array type. (Note that the dimension of an array doesn't enter into the checkcast check, only the component type of the array.)

As a result of executing this instruction, the virtual machine may throw any of the linking errors listed in Chapter 8, "The Linking Model," as possible during resolution of a `CONSTANT_Class_info` entry. If resolution succeeds, but the resolved type cannot be cast to the resolved type, the virtual machine throws a `ClassCastException`.

For more information about the checkcast instruction, see Chapter 15, "Objects and Arrays."

d2f

- Convert double to float

Opcode:

144 (0x90)

Instruction Format:

d2f

Stack:

Before:..., value.word1, value.word2

After:..., result

Description:

The top two words of the operand stack must be a double. To execute the d2f instruction, the Java Virtual Machine pops the double value from the operand stack, converts the double to a float, and pushes the float result.

To convert the double value to float, the Java Virtual Machine first checks to see if the value equals NaN (Not a Number). If so, the float result is also NaN. Else, if the magnitude of the double value is too small to be represented by a float, the float result is a zero of the same sign. Else, if the magnitude of the double value is too large to be represented by a float, the float result is an infinity of the same sign. Otherwise, the virtual machine converts the double value to float zero using IEEE 754 round-to-nearest mode.

Note that this instruction performs a narrowing primitive conversion. Because not all double values are representable by a float, the conversion may result in a loss of magnitude and precision.

For more information about the d2f instruction, see Chapter 11, "Type Conversion."

d2i

- Convert double to int

Opcode:

142 (0x8e)

Instruction Format:

d2i

Stack:

Before:..., value.word1, value.word2

After:..., result

Description:

The top two words of the operand stack must be a double. To execute the d2i instruction, the Java Virtual Machine pops the double value from the operand stack, converts the double to an int, and pushes the int result.

To convert the double value to int, the Java Virtual Machine first checks to see if the value equals NaN (Not a Number). If so, the int result is zero. Else, if the double value is not a positive or negative infinity, the virtual machine rounds the value towards zero using IEEE 754 round-towards-zero mode. If the resulting integral value can be exactly represented by an int, the int result is that integral value. Otherwise, the magnitude of the double value is too

great be represented in an int. If value is positive, the int result is the largest positive integer representable in an int. If value is negative, the int result is the smallest negative integer representable in an int.

Note that this instruction performs a narrowing primitive conversion. Because not all double values are representable by an int, the conversion may result in a loss of magnitude and precision.

For more information about the d2i instruction, see Chapter 11, "Type Conversion."

d2l

- Convert double to long

Opcode:

143 (0x8f)

Instruction Format:

d2l

Stack:

Before:..., value.word1, value.word2

After:..., result.word1, result.word2

Description:

The top two words of the operand stack must be a double. To execute the d2l instruction, the Java Virtual Machine pops the double value from the operand stack, converts the double to a long, and pushes the long result.

To convert the double value to long, the Java Virtual Machine first checks to see if the value equals NaN (Not a Number). If so, the long result is zero. Else, if the double value is not a positive or negative infinity, the virtual machine rounds the value towards zero using IEEE 754 round-towards-zero mode. If the resulting integral value can be exactly represented by a long, the long result is that integral value. Otherwise, the magnitude of the double value is too great be represented in a long. If value is positive, the long result is the largest positive integer representable in a long. If value is negative, the long result is the smallest negative integer representable in an long.

Note that this instruction performs a narrowing primitive conversion. Because not all double values are representable by a long, the conversion may result in a loss of magnitude and precision.

For more information about the d2l instruction, see Chapter 11, "Type Conversion."

dadd

- Add doubles

Opcode:

99 (0x63)

Instruction Format:

dadd

Stack:

Before:..., value1.word1, value1.word2, value2.word1, value2.word2

After:..., result.word1, result.word2

Description:

The top four words of the operand stack must be two doubles, value1 and value2. To execute the dadd instruction, the Java Virtual Machine pops value1 and value2, adds them, and pushes the double result. The result produced by the dadd instruction is governed by the rules of IEEE 754 floating point arithmetic.

For more information about the dadd instruction, see Chapter 14, "Floating Point Arithmetic."

daload

- Load double from array

Opcode:

49 (0x31)

Instruction Format:

daload

Stack:

Before:..., arrayref, index

After:..., value.word1, value.word2

Description:

To execute the daload instruction, the Java Virtual Machine first pops two words from the operand stack. The arrayref word must be a reference that refers to an array of doubles. The index word must be an int. The virtual machine retrieves from the arrayref array the double value specified by index and pushes it onto the operand stack.

If arrayref is null, the Java Virtual Machine throws NullPointerException. Otherwise, if index is not a legal index into the arrayref array, the virtual machine throws ArrayIndexOutOfBoundsException.

For more information about the daload instruction, see Chapter 15, "Objects and Arrays."

dastore

- Store into double array

Opcode:

82 (0x52)

Instruction Format:

dastore

Stack:

Before:..., arrayref, index, value.word1, value.word2

After:...

Description:

To execute the dastore instruction, the Java Virtual Machine first pops four words from the operand stack. The arrayref word must be a reference that refers to an array of doubles. The index word must be an int, and the value words must be a double. The virtual machine stores double value into the arrayref array location specified by index.

If *arrayref* is null, the Java Virtual Machine throws *NullPointerException*. Otherwise, if *index* is not a legal index into the *arrayref* array, the virtual machine throws *ArrayIndexOutOfBoundsException*.

For more information about the *dastore* instruction, see Chapter 15, "Objects and Arrays."

dcmpg

- Compare doubles (1 on NaN)

Opcode:

152 (0x98)

Instruction Format:

dcmpg

Stack:

Before:..., value1.word1, value1.word2, value2.word1, value2.word2

After:..., result

Description:

The top four words of the operand stack must be two doubles, *value1* and *value2*. To execute the *dcmpg* instruction, the Java Virtual Machine pops *value1* and *value2* off the operand stack and compares one against the other. If *value1* equals *value2*, the virtual machine pushes onto the operand stack int result zero. Else, if *value1* is greater than *value2*, the virtual machine pushes onto the operand stack int result one. Otherwise, if *value1* is less than *value2*, the virtual machine pushes onto the operand stack int result negative one. If either *value1* or *value2* equals NaN (Not a Number), the virtual machine pushes onto the operand stack int result one.

The result produced by the *fcmpg* instruction is governed by the rules of IEEE 754 floating point arithmetic. Note that the *dcmpg* instruction differs from the *dcmpl* instruction only in its treatment of NaN. For more information about the *dcmpg* instruction, see Chapter 16, "Control Flow."

dcmpl

- Compare doubles (-1 on NaN)

Opcode:

151 (0x97)

Instruction Format:

dcmpl

Stack:

Before:..., value1.word1, value1.word2, value2.word1, value2.word2

After:..., result

Description:

The top four words of the operand stack must be two doubles, value1 and value2. To execute the *dcmpg* instruction, the Java Virtual Machine pops value1 and value2 off the operand stack and compares one against the other. If value1 equals value2, the virtual machine pushes onto the operand stack int result zero. Else, if value1 is greater than value2, the virtual machine pushes onto the operand stack int result one. Otherwise, if value1 is less than value2, the virtual machine pushes onto the operand stack int result negative one. If either value1 or value2 equals NaN (Not a Number), the virtual machine pushes onto the operand stack int result negative one.

The result produced by the *fcmpl* instruction is governed by the rules of IEEE 754 floating point arithmetic. Note that the *dcmpl* instruction differs from the *dcmpg* instruction only in its treatment of NaN. For more information about the *dcmpl* instruction, see Chapter 16, "Control Flow."

dconst_0

- Push double constant 0.0

Opcode:

14 (0xe)

Instruction Format:

dconst_0

Stack:

Before:...

After:..., <0.0-word1, <0.0-word2

Description:

To execute the *dconst_0* instruction, the Java Virtual Machine pushes the **double** constant 0.0 onto the operand stack.

For more information about the *dconst_0* instruction, see Chapter 10, "Stack Operations."

dconst_1

- Push **double** constant 1.0

Opcode:

15 (0xf)

Instruction Format:

dconst_1

Stack:

Before:...

After:..., <1.0-word1, <1.0-word2

Description:

To execute the *dconst_1* instruction, the Java Virtual Machine pushes the **double** constant 1.0 onto the operand stack.

For more information about the *dconst_1* instruction, see Chapter 10, "Stack Operations."

ddiv

- Divide **doubles**

Opcode:

111 (0x6f)

Instruction Format:

ddiv

Stack:

Before: ..., value1.word1, value1.word2, value2.word1, value2.word2

After: ..., result.word1, result.word2

Description:

The top four words of the operand stack must be two **double**s, value1 and value2. To execute the **ddiv** instruction, the Java Virtual Machine pops value1 and value2, divides value1 by value2 (value1 / value2), and pushes the **double** result. The result produced by the **ddiv** instruction is governed by the rules of IEEE 754 floating point arithmetic.

For more information about the **ddiv** instruction, see Chapter 14, "Floating Point Arithmetic."

dload

- Load **double** from local variable

Opcode:

24 (0x18)

Instruction Format:

dload, index

Stack:

Before: ...

After: ..., value.word1, value.word2

Description:

The index operand, which serves as an 8-bit unsigned index into the local variables of the current frame, must specify the first of two consecutive local variable words that contain a **double**. To execute the **dload** instruction, the Java Virtual Machine pushes onto the operand stack the **double** contained in the two consecutive local variable words specified by index and index + 1.

Note that the **wide** instruction can precede the **lload** instruction, to allow a local variable to be accessed with a 16-bit unsigned offset.

For more information about the **dload** instruction, see Chapter 10, "Stack Operations."

dload_0

- Load **double** from local variable 0

Opcode:

38 (0x26)

Instruction Format:

dload_0

Stack:

Before:...

After:..., value.word1, value.word2

Description:

*The two consecutive local variable words at indexes zero and one must contain a **double**. To execute the **dload_0** instruction, the Java Virtual Machine pushes onto the operand stack the **double** value contained in local variable words zero and one.*

*For more information about the **dload_0** instruction, see Chapter 10, "Stack Operations."*

dload_1

*- Load **double** from local variable 1*

Opcode:

39 (0x27)

Instruction Format:

dload_1

Stack:

Before:...

After:..., value.word1, value.word2

Description:

*The two consecutive local variable words at indexes one and two must contain a **double**. To execute the **dload_1** instruction, the Java Virtual Machine pushes onto the operand stack the **double** value contained in local variable words one and two.*

*For more information about the **dload_1** instruction, see Chapter 10, "Stack Operations."*

dload_2

- Load *double* from local variable 2

Opcode:

40 (0x28)

Instruction Format:

dload_2

Stack:

Before:...

After:..., value.word1, value.word2

Description:

The two consecutive local variable words at indexes two and three must contain a *double*. To execute the *dload_2* instruction, the Java Virtual Machine pushes onto the operand stack the *double* value contained in local variable words two and three.

For more information about the *dload_2* instruction, see Chapter 10, "Stack Operations."

dload_3

- Load *double* from local variable 3

Opcode:

41 (0x29)

Instruction Format:

dload_3

Stack:

Before:...

After:..., value.word1, value.word2

Description:

The two consecutive local variable words at indexes three and four must contain a *double*. To execute

the **dload_3** instruction, the Java Virtual Machine pushes onto the operand stack the **double** value contained in local variable words three and four.

For more information about the **dload_3** instruction, see Chapter 10, "Stack Operations."

dmul

- Multiply **double**s

Opcode:

107 (0x6b)

Instruction Format:

dmul

Stack:

Before: ..., value1.word1, value1.word2, value2.word1, value2.word2

After: ..., result.word1, result.word2

Description:

The top four words of the operand stack must be two **double**s, value1 and value2. To execute the **dmul** instruction, the Java Virtual Machine pops value1 and value2, multiplies them, and pushes the **double** result. The result produced by the **dmul** instruction is governed by the rules of IEEE 754 floating point arithmetic.

For more information about the **dmul** instruction, see Chapter 14, "Floating Point Arithmetic."

dneg

- Negate **double**

Opcode:

119 (0x77)

Instruction Format:

dneg

Stack:

Before: ..., value.word1, value.word2

After:..., result.word1, result.word2

Description:

The top two words of the operand stack must be a **double**. To execute the **dneg** instruction, the Java Virtual Machine pops value, negates it, and pushes the **double** result. The result produced by the **dneg** instruction is governed by the rules of IEEE 754 floating point arithmetic.

Note that the result produced by a **dneg** instruction is not always the same as the number that would be produced by subtracting value from zero with the **dsub** instruction. The range of IEEE 754 floating point numbers includes two zeros, a positive zero and a negative zero. When value is +0.0, the result of the **dneg** instruction is -0.0. By contrast, when subtracting +0.0 from +0.0, the **dsub** instruction yields +0.0.

For more information about the **dneg** instruction, see Chapter 14, "Floating Point Arithmetic."

drem

- Calculate remainder of division of **doubles**

Opcode:

115 (0x73)

Instruction Format:

drem

Stack:

Before:..., value1.word1, value1.word2, value2.word1, value2.word2

After:..., result.word1, result.word2

Description:

The top four words of the operand stack must be two **doubles**, value1 and value2. To execute the **drem** instruction, the Java Virtual Machine pops value1 and value2, calculates the remainder, and pushes the **double** result. The remainder equals value1 - (value1 / value2) * value2, where value1 / value2 is a truncating division, rather than the rounding division required by IEEE 754.

The behavior of **drem** is comparable to that of the C library function **fmod()**. The remainder of two **doubles** can be calculated according to IEEE 754 floating point standard via the **IEEERemainder()** method of class **java.lang.Math**.

For more information about the **drem** instruction, see Chapter 14, "Floating Point Arithmetic."

dreturn

- Return *double* from method

Opcode:

175 (0xaf)

Instruction Format:

dreturn

Stack:

Before: ..., value.word1, value.word2

After: [empty]

Description:

The return type of the returning method must be *double*. The top two words of the operand stack must be a *double*. To execute the *dreturn* instruction, the Java Virtual Machine pops *double* value from the operand stack of the current frame and pushes it onto the operand stack of the invoking method's frame. The virtual machine discards any other words that may still be on the returning method's frame. If the returning method is synchronized, the monitor that was acquired when the method was invoked is released. The invoking method's frame is made current, and the virtual machine continues execution in the invoking method.

For more information about monitors, see Chapter 20, "Thread Synchronization." For more information about the *dreturn* instruction, see Chapter 19, "Method Invocation and Return."

dstore

- Store *double* into local variable

Opcode:

57 (0x39)

Instruction Format:

dstore, index

Stack:

Before: ..., value.word1, value.word2

After: ...

Description:

The index operand must specify a valid 8-bit unsigned index into the local variables of the current frame. The top two words of the operand stack must be a **double**. To execute the **dstore** instruction, the Java Virtual Machine pops the **double** value from the top of the operand stack and stores it into the two consecutive local variable words at indexes *index* and *index + 1*.

Note that the **wide** instruction can precede the **dstore** instruction, to enable a value to be stored into a local variable specified by a 16-bit unsigned offset.

For more information about the **dstore** instruction, see Chapter 10, "Stack Operations."

dstore_0

- Store **double** into local variable 0

Opcode:

71 (0x47)

Instruction Format:

dstore_0

Stack:

Before: ..., value.word1, value.word2

After: ...

Description:

The indexes zero and one must be valid indexes into the local variables of the current stack frame. The top two words on the operand stack must be a **double**. To execute the **dstore_0** instruction, the Java Virtual Machine pops the **double** value from the top of the operand stack and stores it into the two consecutive local variable words at indexes zero and one.

For more information about the **dstore_0** instruction, see Chapter 10, "Stack Operations."

dstore_1

- Store **double** into local variable 1

Opcode:

72 (0x48)

Instruction Format:

dstore_1

Stack:

Before: ..., value.word1, value.word2

After: ...

Description:

*The indexes one and two must be valid indexes into the local variables of the current stack frame. The top two words on the operand stack must be a **double**. To execute the **dstore_1** instruction, the Java Virtual Machine pops the **double** value from the top of the operand stack and stores it into the two consecutive local variable words at indexes one and two.*

*For more information about the **dstore_1** instruction, see Chapter 10, "Stack Operations."*

dstore_2

*- Store **double** into local variable 2*

Opcode:

73 (0x49)

Instruction Format:

dstore_2

Stack:

Before: ..., value.word1, value.word2

After: ...

Description:

*The indexes two and three must be valid indexes into the local variables of the current stack frame. The top two words on the operand stack must be a **double**. To execute the **dstore_2** instruction, the Java Virtual Machine pops the **double** value from the top of the operand stack and stores it into the two consecutive local variable words at indexes two and three.*

*For more information about the **dstore_2** instruction, see Chapter 10, "Stack Operations."*

dstore_3

*- Store **double** into local variable 3*

Opcode:

74 (0x4a)

Instruction Format:

dstore_3

Stack:

Before: ..., value.word1, value.word2

After: ...

Description:

The indexes three and four must be valid indexes into the local variables of the current stack frame. The top two words on the operand stack must be a **double**. To execute the *dstore_3* instruction, the Java Virtual Machine pops the **double** value from the top of the operand stack and stores it into the two consecutive local variable words at indexes three and four.

For more information about the *dstore_3* instruction, see Chapter 10, "Stack Operations."

dsub

- Subtract **doubles**

Opcode:

103 (0x67)

Instruction Format:

dsub

Stack:

Before: ..., value1.word1, value1.word2, value2.word1, value2.word2

After: ..., result.word1, result.word2

Description:

The top four words of the operand stack must be two **doubles**, value1 and value2. To execute the *dsub* instruction, the Java Virtual Machine pops value1 and value2, subtracts value2 from value1 (value1 - value2), and pushes the **double** result. The result produced by the *dsub* instruction is governed by the rules of IEEE 754 floating point arithmetic.

For more information about the **dsub** instruction, see Chapter 14, "Floating Point Arithmetic."

dup

- Duplicate top stack word

Opcode:

89 (0x59)

Instruction Format:

dup

Stack:

Before: ..., word

After: ..., word, word

Description:

To execute the **dup** instruction, the Java Virtual Machine duplicates the top word of the operand stack and pushes the duplicate. This instruction can be used to duplicate any single-word value from the top of the operand stack. It must not be used to duplicate half of a dual word value (**long** or **double**) that occupies the top of the operand stack.

For more information about the **dup** instruction, see Chapter 10, "Stack Operations."

dup_x1

- Duplicate top stack word and put two down

Opcode:

90 (0x5a)

Instruction Format:

dup_x1

Stack:

Before: ..., word2, word1

After: ..., word1, word2, word1

Description:

To execute the **dup_x1** instruction, the Java Virtual Machine duplicates the top word of the operand stack and inserts the duplicate two words down. Both word1 and word2 must be single-word values.

For more information about the **dup_x1** instruction, see Chapter 10, "Stack Operations."

dup_x2

- Duplicate top stack word and put three down

Opcode:

91 (0x5b)

Instruction Format:

dup_x2

Stack:

Before: ..., word3, word2, word1

After: ..., word1, word3, word2, word1

Description:

To execute the **dup_x2** instruction, the Java Virtual Machine duplicates the top word of the operand stack and inserts the duplicate three words down. word1 must be a single-word value. Both word2 and word3 must be single-word values, or together constitute one dual-word value (a **long** or **double**).

For more information about the **dup_x2** instruction, see Chapter 10, "Stack Operations."

dup2

- Duplicate top two stack words

Opcode:

92 (0x5c)

Instruction Format:

dup2

Stack:

Before: ..., word2, word1

After: ..., word2, word1, word2, word1

Description:

To execute the **dup2** instruction, the Java Virtual Machine duplicates the top two words of the operand stack and pushes the duplicate. This instruction may be used to duplicate any dual-word value or any two single-word values that occupy the top of the operand stack. It must not be used if one single-word value and half of a dual-word value (**long** or **double**) occupies the top of the operand stack.

For more information about the **dup2** instruction, see Chapter 10, "Stack Operations."

dup2_x1

- Duplicate top two stack words and put three down

Opcode:

93 (0x5d)

Instruction Format:

dup2_x1

Stack:

Before: ..., word3, word2, word1

After: ..., word2, word1, word3, word2, word1

Description:

To execute the **dup2_x1** instruction, the Java Virtual Machine duplicates the top two words of the operand stack and inserts the duplicate three words down. word3 must be a single-word value. Both word1 and word2 must be single-word values, or together constitute one dual-word value (a **long** or **double**).

For more information about the **dup2_x1** instruction, see Chapter 10, "Stack Operations."

dup2_x2

- Duplicate top two stack words and put four down

Opcode:

94 (0x5e)

Instruction Format:

dup2_x2

Stack:

Before: ..., word4, word3, word2, word1

After: ..., word2, word1, word4, word3, word2, word1

Description:

To execute the **dup2_x2** instruction, the Java Virtual Machine duplicates the top two words of the operand stack and inserts the duplicate four words down. word3 must be a single-word value. Both word1 and word2 must be single-word values, or together constitute one dual-word value (a **long** or **double**). Likewise, both word3 and word4 must be single-word values, or together constitute one dual-word value.

For more information about the **dup2_x2** instruction, see Chapter 10, "Stack Operations."

f2d

- Convert **float** to **double**

Opcode:

141 (0x8d)

Instruction Format:

f2d

Stack:

Before: ..., value

After: ..., result.word1, result.word2

Description:

The top word of the operand stack, value, must be a **float**. To execute the **f2d** instruction, the Java Virtual Machine pops **float** value from the operand stack, converts the **float** to a **double**, and pushes the **double** result.

Note that this instruction performs a widening primitive conversion. Because all **float** values are representable by a **double**, the conversion is exact.

For more information about the **f2d** instruction, see Chapter 11, "Type Conversion."

f2i

- Convert *float* to *int*

Opcode:

139 (0x8b)

Instruction Format:

f2i

Stack:

Before: ..., value

After: ..., result

Description:

The top word of the operand stack, value, must be a *float*. To execute the *f2i* instruction, the Java Virtual Machine pops the *float* value from the operand stack, converts the *float* to an *int*, and pushes the *int* result.

To convert the *float* value to *int*, the Java Virtual Machine first checks to see if the value equals NaN (Not a Number). If so, the *int* result is zero. Else, if the *float* value is not a positive or negative infinity, the virtual machine rounds the value towards zero using IEEE 754 round-towards-zero mode. If the resulting integral value can be exactly represented by an *int*, the *int* result is that integral value. Otherwise, the magnitude of the *float* value is too great to be represented in an *int*. If value is positive, the *int* result is the largest positive integer representable in an *int*. If value is negative, the *int* result is the smallest negative integer representable in an *int*.

Note that this instruction performs a narrowing primitive conversion. Because not all *float* values are representable by an *int*, the conversion may result in a loss of magnitude and precision.

For more information about the *f2i* instruction, see Chapter 11, "Type Conversion."

f2l

- Convert *float* to *long*

Opcode:

140 (0x8c)

Instruction Format:

f2l

Stack:

Before: ..., value

After: ..., result.word1, result.word2

Description:

The top word of the operand stack, value, must be a **float**. To execute the **f2l** instruction, the Java Virtual Machine pops the **float** value from the operand stack, converts the **float** to a **long**, and pushes the **long** result.

To convert the **float** value to **long**, the Java Virtual Machine first checks to see if the value equals NaN (Not a Number). If so, the **long** result is zero. Else, if the **float** value is not a positive or negative infinity, the virtual machine rounds the value towards zero using IEEE 754 round-towards-zero mode. If the resulting integral value can be exactly represented by a **long**, the **long** result is that integral value. Otherwise, the magnitude of the **float** value is too great to be represented in a **long**. If value is positive, the **long** result is the largest positive integer representable in a **long**. If value is negative, the **int** result is the smallest negative integer representable in a **long**.

Note that this instruction performs a narrowing primitive conversion. Because not all **float** values are representable by a **long**, the conversion may result in a loss of magnitude and precision.

For more information about the **f2l** instruction, see Chapter 11, "Type Conversion."

fadd

- Add floats

Opcode:

98 (0x62)

Instruction Format:

fadd

Stack:

Before: ..., value1, value2

After: ..., result

Description:

The top two words of the operand stack, value1 and value2, must be **floats**. To execute the **fadd** instruction, the Java Virtual Machine pops value1 and value2, adds them, and pushes the **float** result.

The result produced by the **fadd** instruction is governed by the rules of IEEE 754 floating point arithmetic.

For more information about the **fadd** instruction, see Chapter 14, "Floating Point Arithmetic."

fa load

- Load **float** from array

Opcode:

48 (0x30)

Instruction Format:

fa load

Stack:

Before: ..., arrayref, index

After: ..., value

Description:

To execute the **fa load** instruction, the Java Virtual Machine first pops two words from the operand stack. The arrayref word must be a **reference** that refers to an array of **floats**. The index word must be an **int**. The virtual machine retrieves from the arrayref array the **float** value specified by index and pushes it onto the operand stack.

If arrayref is **null**, the Java Virtual Machine throws **NullPointerException**. Otherwise, if index is not a legal index into the arrayref array, the virtual machine throws **ArrayIndexOutOfBoundsException**.

For more information about the **fa load** instruction, see Chapter 15, "Objects and Arrays."

fastore

- Store into **float** array

Opcode:

81 (0x51)

Instruction Format:

fastore

Stack:

Before: ..., arrayref, index, value

After: ...

Description:

To execute the **fa load** instruction, the Java Virtual Machine first pops three words from the operand stack. The arrayref word must be a **reference** that refers to an array of **floats**. The index word must be an **int**, and the value word must be a **float**. The virtual machine stores **float** value into the arrayref array location specified by index.

If arrayref is **null**, the Java Virtual Machine throws **NullPointerException**. Otherwise, if index is not a legal index into the arrayref array, the virtual machine throws **ArrayIndexOutOfBoundsException**.

For more information about the **fastore** instruction, see Chapter 15, "Objects and Arrays."

fcmpg

- Compare **floats** (1 on NaN)

Opcode:

150 (0x96)

Instruction Format:

fcmpg

Stack:

Before: ..., value1, value2

After: ..., result

Description:

The top two words of the operand stack, value1 and value2, must be two **floats**. To execute the **fcmpg** instruction, the Java Virtual Machine pops value1 and value2 off the operand stack and compares one against the other. If value1 equals value2, the virtual machine pushes onto the operand stack **int** result zero. Else, if value1 is greater than value2, the virtual machine pushes onto the operand stack **int** result one. Otherwise, if value1 is less than value2, the virtual machine pushes onto the operand stack **int** result negative one. If either value1 or value2 equals NaN (Not a Number), the virtual machine pushes onto the operand stack **int** result one.

The result produced by the **fcmpg** instruction is governed by the rules of IEEE 754 floating point

arithmetic. Note that the ***fcmpg*** instruction differs from the ***fcmpl*** instruction only in its treatment of NaN. For more information about the ***fcmpg*** instruction, see Chapter 16, "Control Flow."

fcmpl

- Compare *floats* (-1 on NaN)

Opcode:

149 (0x95)

Instruction Format:

fcmpl

Stack:

Before: ..., value1, value2

After: ..., result

Description:

The top two words of the operand stack, value1 and value2, must be two *floats*. To execute the ***fcmpl*** instruction, the Java Virtual Machine pops value1 and value2 off the operand stack and compares one against the other. If value1 equals value2, the virtual machine pushes onto the operand stack *int* result zero. Else, if value1 is greater than value2, the virtual machine pushes onto the operand stack *int* result one. Otherwise, if value1 is less than value2, the virtual machine pushes onto the operand stack *int* result negative one. If either value1 or value2 equals NaN (Not a Number), the virtual machine pushes onto the operand stack *int* result negative one.

The result produced by the ***fcmpl*** instruction is governed by the rules of IEEE 754 floating point arithmetic. Note that the ***fcmpl*** instruction differs from the ***fcmpg*** instruction only in its treatment of NaN. For more information about the ***fcmpl*** instruction, see Chapter 16, "Control Flow."

fconst_0

- Push *float* constant 0.0

Opcode:

11 (0xb)

Instruction Format:

fconst_0

Stack:

Before:...

After:..., <0.0

Description:

*To execute the **fconst_0** instruction, the Java Virtual Machine pushes the **float** constant 0.0 onto the operand stack.*

*For more information about the **fconst_0** instruction, see Chapter 10, "Stack Operations."*

fconst_1

*- Push **float** constant 1.0*

Opcode:

12 (0xc)

Instruction Format:

fconst_1

Stack:

Before:...

After:..., <1.0

Description:

*To execute the **fconst_1** instruction, the Java Virtual Machine pushes the **float** constant 1.0 onto the operand stack.*

*For more information about the **fconst_1** instruction, see Chapter 10, "Stack Operations."*

fconst_2

*- Push **float** constant 2.0*

Opcode:

13 (0xd)

Instruction Format:

fconst_2

Stack:

Before:...

After:..., <2.0

Description:

To execute the *fconst_2* instruction, the Java Virtual Machine pushes the *float* constant 2.0 onto the operand stack.

For more information about the *fconst_2* instruction, see Chapter 10, "Stack Operations."

fdiv

- Divide *floats*

Opcode:

110 (0x6e)

Instruction Format:

fdiv

Stack:

Before:..., value1, value2

After:..., result

Description:

The top two words of the operand stack, *value1* and *value2*, must be *floats*. To execute the *fdiv* instruction, the Java Virtual Machine pops *value1* and *value2*, divides *value1* by *value2* (*value1* / *value2*), and pushes the *float* result. The result produced by the *fdiv* instruction is governed by the rules of IEEE 754 floating point arithmetic.

For more information about the *fdiv* instruction, see Chapter 14, "Floating Point Arithmetic."

fload

- Load *float* from local variable

Opcode:

23 (0x17)

Instruction Format:

fload, index

Stack:

Before:...

After:..., value

Description:

The index operand, which serves as an 8-bit unsigned index into the local variables of the current frame, must specify a local variable word that contains an *float*. To execute the *fload* instruction, the Java Virtual Machine pushes onto the operand stack the *float* contained in the local variable word specified by index.

Note that the *wide* instruction can precede the *fload* instruction, to allow a local variable to be accessed with a 16-bit unsigned offset.

For more information about the *fload* instruction, see Chapter 10, "Stack Operations."

fload_0

- Load *float* from local variable 0

Opcode:

34 (0x22)

Instruction Format:

fload_0

Stack:

Before:...

After:..., value

Description:

The local variable word at index zero must contain a *float*. To execute the *fload_0* instruction, the Java Virtual Machine pushes onto the operand stack the *float* value contained in local variable word zero.

For more information about the *fload_0* instruction, see Chapter 10, "Stack Operations."

fload_1

- Load *float* from local variable 1

Opcode:

35 (0x23)

Instruction Format:

fload_1

Stack:

Before:...

After:..., value

Description:

The local variable word at index one must contain a *float*. To execute the *fload_1* instruction, the Java Virtual Machine pushes onto the operand stack the *float* value contained in local variable word one.

For more information about the *fload_1* instruction, see Chapter 10, "Stack Operations."

fload_2

- Load *float* from local variable 2

Opcode:

36 (0x24)

Instruction Format:

fload_2

Stack:

Before:...

After:..., value

Description:

The local variable word at index two must contain a *float*. To execute the *fload_2* instruction, the Java Virtual Machine pushes onto the operand stack the *float* value contained in local variable word two.

For more information about the *fload_2* instruction, see Chapter 10, "Stack Operations."

fload_3

- Load *float* from local variable 3

Opcode:

37 (0x25)

Instruction Format:

fload_3

Stack:

Before:...

After:..., value

Description:

The local variable word at index three must contain a *float*. To execute the *fload_3* instruction, the Java Virtual Machine pushes onto the operand stack the *float* value contained in local variable word three.

For more information about the *fload_3* instruction, see Chapter 10, "Stack Operations."

fmul

- Multiply *floats*

Opcode:

106 (0x6a)

Instruction Format:

fmul

Stack:

Before:..., value1, value2

After:..., result

Description:

The top two words of the operand stack, *value1* and *value2*, must be **floats**. To execute the **fmul** instruction, the Java Virtual Machine pops *value1* and *value2*, multiplies them, and pushes the **float** result. The result produced by the **fmul** instruction is governed by the rules of IEEE 754 floating point arithmetic.

For more information about the **fmul** instruction, see Chapter 14, "Floating Point Arithmetic."

fneg

- Negate **float**

Opcode:

118 (0x76)

Instruction Format:

fneg

Stack:

Before:..., *value*

After:..., result

Description:

The top word of the operand stack, *value*, must be a **float**. To execute the **fneg** instruction, the Java Virtual Machine pops *value*, negates it, and pushes the **float** result. The result produced by the **fneg** instruction is governed by the rules of IEEE 754 floating point arithmetic.

Note that the result produced by an **fneg** instruction is not always the same as the number that would be produced by subtracting *value* from zero with the **fsub** instruction. The range of IEEE 754 floating point numbers includes two zeros, a positive zero and a negative zero. When *value* is +0.0, the result of the **fneg** instruction is -0.0. By contrast, when subtracting +0.0 from +0.0, the **fsub** instruction yields +0.0.

For more information about the **fneg** instruction, see Chapter 14, "Floating Point Arithmetic."

frem

- Calculate remainder of division of **floats**

Opcode:

114 (0x72)

Instruction Format:

frem

Stack:

Before: ..., value1, value2

After: ..., result

Description:

The top two words of the operand stack, value1 and value2, must be *floats*. To execute the *frem* instruction, the Java Virtual Machine pops value1 and value2, calculates the remainder, and pushes the *float* result. The remainder equals $\text{value1} - (\text{value1} / \text{value2}) * \text{value2}$, where $\text{value1} / \text{value2}$ is a truncating division, rather than the rounding division required by IEEE 754.

The behavior of *frem* is comparable to that of the C library function *fmod()*. The remainder of two *floats* can be calculated according to IEEE 754 floating point standard via the *IEEERemainder()* method of class *java.lang.Math*.

For more information about the *frem* instruction, see Chapter 14, "Floating Point Arithmetic."

freturn

- Return *float* from method

Opcode:

174 (0xae)

Instruction Format:

freturn

Stack:

Before: ..., value

After: [empty]

Description:

The return type of the returning method must be *float*. The top word of the operand stack, value, must be a *float*. To execute the *freturn* instruction, the Java Virtual Machine pops *float* value from the operand stack of the current frame and pushes it onto the operand stack of the invoking method.

frame. The virtual machine discards any other words that may still be on the returning method's frame. If the returning method is synchronized, the monitor that was acquired when the method was invoked is released. The invoking method's frame is made current, and the virtual machine continues execution in the invoking method.

For more information about monitors, see Chapter 20, "Thread Synchronization." For more information about the *freturn* instruction, see Chapter 19, "Method Invocation and Return."

fstore

- Store *float* into local variable

Opcode:

56 (0x38)

Instruction Format:

fstore, index

Stack:

Before: ..., value

After: ...

Description:

The index operand must specify a valid 8-bit unsigned index into the local variables of the current frame. The value on the top of the operand stack must be a *float*. To execute the *fstore* instruction, the Java Virtual Machine pops the *float* value from the top of the operand stack and stores it into the local variable word specified by index.

Note that the *wide* instruction can precede the *fstore* instruction, to enable a value to be stored into a local variable specified by a 16-bit unsigned offset.

For more information about the *fstore* instruction, see Chapter 10, "Stack Operations."

fstore_0

- Store *float* into local variable 0

Opcode:

67 (0x43)

Instruction Format:

fstore_0

Stack:

Before: ..., value

After: ...

Description:

*The index zero must be a valid index into the local variables of the current stack frame, and the value word on the top of the operand stack must be a **float**. To execute the **fstore_0** instruction, the Java Virtual Machine pops the **float** value from the top of the operand stack and stores it into the local variable word at index zero.*

*For more information about the **fstore_0** instruction, see Chapter 10, "Stack Operations."*

fstore_1

*- Store **float** into local variable 1*

Opcode:

68 (0x44)

Instruction Format:

fstore_1

Stack:

Before: ..., value

After: ...

Description:

*The index one must be a valid index into the local variables of the current stack frame, and the value word on the top of the operand stack must be a **float**. To execute the **fstore_1** instruction, the Java Virtual Machine pops the **float** value from the top of the operand stack and stores it into the local variable word at index one.*

*For more information about the **fstore_1** instruction, see Chapter 10, "Stack Operations."*

fstore_2

*- Store **float** into local variable 2*

Opcode:

69 (0x45)

Instruction Format:

fstore_2

Stack:

Before: ..., value

After: ...

Description:

The index two must be a valid index into the local variables of the current stack frame, and the value word on the top of the operand stack must be a *float*. To execute the *fstore_2* instruction, the Java Virtual Machine pops the *float* value from the top of the operand stack and stores it into the local variable word at index two.

For more information about the *fstore_2* instruction, see Chapter 10, "Stack Operations."

fstore_3

- Store *float* into local variable 3

Opcode:

70 (0x46)

Instruction Format:

fstore_3

Stack:

Before: ..., value

After: ...

Description:

The index three must be a valid index into the local variables of the current stack frame, and the value word on the top of the operand stack must be a *float*. To execute the *fstore_3* instruction, the Java Virtual Machine pops the *float* value from the top of the operand stack and stores it into the local variable word at index three.

For more information about the *fstore_3* instruction, see Chapter 10, "Stack Operations."

fsub

- Subtract *floats*

Opcode:

102 (0x66)

Instruction Format:

fsub

Stack:

Before: ..., value1, value2

After: ..., result

Description:

The top two words of the operand stack, value1 and value2, must be *floats*. To execute the *fsub* instruction, the Java Virtual Machine pops value1 and value2, subtracts value2 from value1 (value1 - value2), and pushes the *float* result. The result produced by the *fsub* instruction is governed by the rules of IEEE 754 floating point arithmetic.

For more information about the *fsub* instruction, see Chapter 14, "Floating Point Arithmetic."

getfield

- Fetch field from object

Opcode:

180 (0xb4)

Instruction Format:

getfield, indexbyte1, indexbyte2

Stack:

Before: ..., objectref

After: ..., value

or

Before: ..., objectref

After: ..., value.word1, value.word2

Description:

The top word of the stack, objectref, must be a **reference**. To execute the **getfield** instruction, the Java Virtual Machine first forms an unsigned 16-bit index into the constant pool by calculating $(\text{indexbyte1} \ll 8) \mid \text{indexbyte2}$. The virtual machine then looks up the constant pool entry specified by the calculated index. The constant pool entry at that index must be a **CONSTANT_Fieldref_info** entry. If it hasn't already, the virtual machine resolves the entry, which yields the field's width and the field's offset from the beginning of the object image. If the resolution is successful, the virtual machine pops the objectref and fetches the field from the object pointed to by objectref. If the type of the field is byte, short, or char, the virtual machine sign-extends the field's value to an int and pushes the single-word int value onto the operand stack. Else, if the type is int, boolean, float, or reference, the virtual machine pushes the single-word value onto the operand stack. Otherwise, the type is long or double, and the virtual machine pushes the dual-word value onto the operand stack.

As a result of executing this instruction, the virtual machine may throw any of the linking errors listed in Chapter 8, "The Linking Model," as possible during resolution of a **CONSTANT_Fieldref_info** entry. As part of the process of resolving the **CONSTANT_Fieldref_info** entry, the virtual machine checks whether the field's access permission enables the current class to access the field. If the field is protected, the virtual machine makes certain the field is a member of the either the current class or a superclass of the current class, and that the class of the object pointed to by objectref is either the current class or a subclass of the current class. If not (or if there is any other access permission problem), the virtual machine throws **IllegalAccessError**. Else, if the field exists and is accessible from the current class, but the field is static, the virtual machine throws **IncompatibleClassChangeError**. Otherwise, if objectref is null, the virtual machine throws **NullPointerException**.

For more information about the **getfield** instruction, see Chapter 15, "Objects and Arrays."

getstatic

- Fetch static field from class

Opcod

178 (0xb2)

Instruction Format:

`getstatic, indexbyte1, indexbyte2`

Stack:

Before:...

After:..., value

or

Before:...

After:..., value.word1, value.word2

Description:

To execute the `getstatic` instruction, the Java Virtual Machine first forms an unsigned 16-bit index into the constant pool by calculating $(\text{indexbyte1} \ll 8) \mid \text{indexbyte2}$. The virtual machine then looks up the constant pool entry specified by the calculated index. The constant pool entry at that index must be a `CONSTANT_Fieldref_info` entry. If it hasn't already, the virtual machine resolves the entry. If the resolution is successful, the virtual machine fetches the value of the static field. If the type of the field is `byte`, `short`, or `char`, the virtual machine sign-extends the field's value to `int` and pushes the single-word `int` value onto the operand stack. Else, if the type is `int`, `boolean`, `float`, or `reference` the virtual machine pushes the single-word value onto the operand stack. Otherwise, the type is `long` or `double`, and the virtual machine pushes the dual-word value onto the operand stack.

As a result of executing this instruction, the virtual machine may throw any of the linking errors listed in Chapter 8, "The Linking Model," as possible during resolution of a `CONSTANT_Fieldref_info` entry. As part of the process of resolving the `CONSTANT_Fieldref_info` entry, the virtual machine checks whether the field's access permission enables the current class to access the field. If the field is protected, the virtual machine makes certain the field is a member of the either the current class or a superclass of the current class. If not (or if there is any other access permission problem), the virtual machine throws `IllegalAccessException`. Else, if the field exists and is accessible from the current class, but the field is not static, the virtual machine throws `IncompatibleClassChangeError`.

For more information about the `getstatic` instruction, see Chapter 15, "Objects and Arrays."

goto

- Branch always

Opcode:

167 (0xa7)

Instruction Format:

goto, branchbyte1, branchbyte2

Stack:

No change

Description:

To execute the goto instruction, the Java Virtual Machine forms a signed 16-bit offset by calculating (branchbyte1 8) | branchbyte2. The virtual machine then calculates a target (program counter) address by adding the calculated offset to the address of the goto opcode. The target address must be the address of an opcode within the same method as the goto opcode. The virtual machine jumps to the target address and continues execution there.

For more information about the goto instruction, see Chapter 16, "Control Flow."

goto_w

- Branch always (wide index)

Opcode:

200 (0xc8)

Instruction Format:

goto_w = 200, branchbyte1, branchbyte2, branchbyte3, branchbyte4

Stack:

No change

Description:

To execute the goto_w instruction, the Java Virtual Machine forms a

signed 32-bit offset by calculating $(\text{branchbyte1} \ll 24) \mid (\text{branchbyte2} \ll 16) \mid (\text{branchbyte3} \ll 8) \mid \text{branchbyte4}$. The virtual machine then calculates a target (program counter) address by adding the calculated offset to the address of the `goto_w` opcode. The target address must be the address of an opcode within the same method as the `goto_w` opcode. The virtual machine jumps to the target address and continues execution there.

Note that despite the 32-bit offset of the `goto_w` instruction, Java methods are currently (in both the 1.0 and 1.1 releases) limited to 65,535 bytes by three items in the Java class file format: the sizes of the indexes in the `LineNumberTable` attribute, the `LocalVariableTable` attribute, and the `Code` attribute's `exception_table` item. According to the Java Virtual Machine specification, the 65,536 byte limit to Java methods may be raised in a future release. For more information about the `goto_w` instruction, see Chapter 16, "Control Flow."

i2b

- Convert `int` to `byte`

Opcode:

145 (0x91)

Instruction Format:

i2b

Stack:

Before:..., `value`

After:..., `result`

Description:

The top word of the operand stack, `value`, must be an `int`. To execute the *i2b* instruction, the Java Virtual Machine pops `int` value from the operand stack, truncates the `int` to a `byte`, sign-extends the result back to an `int`, and pushes the `int` result.

Note that this instruction performs a narrowing primitive conversion. As a result of this conversion, magnitude information may be lost and the sign bit may change.

For more information about the *i2b* instruction, see Chapter 11, "Type

Conversion."

i2c

- Convert *int* to *char*

Opcode:

146 (0x92)

Instruction Format:

i2c

Stack:

Before:..., value

After:..., result

Description:

The top word of the operand stack, value, must be an int. To execute the i2c instruction, the Java Virtual Machine pops int value from the operand stack, truncates the int to a char, zero-extends the result back to an int, and pushes the int result.

Note that this instruction performs a narrowing primitive conversion. As a result of this conversion, magnitude information may be lost and, as the result is always positive, the sign bit may change.

For more information about the i2c instruction, see Chapter 11, "Type Conversion."

i2d

- Convert *int* to *double*

Opcode:

135 (0x87)

Instruction Format:

i2d

Stack:

Before:..., value

After:..., result.word1, result.word2

Description:

The top word of the operand stack, value, must be an int. To execute the i2d instruction, the Java Virtual Machine pops int value from the operand stack, sign-extends the int to a double, and pushes the double result.

Note that this instruction performs a widening primitive conversion. Because all int values are representable by a double, the conversion is exact.

For more information about the i2d instruction, see Chapter 11, "Type Conversion."

i2f

- Convert int to float

Opcode:

134 (0x86)

Instruction Format:

i2f

Stack:

Before:..., value

After:..., result

Description:

The top word of the operand stack, value, must be an int. To execute the i2f instruction, the Java Virtual Machine pops int value from the operand stack, converts the int to a float using the IEEE round-to-nearest mode, and pushes the float result.

Note that this instruction performs a widening primitive conversion. Because not all int values are exactly representable by a float, the conversion may result in a loss of precision.

For more information about the i2f instruction, see Chapter 11, "Type

Conversion."

i2l

- Convert int to long

Opcode:

133 (0x85)

Instruction Format:

i2l

Stack:

Before:..., value

After:..., result.word1, result.word2

Description:

The top word of the operand stack, value, must be an int. To execute the i2l instruction, the Java Virtual Machine pops int value from the operand stack, sign-extends the int to a long, and pushes the long result.

Note that this instruction performs a widening primitive conversion. Because all int values are representable by a long, the conversion is exact.

For more information about the i2l instruction, see Chapter 11, "Type Conversion."

i2s

- Convert int to short

Opcode:

147 (0x93)

Instruction Format:

i2s

Stack:

Before:..., value

After:..., result

Description:

The top word of the operand stack, value, must be an int. To execute the i2s instruction, the Java Virtual Machine pops int value from the operand stack, truncates the int to a short, sign-extends the result back to an int, and pushes the int result.

Note that this instruction performs a narrowing primitive conversion. As a result of this conversion, magnitude information may be lost and the sign bit may change.

For more information about the i2s instruction, see Chapter 11, "Type Conversion."

iadd

- Add ints

Opcode:

96 (0x60)

Instruction Format:

iadd

Stack:

Before:..., value1, value2

After:..., result

Description:

The top two words of the operand stack, value1 and value2, must be ints. To execute the iadd instruction, the Java Virtual Machine pops value1 and value2, adds them, and pushes the int result. If overflow occurs, result is the 32 lowest order bits of the true mathematical result represented in a sufficiently wide two's-complement format, and the sign of result is different from that of the true mathematical result.

For more information about the iadd instruction, see Chapter 12, "Integer Arithmetic."

iaload

- Load int from array

Opcode:

46 (0x2e)

Instruction Format:

iaload

Stack:

Before:..., arrayref, index

After:..., value

Description:

To execute the iaload instruction, the Java Virtual Machine first pops two words from the operand stack. The arrayref word must be a reference that refers to an array of ints. The index word must be an int. The virtual machine retrieves from the arrayref array the int value specified by index and pushes it onto the operand stack.

If arrayref is null, the Java Virtual Machine throws NullPointerException. Otherwise, if index is not a legal index into the arrayref array, The virtual machine throws ArrayIndexOutOfBoundsException.

For more information about the iaload instruction, see Chapter 15, "Objects and Arrays."

iand

- Perform boolean AND on ints

Opcode:

126 (0x7e)

Instruction Format:

iand

Stack:

Before:..., value1, value2

After:..., result

Description:

The top two words of the operand stack, value1 and value2, must be ints. To execute the iand instruction, the Java Virtual Machine pops value1 and value2, bitwise ANDs them, and pushes the int result.

For more information about the iand instruction, see Chapter 13, "Logic."

iastore

- Store into int array

Opcode:

79 (0x4f)

Instruction Format:

iastore

Stack:

Before:..., arrayref, index, value

After:...

Description:

To execute the iastore instruction, the Java Virtual Machine first pops three words from the operand stack. The arrayref word must be a reference that refers to an array of ints. The index and value words must be ints. The virtual machine stores int value into the arrayref array location specified by index.

If arrayref is null, the Java Virtual Machine throws NullPointerException. Otherwise, if index is not a legal index into the arrayref array, The virtual machine throws ArrayIndexOutOfBoundsException.

For more information about the iastore instruction, see Chapter 15, "Objects and Arrays."

iconst_0

- Push int constant 0

Opcode:

3 (0x3)

Instruction Format:

iconst_0

Stack:

Before:...

After:..., 0

Description:

*To execute the *iconst_0* instruction, the Java Virtual Machine pushes the int constant 0 onto the operand stack.*

*For more information about the *iconst_0* instruction, see Chapter 10, "Stack Operations."*

iconst_1

- Push int constant 1

Opcode:

4 (0x4)

Instruction Format:

iconst_1

Stack:

Before:...

After:..., 1

Description:

*To execute the *iconst_1* instruction, the Java Virtual Machine pushes the int constant 1 onto the operand stack.*

*For more information about the *iconst_1* instruction, see Chapter 10,*

"Stack Operations."

iconst_2

- Push int constant 2

Opcode:

5 (0x5)

Instruction Format:

iconst_2

Stack:

Before:...

After:..., 2

Description:

To execute the `iconst_2` instruction, the Java Virtual Machine pushes the int constant 2 onto the operand stack.

For more information about the `iconst_2` instruction, see Chapter 10, "Stack Operations."

iconst_3

- Push int constant 3

Opcode:

6 (0x6)

Instruction Format:

iconst_3

Stack:

Before:...

After:..., 3

Description:

To execute the `iconst_3` instruction, the Java Virtual Machine pushes the int constant 3 onto the operand stack.

For more information about the `iconst_3` instruction, see Chapter 10, "Stack Operations."

`iconst_4`

- Push int constant 4

Opcode:

7 (0x7)

Instruction Format:

`iconst_4`

Stack:

Before:...

After:..., 4

Description:

To execute the `iconst_4` instruction, the Java Virtual Machine pushes the int constant 4 onto the operand stack.

For more information about the `iconst_4` instruction, see Chapter 10, "Stack Operations."

`iconst_5`

- Push int constant 5

Opcode:

8 (0x8)

Instruction Format:

`iconst_5`

Stack:

Before:...

After:..., 5

Description:

To execute the `iconst_5` instruction, the Java Virtual Machine pushes the int constant 5 onto the operand stack.

For more information about the `iconst_5` instruction, see Chapter 10, "Stack Operations."

`iconst_m1`

- Push int constant -1

Opcode:

2 (0x2)

Instruction Format:

`iconst_m1`

Stack:

Before:...

After:..., -1

Description:

To execute the `iconst_m1` instruction, the Java Virtual Machine pushes the int constant -1 onto the operand stack.

For more information about the `iconst_m1` instruction, see Chapter 10, "Stack Operations."

`idiv`

- Divide ints

Opcode:

108 (0x6c)

Instruction Format:

`idiv`

Stack:

Before:..., value1, value2

After:..., result

Description:

The top two words of the operand stack, value1 and value2, must be ints. To execute the idiv instruction, the Java Virtual Machine pops value1 and value2, integer divides value1 by value2 (value1 / value2), and pushes the int result.

Integer division rounds the magnitude of the true mathematical quotient towards zero to the nearest integer. If the magnitude of the denominator is greater than that of the numerator, the int result is zero. Else, with one special exception, the sign of result is positive if the signs of the numerator and denominator are the same, negative if they are different. The exception to this rule is when the numerator is the smallest negative integer that can be represented by an int and the denominator is -1. For this division, the true mathematical result is one greater than the largest positive integer that can be represented by an int. As a consequence, the division overflows and the result is equal to the numerator.

If value2 (the denominator) is zero, the Java Virtual Machine throws ArithmeticException.

For more information about the idiv instruction, see Chapter 12, "Integer Arithmetic."

ifeq

- Branch if equal to 0

Opcode:

153 (0x99)

Instruction Format:

ifeq, branchbyte1, branchbyte2

Stack:

Before:..., value

After:...

Description:

The top word of the operand stack, value, must be an int. To execute the ifeq instruction, the Java Virtual Machine pops value off the operand stack and compares it against zero. If value equals zero, The virtual machine forms a signed 16-bit offset by calculating $(\text{branchbyte1} \ll 8) \mid \text{branchbyte2}$. The virtual machine then calculates a target (program counter) address by adding the calculated offset to the address of the ifeq opcode. The target address must be the address of an opcode within the same method as the ifeq opcode. The virtual machine jumps to the target address and continues execution there. Otherwise, if value does not equal zero, The virtual machine does not take the jump. It simply continues execution at the instruction immediately following the ifeq instruction.

For more information about the ifeq instruction, see Chapter 16, "Control Flow."

ifge

- Branch if greater than or equal to 0

Opcode:

156 (0x9c)

Instruction Format:

ifge, branchbyte1, branchbyte2

Stack:

Before:..., value

After:...

Description:

The top word of the operand stack, value, must be an int. To execute the ifge instruction, the Java Virtual Machine pops value off the operand stack and compares it against zero. If value is greater than or equal to zero, The virtual machine forms a signed 16-bit offset by calculating $(\text{branchbyte1} \ll 8) \mid \text{branchbyte2}$. The virtual machine then calculates a target (program counter) address by adding the calculated offset to the address of the ifge opcode. The target address must be the address of an opcode within the same method as the ifge opcode. The virtual machine jumps to the target address and continues execution there. Otherwise, if value is not greater than or

equal to zero, The virtual machine does not take the jump. It simply continues execution at the instruction immediately following the ifge instruction.

For more information about the ifge instruction, see Chapter 16, "Control Flow."

ifgt

- Branch if greater than 0

Opcode:

157 (0x9d)

Instruction Format:

ifgt, branchbyte1, branchbyte2

Stack:

Before:..., value

After:...

Description:

The top word of the operand stack, value, must be an int. To execute the ifgt instruction, the Java Virtual Machine pops value off the operand stack and compares it against zero. If value is greater than zero, The virtual machine forms a signed 16-bit offset by calculating (branchbyte1 8) | branchbyte2. The virtual machine then calculates a target (program counter) address by adding the calculated offset to the address of the ifgt opcode. The target address must be the address of an opcode within the same method as the ifgt opcode. The virtual machine jumps to the target address and continues execution there. Otherwise, if value is not greater than zero, The virtual machine does not take the jump. It simply continues execution at the instruction immediately following the ifgt instruction.

For more information about the ifgt instruction, see Chapter 16, "Control Flow."

ifle

- Branch if less than or equal to 0

Opcode:

158 (0x9e)

Instruction Format:

ifle, branchbyte1, branchbyte2

Stack:

Before:..., value

After:...

Description:

The top word of the operand stack, value, must be an int. To execute the ifle instruction, the Java Virtual Machine pops value off the operand stack and compares it against zero. If value is less than or equal to zero, The virtual machine forms a signed 16-bit offset by calculating (branchbyte1 8) | branchbyte2. The virtual machine then calculates a target (program counter) address by adding the calculated offset to the address of the ifle opcode. The target address must be the address of an opcode within the same method as the ifle opcode. The virtual machine jumps to the target address and continues execution there. Otherwise, if value is not less than or equal to zero, The virtual machine does not take the jump. It simply continues execution at the instruction immediately following the ifle instruction.

For more information about the ifle instruction, see Chapter 16, "Control Flow."

iflt

- Branch if less than 0

Opcode:

155 (0x9b)

Instruction Format:

iflt, branchbyte1, branchbyte2

Stack:

Before:..., value

After:...

Description:

The top word of the operand stack, value, must be an int. To execute the iflt instruction, the Java Virtual Machine pops value off the operand stack and compares it against zero. If value is less than zero, The virtual machine forms a signed 16-bit offset by calculating (branchbyte1 8) | branchbyte2. The virtual machine then calculates a target (program counter) address by adding the calculated offset to the address of the iflt opcode. The target address must be the address of an opcode within the same method as the iflt opcode. The virtual machine jumps to the target address and continues execution there. Otherwise, if value is not less than zero, The virtual machine does not take the jump. It simply continues execution at the instruction immediately following the iflt instruction.

For more information about the iflt instruction, see Chapter 16, "Control Flow."

ifne

- Branch if not equal to 0

Opcode:

154 (0x9a)

Instruction Format:

ifne, branchbyte1, branchbyte2

Stack:

Before:..., value

After:...

Description:

The top word of the operand stack, value, must be an int. To execute the ifne instruction, the Java Virtual Machine pops value off the operand stack and compares it against zero. If value does not equal zero, The virtual machine forms a signed 16-bit offset by calculating (branchbyte1 8) | branchbyte2. The virtual machine then calculates a target (program counter) address by adding the calculated offset to the address of the ifne opcode. The target address must be the address of an opcode within the same method as the ifne opcode. The virtual machine jumps to the target address and continues execution there. Otherwise, if value does equal zero, The virtual machine does

not take the jump. It simply continues execution at the instruction immediately following the ifne instruction.

For more information about the ifne instruction, see Chapter 16, "Control Flow."

ifnonnull

- Branch if not null

Opcode:

199 (0xc7)

Instruction Format:

ifnonnull, branchbyte1, branchbyte2

Stack:

Before:..., value

After:...

Description:

The top word of the operand stack, value, must be a reference. To execute the ifnonnull instruction, the Java Virtual Machine pops value off the operand stack and compares it against null. If value is not null, The virtual machine forms a signed 16-bit offset by calculating $(\text{branchbyte1} \ll 8) \mid \text{branchbyte2}$. The virtual machine then calculates a target (program counter) address by adding the calculated offset to the address of the ifnonnull opcode. The target address must be the address of an opcode within the same method as the ifnonnull opcode. The virtual machine jumps to the target address and continues execution there. Otherwise, if value is null, The virtual machine does not take the jump. It simply continues execution at the instruction immediately following the ifnonnull instruction.

For more information about the ifnonnull instruction, see Chapter 16, "Control Flow."

ifnull

- Branch if null

Opcode:

198 (0xc6)

Instruction Format:

ifnull, branchbyte1, branchbyte2

Stack:

Before:..., value

After:...

Description:

The top word of the operand stack, value, must be a reference. To execute the ifnull instruction, the Java Virtual Machine pops value off the operand stack and compares it against null. If value is null, The virtual machine forms a signed 16-bit offset by calculating (branchbyte1 8) | branchbyte2. The virtual machine then calculates a target (program counter) address by adding the calculated offset to the address of the ifnull opcode. The target address must be the address of an opcode within the same method as the ifnull opcode. The virtual machine jumps to the target address and continues execution there. Otherwise, if value is not null, The virtual machine does not take the jump. It simply continues execution at the instruction immediately following the ifnull instruction.

For more information about the ifnull instruction, see Chapter 16, "Control Flow."

if_acmpeq

- Branch if object references are equal

Opcode:

165 (0xa5)

Instruction Format:

if_acmpeq, branchbyte1, branchbyte2

Stack:

Before:..., value1, value2

After:...

Description:

The top two words of the operand stack, value1 and value2, must be references. To execute the `if_acmpeq` instruction, the Java Virtual Machine pops value1 and value2 off the operand stack and compares one against the other. If value1 equals value2 (in other words, if they both point to exactly the same object or are both null), The virtual machine forms a signed 16-bit offset by calculating $(\text{branchbyte1} \ll 8) \mid \text{branchbyte2}$. The virtual machine then calculates a target (program counter) address by adding the calculated offset to the address of the `if_acmpeq` opcode. The target address must be the address of an opcode within the same method as the `if_acmpeq` opcode. The virtual machine jumps to the target address and continues execution there. Otherwise, if value does not equal value2, The virtual machine does not take the jump. It simply continues execution at the instruction immediately following the `if_acmpeq` instruction.

For more information about the `if_acmpeq` instruction, see Chapter 16, "Control Flow."

`if_acmpne`

- Branch if object references not equal

Opcode:

166 (0xa6)

Instruction Format:

`if_acmpne`, branchbyte1, branchbyte2

Stack:

Before:..., value1, value2

After:...

Description:

The top two words of the operand stack, value1 and value2, must be references. To execute the `if_acmpne` instruction, the Java Virtual Machine pops value1 and value2 off the operand stack and compares one against the other. If value1 does not equal value2 (in other words, if they don't both point to exactly the same object and they aren't both null), The virtual machine forms a signed 16-bit offset by calculating $(\text{branchbyte1} \ll 8) \mid \text{branchbyte2}$. The virtual machine then calculates a target (program counter) address by adding the

calculated offset to the address of the `if_acmpne` opcode. The target address must be the address of an opcode within the same method as the `if_acmpne` opcode. The virtual machine jumps to the target address and continues execution there. Otherwise, if value equals value2, The virtual machine does not take the jump. It simply continues execution at the instruction immediately following the `if_acmpne` instruction.

For more information about the `if_acmpne` instruction, see Chapter 16, "Control Flow."

`if_icmpeq`

- Branch if ints equal

Opcode:

159 (0x9f)

Instruction Format:

`if_icmpeq`, branchbyte1, branchbyte2

Stack:

Before:..., value1, value2

After:...

Description:

The top two words of the operand stack, value1 and value2, must be ints. To execute the `if_icmpeq` instruction, the Java Virtual Machine pops value1 and value2 off the operand stack and compares one against the other. If value1 equals value2, The virtual machine forms a signed 16-bit offset by calculating $(\text{branchbyte1} \ll 8) \mid \text{branchbyte2}$. The virtual machine then calculates a target (program counter) address by adding the calculated offset to the address of the `if_icmpeq` opcode. The target address must be the address of an opcode within the same method as the `if_icmpeq` opcode. The virtual machine jumps to the target address and continues execution there. Otherwise, if value does not equal value2, The virtual machine does not take the jump. It simply continues execution at the instruction immediately following the `if_icmpeq` instruction.

For more information about the `if_icmpeq` instruction, see Chapter 16, "Control Flow."

`if_icmpge`

- Branch if int greater than or equal to other int

Opcode:

162 (0xa2)

Instruction Format:

if_icmpge, branchbyte1, branchbyte2

Stack:

Before:..., value1, value2

After:...

Description:

The top two words of the operand stack, value1 and value2, must be ints. To execute the *if_icmpge* instruction, the Java Virtual Machine pops value1 and value2 off the operand stack and compares one against the other. If value1 is greater than or equal to value2, The virtual machine forms a signed 16-bit offset by calculating (branchbyte1 8) | branchbyte2. The virtual machine then calculates a target (program counter) address by adding the calculated offset to the address of the *if_icmpge* opcode. The target address must be the address of an opcode within the same method as the *if_icmpge* opcode. The virtual machine jumps to the target address and continues execution there. Otherwise, if value is not greater than or equal to value2, The virtual machine does not take the jump. It simply continues execution at the instruction immediately following the *if_icmpge* instruction.

For more information about the *if_icmpge* instruction, see Chapter 16, "Control Flow."

if_icmpgt

- Branch if int greater than other int

Opcode:

163 (0xa3)

Instruction Format:

if_icmpgt, branchbyte1, branchbyte2

Stack:

Before:..., value1, value2

After:...

Description:

The top two words of the operand stack, value1 and value2, must be ints. To execute the if_icmpgt instruction, the Java Virtual Machine pops value1 and value2 off the operand stack and compares one against the other. If value1 is greater than value2, The virtual machine forms a signed 16-bit offset by calculating (branchbyte1 8) | branchbyte2. The virtual machine then calculates a target (program counter) address by adding the calculated offset to the address of the if_icmpgt opcode. The target address must be the address of an opcode within the same method as the if_icmpgt opcode. The virtual machine jumps to the target address and continues execution there. Otherwise, if value is not greater than value2, The virtual machine does not take the jump. It simply continues execution at the instruction immediately following the if_icmpgt instruction.

For more information about the if_icmpgt instruction, see Chapter 16, "Control Flow."

if_icmple

- Branch if int less than or equal to other int

Opcode:

164 (0xa4)

Instruction Format:

if_icmple, branchbyte1, branchbyte2

Stack:

Before:..., value1, value2

After:...

Description:

The top two words of the operand stack, value1 and value2, must be ints. To execute the if_icmple instruction, the Java Virtual Machine pops value1 and value2 off the operand stack and compares one against the other. If value1 is less than or equal to value2, The virtual machine forms a signed 16-bit offset by calculating (branchbyte1 8) |

branchbyte2. The virtual machine then calculates a target (program counter) address by adding the calculated offset to the address of the if_icmple opcode. The target address must be the address of an opcode within the same method as the if_icmple opcode. The virtual machine jumps to the target address and continues execution there. Otherwise, if value is not less than or equal to value2, The virtual machine does not take the jump. It simply continues execution at the instruction immediately following the if_icmple instruction.

For more information about the if_icmple instruction, see Chapter 16, "Control Flow."

if_icmplt

- Branch if int less than other int

Opcode:

161 (0xa1)

Instruction Format:

if_icmplt, branchbyte1, branchbyte2

Stack:

Before:..., value1, value2

After:...

Description:

The top two words of the operand stack, value1 and value2, must be ints. To execute the if_icmplt instruction, the Java Virtual Machine pops value1 and value2 off the operand stack and compares one against the other. If value1 is less than value2, The virtual machine forms a signed 16-bit offset by calculating (branchbyte1 8) | branchbyte2. The virtual machine then calculates a target (program counter) address by adding the calculated offset to the address of the if_icmplt opcode. The target address must be the address of an opcode within the same method as the if_icmplt opcode. The virtual machine jumps to the target address and continues execution there. Otherwise, if value is not less than value2, The virtual machine does not take the jump. It simply continues execution at the instruction immediately following the if_icmplt instruction.

For more information about the if_icmplt instruction, see Chapter 16, "Control Flow."

if_icmpne

- Branch if ints not equal

Opcode:

160 (0xa0)

Instruction Format:

if_icmpne, *branchbyte1*, *branchbyte2*

Stack:

Before:..., *value1*, *value2*

After:...

Description:

The top two words of the operand stack, *value1* and *value2*, must be ints. To execute the *if_icmpne* instruction, the Java Virtual Machine pops *value1* and *value2* off the operand stack and compares one against the other. If *value1* does not equal *value2*, The virtual machine forms a signed 16-bit offset by calculating $(\text{branchbyte1} \ll 8) \mid \text{branchbyte2}$. The virtual machine then calculates a target (program counter) address by adding the calculated offset to the address of the *if_icmpne* opcode. The target address must be the address of an opcode within the same method as the *if_icmpne* opcode. The virtual machine jumps to the target address and continues execution there. Otherwise, if *value1* equals *value2*, The virtual machine does not take the jump. It simply continues execution at the instruction immediately following the *if_icmpne* instruction.

For more information about the *if_icmpne* instruction, see Chapter 16, "Control Flow."

iinc

- Increment int local variable by constant

Opcode:

132 (0x84)

Instruction Format:

iinc, *index*, *const*

Stack:

No change

Description:

The index operand must specify a valid 8-bit unsigned index into the local variables of the current frame. To execute the `iinc` instruction, the Java Virtual Machine adds the 8-bit signed increment `const` to the local variable word specified by `index`.

Note that the `wide` instruction can precede the `iinc` instruction, to enable a local variable specified by a 16-bit unsigned offset to be incremented by a signed 16-bit constant.

For more information about the `iinc` instruction, see Chapter 12, "Integer Arithmetic."

`iload`

- Load `int` from local variable

Opcode:

21 (0x15)

Instruction Format:

`iload, index`

Stack:

Before:...

After:..., value

Description:

The index operand, which serves as an 8-bit unsigned index into the local variables of the current frame, must specify a local variable word that contains an `int`. To execute the `iload` instruction, the Java Virtual Machine pushes onto the operand stack the `int` value contained in the local variable word specified by `index`.

Note that the `wide` instruction can precede the `iload` instruction, to allow a local variable to be accessed with a 16-bit unsigned offset.

For more information about the `iload` instruction, see Chapter 10,

"Stack Operations."

iload_0

- Load int from local variable 0

Opcode:

26 (0x1a)

Instruction Format:

iload_0

Stack:

Before:...

After:..., value

Description:

The local variable word at index zero must contain an int. To execute the `iload_0` instruction, the Java Virtual Machine pushes onto the operand stack the int value contained in the local variable word zero.

For more information about the `iload_0` instruction, see Chapter 10, "Stack Operations."

iload_1

- Load int from local variable 1

Opcode:

27 (0x1b)

Instruction Format:

iload_1

Stack:

Before:...

After:..., value

Description:

The local variable word at index one must contain an int. To execute the `iload_1` instruction, the Java Virtual Machine pushes onto the operand stack the int contained in the local variable word one.

For more information about the `iload_1` instruction, see Chapter 10, "Stack Operations."

`iload_2`

- Load int from local variable 2

Opcode:

28 (0x1c)

Instruction Format:

`iload_2`

Stack:

Before:...

After:..., value

Description:

The local variable word at index two must contain an int. To execute the `iload_2` instruction, the Java Virtual Machine pushes onto the operand stack the int value contained in the local variable word two.

For more information about the `iload_2` instruction, see Chapter 10, "Stack Operations."

`iload_3`

- Load int from local variable 3

Opcode:

29 (0x1d)

Instruction Format:

`iload_3`

Stack:

Before:...

After:..., value

Description:

The local variable word at index three must contain an int. To execute the `iload_3` instruction, the Java Virtual Machine pushes onto the operand stack the int value contained in the local variable word three.

For more information about the `iload_3` instruction, see Chapter 10, "Stack Operations."

`imul`

- Multiply ints

Opcode:

104 (0x68)

Instruction Format:

`imul`

Stack:

Before:..., value1, value2

After:..., result

Description:

The top two words of the operand stack, value1 and value2, must be ints. To execute the `imul` instruction, the Java Virtual Machine pops value1 and value2, multiplies them, and pushes the int result. If overflow occurs, result is the 32 lowest order bits of the true mathematical result represented in a sufficiently wide two's-complement format, and the sign of result may be different from that of the true mathematical result.

For more information about the `imul` instruction, see Chapter 12, "Integer Arithmetic."

`ineg`

- Negate *int*

Opcode:

116 (0x74)

Instruction Format:

ineg

Stack:

Before:..., *value*

After:..., *result*

Description:

The top word of the operand stack, *value*, must be an *int*. To execute the *ineg* instruction, the Java Virtual Machine pops *value*, negates it, and pushes the *int* result.

The result produced by an *ineg* instruction is the same number that would be produced by subtracting *value* from zero with the *isub* instruction. As a consequence, when *value* is the smallest negative integer that can be represented by an *int*, the negation overflows. For this negation, the true mathematical result is one greater than the largest positive integer that can be represented by an *int*, and the actual result is equal to *value* with no change in sign.

For more information about the *ineg* instruction, see Chapter 12, "Integer Arithmetic."

instanceof

- Determine if an object is of given type

Opcode:

193 (0xc1)

Instruction Format:

instanceof, *indexbyte1*, *indexbyte2*

Stack:

Before:..., *objectref*

After:..., result

Description:

The top word of the stack, *objectref*, must be a reference. To execute the *instanceof* instruction, the Java Virtual Machine first forms an unsigned 16-bit index into the constant pool by calculating $(\text{indexbyte1} \ll 8) \mid \text{indexbyte2}$. The virtual machine then looks up the constant pool entry specified by the calculated index. The constant pool entry at that index must be a *CONSTANT_Class_info* entry. If it hasn't already, The virtual machine resolves the entry. The entry may be a class, interface, or array type. If *objectref* is not null and the object pointed to by *objectref* is an "instance of" the resolved type, The virtual machine pushes an int one result onto the operand stack. Otherwise, The virtual machine pushes an int zero result onto the operand stack.

To determine whether the object pointed to by *objectref* is an "instance of" the resolved type, The virtual machine first determines whether the object is a class instance or array. (It can't be an interface instance, because interfaces can't be instantiated.) If it is a class instance, and the resolved type is a class, not an interface, the object is "an instance" of the resolved class if the object's class is the resolved class or a subclass of the resolved class. Else, if it is a class instance, and the resolved type is an interface, not a class, the object is "an instance" of the resolved interface if the object's class implements the resolved interface. Otherwise, the object is an array. If the resolved type is a class, it must be *java.lang.Object*. Else, if the resolved type is an array of primitive types, the object must be an array of the same primitive type. Otherwise, the resolved type must be an array with a component type of some reference type, and the object must be an array with a component type that is an "instance of" the component type of the resolved array type. (Note that the dimension of an array doesn't enter into the *instanceof* check, only the component type of the array.)

As a result of executing this instruction, The virtual machine may throw any of the linking errors listed in Chapter 8, "The Linking Model," as possible during resolution of a *CONSTANT_Class_info* entry.

For more information about the *instanceof* instruction, see Chapter 15, "Objects and Arrays."

invokeinterface

- Invoke interface method

Opcode:

185 (0xb9)

Instruction Format:

invokeinterface, indexbyte1, indexbyte2, nargs, 0

Stack:

Before:..., objectref, [arg1, [arg2 ...]]

After:...

Description:

To execute the invokeinterface instruction, the Java Virtual Machine first forms an unsigned 16-bit index into the constant pool by calculating (indexbyte1 8) | indexbyte2. The virtual machine then looks up the constant pool entry specified by the calculated index. The constant pool entry at that index must be a CONSTANT_InterfaceMethodref_info entry. If it hasn't already, The virtual machine resolves the entry. The resolved method's descriptor must exactly match the descriptor of one of the methods declared in the resolved interface. The method must not be an instance initialization method, "<init", or a class initialization method, "<clinit."

The nargs operand is an unsigned byte that indicates the number of words of parameters required by the method being invoked (including the hidden this reference). The operand stack must contain nargs - 1 words of parameters and the objectref word. The parameter words must match the order and type of parameters required by the resolved method. The objectref word, the reference to the object upon which to invoke the instance method, must be a reference. If the resolution is successful, The virtual machine pops nargs - 1 parameter words and objectref.

To invoke the method, The virtual machine retrieves the direct reference to the instance method to invoke from a method table. It locates the method table for the class of object pointed to by objectref and searches through it for a method with a name and descriptor that matches exactly the name and descriptor of the resolved method. (If the object's class is an array type, The virtual machine uses the method table for class java.lang.Object. An array type, of course, can only implement an interface if Object itself implements the interface.)

If the method is synchronized, the Java Virtual Machine, on behalf of the current thread, acquires the monitor associated with objectref.

If the method to invoke is not native, The virtual machine creates a new stack frame for the method and pushes the new stack frame onto the current thread's Java stack. The virtual machine then places the objectref word and nargs - 1 parameter words that it popped from the operand stack of the calling method's frame into the local variables of the new stack frame. It places objectref into local variable position zero, arg1 into local variable position one, and so on. (objectref is the hidden this reference passed to all instance methods.) The virtual machine makes the new stack frame current, sets the program counter to the address of the first instruction in the new method, and continues execution there.

If the method to invoke is native, The virtual machine invokes the native method in an implementation-dependent manner.

As a result of executing this instruction, The virtual machine may throw any of the linking errors listed in Chapter 8, "The Linking Model," as possible during resolution of a CONSTANT_InterfaceMethodref_info entry. If no method exists in the object's class with the required name and descriptor, The virtual machine throws IncompatibleClassChangeError. Else, if the method exists but is static, The virtual machine throws IncompatibleClassChangeError. Else, if the method exists, but is not public, the virtual machine throws IllegalAccessException. Else, if the method is abstract, The virtual machine throws AbstractMethodError. Else, if the method is native and the native implementation of the method can't be loaded or linked, The virtual machine throws UnsatisfiedLinkError. Otherwise, if objectref is null, The virtual machine throws NullPointerException.

For more information about the invokeinterface instruction, see Chapter 19, "Method Invocation and Return."

invokespecial

- Invoke instance method with special handling for private, superclass, and instance initialization methods

Opcode:

183 (0xb7)

Instruction Format:

invokespecial, indexbyte1, indexbyte2

Stack:

Before:..., objectref, [arg1, [arg2 ...]]

After:...

Description:

To execute the `invokespecial` instruction, the Java Virtual Machine first forms an unsigned 16-bit index into the constant pool by calculating $(\text{indexbyte1} \ll 8) \mid \text{indexbyte2}$. The virtual machine then looks up the constant pool entry specified by the calculated index. The constant pool entry at that index must be a `CONSTANT_Methodref_info` entry. If it hasn't already, the virtual machine resolves the entry, which yields a direct reference to the method's data, including the number of words of parameters to the method, `nargs`. The resolved method's descriptor must exactly match the descriptor of one of the methods declared in the resolved class. The method must not be a class initialization method, "`<clinit>`."

The `invokespecial` instruction is used to invoke three special kinds of instance methods: superclass methods, private methods, and instance initialization methods. `invokespecial` contrasts with `invokevirtual` in the way it invokes an instance method. Whereas `invokevirtual` always selects the method to invoke at run-time based on the class of the object (dynamic binding), `invokespecial` normally (with one exception) selects the method to invoke at compile-time based on the type of the reference (static binding).

The one exception to `invokespecial`'s pure static binding behavior occurs if

- o the resolved method is not private and is not an instance initialization method,
- o the resolved method's class is a superclass of the current method's class,
- o and the `ACC_SUPER` flag is set in the current method's class.

In this situation, the Java Virtual Machine dynamically selects at run-time the method to invoke by finding the method in the closest superclass that has a descriptor that exactly matches the resolved method's descriptor, irrespective of the class of the resolved method. In the majority of cases, the class of the selected method will most likely be the class of the resolved method anyway, but it is possible that it could be some other class.

For example, imagine you create an inheritance hierarchy of three classes: `Animal`, `Dog`, and `CockerSpaniel`. Assume class `Dog` extends class `Animal`, class `CockerSpaniel` extends class `Dog`, and that a method defined in `CockerSpaniel` uses `invokespecial` to invoke a non-private superclass method named `walk()`. Assume also that when you

compiled CockerSpaniel, the compiler set the ACC_SUPER flag. In addition, assume that when you compiled CockerSpaniel, class Animal defined a walk() method, but Dog didn't. In that case, the symbolic reference from CockerSpaniel to the walk() method would give Animal as its class. When the invokespecial instruction in CockerSpaniel's method is executed, The virtual machine would dynamically select and invoke Animal's walk() method.

Now imagine that later, you added a walk() method to Dog, and recompiled Dog, but didn't recompile CockerSpaniel. CockerSpaniel's symbolic reference to the superclass walk() method still claims Animal as its class, even though there is now an implementation of walk() in Dog's class file. Nevertheless, when the invokespecial instruction in CockerSpaniel's method is executed, The virtual machine would dynamically select and invoke Dog's implementation of the walk() method.

This special (not static binding) treatment of superclass invocations was the motivation for adding the ACC_SUPER flag to the class access_flags item of class files, and for changing the name of this opcode from its original name, invokenonvirtual, to its current name, invokespecial. If the CockerSpaniel class of the previous example had been compiled by an old compiler that didn't set the ACC_SUPER flag, The virtual machine would invoke Animal's implementation of walk() regardless of whether Dog declared a walk(). As mentioned in Chapter 6, "The Java Class File," all new Java compilers should set the ACC_SUPER flag in every class file they generate.

If the resolved method is an instance initialization method, "<init," the method must be invoked only once on each uninitialized (except to default initial values) object. In addition, an instance initialization method must be invoked on each uninitialized object before the first backwards branch. In other words, the bytecodes of a method need not call an <init method with invokespecial right after a new instruction. Other instructions could intervene between the new and the invokespecial for the <init, but none of those instructions may branch backwards (such as a goto to the beginning of the method).

The operand stack must contain nargs - 1 words of parameters and the objectref word. The parameter words must match the order and type of parameters required by the resolved method. The objectref word, the reference to the object upon which to invoke the instance method, must be a reference. If the resolution is successful, The virtual machine pops nargs - 1 parameter words and objectref.

If the method is synchronized, the Java Virtual Machine, on behalf of the current thread, acquires the monitor associated with objectref.

If the method to invoke is not native, The virtual machine creates a new stack frame for the method and pushes the new stack frame onto

the current thread's Java stack. The virtual machine then places the objectref word and nargs - 1 parameter words that it popped from the operand stack of the calling method's frame into the local variables of the new stack frame. It places objectref into local variable position zero, arg1 into local variable position one, and so on. (objectref is the hidden this reference passed to all instance methods.) The virtual machine makes the new stack frame current, sets the program counter to the address of the first instruction in the new method, and continues execution there.

If the method to invoke is native, The virtual machine invokes the native method in an implementation-dependent manner.

As a result of executing this instruction, The virtual machine may throw any of the linking errors listed in Chapter 8, "The Linking Model," as possible during resolution of a CONSTANT_Methodref_info entry. As part of the process of resolving the CONSTANT_Methodref_info entry, The virtual machine checks whether the method's access permission enables the current class to access the method. If the method is protected, The virtual machine makes certain the method is a member of the either the current class or a superclass of the current class, and that the class of the object pointed to by objectref is either the current class or a subclass of the current class. If not (or if there is any other access permission problem), The virtual machine throws `IllegalAccessException`. Else, if the method exists and is accessible from the current class, but the method is static, The virtual machine throws `IncompatibleClassChangeError`. Else, if the method is abstract, The virtual machine throws `AbstractMethodError`. Else, if the method is native and the native implementation of the method can't be loaded or linked, The virtual machine throws `UnsatisfiedLinkError`. Otherwise, if objectref is null, The virtual machine throws `NullPointerException`.

For more information about the invokespecial instruction, see Chapter 19, "Method Invocation and Return."

invokestatic

- Invoke a class (static) method

Opcod:

184 (0xb8)

Instruction Format:

invokestatic, indexbyte1, indexbyte2

Stack:

Before:..., [arg1, [arg2 ...]

After:...

Description:

To execute the `invokestatic` instruction, the Java Virtual Machine first forms an unsigned 16-bit index into the constant pool by calculating $(\text{indexbyte1} \ll 8) \mid \text{indexbyte2}$. The virtual machine then looks up the constant pool entry specified by the calculated index. The constant pool entry at that index must be a `CONSTANT_Methodref_info` entry. If it hasn't already, The virtual machine resolves the entry, which yields a direct reference to the method's data, including the number of words of parameters to the method, `nargs`. The resolved method's descriptor must exactly match the descriptor of one of the methods declared in the resolved class. The method must not be an instance initialization method, "`<init>`", or a class initialization method, "`<clinit>`".

The operand stack must contain `nargs` words of parameters. The parameter words must match the order and type of parameters required by the resolved method. If the resolution is successful, The virtual machine pops the `nargs` parameter words.

If the method is synchronized, the Java Virtual Machine, on behalf of the current thread, acquires the monitor associated with the Class instance that represents the resolved method's class.

If the method to invoke is not native, The virtual machine creates a new stack frame for the method and pushes the new stack frame onto the current thread's Java stack. The virtual machine then places the `nargs` parameter words that it popped from the operand stack of the calling method's frame into the local variables of the new stack frame. It places `arg1` into local variable position zero, `arg2` into local variable position one, and so on. The virtual machine makes the new stack frame current, sets the program counter to the address of the first instruction in the new method, and continues execution there.

If the method to invoke is native, The virtual machine invokes the native method in an implementation-dependent manner.

As a result of executing this instruction, The virtual machine may throw any of the linking errors listed in Chapter 8, "The Linking Model," as possible during resolution of a `CONSTANT_Methodref_info` entry. As part of the process of resolving the `CONSTANT_Methodref_info` entry, The virtual machine checks whether the

method's access permission enables the current class to access the method. If the method is protected, The virtual machine makes certain the method is a member of the either the current class or a superclass of the current class. If not (or if there is any other access permission problem), The virtual machine throws `IllegalAccessError`. Else, if the method exists and is accessible from the current class, but the method is not static, The virtual machine throws `IncompatibleClassChangeError`. Else, if the method is abstract, The virtual machine throws `AbstractMethodError`. Otherwise, if the method is native and the native implementation of the method can't be loaded or linked, The virtual machine throws `UnsatisfiedLinkError`.

For more information about the `invokestatic` instruction, see Chapter 19, "Method Invocation and Return."

`invokevirtual`

- Invoke instance method, dispatch based on an object's class at run-time

Opcode:

182 (0xb6)

Instruction Format:

`invokevirtual, indexbyte1, indexbyte2`

Stack:

Before:..., `objectref`, [`arg1`, [`arg2 ...`]]

After:...

Description:

To execute the `invokevirtual` instruction, the Java Virtual Machine first forms an unsigned 16-bit index into the constant pool by calculating $(\text{indexbyte1} \ll 8) \mid \text{indexbyte2}$. The virtual machine then looks up the constant pool entry specified by the calculated index. The constant pool entry at that index must be a `CONSTANT_Methodref_info` entry. If it hasn't already, The virtual machine resolves the entry, which yields the method's method table index, index, and the number of words of parameters to the method, `nargs`. The resolved method's descriptor must exactly match the descriptor of one of the methods declared in the resolved class. The method must not be an instance initialization method, "`<init`", or a class initialization method, "`<clinit`".

The operand stack must contain `nargs - 1` words of parameters and the `objectref` word. The parameter words must match the order and type of parameters required by the resolved method. The `objectref` word, the reference to the object upon which to invoke the instance method, must be a reference. If the resolution is successful, The virtual machine pops `nargs - 1` parameter words and `objectref`.

To invoke the method, The virtual machine retrieves the direct reference to the instance method to invoke from a method table. It locates the method table for the class of object pointed to by `objectref` and looks up the direct reference that occupies method table position index. (If the object's class is an array type, The virtual machine uses the method table for class `java.lang.Object`.)

If the method is synchronized, the Java Virtual Machine, on behalf of the current thread, acquires the monitor associated with `objectref`.

If the method to invoke is not native, The virtual machine creates a new stack frame for the method and pushes the new stack frame onto the current thread's Java stack. The virtual machine then places the `objectref` word and `nargs - 1` parameter words that it popped from the operand stack of the calling method's frame into the local variables of the new stack frame. It places `objectref` into local variable position zero, `arg1` into local variable position one, and so on. (`objectref` is the hidden this reference passed to all instance methods.) The virtual machine makes the new stack frame current, sets the program counter to the address of the first instruction in the new method, and continues execution there.

If the method to invoke is native, The virtual machine invokes the native method in an implementation-dependent manner.

As a result of executing this instruction, The virtual machine may throw any of the linking errors listed in Chapter 8, "The Linking Model," as possible during resolution of a `CONSTANT_Methodref_info` entry. As part of the process of resolving the `CONSTANT_Methodref_info` entry, The virtual machine checks whether the method's access permission enables the current class to access the method. If the method is protected, The virtual machine makes certain the method is a member of the either the current class or a superclass of the current class, and that the class of the object pointed to by `objectref` is either the current class or a subclass of the current class. If not (or if there is any other access permission problem), The virtual machine throws `IllegalAccessError`. Else, if the method exists and is accessible from the current class, but the method is static, The virtual machine throws `IncompatibleClassChangeError`. Else, if the method is abstract, The virtual machine throws `AbstractMethodError`. Else, if the method is native and the native implementation of the method can't be loaded or linked, The virtual machine throws `UnsatisfiedLinkError`. Otherwise,

*if objectref is null, The virtual machine throws
NullPointerException.*

For more information about the invokevirtual instruction, see Chapter 19, "Method Invocation and Return."

ior

- Perform boolean OR on ints

Opcode:

128 (0x80)

Instruction Format:

ior

Stack:

Before:..., value1, value2

After:..., result

Description:

The top two words of the operand stack, value1 and value2, must be ints. To execute the ior instruction, the Java Virtual Machine pops value1 and value2, bitwise ORs them, and pushes the int result.

For more information about the ior instruction, see Chapter 13, "Logic."

irem

- Calculate remainder of division of ints

Opcode:

112 (0x70)

Instruction Format:

irem

Stack:

Before:..., value1, value2

After:..., result

Description:

*The top two words of the operand stack, value1 and value2, must be ints. To execute the irem instruction, the Java Virtual Machine pops value1 and value2, calculates the integer remainder, and pushes the int result. The integer remainder equals value1 - (value1 / value2) * value2.*

The irem instruction implements Java's remainder operator: %. The irem behaves such that the Java expression shown below is always true, where n and d are any two ints:

begin

*$(n/d)*d + (n\%d) == n$*

end

This behavior means that the result of an irem instruction always takes the same sign as the numerator, which is popped off the operand stack as value1.

If value2 (the denominator) is zero, the Java Virtual Machine throws ArithmeticException.

For more information about the irem instruction, see Chapter 12, "Integer Arithmetic."

ireturn

- Return int from method

Opcode:

172 (0xac)

Instruction Format:

ireturn

Stack:

Before:..., value

After:[empty]

Description:

The return type of the returning method must be byte, short, int, or char. The top word of the operand stack, value, must be an int. To execute the ireturn instruction, the Java Virtual Machine pops int value from the operand stack of the current frame and pushes it onto the operand stack of the invoking method's frame. The virtual machine discards any other words that may still be on the returning method's frame. If the returning method is synchronized, the monitor that was acquired when the method was invoked is released. The invoking method's frame is made current, and The virtual machine continues execution in the invoking method.

For more information about the ireturn instruction, see Chapter 19, "Method Invocation and Return." For more information about monitors, see Chapter 20, "Thread Synchronization."

ishl

- Perform left shift on int

Opcode:

120 (0x78)

Instruction Format:

ishl

Stack:

Before:..., value1, value2

After:..., result

Description:

The top two words of the operand stack, value1 and value2, must be ints. To execute the ishl instruction, the Java Virtual Machine pops value1 and value2, shifts value1 left by the number of bits specified in the 5 lowest order bits of value2 (from 0 to 31 bit positions), and pushes the int result.

For more information about the ishl instruction, see Chapter 13, "Logic."

ishr

- Perform arithmetic right shift on *int*

Opcode:

122 (0x7a)

Instruction Format:

ishr

Stack:

Before:..., *value1*, *value2*

After:..., *result*

Description:

The top two words of the operand stack, *value1* and *value2*, must be *ints*. To execute the *ishr* instruction, the Java Virtual Machine pops *value1* and *value2*, shifts *value1* right with sign extension by the number of bits specified in the 5 lowest order bits of *value2* (from 0 to 31 bit positions), and pushes the *int* result.

For more information about the *ishr* instruction, see Chapter 13, "Logic."

istore

- Store *int* into local variable

Opcode:

54 (0x36)

Instruction Format:

istore, *index*

Stack:

Before:..., *value*

After:...

Description:

The *index* operand must specify a valid 8-bit unsigned index into the

local variables of the current frame. The value word on the top of the operand stack must be an int. To execute the `istore` instruction, the Java Virtual Machine pops the int value from the top of the operand stack and stores it into the local variable word specified by index.

Note that the `wide` instruction can precede the `istore` instruction, to enable a value to be stored into a local variable specified by a 16-bit unsigned offset.

For more information about the `istore` instruction, see Chapter 10, "Stack Operations."

`istore_0`

- Store int into local variable 0

Opcode:

59 (0x3b)

Instruction Format:

`istore_0`

Stack:

Before:..., value

After:...

Description:

The index zero must be a valid index into the local variables of the current stack frame, and the value word on the top of the operand stack must be an int. To execute the `istore_0` instruction, the Java Virtual Machine pops the int value from the top of the operand stack and stores it into the local variable word at index zero.

For more information about the `istore_0` instruction, see Chapter 10, "Stack Operations."

`istore_1`

- Store int into local variable 1

Opcode:

60 (0x3c)

Instruction Format:

istore_1

Stack:

Before:..., value

After:...

Description:

*The index one must be a valid index into the local variables of the current stack frame, and the value word on the top of the operand stack must be an int. To execute the *istore_1* instruction, the Java Virtual Machine pops the int value from the top of the operand stack and stores it into the local variable word at index one.*

*For more information about the *istore_1* instruction, see Chapter 10, "Stack Operations."*

istore_2

- Store int into local variable 2

Opcode:

61 (0x3d)

Instruction Format:

istore_2

Stack:

Before:..., value

After:...

Description:

*The index two must be a valid index into the local variables of the current stack frame, and the value word on the top of the operand stack must be an int. To execute the *istore_2* instruction, the Java Virtual Machine pops the int value from the top of the operand stack and stores it into the local variable word at index two.*

For more information about the `istore_2` instruction, see Chapter 10, "Stack Operations."

`istore_3`

- Store int into local variable 3

Opcode:

62 (0x3e)

Instruction Format:

`istore_3`

Stack:

Before:..., value

After:...

Description:

The index three must be a valid index into the local variables of the current stack frame, and the value word on the top of the operand stack must be an int. To execute the `istore_3` instruction, the Java Virtual Machine pops the int value from the top of the operand stack and stores it into the local variable word at index three.

For more information about the `istore_3` instruction, see Chapter 10, "Stack Operations."

`isub`

- Subtract ints

Opcode:

100 (0x64)

Instruction Format:

`isub`

Stack:

Before:..., value1, value2

After:..., result

Description:

The top two words of the operand stack, value1 and value2, must be ints. To execute the isub instruction, the Java Virtual Machine pops value1 and value2, subtracts value2 from value1 (value1 - value2), and pushes the int result. If overflow occurs, result is the 32 lowest order bits of the true mathematical result represented in a sufficiently wide two's-complement format, and the sign bit of result may be different from the true mathematical result.

For more information about the isub instruction, see Chapter 12, "Integer Arithmetic."

iushr

- Perform logical right shift on int

Opcode:

124 (0x7c)

Instruction Format:

iushr

Stack:

Before:..., value1, value2

After:..., result

Description:

The top two words of the operand stack, value1 and value2, must be ints. To execute the iushr instruction, the Java Virtual Machine pops value1 and value2, shifts value1 right with zero extension by the number of bits specified in the 5 lowest order bits of value2 (from 0 to 31 bit positions), and pushes the int result.

For more information about the iushr instruction, see Chapter 13, "Logic."

ixor

- Perform boolean XOR on ints

Opcode:

130 (0x82)

Instruction Format:

ixor

Stack:

Before:..., value1, value2

After:..., result

Description:

The top two words of the operand stack, value1 and value2, must be ints. To execute the ixor instruction, the Java Virtual Machine pops value1 and value2, bitwise exclusive ORs them, and pushes the int result.

For more information about the ixor instruction, see Chapter 13, "Logic."

jsr

- Jump to subroutine

Opcode:

168 (0xa8)

Instruction Format:

jsr, branchbyte1, branchbyte2

Stack:

Before:...

After:..., address

Description:

To execute the jsr instruction, the Java Virtual Machine pushes the (program counter) address of the opcode immediately following the jsr instruction, the address word, onto the operand stack. The virtual machine then forms a signed 16-bit offset by calculating (branchbyte1

8) | branchbyte2. It calculates a target (program counter) address by adding the calculated offset to the address of the jsr opcode. The target address must be the address of an opcode within the same method as the jsr opcode. The virtual machine jumps to the target address and continues execution there.

For more information about the jsr instruction, see Chapter 18, "Finally Clauses."

jsr_w

- Jump to subroutine (wide index)

Opcode:

201 (0xc9)

Instruction Format:

jsr_w, branchbyte1, branchbyte2, branchbyte3, branchbyte4

Stack:

Before:...

After:..., address

Description:

To execute the jsr_w instruction, the Java Virtual Machine pushes the (program counter) address of the opcode immediately following the jsr_w instruction, the address word, onto the operand stack. The virtual machine then forms a signed 32-bit offset by calculating (branchbyte1 24) | (branchbyte2 16) | (branchbyte3 8) | branchbyte4. The virtual machine then calculates a target (program counter) address by adding the calculated offset to the address of the jsr_w opcode. The target address must be the address of an opcode within the same method as the jsr_w opcode. The virtual machine jumps to the target address and continues execution there.

Note that despite the 32-bit offset of the jsr_w instruction, Java methods are currently (in both the 1.0 and 1.1 releases) limited to 65,535 bytes by three items in the Java class file format: the sizes of the indexes in the LineNumberTable attribute, the LocalVariableTable attribute, and the Code attribute's exception_table item. According to the Java Virtual Machine specification, the 65,536 byte limit to Java methods may be raised in a future release. For more information about the jsr_w instruction, see Chapter 18, "Finally Clauses."

l2d

- Convert long to double

Opcode:

138 (0x8a)

Instruction Format:

l2d

Stack:

Before:..., value.word1, value.word2

After:..., result.word1, result.word2

Description:

The top two words of the operand stack must be a long. To execute the l2d instruction, the Java Virtual Machine pops the long value from the operand stack, converts the long to a double using the IEEE round-to-nearest mode, and pushes the double result.

Note that this instruction performs a widening primitive conversion. Because not all long values are exactly representable by a double, the conversion may result in a loss of precision.

For more information about the l2d instruction, see Chapter 11, "Type Conversion."

l2f

- Convert long to float

Opcode:

137 (0x89)

Instruction Format:

l2f

Stack:

Before:..., value.word1, value.word2

After:..., result

Description:

The top two words of the operand stack must be a long. To execute the l2f instruction, the Java Virtual Machine pops the long value from the operand stack, converts the long to a float using the IEEE round-to-nearest mode, and pushes the float result.

Note that this instruction performs a widening primitive conversion. Because not all long values are exactly representable by a float, the conversion may result in a loss of precision.

For more information about the l2f instruction, see Chapter 11, "Type Conversion."

l2i

- Convert long to int

Opcode:

136 (0x88)

Instruction Format:

l2i

Stack:

Before:..., value.word1, value.word2

After:..., result

Description:

The top two words of the operand stack must be a long. To execute the l2i instruction, the Java Virtual Machine pops long value from the operand stack, truncates the long to a int, and pushes the int result.

Note that this instruction performs a narrowing primitive conversion. As a result of this conversion, magnitude information may be lost and the sign bit may change.

For more information about the l2i instruction, see Chapter 11, "Type Conversion."

ladd

- Add longs

Opcode:

97 (0x61)

Instruction Format:

ladd

Stack:

Before:..., value1.word1, value1.word2, value2.word1, value2.word2

After:..., result.word1, result.word2

Description:

The top four words of the operand stack must be two longs, value1 and value2. To execute the ladd instruction, the Java Virtual Machine pops value1 and value2, adds them, and pushes the long result. If overflow occurs, result is the 64 lowest order bits of the true mathematical result represented in a sufficiently wide two's-complement format, and the sign of result is different from that of the true mathematical result.

For more information about the ladd instruction, see Chapter 12, "Integer Arithmetic."

laload

- Load long from array

Opcode:

47 (0x2f)

Instruction Format:

laload

Stack:

Before:..., arrayref, index

After:..., value.word1, value.word2

Description:

To execute the `laload` instruction, the Java Virtual Machine first pops two words from the operand stack. The `arrayref` word must be a reference that refers to an array of longs. The `index` word must be an `int`. The virtual machine retrieves from the `arrayref` array the long value specified by `index` and pushes it onto the operand stack.

If `arrayref` is null, the Java Virtual Machine throws `NullPointerException`. Otherwise, if `index` is not a legal index into the `arrayref` array, The virtual machine throws `ArrayIndexOutOfBoundsException`.

For more information about the `laload` instruction, see Chapter 15, "Objects and Arrays."

land

- Perform boolean AND on longs

Opcode:

127 (0x7f)

Instruction Format:

`land`

Stack:

Before:..., `value1.word1`, `value1.word2`, `value2.word1`, `value2.word2`

After:..., `result.word1`, `result.word2`

Description:

The top four words of the operand stack must be two longs, `value1` and `value2`. To execute the `land` instruction, the Java Virtual Machine pops `value1` and `value2`, bitwise ANDs them, and pushes the long result.

For more information about the `land` instruction, see Chapter 13, "Logic."

lastore

- Store into long array

Opcode:

80 (0x50)

Instruction Format:

lastore

Stack:

Before:..., arrayref, index, value.word1, value.word2

After:...

Description:

To execute the laload instruction, the Java Virtual Machine first pops four words from the operand stack. The arrayref word must be a reference that refers to an array of longs. The index word must be an int, and the value words must be a long. The virtual machine stores long value into the arrayref array location specified by index.

If arrayref is null, the Java Virtual Machine throws NullPointerException. Otherwise, if index is not a legal index into the arrayref array, The virtual machine throws ArrayIndexOutOfBoundsException.

For more information about the lastore instruction, see Chapter 15, "Objects and Arrays."

lcmp

- Compare longs

Opcode:

148 (0x94)

Instruction Format:

lcmp

Stack:

Before:..., value1.word1, value1.word2, value2.word1, value2.word2

After:..., result

Description:

The top four words of the operand stack must be two longs, value1 and value2. To execute the `lcmp` instruction, the Java Virtual Machine pops value1 and value2 off the operand stack and compares one against the other. If value1 equals value2, The virtual machine pushes onto the operand stack int result zero. Else, if value1 is greater than value2, The virtual machine pushes onto the operand stack int result one. Otherwise, if value1 is less than value2, The virtual machine pushes onto the operand stack int result negative one.

For more information about the `lcmp` instruction, see Chapter 16, "Control Flow."

`lconst_0`

- Push long constant 0

Opcode:

9 (0x9)

Instruction Format:

`lconst_0`

Stack:

Before:...

After:..., <0-word1, <0-word2

Description:

To execute the `lconst_0` instruction, the Java Virtual Machine pushes the long constant 0 onto the operand stack.

For more information about the `lconst_0` instruction, see Chapter 10, "Stack Operations."

`lconst_1`

- Push long constant 1

Opcode:

10 (0xa)

Instruction Format:

`lconst_1`

Stack:

Before:...

After:..., <1-word1, <1-word2

Description:

To execute the `lconst_1` instruction, the Java Virtual Machine pushes the long constant 1 onto the operand stack.

For more information about the `lconst_1` instruction, see Chapter 10, "Stack Operations."

`ldc`

- Push item from constant pool

Opcode:

`18 (0x12)`

Instruction Format:

`ldc, index`

Stack:

Before:...

After:..., item

Description:

The index operand must be a valid unsigned 8-bit index into the current constant pool. To execute the `ldc` instruction, the Java Virtual Machine first looks up the constant pool entry specified by the index operand. At constant pool entry index, The virtual machine must find either a `CONSTANT_Integer_info`, `CONSTANT_Float_info`, or `CONSTANT_String_info` entry. If it hasn't already, The virtual machine resolves the entry. If the entry is a `CONSTANT_Integer_info`, The virtual machine pushes the int value represented by the entry onto the operand stack. Else, if the entry is a `CONSTANT_Float_info`, The virtual machine pushes the float value represented by the entry onto

the operand stack. Otherwise, the entry is a `CONSTANT_String_info` entry, and The virtual machine pushes a reference to the interned String object that was produced by the process of resolving the entry onto the operand stack.

Note that the `ldc_w` instruction performs the same function, but offers a wide (16-bit) constant pool index.

For more information about the `ldc` instruction, see Chapter 10, "Stack Operations."

`ldc_w`

- Push item from constant pool (wide index)

Opcode:

19 (0x13)

Instruction Format:

`ldc_w, indexbyte1, indexbyte2`

Stack:

Before:...

After:..., item

Description:

To execute the `ldc_w` instruction, the Java Virtual Machine first forms an unsigned 16-bit index into the constant pool by calculating $(\text{indexbyte1} \ll 8) \mid \text{indexbyte2}$. The virtual machine then looks up the constant pool entry specified by the calculated index. The constant pool entry at that index must be either a `CONSTANT_Integer_info`, a `CONSTANT_Float_info`, or a `CONSTANT_String_info` entry. If it hasn't already, The virtual machine resolves the entry. If the entry is a `CONSTANT_Integer_info`, The virtual machine pushes the int value represented by the entry onto the operand stack. Else, if the entry is a `CONSTANT_Float_info`, The virtual machine pushes the float value it represents onto the operand stack. Otherwise, the entry is a `CONSTANT_String_info`, and The virtual machine pushes a reference to the interned String object that was produced by the process of resolving the entry onto the operand stack.

Note that the `ldc_w` instruction performs the same function as `ldc`, but offers a wide (16-bit) constant pool index as opposed to `ldc`'s 8-

bit constant pool index.

For more information about the ldc_w instruction, see Chapter 10, "Stack Operations."

ldc2_w

- Push long or double from constant pool (wide index)

Opcode:

20 (0x14)

Instruction Format:

ldc2_w, indexbyte1, indexbyte2

Stack:

Before:...

After:..., item.word1, item.word2

Description:

To execute the ldc2_w instruction, the Java Virtual Machine first forms an unsigned 16-bit index into the constant pool by calculating (indexbyte1 8) | indexbyte2. The virtual machine then looks up the constant pool entry specified by the calculated index. The constant pool entry at that index must be either a CONSTANT_Long_info or CONSTANT_Double_info entry. The virtual machine resolves the entry. If the entry is a CONSTANT_Long_info, The virtual machine pushes the long value represented by the entry onto the operand stack. Otherwise, it is a CONSTANT_Double_info entry, and The virtual machine pushes the double value represented by the entry onto the operand stack.

Note that there is no "ldc2" instruction that performs the same function as ldc2_w but with an 8-bit constant pool index. All two-word constants must be retrieved from the constant pool with an ldc2_w instruction, which has a wide (16-bit) constant pool index.

For more information about the ldc2_w instruction, see Chapter 10, "Stack Operations."

ldiv

- Divide longs

Opcode:

109 (0x6d)

Instruction Format:

`ldiv`

Stack:

Before:..., value1.word1, value1.word2, value2.word1, value2.word2

After:..., result.word1, result.word2

Description:

The top four words of the operand stack must be two longs, value1 and value2. To execute the `ldiv` instruction, the Java Virtual Machine pops value1 and value2, integer divides value1 by value2 (value1 / value2), and pushes the long result.

Integer division rounds the magnitude of the true mathematical quotient towards zero to the nearest integer. If the magnitude of the denominator is greater than that of the numerator, the long result is zero. Else, with one special exception, the sign of result is positive if the signs of the numerator and denominator are the same, negative if they are different. The exception to this rule is when the numerator is the smallest negative integer that can be represented by a long and the denominator is -1. For this division, the true mathematical result is one greater than the largest positive integer that can be represented by a long. As a consequence, the division overflows and the result is equal to the numerator.

If value2 (the denominator) is zero, the Java Virtual Machine throws `ArithmeticException`.

For more information about the `ldiv` instruction, see Chapter 12, "Integer Arithmetic."

`lload`

- Load long from local variable

Opcode:

22 (0x16)

Instruction Format:

`lload, index`

Stack:

Before:...

After:..., value.word1, value.word2

Description:

The index operand, which serves as an 8-bit unsigned index into the local variables of the current frame, must specify the first of two consecutive local variable words that contain a long. To execute the lload instruction, the Java Virtual Machine pushes onto the operand stack the long contained in the two consecutive local variable words specified by index and index + 1.

Note that the wide instruction can precede the lload instruction, to allow a local variable to be accessed with a 16-bit unsigned offset.

For more information about the lload instruction, see Chapter 10, "Stack Operations."

lload_0

- Load long from local variable 0

Opcode:

30 (0x1e)

Instruction Format:

`lload_0`

Stack:

Before:...

After:..., value.word1, value.word2

Description:

The two consecutive local variable words at indexes zero and one must contain a long. To execute the lload_0 instruction, the Java Virtual Machine pushes onto the operand stack the long value contained in local variable words zero and one.

For more information about the lload_0 instruction, see Chapter 10, "Stack Operations."

lload_1

- Load long from local variable 1

Opcode:

31 (0x1f)

Instruction Format:

lload_1

Stack:

Before:...

After:..., value.word1, value.word2

Description:

The two consecutive local variable words at indexes one and two must contain a long. To execute the lload_1 instruction, the Java Virtual Machine pushes onto the operand stack the long value contained in local variable words one and two.

For more information about the lload_1 instruction, see Chapter 10, "Stack Operations."

lload_2

- Load long from local variable 2

Opcode:

32 (0x20)

Instruction Format:

lload_2

Stack:

Before:...

After:..., value.word1, value.word2

Description:

The two consecutive local variable words at indexes two and three must contain a long. To execute the `lload_2` instruction, the Java Virtual Machine pushes onto the operand stack the long value contained in local variable words two and three.

For more information about the `lload_2` instruction, see Chapter 10, "Stack Operations."

`lload_3`

- Load long from local variable 3

Opcode:

33 (0x21)

Instruction Format:

`lload_3`

Stack:

Before:...

After:..., value.word1, value.word2

Description:

The two consecutive local variable words at indexes three and four must contain a long. To execute the `lload_3` instruction, the Java Virtual Machine pushes onto the operand stack the long value contained in local variable words three and four.

For more information about the `lload_3` instruction, see Chapter 10, "Stack Operations."

`lmul`

- Multiply longs

Opcode:

105 (0x69)

Instruction Format:

`lmul`

Stack:

Before:..., value1.word1, value1.word2, value2.word1, value2.word2

After:..., result.word1, result.word2

Description:

The top four words of the operand stack must be two longs, value1 and value2. To execute the `lmul` instruction, the Java Virtual Machine pops value1 and value2, multiplies them, and pushes the long result. If overflow occurs, result is the 64 lowest order bits of the true mathematical result represented in a sufficiently wide two's-complement format, and the sign of result may be different from that of the true mathematical result.

For more information about the `lmul` instruction, see Chapter 12, "Integer Arithmetic."

lneg

- Negate long

Opcode:

117 (0x75)

Instruction Format:

`lneg`

Stack:

Before:..., value.word1, value.word2

After:..., result.word1, result.word2

Description:

The top two words of the operand stack must be a long, value. To execute the `lneg` instruction, the Java Virtual Machine pops value, negates it, and pushes the long result.

The result produced by an `lneg` instruction is the same number that would be produced by subtracting value from zero with the `lsub` instruction. As a consequence, when value is the smallest negative

integer that can be represented by an long, the negation overflows. For this negation, the true mathematical result is one greater than the largest positive integer that can be represented by an long, and the actual result is equal to value with no change in sign.

For more information about the lneg instruction, see Chapter 12, "Integer Arithmetic."

lookupswitch

- Access jump table by key match and jump

Opcode:

171 (0xab)

Instruction Format:

lookupswitch, ...0-3 byte pad..., defaultbyte1, defaultbyte2, defaultbyte3, defaultbyte4, npairs1, npairs2, npairs3, npairs4, ..match-offset pairs..

Stack:

Before:..., key

After:...

Description:

The lookupswitch opcode is followed by zero to three bytes of padding-- enough so that the byte immediately following the padding starts at an address that is a multiple of four bytes from the beginning of the method. Each padding byte is a zero. Immediately following the padding is a signed 32-bit default branch offset, default. Following default is npairs, a signed count of the number of case value/branch offset pairs embedded in this lookupswitch instruction. The value of npairs must be greater than or equal to zero. Following npairs are the case value/branch offset pairs themselves. For each pair, the signed 32-bit case value, match, precedes the signed 32-bit branch offset, offset. The virtual machine calculates all of these signed 32-bit values from the four individual bytes as (byte1 24) | (byte2 16) | (byte3 8) | byte4.

The top word of the operand stack, key, must be an int. To execute the lookupswitch instruction, the Java Virtual Machine pops key off the operand stack and compares it to the match values. If the key is equal to one of the match values, The virtual machine calculates a target (program counter) address by adding the signed offset that corresponds to the

matching match value to the address of the lookupswitch opcode. The target address must be the address of an opcode within the same method as the lookupswitch opcode. The virtual machine jumps to the target address and continues execution there.

For more information about the lookupswitch instruction, see Chapter 16, "Control Flow."

lor

- Perform boolean OR on longs

Opcode:

129 (0x81)

Instruction Format:

lor

Stack:

Before:..., value1.word1, value1.word2, value2.word1, value2.word2

After:..., result.word1, result.word2

Description:

The top four words of the operand stack must be two longs, value1 and value2. To execute the lor instruction, the Java Virtual Machine pops value1 and value2, bitwise ORs them, and pushes the long result.

For more information about the lor instruction, see Chapter 13, "Logic."

lrem

- Calculate remainder of division of longs

Opcode:

113 (0x71)

Instruction Format:

lrem

Stack:

Before:..., value1.word1, value1.word2, value2.word1, value2.word2

After:..., result.word1, result.word2

Description:

*The top four words of the operand stack must be two longs, value1 and value2. To execute the lrem instruction, the Java Virtual Machine pops value1 and value2, calculates the integer remainder, and pushes the long result. The integer remainder equals value1 - (value1 / value2) * value2.*

The lrem instruction implements Java's remainder operator, %, on longs. The lrem behaves such that the Java expression shown below is always true, where n and d are any two longs:

begin

*$(n/d)*d + (n\%d) == n$*

end

This behavior means that the result of an lrem instruction always takes the same sign as the numerator, which is popped off the operand stack as value1.

If value2 (the denominator) is zero, the Java Virtual Machine throws ArithmeticException.

For more information about the lrem instruction, see Chapter 12, "Integer Arithmetic."

lreturn

- Return long from method

Opcod:

173 (0xad)

Instruction Format:

lreturn

Stack:

Before:..., value.word1, value.word2

After:[empty]

Description:

The return type of the returning method must be long. The top two words of the operand stack must be a long. To execute the lreturn instruction, the Java Virtual Machine pops the long value from the operand stack of the current frame and pushes it onto the operand stack of the invoking method's frame. The virtual machine discards any other words that may still be on the returning method's frame. If the returning method is synchronized, the monitor that was acquired when the method was invoked is released. The invoking method's frame is made current, and The virtual machine continues execution in the invoking method.

For more information about monitors, see Chapter 20, "Thread Synchronization." For more information about the lreturn instruction, see Chapter 19, "Method Invocation and Return."

lshl

- Perform left shift on long

Opcode:

121 (0x79)

Instruction Format:

lshl

Stack:

Before:..., value1.word1, value1.word2, value2

After:..., result.word1, result.word2

Description:

The top word of the operand stack, value2, must be an int. The next two words down must be a long, value1. To execute the lshl instruction, the Java Virtual Machine pops value1 and value2, shifts value1 left by the number of bits specified in the 6 lowest order bits of value2 (from 0 to 63 bit positions), and pushes the long result.

For more information about the lshl instruction, see Chapter 13, "Logic."

lshr

- Perform arithmetic right shift on long

Opcode:

123 (0x7b)

Instruction Format:

lshr

Stack:

Before:..., value1.word1, value1.word2, value2

After:..., result.word1, result.word2

Description:

The top word of the operand stack, *value2*, must be an int. The next two words down must be a long, *value1*. To execute the *lshr* instruction, the Java Virtual Machine pops *value1* and *value2*, shifts *value1* right with sign extension by the number of bits specified in the 6 lowest order bits of *value2* (from 0 to 63 bit positions), and pushes the long result.

For more information about the *lshr* instruction, see Chapter 13, "Logic."

lstore

- Store long into local variable

Opcode:

55 (0x37)

Instruction Format:

lstore, index

Stack:

Before:..., value.word1, value.word2

After:...

Description:

The index operand must specify a valid 8-bit unsigned index into the local variables of the current frame. The top two words of the operand stack must be a long. To execute the `lstore` instruction, the Java Virtual Machine pops the long value from the top of the operand stack and stores it into the two consecutive local variable words at indexes `index` and `index + 1`.

Note that the wide instruction can precede the `lstore` instruction, to enable a value to be stored into a local variable specified by a 16-bit unsigned offset.

For more information about the `lstore` instruction, see Chapter 10, "Stack Operations."

`lstore_0`

- Store long into local variable 0

Opcode:

63 (0x3f)

Instruction Format:

`lstore_0`

Stack:

Before:..., `value.word1`, `value.word2`

After:...

Description:

The indexes zero and one must be valid indexes into the local variables of the current stack frame. The top two words on the operand stack must be a long. To execute the `lstore_0` instruction, the Java Virtual Machine pops the long value from the top of the operand stack and stores it into the two consecutive local variable words at indexes zero and one.

For more information about the `lstore_0` instruction, see Chapter 10, "Stack Operations."

`lstore_1`

- Store long into local variable 1

Opcode:

64 (0x40)

Instruction Format:

`lstore_1`

Stack:

Before:..., `value.word1`, `value.word2`

After:...

Description:

The indexes one and two must be valid indexes into the local variables of the current stack frame. The top two words on the operand stack must be a long. To execute the `lstore_1` instruction, the Java Virtual Machine pops the long value from the top of the operand stack and stores it into the two consecutive local variable words at indexes one and two.

For more information about the `lstore_1` instruction, see Chapter 10, "Stack Operations."

`lstore_2`

- Store long into local variable 2

Opcode:

65 (0x41)

Instruction Format:

`lstore_2`

Stack:

Before:..., `value.word1`, `value.word2`

After:...

Description:

The indexes two and three must be valid indexes into the local variables of the current stack frame. The top two words on the operand stack must be a long. To execute the lstore_2 instruction, the Java Virtual Machine pops the long value from the top of the operand stack and stores it into the two consecutive local variable words at indexes two and three.

For more information about the lstore_2 instruction, see Chapter 10, "Stack Operations."

lstore_3

- Store long into local variable 3

Opcode:

66 (0x42)

Instruction Format:

lstore_3

Stack:

Before:..., value.word1, value.word2

After:...

Description:

The indexes three and four must be valid indexes into the local variables of the current stack frame. The top two words on the operand stack must be a long. To execute the lstore_3 instruction, the Java Virtual Machine pops the long value from the top of the operand stack and stores it into the two consecutive local variable words at indexes three and four.

For more information about the lstore_3 instruction, see Chapter 10, "Stack Operations."

lsub

- Subtract longs

Opcode:

101 (0x65)

Instruction Format:

`lsub`

Stack:

Before:..., `value1.word1`, `value1.word2`, `value2.word1`, `value2.word2`

After:..., `result.word1`, `result.word2`

Description:

The top four words of the operand stack must be two longs, `value1` and `value2`. To execute the `lsub` instruction, the Java Virtual Machine pops `value1` and `value2`, subtracts `value2` from `value1` (`value1 - value2`), and pushes the long result. If overflow occurs, result is the 64 lowest order bits of the true mathematical result represented in a sufficiently wide two's-complement format, and the sign bit of result is different from the true mathematical result.

For more information about the `lsub` instruction, see Chapter 12, "Integer Arithmetic."

lushr

- Perform logical right shift on long

Opcode:

125 (0x7d)

Instruction Format:

`lushr`

Stack:

Before:..., `value1.word1`, `value1.word2`, `value2.word1`, `value2.word2`

After:..., `result.word1`, `result.word2`

Description:

The top word of the operand stack, `value2`, must be an int. The next two words down must be a long, `value1`. To execute the `lushr` instruction, the Java Virtual Machine pops `value1` and `value2`, shifts `value1` right with zero extension by the number of bits specified in the 6 lowest order bits of `value2` (from 0 to 63 bit positions), and

pushes the long result.

For more information about the lushr instruction, see Chapter 13, "Logic."

lxor

- Perform boolean XOR on longs

Opcode:

131 (0x83)

Instruction Format:

lxor

Stack:

Before:..., value1.word1, value1.word2, value2.word1, value2.word2

After:..., result.word1, result.word2

Description:

The top four words of the operand stack must be two longs, value1 and value2. To execute the lxor instruction, the Java Virtual Machine pops value1 and value2, bitwise exclusive ORs them, and pushes the long result.

For more information about the lxor instruction, see Chapter 13, "Logic."

monitorenter

- Enter and acquire object monitor

Opcode:

194 (0xc2)

Instruction Format:

monitorenter

Stack:

Before:..., objectref

After:...

Description:

The top word of the operand stack, objectref, must be a reference. To execute the monitorenter instruction, the Java Virtual Machine pops objectref and, on behalf of the current thread, acquires the monitor associated with objectref.

If the objectref word is null, The virtual machine throws NullPointerException.

For more information about the monitorenter instruction, see Chapter 20, "Thread Synchronization."

monitorexit

- Release and exit object monitor

Opcode:

195 (0xc3)

Instruction Format:

monitorexit

Stack:

Before:..., objectref

After:...

Description:

The top word of the operand stack, objectref, must be a reference. To execute the monitorenter instruction, the Java Virtual Machine pops objectref and, on behalf of the current thread, releases and exits the monitor associated with objectref.

If the objectref word is null, The virtual machine throws NullPointerException.

For more information about the monitorexit instruction, see Chapter 20, "Thread Synchronization."

multianewarray

- Allocate new multi-dimensional array

Opcode:

197 (0xc5)

Instruction Format:

multianewarray, indexbyte1, indexbyte2, dimensions

Stack:

Before:..., count1, [count2,...]

After:arrayref

Description:

The dimensions operand, an unsigned byte that indicates the number of dimensions in the array to create, must be greater than or equal to one. The top dimensions words of the operand stack, count1, [count2,...], must contain ints that have nonnegative values. Each of these words gives the number of elements in one dimension of the array. For example, consider an array declared as:

begin cc.

int[][][] example = new int[3][4][5];

end

For this array, dimensions would be equal to three and the operand stack would contain three "count" words. The count1 word would be three, count2 four, and count3 five.

To execute the multianewarray instruction, the Java Virtual Machine first forms an unsigned 16-bit index into the constant pool by calculating (indexbyte1 8) | indexbyte2. The virtual machine then looks up the constant pool entry specified by the calculated index. The constant pool entry at that index must be a CONSTANT_Class_info entry. If it hasn't already, The virtual machine resolves the entry. The entry must be an array type of dimensionality greater than or equal to dimensions.

If the resolution is successful, the Java Virtual Machine creates the dimensions-dimensional array on the heap. The virtual machine initializes the elements of the first-dimension array with references to the second-dimension arrays; it initializes the elements of each of the second-dimension array with references to the third-dimension

arrays, and so on. The virtual machine initializes the elements of the last-dimension arrays with their default initial values. Lastly, the Java Virtual Machine pushes a reference to the new dimensions-dimensional array onto the operand stack.

As a result of executing this instruction, The virtual machine may throw any of the linking errors listed in Chapter 8, "The Linking Model," as possible during resolution of a `CONSTANT_Class_info` entry. If resolution succeeds and the array's component type is a reference type, but the current class does not have permission to access that reference type, The virtual machine throws `IllegalAccessException`. Otherwise, if any of the counts popped off of the operand stack have a negative value, the Java Virtual Machine throws `NegativeArraySizeException`.

Note that the dimensionality of the array type specified in the resolved `CONSTANT_Class_info` entry need not be equal to the dimensions operand--it can also be greater than dimensions. For example, the name of the array class for a three-dimensional array of ints is `"[[[I`. The actual constant pool entry referenced by a `multianewarray` instruction that creates a three-dimensional array of ints must have at least three dimensions, but may have more. For instance, resolved array types of `"[[[I`, `"[[[[I`, and `"[[[[[[I` would all yield three-dimensional arrays of int so long as the dimensions operand is three. This flexibility in specifying multidimensional array types to the `multianewarray` instruction can reduce the number of entries required in the constant pool of some classes.

For more information about the `multianewarray` instruction, see Chapter 15, "Objects and Arrays."

new

- Create a new object

Opcode:

187 (0xbb)

Instruction Format:

`new, indexbyte1, indexbyte2`

Stack:

Before:...

After:..., objectref

Description:

To execute the new instruction, the Java Virtual Machine first forms an unsigned 16-bit index into the constant pool by calculating $(\text{indexbyte1} \ll 8) \mid \text{indexbyte2}$. The virtual machine then looks up the constant pool entry specified by the calculated index. The constant pool entry at that index must be a `CONSTANT_Class_info` entry. If it hasn't already, The virtual machine resolves the entry. The entry must be a class type, not an interface or array type. The virtual machine allocates sufficient memory from the heap for the new object's image, and sets the object's instance variables to their default initial values. Lastly, The virtual machine pushes `objectref`, a reference to the new object, onto the operand stack.

As a result of executing this instruction, The virtual machine may throw any of the linking errors listed in Chapter 8, "The Linking Model," as possible during resolution of a `CONSTANT_Class_info` entry. If resolution succeeds, but the resolved type is an interface, abstract class, or array, The virtual machine throws an `InstantiationError`. Else, if the current class doesn't have permission to access the resolved class, The virtual machine throws an `IllegalAccessException`.

For more information about the new instruction, see Chapter 15, "Objects and Arrays."

newarray

- Allocate new array of primitive type components

Opcode:

188 (0xbc)

Instruction Format:

`newarray, atype`

Stack:

Before:..., count

After:arrayref

Description:

The top word of the operand stack, `count`, must be an `int`. The `atype` operand, which is used to indicate the array type, must take one of

the values shown in Table A-2. To execute the `newarray` instruction, the Java Virtual Machine pops `count` and creates on the heap an array of size `count` of the primitive type specified by `atype`. The virtual machine initializes each array element to its default initial value and pushes `arrayref`, a reference to the new array, onto the operand stack.

Table A-2. Values for `atype`

2 columns

Array Type `atype`

`T_BOOLEAN`4

`T_CHAR`5

`T_FLOAT`6

`T_DOUBLE`7

`T_BYTE`8

`T_SHORT`9

`T_INT`10

`T_LONG`11

end table

If `count` is less than zero, the Java Virtual Machine throws `NegativeArraySizeException`.

For more information about the `newarray` instruction, see Chapter 15, "Objects and Arrays."

`nop`

- Do nothing

Opcode:

0 (0x0)

Instruction Format:

`nop`

Stack:

No change

Description:

To execute the nop instruction, the Java Virtual Machine takes a coffee break.

For more information about the nop instruction, see Chapter 10, "Stack Operations."

pop

- Pop top stack word

Opcode:

87 (0x57)

Instruction Format:

pop

Stack:

Before:..., word

After:...

Description:

To execute the pop instruction, the Java Virtual Machine pops the top word from the operand stack. This instruction can be used to pop any single-word value from the top of the operand stack. It must not be used to remove half of a dual word value (long or double) from the top of the operand stack.

For more information about the pop instruction, see Chapter 10, "Stack Operations."

pop2

- Pop top two stack words

Opcode:

88 (0x58)

Instruction Format:

pop2

Stack:

Before:..., word2, word1

After:...

Description:

To execute the pop2 instruction, the Java Virtual Machine pops the top two words from the operand stack. This instruction can be used to pop any dual-word value from the top of the operand stack, or any two single-word values. It must not be used to remove one single-word value and half of a dual word value (long or double) from the top of the operand stack.

For more information about the pop2 instruction, see Chapter 10, "Stack Operations."

putfield

- Set field in object

Opcode:

181 (0xb5)

Instruction Format:

putfield, indexbyte1, indexbyte2

Stack:

Before:..., objectref, value

After:...

or

Before:..., objectref, value.word1, value.word2

After:...

Description:

To execute the `putfield` instruction, the Java Virtual Machine first forms an unsigned 16-bit index into the constant pool by calculating $(\text{indexbyte1} \ll 8) \mid \text{indexbyte2}$. The virtual machine then looks up the constant pool entry specified by the calculated index. The constant pool entry at that index must be a `CONSTANT_Fieldref_info` entry. If it hasn't already, the virtual machine resolves the entry, which yields the field's width and the field's offset from the beginning of the object image.

The type of the single or double-word value occupying the top of the stack must be compatible with the descriptor of the resolved field. If the resolved field's descriptor is `byte`, `short`, `char`, `boolean`, or `int`, the type of value must be `int`. If the resolved field's descriptor is `long`, the type of value must be `long`. If the resolved field's descriptor is `float`, the type of value must be `float`. If the resolved field's descriptor is `double`, the type of value must be `double`. If the resolved field's descriptor is a reference type, the type of value must be reference and must be assignment-compatible the resolved descriptor type. The `objectref` word must be a reference.

The virtual machine pops value and `objectref` and assigns value to the appropriate field in the object pointed to by `objectref`.

As a result of executing this instruction, the virtual machine may throw any of the linking errors listed in Chapter 8, "The Linking Model," as possible during resolution of a `CONSTANT_Fieldref_info` entry. As part of the process of resolving the `CONSTANT_Fieldref_info` entry, the virtual machine checks whether the field's access permission enables the current class to access the field. If the field is protected, the virtual machine makes certain the field is a member of either the current class or a superclass of the current class, and that the class of the object pointed to by `objectref` is either the current class or a subclass of the current class. If not (or if there is any other access permission problem), the virtual machine throws `IllegalAccessException`. Else, if the field exists and is accessible from the current class, but the field is static, the virtual machine throws `IncompatibleClassChangeError`. Otherwise, if `objectref` is null, the virtual machine throws `NullPointerException`.

For more information about the `putfield` instruction, see Chapter 15, "Objects and Arrays."

`putstatic`

- Set static field in class

Opcode:

179 (0xb3)

Instruction Format:

putstatic, indexbyte1, indexbyte2

Stack:

Before:..., value

After:...

or

Before:..., value.word1, value.word2

After:...

Description:

To execute the *putstatic* instruction, the Java Virtual Machine first forms an unsigned 16-bit index into the constant pool by calculating $(\text{indexbyte1} \ll 8) \mid \text{indexbyte2}$. The virtual machine then looks up the constant pool entry specified by the calculated index. The constant pool entry at that index must be a *CONSTANT_Fieldref_info* entry. If it hasn't already, The virtual machine resolves the entry.

The type of the single or double-word value occupying the top of the stack must be compatible with the descriptor of the resolved field. If the resolved field's descriptor is *byte*, *short*, *char*, *boolean*, or *int*, the type of value must be *int*. If the resolved field's descriptor is *long*, the type of value must be *long*. If the resolved field's descriptor is *float*, the type of value must be *float*. If the resolved field's descriptor is *double*, the type of value must be *double*. If the resolved field's descriptor is a reference type, the type of value must be reference and must be assignment-compatible the resolved descriptor type.

The virtual machine pops and assigns value to the appropriate static field.

As a result of executing this instruction, The virtual machine may throw any of the linking errors listed in Chapter 8, "The Linking Model," as possible during resolution of a *CONSTANT_Fieldref_info* entry. As part of the process of resolving the *CONSTANT_Fieldref_info* entry, The virtual machine checks whether the field's access permission enables the current class to access the field. If the field is protected, The virtual machine makes certain the field is a member of the either the current class or a superclass of the current class. If not (or if there is any other access permission problem), The virtual machine throws *IllegalAccessError*. Else, if the field

exists and is accessible from the current class, but the field is not static, The virtual machine throws `IncompatibleClassChangeError`.

For more information about the `putstatic` instruction, see Chapter 15, "Objects and Arrays."

`ret`

- Return from subroutine

`Opcode:`

169 (0xa9)

`Instruction Format:`

ret, index

`Stack:`

No change

`Description:`

The index operand is an 8-bit unsigned offset into the local variables. The local variable word specified by index must be a returnAddress. To execute the `ret` instruction, the Java Virtual Machine sets the program counter to the returnAddress value stored in local variable index, and continues execution there. (In other words, The virtual machine jumps to the returnAddress.)

Note that the wide instruction can precede the `ret` instruction, to allow a local variable to be accessed with a 16-bit unsigned offset.

For more information about the `ret` instruction, see Chapter 18, "Finally Clauses."

`return`

- Return (void) from method

`Opcode:`

177 (0xb1)

`Instruction Format:`

return

Stack:

Before:...

After:[empty]

Description:

The return type of the returning method must be void. To execute the return instruction, the Java Virtual Machine discards any words that may still be on the returning method's frame. If the returning method is synchronized, the monitor that was acquired when the method was invoked is released. The invoking method's frame is made current, and the virtual machine continues execution in the invoking method.

For more information about monitors, see Chapter 20, "Thread Synchronization." For more information about the return instruction, see Chapter 19, "Method Invocation and Return."

aload

- Load short from array

Opcode:

53 (0x35)

Instruction Format:

aload

Stack:

Before:..., arrayref, index

After:..., value

Description:

To execute the aload instruction, the Java Virtual Machine first pops two words from the operand stack. The arrayref word must be a reference that refers to an array of shorts. The index word must be an int. The virtual machine retrieves from the arrayref array the short value specified by index, sign-extends it to an int, and pushes it onto the operand stack.

If arrayref is null, the Java Virtual Machine throws NullPointerException. Otherwise, if index is not a legal index into

*the arrayref array, The virtual machine throws
ArrayIndexOutOfBoundsException.*

*For more information about the saload instruction, see Chapter 15,
"Objects and Arrays."*

sastore

- Store into short array

Opcode:

86 (0x56)

Instruction Format:

sastore

Stack:

Before:..., array, index, value

After:...

Description:

To execute the sastore instruction, the Java Virtual Machine first pops three words from the operand stack. The arrayref word must be a reference that refers to an array of shorts. The index and value words must be ints. The virtual machine truncates the int value to a short and stores it into the arrayref array location specified by index.

*If arrayref is null, the Java Virtual Machine throws
NullPointerException. Otherwise, if index is not a legal index into
the arrayref array, The virtual machine throws
ArrayIndexOutOfBoundsException.*

*For more information about the sastore instruction, see Chapter 15,
"Objects and Arrays."*

sipush

- Push 16-bit signed integer

Opcode:

17 (0x11)

Instruction Format:

sipush, byte1, byte2

Stack:

Before:...

After:..., value

Description:

To execute the sipush instruction, the Java Virtual Machine first forms an intermediate 16-bit signed integer from byte1 and byte2 by calculating (byte1 8) | byte2. The virtual machine then sign-extends the intermediate 16-bit signed integer to an int, and pushes the resulting int value onto the operand stack.

For more information about the sipush instruction, see Chapter 10, "Stack Operations."

swap

- Swap top two stack words

Opcode:

95 (0x5f)

Instruction Format:

swap

Stack:

Before:..., word2, word1

After:..., word1, word2

Description:

To execute the swap instruction, the Java Virtual Machine swaps the top two words of the operand stack. Both word1 and word2 must be single-word values.

For more information about the swap instruction, see Chapter 10, "Stack Operations."

tableswitch

- Access jump table by index and jump

Opcode:

170 (0xaa)

Instruction Format:

tableswitch, ...0-3 byte pad..., defaultbyte1, defaultbyte2, defaultbyte3, defaultbyte4, lowbyte1, lowbyte2, lowbyte3, lowbyte4, highbyte1, highbyte2, highbyte3, highbyte4, ...jump offsets...

Stack:

Before:..., index

After:...

Description:

The tableswitch opcode is followed by zero to three bytes of padding-- enough so that the byte immediately following the padding starts at an address that is a multiple of four bytes from the beginning of the method. Each padding byte is a zero. Immediately following the padding is a signed 32-bit default branch offset, default. Following default are two signed 32-bit endpoints of the range of case values embedded in this tableswitch instruction: low, the low endpoint, and high, the high endpoint. The value of low must be less than or equal to the value of high. Following low and high are high - low + 1 signed 32-bit branch offsets--one branch offset for high, one for low, and one for each integer case value in between high and low. These offsets serve as a zero-based jump table. The branch offset for low, the first entry in the jump table, immediately follows the high endpoint value. The virtual machine calculates all of these signed 32-bit values from the four individual bytes as (byte1 24) | (byte2 16) | (byte3 8) | byte4.

The top word of the operand stack, index, must be an int. To execute the tableswitch instruction, the Java Virtual Machine pops index off the operand stack, compares it to low and high, and selects a branch offset. If index is less than low and greater than high, The virtual machine selects default as the branch offset. Else, The virtual machine selects the branch offset at position index - low in the jump table. Once it has selected a branch offset, The virtual machine calculates a target (program counter) address by adding the signed branch offset to the address of the tableswitch opcode. The target address must be the address of an opcode within the same method as

the `tableswitch` opcode. The virtual machine jumps to the target address and continues execution there.

For more information about the `tableswitch` instruction, see Chapter 16, "Control Flow."

wide

- Extend a local variable index with additional bytes

Opcode:

196 (0xc4)

Instruction Format:

`wide, <opcode, indexbyte1, indexbyte2`

or

`wide, iinc, indexbyte1, indexbyte2, constbyte1, constbyte2`

Stack:

To see how the stack changes when an instruction is modified by `wide`, see the entry for the unmodified instruction. The stack change for an instruction modified by `wide` is identical to the stack change for that same instruction unmodified.

Description:

The `wide` opcode modifies instructions that reference the local variables, extending the modified instruction's unsigned 8-bit local variable index to an unsigned 16-bit index. As shown above, a `wide` instruction comes in two formats. When the `wide` opcode modifies `iload`, `fload`, `aload`, `lload`, `dload`, `istore`, `fstore`, `astore`, `lstore`, `dstore`, or `ret`, the instruction has the first format. When the `wide` opcode modifies `iinc`, the instruction has the second format.

To execute any `wide` instruction, the Java Virtual Machine first forms an unsigned 16-bit index into the local variables by calculating $(\text{indexbyte1} \ll 8) \mid \text{indexbyte2}$. Irrespective of the opcode that `wide` modifies, the calculated index must be a valid index into the local variables of the current frame. If the `wide` opcode modifies `lload`, `dload`, `lstore`, or `dstore`, then one more than the calculated index must also be a valid index into the local variables. Given a valid unsigned 16-bit (wide) local variable index, The virtual machine executes the modified instruction using the wide index.

When a wide opcode modifies an `iinc` instruction, the format of the `iinc` instruction changes. An unmodified `iinc` instruction has two operands, an 8-bit unsigned local variable index, and an 8-bit signed increment `const`. When modified by `wide`, however, both the `iinc` instruction's local variable index and its increment are extended by an extra byte. To execute this instruction, The virtual machine forms a signed 16-bit increment by calculating $(\text{constbyte1} \ll 8) \mid \text{constbyte2}$.

Note that from the perspective of the bytecode verifier, the opcode modified by `wide` is seen as an operand to the wide opcode. The bytecodes are not allowed to treat opcodes modified by `wide` independently. For example, it is illegal for a `goto` instruction to jump directly to an opcode modified by `wide`.

For more information about the wide instruction, see Chapter 10, "Stack Operations."

Appendix B

Opcode Mnemonic by Function Group

Stack Operations (Chapter 10)

Instructions that push a constant onto the stack

<code>aconst_null</code>	Push <code>null</code> object reference
<code>iconst_m1</code>	Push <code>int</code> constant -1
<code>iconst_0</code>	Push <code>int</code> constant 0
<code>iconst_1</code>	Push <code>int</code> constant 1
<code>iconst_2</code>	Push <code>int</code> constant 2
<code>iconst_3</code>	Push <code>int</code> constant 3
<code>iconst_4</code>	Push <code>int</code> constant 4
<code>iconst_5</code>	Push <code>int</code> constant 5
<code>lconst_0</code>	Push <code>long</code> constant 0
<code>lconst_1</code>	Push <code>long</code> constant 1
<code>fconst_0</code>	Push <code>float</code> constant 0.0
<code>fconst_1</code>	Push <code>float</code> constant 1.0
<code>fconst_2</code>	Push <code>float</code> constant 2.0
<code>dconst_0</code>	Push <code>double</code> constant 0.0
<code>dconst_1</code>	Push <code>double</code> constant 1.0
<code>bipush</code>	Push 8-bit signed integer
<code>sipush</code>	Push 16-bit signed integer
<code>ldc</code>	Push item from constant pool
<code>ldcw</code>	Push item from constant pool (wide index)
<code>ldc2_w</code>	Push <code>long</code> or <code>double</code> from constant pool (wide index)

Instructions that load a local variable onto the stack

<code>iload</code>	Load <code>int</code> from local variable
<code>lload</code>	Load <code>long</code> from local variable

<code>fload</code>	Load <code>float</code> from local variable
<code>dload</code>	Load <code>double</code> from local variable
<code>aload</code>	Load <code>reference</code> from local variable
<code>iload_0</code>	Load <code>int</code> from local variable 0
<code>iload_1</code>	Load <code>int</code> from local variable 1
<code>iload_2</code>	Load <code>int</code> from local variable 2
<code>iload_3</code>	Load <code>int</code> from local variable 3
<code>lload_0</code>	Load <code>long</code> from local variable 0
<code>lload_1</code>	Load <code>long</code> from local variable 1
<code>lload_2</code>	Load <code>long</code> from local variable 2
<code>lload_3</code>	Load <code>long</code> from local variable 3
<code>fload_0</code>	Load <code>float</code> from local variable 0
<code>fload_1</code>	Load <code>float</code> from local variable 1
<code>fload_2</code>	Load <code>float</code> from local variable 2
<code>fload_3</code>	Load <code>float</code> from local variable 3
<code>dload_0</code>	Load <code>double</code> from local variable 0
<code>dload_1</code>	Load <code>double</code> from local variable 1
<code>dload_2</code>	Load <code>double</code> from local variable 2
<code>dload_3</code>	Load <code>double</code> from local variable 3
<code>aload_0</code>	Load <code>reference</code> from local variable 0
<code>aload_1</code>	Load <code>reference</code> from local variable 1
<code>aload_2</code>	Load <code>reference</code> from local variable 2
<code>aload_3</code>	Load <code>reference</code> from local variable 3
<code>iaload</code>	Load <code>int</code> from array
<code>laload</code>	Load <code>long</code> from array
<code>faload</code>	Load <code>float</code> from array
<code>daload</code>	Load <code>double</code> from array
<code>aaload</code>	Load <code>reference</code> from array
<code>baload</code>	Load <code>byte</code> or <code>boolean</code> from array
<code>caload</code>	Load <code>char</code> from array
<code>saload</code>	Load <code>short</code> from array

Instructions that store a value from the stack into a local variable

<code>istore</code>	Store <code>int</code> into local variable
<code>lstore</code>	Store <code>long</code> into local variable
<code>fstore</code>	Store <code>float</code> into local variable
<code>dstore</code>	Store <code>double</code> into local variable
<code>astore</code>	Store <code>reference</code> or <code>returnAddress</code> into local variable

istore_0 Store int into local variable 0
istore_1 Store int into local variable 1
istore_2 Store int into local variable 2
istore_3 Store int into local variable 3
lstore_0 Store long into local variable 0
lstore_1 Store long into local variable 1
lstore_2 Store long into local variable 2
lstore_3 Store long into local variable 3
fstore_0 Store float into local variable 0
fstore_1 Store float into local variable 1
fstore_2 Store float into local variable 2
fstore_3 Store float into local variable 3
dstore_0 Store double into local variable 0
dstore_1 Store double into local variable 1
dstore_2 Store double into local variable 2
dstore_3 Store double into local variable 3
astore_0 Store reference or returnAddress into local variable 0
astore_1 Store reference or returnAddress into local variable 1
astore_2 Store reference or returnAddress into local variable 2
astore_3 Store reference or returnAddress into local variable 3
iastore Store into int array
lastore Store into long array
fastore Store into float array
dastore Store into double array
aastore Store into reference array
bastore Store into byte or boolean array
castore Store into char array
sastore Store into short array

The wide instruction

wide Extend a local variable index with additional bytes

Generic (typeless) stack operations

nop Do nothing
pop Pop top stack word
pop2 Pop top two stack words
dup Duplicate top stack word
dup_x1 Duplicate top stack word and put two down

dup_x2	Duplicate top stack word and put three down
dup2	Duplicate top two stack words
dup2_x1	Duplicate top two stack words and put two down
dup2_x2	Duplicate top two stack words and put three down
swap	Swap top two stack words

Type Conversion (Chapter 11)

i2l	Convert int to long
i2f	Convert int to float
i2d	Convert int to double
l2i	Convert long to int
l2f	Convert long to float
l2d	Convert long to double
f2i	Convert float to int
f2l	Convert float to long
f2d	Convert float to double
d2i	Convert double to int
d2l	Convert double to long
d2f	Convert double to float
i2b	Convert int to byte
i2c	Convert int to char
i2s	Convert int to short

Integer Arithmetic (Chapter 12)

iadd	Add ints
ladd	Add longs
isub	Subtract ints
lsub	Subtract longs
imul	Multiply ints
lmul	Multiply longs
idiv	Divide ints
ldiv	Divide longs
irem	Calculate remainder of division of ints
lrem	Calculate remainder of division of longs
ineg	Negate int
lneg	Negate long
iinc	Increment int local variable by constant

Logic (Chapter 13)

Shift operations

ishl	Perform left shift on <code>int</code>
lshl	Perform left shift on <code>long</code>
ishr	Perform arithmetic right shift on <code>int</code>
lshr	Perform arithmetic right shift on <code>long</code>
iushr	Perform logical right shift on <code>int</code>
lushr	Perform logical right shift on <code>long</code>

Bitwise boolean operations

iand	Perform boolean AND on <code>ints</code>
land	Perform boolean AND on <code>longs</code>
ior	Perform boolean OR on <code>ints</code>
lor	Perform boolean OR on <code>longs</code>
ixor	Perform boolean XOR on <code>ints</code>
lxor	Perform boolean XOR on <code>longs</code>

Floating Point Arithmetic (Chapter 14)

fadd	Add <code>floats</code>
dadd	Add <code>doubles</code>
fsub	Subtract <code>floats</code>
dsub	Subtract <code>doubles</code>
fmul	Multiply <code>floats</code>
dmul	Multiply <code>doubles</code>
fdiv	Divide <code>floats</code>
ddiv	Divide <code>doubles</code>
frem	Calculate remainder of division of <code>floats</code>
drem	Calculate remainder of division of <code>doubles</code>
fneg	Negate <code>float</code>
dneg	Negate <code>double</code>

Objects and Arrays (Chapter 15)

Instructions that deal with objects

new	Create a new object
checkcast	Make sure object is of a given type

getfield	Fetch field from object
putfield	Set field in object
getstatic	Fetch static field from class
putstatic	Set static field in class
instanceof	Determine if an object is of a given type

Instructions that deal with arrays

newarray	Allocate new array of primitive type components
anewarray	Allocate new array of reference type components
arraylength	Get length of an array
multianewarray	Allocate a new multi-dimensional array

Control Flow (Chapter 16)

Conditional branch instructions

ifeq	Branch if equal to 0
ifne	Branch if not equal to 0
iflt	Branch if less than 0
ifge	Branch if greater than or equal to 0
ifgt	Branch if greater than 0
ifle	Branch if less than or equal to 0
if_icmpeq	Branch if <code>ints</code> equal
if_icmpne	Branch if <code>ints</code> not equal
if_icmplt	Branch if <code>int</code> less than other <code>int</code>
if_icmpge	Branch if <code>int</code> greater than or equal to other <code>int</code>
if_icmpgt	Branch if <code>int</code> greater than other <code>int</code>
if_icmple	Branch if <code>int</code> less than or equal to other <code>int</code>
ifnull	Branch if <code>null</code>
ifnonnull	Branch if not <code>null</code>
if_acmpeq	Branch if object references are equal
if_acmpne	Branch if object references not equal

Comparison instructions

lcmp	Compare <code>longs</code>
fcmpl	Compare <code>floats</code> (-1 on NaN)
fcmpg	Compare <code>floats</code> (1 on NaN)
dcmpl	Compare <code>doubles</code> (-1 on NaN)
dcmpg	Compare <code>doubles</code> (1 on NaN)

Unconditional branch instructions

goto	Branch always
goto_w	Branch always (wide index)

Table jumping instructions

tableswitch	Access jump table by index and jump
lookupswitch	Access jump table by key match and jump

Exceptions (Chapter 17)

athrow	Throw exception or error
--------	--------------------------

Finally Clauses (Chapter 18)

jsr	Jump to subroutine
jsr_w	Jump to subroutine (wide index)
ret	Return from subroutine

Method Invocation and Return (Chapter 19)

Method invocation instructions

invokevirtual	Invoke instance method, dispatch based on run-time type
invokespecial	Invoke instance method, dispatching based on compile-time type
invokestatic	Invoke a class (static) method
invokeinterface	Invoke interface method

Method return instructions

ireturn	Return int from method
lreturn	Return long from method
freturn	Return float from method
dreturn	Return double from method
areturn	Return reference from method
return	Return (void) from method

Thread Synchronization (Chapter 20)

monitorenter	Enter and acquire object monitor
--------------	----------------------------------

`monitorexit`

Release and exit object monitor

Appendix C

Opcode Mnemonics by Opcode

Standard Opcodes

0	0x00	nop
1	0x01	aconst_null
2	0x02	iconst_m1
3	0x03	iconst_0
4	0x04	iconst_1
5	0x05	iconst_2
6	0x06	iconst_3
7	0x07	iconst_4
8	0x08	iconst_5
9	0x09	lconst_0
10	0x0a	lconst_1
11	0x0b	fconst_0
12	0x0c	fconst_1
13	0x0d	fconst_2
14	0x0e	dconst_0
15	0x0f	dconst_1
16	0x10	bipush
17	0x11	sipush
18	0x12	ldc
19	0x13	ldc_w
20	0x14	ldc2_w
21	0x15	iload
22	0x16	lload
23	0x17	fload
24	0x18	dload
25	0x19	aload
26	0x1a	iload_0

27	0x1b	iload_1
28	0x1c	iload_2
29	0x1d	iload_3
30	0x1e	lload_0
31	0x1f	lload_1
32	0x20	lload_2
33	0x21	lload_3
34	0x22	fload_0
35	0x23	fload_1
36	0x24	fload_2
37	0x25	fload_3
38	0x26	dload_0
39	0x27	dload_1
40	0x28	dload_2
41	0x29	dload_3
42	0x2a	aload_0
43	0x2b	aload_1
44	0x2c	aload_2
45	0x2d	aload_3
46	0x2e	iaload
47	0x2f	laload
48	0x30	faload
49	0x31	daload
50	0x32	aaload
51	0x33	baload
52	0x34	caload
53	0x35	saload
54	0x36	istore
55	0x37	lstore
56	0x38	fstore
57	0x39	dstore
58	0x3a	astore
59	0x3b	istore_0
60	0x3c	istore_1
61	0x3d	istore_2
62	0x3e	istore_3
63	0x3f	lstore_0
64	0x40	lstore_1
65	0x41	lstore_2
66	0x42	lstore_3
67	0x43	fstore_0

68	0x44	fstore_1
69	0x45	fstore_2
70	0x46	fstore_3
71	0x47	dstore_0
72	0x48	dstore_1
73	0x49	dstore_2
74	0x4a	dstore_3
75	0x4b	astore_0
76	0x4c	astore_1
77	0x4d	astore_2
78	0x4e	astore_3
79	0x4f	iastore
80	0x50	lastore
81	0x51	fastore
82	0x52	dastore
83	0x53	aastore
84	0x54	bastore
85	0x55	castore
86	0x56	sastore
87	0x57	pop
88	0x58	pop2
89	0x59	dup
90	0x5a	dup_x1
91	0x5b	dup_x2
92	0x5c	dup2
93	0x5d	dup2_x1
94	0x5e	dup2_x2
95	0x5f	swap
96	0x60	iadd
97	0x61	ladd
98	0x62	fadd
99	0x63	dadd
100	0x64	isub
101	0x65	lsub
102	0x66	fsub
103	0x67	dsub
104	0x68	imul
105	0x69	lmul
106	0x6a	fmul
107	0x6b	dmul
108	0x6c	idiv

109	0x6d	ldiv
110	0x6e	fdiv
111	0x6f	ddiv
112	0x70	irem
113	0x71	lrem
114	0x72	frem
115	0x73	drem
116	0x74	ineg
117	0x75	lneg
118	0x76	fneg
119	0x77	dneg
120	0x78	ishl
121	0x79	lshl
122	0x7a	ishr
123	0x7b	lshr
124	0x7c	iushr
125	0x7d	lushr
126	0x7e	iand
127	0x7f	land
128	0x80	ior
129	0x81	lor
130	0x82	ixor
131	0x83	lxor
132	0x84	iinc
133	0x85	i2l
134	0x86	i2f
135	0x87	i2d
136	0x88	l2i
137	0x89	l2f
138	0x8a	l2d
139	0x8b	f2i
140	0x8c	f2l
141	0x8d	f2d
142	0x8e	d2i
143	0x8f	d2l
144	0x90	d2f
145	0x91	i2b
146	0x92	i2c
147	0x93	i2s
148	0x94	lcmp
149	0x95	fcmpl

150	0x96	fcmpg
151	0x97	dcmpl
152	0x98	dcmpg
153	0x99	ifeq
154	0x9a	ifne
155	0x9b	iflt
156	0x9c	ifge
157	0x9d	ifgt
158	0x9e	ifle
159	0x9f	if_icmpeq
160	0xa0	if_icmpne
161	0xa1	if_icmplt
162	0xa2	if_icmpge
163	0xa3	if_icmpgt
164	0xa4	if_icmple
165	0xa5	if_acmpeq
166	0xa6	if_acmpne
167	0xa7	goto
168	0xa8	jsr
169	0xa9	ret
170	0xaa	tableswitch
171	0xab	lookupswitch
172	0xac	ireturn
173	0xad	lreturn
174	0xae	freturn
175	0xaf	dreturn
176	0xb0	areturn
177	0xb1	return
178	0xb2	getstatic
179	0xb3	putstatic
180	0xb4	getfield
181	0xb5	putfield
182	0xb6	invokevirtual
183	0xb7	invokespecial
184	0xb8	invokestatic
185	0xb9	invokeinterface
187	0xbb	new
188	0xbc	newarray
189	0xbd	anewarray
190	0xbe	arraylength
191	0xbf	athrow

192	0xc0	checkcast
193	0xc1	instanceof
194	0xc2	monitorenter
195	0xc3	monitorexit
196	0xc4	wide
197	0xc5	multianewarray
198	0xc6	ifnull
199	0xc7	ifnonnull
200	0xc8	goto_w
201	0xc9	jsr_w

Quick Opcodes

203	0xcb	ldc_quick
204	0xcc	ldc_w_quick
205	0xcd	ldc2_w_quick
206	0xce	getfield_quick
207	0xcf	putfield_quick
208	0xd0	getfield2_quick
209	0xd1	putfield2_quick
210	0xd2	getstatic_quick
211	0xd3	putstatic_quick
212	0xd4	getstatic2_quick
213	0xd5	putstatic2_quick
214	0xd6	invokevirtual_quick
215	0xd7	invokenonvirtual_quick
216	0xd8	invokesuper_quick
217	0xd9	invokestatic_quick
218	0xda	invokeinterface_quick
219	0xdb	invokevirtualobject_quick
221	0xdd	new_quick
222	0xde	anewarray_quick
223	0xdf	multianewarray_quick
224	0xe0	checkcast_quick
225	0xe1	instanceof_quick
226	0xe2	invokevirtual_quick_w
227	0xe3	getfield_quick_w
228	0xe4	putfield_quick_w

Reserved Opcodes

202	0xca	breakpoint
254	0xfe	impdep1
255	0xff	impdep2